

Machine Learning for OpenCV 4

Second Edition

Intelligent algorithms for building image processing apps using
OpenCV 4, Python, and scikit-learn



Packt

www.packt.com

Aditya Sharma, Vishwesh Ravi Shrimali
and Michael Beyeler

Machine Learning for OpenCV 4 *Second Edition*

Intelligent algorithms for building image processing apps using
OpenCV 4, Python, and scikit-learn

Aditya Sharma
Vishwesh Ravi Shrimali

Michael Beyeler



BIRMINGHAM - MUMBAI

Machine Learning for OpenCV 4 Second Edition

Copyright © 2019 Packt Publishing

All rights reserved. No part of this book may be reproduced, stored in a retrieval system, or transmitted in any form or by any means, without the prior written permission of the publisher, except in the case of brief quotations embedded in critical articles or reviews.

Every effort has been made in the preparation of this book to ensure the accuracy of the information presented. However, the information contained in this book is sold without warranty, either express or implied. Neither the authors, nor Packt Publishing or its dealers and distributors, will be held liable for any damages caused or alleged to have been caused directly or indirectly by this book. ...



Packt.com

Subscribe to our online digital library for full access to over 7,000 books and videos, as well as industry leading tools to help you plan your personal development and advance your career. For more information, please visit our website.

Why subscribe?

- Spend less time learning and more time coding with practical eBooks and Videos from over 4,000 industry professionals
- Improve your learning with Skill Plans built especially for you
- Get a free eBook or video every month
- Fully searchable for easy access to vital information
- Copy and paste, print, and bookmark content

Did you know that Packt offers eBook versions of every book published, with PDF and ePub files available? You can upgrade to the eBook version at www.packt.com and as a print book customer, you are entitled to a discount on the eBook copy. Get in touch with us at customercare@packtpub.com for more details.

At www.packt.com, you can also read a collection of free technical articles, sign up for a range ...

Contributors

About the authors

Aditya Sharma is a senior engineer at Robert Bosch working on solving real-world autonomous computer vision problems. At Robert Bosch, he also secured first place at an AI hackathon 2019. He has been associated with some of the premier institutes of India, including IIT Mandi and IIIT Hyderabad. At IIT, he published papers on medical imaging using deep learning at ICIP 2019 and MICCAI 2019. At IIIT, his work revolved around document image super-resolution.

He is a motivated writer and has written many articles on machine learning and deep learning for DataCamp and LearnOpenCV. Aditya runs his own YouTube channel and has contributed as a speaker at the NCVPRIPG conference (2017) and Aligarh Muslim University for a workshop ...

About the reviewers

Wilson Choo is a deep learning engineer working on deep learning modeling research. He has a deep interest in creating applications that implement deep learning, computer vision, and machine learning.

His past work includes the validation and benchmarking of Intel OpenVINO Toolkit algorithms, as well as custom Android OS validation. He has experience in integrating deep learning applications in different hardware and OSes. His native programming languages are Java, Python, and C++.

Robert B. Fisher has a PhD from the University of Edinburgh, where he also served as a college dean of research. He is currently the industrial liaison committee chair for the International Association for Pattern Recognition. His research covers topics mainly in high-level computer vision and 3D video analysis, which has led to 5 books and 300 peer-reviewed scientific articles or book chapters (Google H-index: 46). Most recently, he has been the coordinator of an EC-funded project that's developing a gardening robot. He has developed several online computer vision resources with over 1 million hits. He is a fellow of the International Association for Pattern Recognition and the British Machine Vision Association.

Packt is searching for authors like you

If you're interested in becoming an author for Packt, please visit authors.packtpub.com and apply today. We have worked with thousands of developers and tech professionals, just like you, to help them share their insight with the global tech community. You can make a general application, apply for a specific hot topic that we are recruiting an author for, or submit your own idea.

What this book covers

Chapter 1, *A Taste of Machine Learning*, starts us off with installing the required software and Python modules for this book.

Chapter 2, *Working with Data in OpenCV*, takes a look at some basic OpenCV functions.

Chapter 3, *First Steps in Supervised Learning*, will cover the basics of supervised learning methods in machine learning. We will have a look at some examples of supervised learning methods using OpenCV and the scikit-learn library in Python.

Chapter 4, *Representing Data and Engineering Features*, will cover concepts such as feature detection and feature recognition using ORB in OpenCV. We will also try to understand important concepts such as the curse of dimensionality.

Chapter 5, *Using Decision Trees to Make a Medical Diagnosis*, will introduce decision trees and important concepts related to them, including the depth of trees and techniques such as pruning. We will also cover a practical application of predicting breast cancer diagnoses using decision trees.

Chapter 6, *Detecting Pedestrians with Support Vector Machines*, will start off with an introduction to support vector machines and how they can be implemented in OpenCV. We will also cover an application of pedestrian detection using OpenCV.

Chapter 7, *Implementing a Spam Filter with Bayesian Learning*, will discuss techniques such as the Naive Bayes algorithm, multinomial

Naive Bayes, and more, as well as how they can be implemented. Finally, we will build a machine learning application to classify data into spam and ham.

[Chapter 8](#), *Discovering Hidden Structures with Unsupervised Learning*, will be our first introduction to the second class of machine learning algorithms—unsupervised learning. We will discuss techniques such as clustering using k-nearest neighbors, k-means, and more.

[Chapter 9](#), *Using Deep Learning to Classify Handwritten Digits*, will introduce deep learning techniques and we will see how we can use deep neural networks to classify images from the MNIST dataset.

[Chapter 10](#), *Ensemble Methods for Classification*, will cover topics such as random forest, bagging, and boosting for classification purposes.

[Chapter 11](#), *Selecting the Right Model with Hyperparameter Tuning*, will go over the process of selecting the optimum set of parameters in various machine learning methods in order to improve the performance of a model.

[Chapter 12](#), *Using OpenVINO with OpenCV*, will introduce OpenVINO Toolkit, which was introduced in OpenCV 4.0. We will also go over how we can use it in OpenCV using image classification as an example.

[Chapter 13](#), *Conclusion*, will provide a summary of the major topics that we have covered in the book and talk about what you can do next.

Table of Contents

[Title Page](#)

[Copyright and Credits](#)

[Machine Learning for OpenCV 4 Second Edition](#)

[About Packt](#)

[Why subscribe?](#)

[Contributors](#)

[About the authors](#)

[About the reviewers](#)

[Packt is searching for authors like you](#)

[Preface](#)

[Who this book is for](#)

[What this book covers](#)

[To get the most out of this book](#)

[Download the example code files](#)

[Download the color images](#)

[Conventions used](#)

[Get in touch](#)

[Reviews](#)

[1. Section 1: Fundamentals of Machine Learning and OpenCV](#)

[1. A Taste of Machine Learning](#)

[Technical requirements](#)

[Getting started with machine learning](#)

[Problems that machine learning can solve](#)

[Getting started with Python](#)

[Getting started with OpenCV](#)

[Installation](#)

[Getting the latest code for this book](#)

Preface

As the world changes and humans build smarter and better machines, the demand for machine learning and computer vision experts increases. Machine learning, as the name suggests, is the process of a machine learning to make predictions given a certain set of parameters as input. Computer vision, on the other hand, gives a machine vision; that is, it makes the machine aware of visual information. When you combine these technologies, you get a machine that can use visual data to make predictions, which brings machines one step closer to having human capabilities. When you add deep learning to it, the machine can even surpass human capabilities in terms of making predictions. This might seem far-fetched, but with AI systems taking over decision-based systems, this has actually become a reality. You have AI cameras, AI monitors, AI sound systems, AI-powered processors, and more. We cannot promise you that you will be able to build an AI camera after reading this book, but we do intend to provide you with the tools necessary for you to do so. The most powerful tool that we are going to introduce is the OpenCV library, which is the world's largest computer vision library. Even though its use in machine learning is not very common, we have provided some examples and concepts on how it can be used for machine learning. We have gone with a hands-on approach in this book and we recommend that you try out every single piece of code present in this book to build an application that showcases your knowledge. The world is changing and this book is our way of helping young minds change it for the better.

Who this book is for

We have tried to explain all the concepts from scratch to make the book suitable for beginners as well as advanced readers. We recommend that readers have some basic knowledge of Python programming, but it's not mandatory. Whenever you encounter some Python syntax that you are not able to understand, make sure you look it up on the internet. *Help is always provided to those who look for it.*

To get the most out of this book

We recommend that you go through any good Python programming book or online tutorials or videos, if you are a beginner in Python. You can also have a look at DataCamp (<http://www.datacamp.com>) to learn Python using interactive lessons.

We also recommend that you learn some basic concepts about the Matplotlib library in Python. You can try out this tutorial for that: <https://www.datacamp.com/community/tutorials/matplotlib-tutorial-python>.

You don't need to have anything installed on your system for this book before starting it. We will cover all the installation steps in the first chapter.

Download the example code files

You can download the example code files for this book from your account at www.packt.com. If you purchased this book elsewhere, you can visit www.packtpub.com/support and register to have the files emailed directly to you.

You can download the code files by following these steps:

1. Log in or register at www.packt.com.
2. Select the Support tab.
3. Click on Code Downloads.
4. Enter the name of the book in the Search box and follow the onscreen instructions.

Once the file is downloaded, please make sure that you unzip or extract the folder using the latest version of:

- WinRAR/7-Zip for Windows
- Zipeg/iZip/UnRarX for Mac
- 7-Zip/PeaZip for Linux

The code bundle for the book is also hosted on GitHub at <https://github.com/PacktPublishing/Machine-Learning-for-OpenCV-Second-Edition> ...

Download the color images

We also provide a PDF file that has color images of the screenshots/diagrams used in this book. You can download it here: http://www.packtpub.com/sites/default/files/downloads/9781789536300_ColorImages.pdf.

Conventions used

There are a number of text conventions used throughout this book.

CodeInText: Indicates code words in text, database table names, folder names, filenames, file extensions, pathnames, dummy URLs, user input, and Twitter handles. Here is an example: "We can set both `max_samples<1.0` and `max_features<1.0` to implement the random patches method."

A block of code is set as follows:

```
In [1]: from sklearn.ensemble import BaggingClassifier... from  
sklearn.neighbors import KNeighborsClassifier... bag_knn =  
BaggingClassifier(KNeighborsClassifier(), ... n_estimators=10)
```

Any command-line input or output is written as follows:

```
$ conda install package_name
```

Bold: Indicates a new term, an important word, or words that you see onscreen.

Warnings ...

Get in touch

Feedback from our readers is always welcome.

General feedback: If you have questions about any aspect of this book, mention the book title in the subject of your message and email us at customercare@packtpub.com.

Errata: Although we have taken every care to ensure the accuracy of our content, mistakes do happen. If you have found a mistake in this book, we would be grateful if you would report this to us. Please visit www.packtpub.com/support/errata, selecting your book, clicking on the Errata Submission Form link, and entering the details.

Piracy: If you come across any illegal copies of our works in any form on the internet, we would be grateful if you would provide us with the location address or website name. Please contact us at copyright@packt.com with a link to the material.

If you are interested in becoming an author: If there is a topic that you have expertise in and you are interested in either writing or contributing to a book, please visit authors.packtpub.com.

Reviews

Please leave a review. Once you have read and used this book, why not leave a review on the site that you purchased it from? Potential readers can then see and use your unbiased opinion to make purchase decisions, we at Packt can understand what you think about our products, and our authors can see your feedback on their book. Thank you!

For more information about Packt, please visit packt.com.

Section 1: Fundamentals of Machine Learning and OpenCV

In the very first section of this book, we will go over the basics of machine learning and OpenCV, starting with installing the required libraries, and then moving on to basic OpenCV functions, the basics of supervised learning and their applications, and finally, feature detection and recognition using OpenCV.

This section includes the following chapters:

- [Chapter 1, A Taste of Machine Learning](#)
- [Chapter 2, Working with Data in OpenCV](#)
- [Chapter 3, First Steps in Supervised Learning](#)
- [Chapter 4, Representing Data and Engineering Features](#)

A Taste of Machine Learning

So, you have decided to enter the field of machine learning. That's great!

Nowadays, machine learning is all around us—from protecting our email, to automatically tagging our friends in pictures, to predicting what movies we like. As a form of artificial intelligence, machine learning enables computers to learn through experience; to make predictions about the future using collected data from the past. On top of that, computer vision is one of today's most exciting application fields of machine learning, with deep learning and convolutional neural networks driving innovative systems such as self-driving cars and Google's DeepMind.

However, fret not; your application does not need to be as large-scale or world-changing ...

Technical requirements

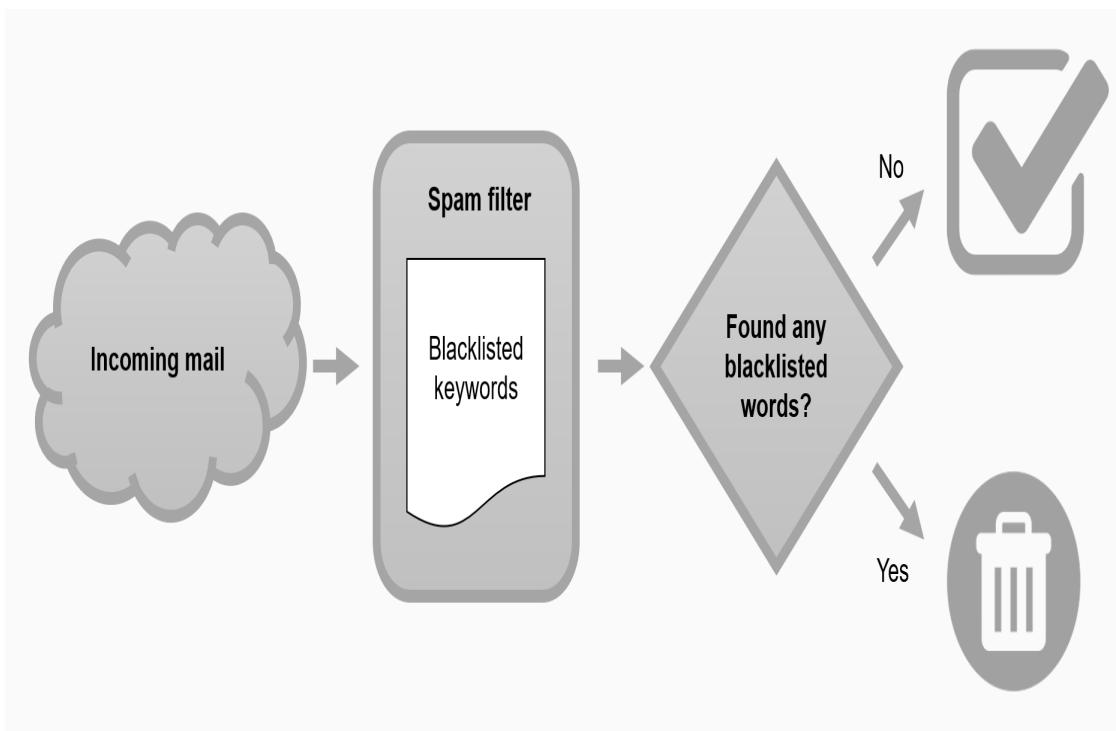
You can refer to the code for this chapter at the following link: <https://github.com/PacktPublishing/Machine-Learning-for-OpenCV-Second-Edition/tree/master/Chapter01>.

Here is a short summary of the software and hardware requirements:

- OpenCV version 4.1.x (4.1.0 or 4.1.1 will both work just fine).
- Python version 3.6 (any Python version 3.x will be fine).
- Anaconda Python 3 for installing Python and the required modules.
- You can use any OS—macOS, Windows, and Linux-based OS—with this book. We recommend you have at least 4 GB RAM in your system.
- You don't need to have a GPU to run the code provided with this book.

Getting started with machine learning

Machine learning has been around for at least 60 years. Growing out of the quest for artificial intelligence, early machine learning systems inferred the hand-coded rules of `if...else` statements to process data and make decisions. Think of a spam filter whose job is to parse incoming emails and move unwanted messages to a spam folder as shown here in the following diagram:



We could come up with a blacklist of words that, whenever they show up in a message, would mark an email as spam. This is a simple example of a hand-coded expert system. (We will build a smarter one in [Chapter 7, Implementing a Spam ...](#)

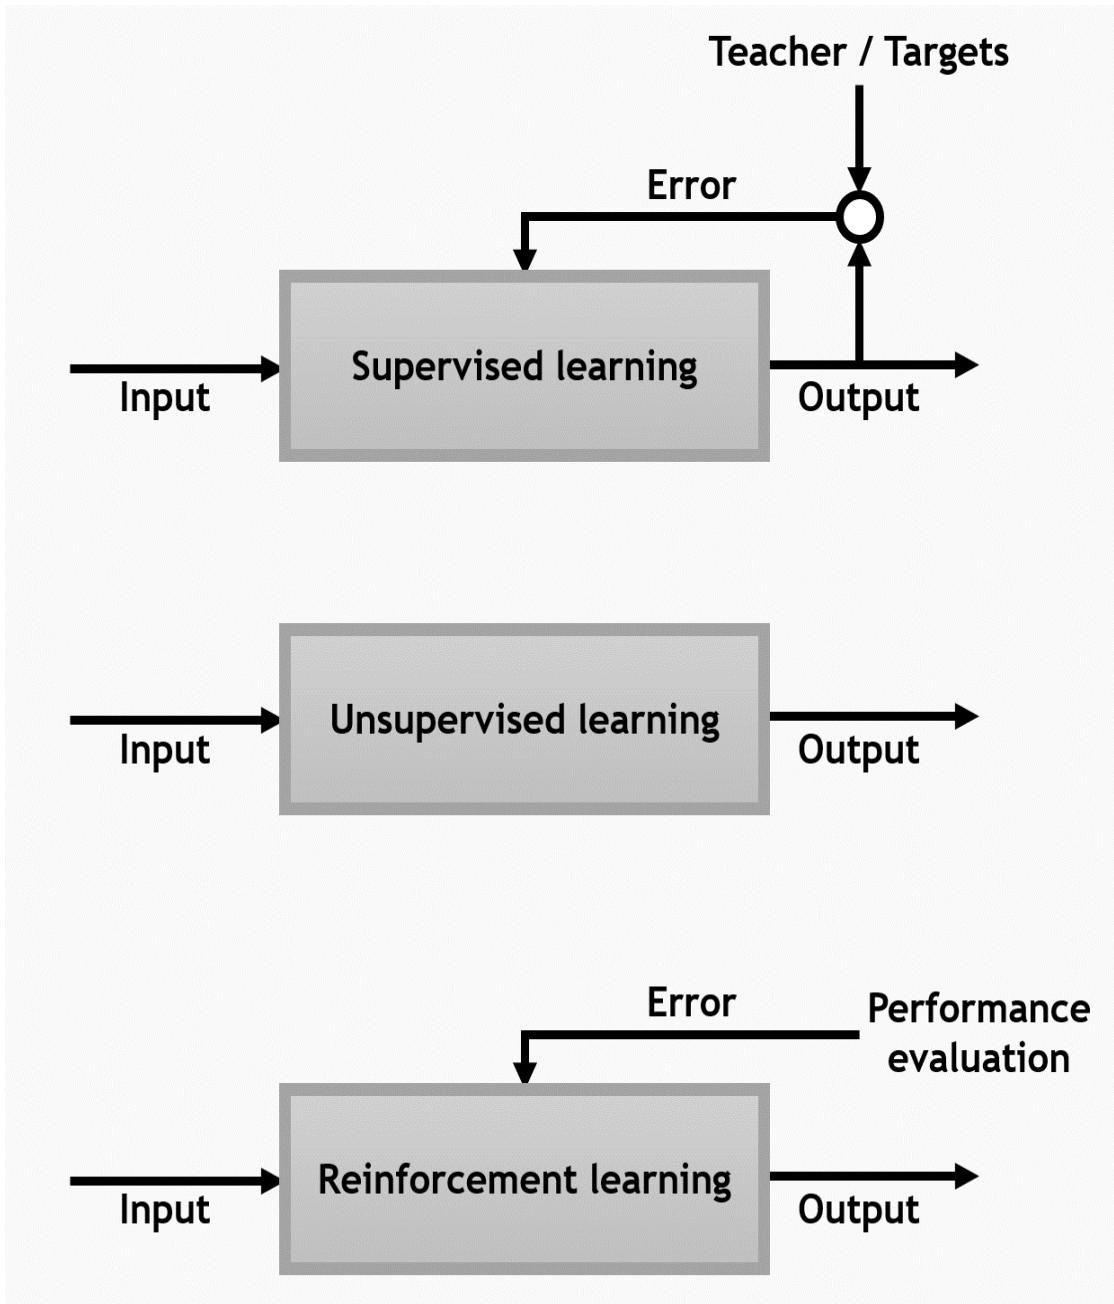
Problems that machine learning can solve

Most machine learning problems belong to one of the following three main categories:

- In supervised learning, we have what is referred to as the label for a data point. Now, this can be the class of an object that is captured in the image, a bounding box around a face, the digit present in the image, or anything else. Think of it as a teacher who teaches but also tells you what the correct answer is to a problem. Now, the student can try to devise a model or an equation that takes into account all the problems and their correct answers and finds out the answer to a problem that does (or does not) have a correct answer. The data that goes into learning the model is called the **training data** and the data on which the process/model is tested is called **test data**. These predictions come in two flavors, such as identifying new photos with the correct animal (called a **classification problem**) or assigning accurate sale prices to other used cars (called a **regression problem**). Don't worry if this seems a little over your head for now—we will have the entirety of the book to nail down the details.

- In unsupervised learning, data points have no labels associated with them ([Chapter 8, Discovering Hidden Structures with Unsupervised Learning](#)). Think of it like a class where the instructor gives you a jumbled puzzle and leaves it up to you to figure out what to do. Here, the most common result is the clusters, which contain objects with similar characteristics. It can also result in different ways of looking at higher dimensional data (complex data) so that it appears simpler.
- Reinforcement learning is about maximizing a reward in a problem. So, if the teacher gives you a candy for every correct answer and punishes you for every incorrect answer, he/she is reinforcing the concepts by making you increase the number of times you receive candies rather than the number of times you are subjected to a punishment.

These three main categories are illustrated in the following diagram:



Now that we have covered the main machine learning categories, let's go over some concepts in Python that will prove very useful along the journey of this book.

Getting started with Python

Python has become the common language for many data science and machine learning applications, thanks to its great number of open source libraries for processes such as data loading, data visualization, statistics, image processing, and natural language processing. One of the main advantages of using Python is the ability to interact directly with the code, using a Terminal or other tools such as the **Jupyter Notebook**, which we'll look at shortly.

If you have mostly been using OpenCV in combination with C++, I would strongly suggest that you switch to Python, at least for the purpose of studying this book. This decision has not been made out of spite! Quite the contrary: I have done my fair share of C/C++ programming—especially ...

Getting started with OpenCV

Being the avid user of OpenCV that I believe you are, I probably don't have to convince you about the power of OpenCV.

Built to provide a common infrastructure for computer vision applications, OpenCV has become a comprehensive set of both classic and state-of-the-art computer vision and machine learning algorithms. According to their own documentation, OpenCV has a user community of more than 47,000 people and has been downloaded over seven million times. That's pretty impressive! As an open source project, it is very easy for researchers, businesses, and government bodies to utilize and modify already available code.

That being said, a number of open source machine learning libraries have popped up as part of the recent machine learning boom that provide far more functionality than OpenCV. A prominent example is **scikit-learn**, which provides a number of state-of-the-art machine learning algorithms as well as a wealth of online tutorials and code snippets. As OpenCV was developed mainly to provide computer vision algorithms, its machine learning functionality is restricted to a single module, called `m1`. As we will see in this book, OpenCV still provides a number of state-of-the-art algorithms, but sometimes lacks a bit in functionality. In these rare cases, instead of reinventing the wheel, we will simply use scikit-learn for our purposes.

Last but not least, installing OpenCV using the Python Anaconda distribution is essentially a one-liner as we'll see in the following sections.

If you are a more advanced user who wants to build real-time applications, OpenCV's algorithms are well-optimized for this task, and Python provides

*several ways to speed up computations where it is necessary (using, for example, **Cython** or parallel processing libraries such as **joblib** or **dask**).*

Installation

Before we get started, let's make sure that we have all the tools and libraries installed that are necessary to create a fully functioning data science environment. After downloading the latest code for this book from GitHub, we are going to install the following software:

- Python's Anaconda distribution, based on Python 3.6 or higher
- OpenCV 4.1
- Some supporting packages

*Don't feel like installing stuff? You can also visit <https://mybinder.org/v2/gh/PacktPublishing/Machine-Learning-for-OpenCV-Second-Edition/master>, where you will find all the code for this book in an interactive, executable environment and 100% free and open source, thanks to the **Binder** project.*

Getting the latest code for this book

You can get the latest code for this book from GitHub: <https://github.com/PacktPublishing/Machine-Learning-for-OpenCV-Second-Edition>. You can either download a `.zip` package (beginners) or clone the repository using Git (intermediate users).

Git is a version control system that allows you to track changes in files and collaborate with others on your code. In addition, the web platform GitHub makes it easy for people to share their code with you on a public server. As I make improvements to the code, you can easily update your local copy, file bug reports, or suggest code changes.

If you choose to go with git, the first step is to make sure it is installed (<https://git-scm.com/downloads>).

Then, open a Terminal (or Command Prompt, as it is called in Windows):

- On Windows 10, right-click on the Start Menu button, and select Command Prompt.
- On macOS X, press *Cmd + Space* to open spotlight search, then type `terminal`, and hit *Enter*.
- On Ubuntu, Linux/Unix and friends, press *Ctrl + Alt + T*. On Red Hat, right-click on the desktop and choose Open Terminal from the menu.

Navigate to a directory where you want the code downloaded:

```
cd Desktop
```

Then you can grab a local copy of the latest code by typing the following:

```
git clone https://github.com/PacktPublishing/Machine-Learning-for-  
OpenCV-Second-Edition.git
```

This will download the latest code in a folder called `opencv-ML`.

After a while, the code might change online. In that case, you can update your local copy by running the following command from within the `opencv-ML` directory:

```
git pull origin master
```

Getting to grips with Python's Anaconda distribution

Anaconda is a free Python distribution developed by Continuum Analytics that is made for scientific computing. It works across Windows, Linux, and macOS X platforms and is free, even for commercial use. However, the best thing about it is that it comes with a number of preinstalled packages that are essential for data science, math, and engineering. These packages include the following:

- **NumPy**: A fundamental package for scientific computing in Python that provides functionality for multidimensional arrays, high-level mathematical functions, and pseudo-random number generators
- **SciPy**: A collection of functions for scientific computing in Python that provides advanced linear algebra routines, ...

Installing OpenCV in a conda environment

We will carry out the following steps to install OpenCV:

1. In a Terminal, navigate to the directory where you downloaded the following code:

```
$ cd Desktop/OpenCV-ML
```

2. Then, run the following command to create a conda environment based on Python 3.6, which will also install all the necessary packages listed in the `environment.yml` file (available in the GitHub repository) in one fell swoop:

```
$ conda create env -f environment.yml
```

3. You can also have a look at the following `environment.yml` file:

```
name: OpenCV-ML
channels:
  - conda-forge
dependencies:
  - python==3.6
  - numpy==1.15.4
  - scipy==1.1.0
```

```
- scikit-learn==0.20.1
- matplotlib
- jupyter==1.0
- notebook==5.7.4
- pandas==0.23.4
- theano
- keras==2.2.4
- mkl-service==1.1.2
- pip
- pip:
  - opencv-contrib-python==4.1.0.25
```

Notice that the environment's name will be `opencv-ML`. This code will use the `conda-forge` channel to download all the `conda` based dependencies and use `pip` to install `opencv` 4.0 (along with `opencv_contrib`).

4. To activate the environment, type one of the following, depending on your platform:

```
$ source activate OpenCV-ML    # on Linux / Mac OS X
$ activate OpenCV-ML          # on Windows
```

5. When we close the Terminal, the session will be deactivated —so we will have to run this last command again the next time we open a new Terminal. We can also deactivate the environment by hand:

```
$ source deactivate  # on Linux / Mac OS X
$ deactivate        # on Windows
```

And done! Let's verify whether all this installation was successful or not.

Verifying the installation

It's a good idea to double-check our installation. While our terminal is still open, we start IPython, which is an interactive shell to run Python commands:

```
$ ipython
```

Next, make sure that you are running (at least) Python 3.6 and not Python 2.7. You might see the version number displayed in IPython's welcome message. If not, you can run the following commands:

```
In [1]: import sys...      print(sys.version)      3.6.0 | packaged by  
conda-forge | (default, Feb 9 2017, 14:36:55) [GCC 4.8.2 20140120 (Red  
Hat 4.8.2-15)]
```

Now try to import OpenCV as follows:

```
In [2]: import cv2
```

You should get no error messages. Then, try to find out the version number like so:

```
In [3]: cv2.__version__out[3]: '4.0.0'
```

Make sure that OpenCV's version ...

Getting a glimpse of OpenCV's ml module

Starting with OpenCV 3.1, all machine learning related functions in OpenCV have been grouped into the `ml` module. This has been the case for the C++ API for quite some time. You can get a glimpse of what's to come by displaying all functions in the `ml` module:

```
In [4]: dir(cv2.ml)
Out[4]: ['ANN_MLP_ANNEAL',
 'ANN_MLP_BACKPROP',
 'ANN_MLP_GAUSSIAN',
 'ANN_MLP_IDENTITY',
 'ANN_MLP_LEAKYRELU',
 'ANN_MLP_NO_INPUT_SCALE',
 'ANN_MLP_NO_OUTPUT_SCALE',
 ...
 '__spec__']
```

If you have installed an older version of OpenCV, the `ml` module might not be present. For example, the k-nearest neighbor algorithm (which we will talk about in [chapter 3](#), First Steps in Supervised Learning) used to be called `cv2.KNearest()` but is now called `cv2.ml.KNearest_create()`. In order to avoid confusion throughout the book, I recommend using OpenCV 4.0.

This is all good but you will be wondering by now why should you even learn machine learning, and what are its applications? Let's answer this question in the next section.

Applications of machine learning

Machine learning, artificial intelligence, deep learning, and data science are four terms that I believe are going to change the way we have always looked at things. Let's see if I can convince you why I believe so.

From making a computer learn how to play Go and defeat the world champion of the very same game to using the same branch to detect whether a person has a tumor or not just by seeing their brain's CT Scan, machine learning has left its mark in every single domain. One of the projects that I worked on was using machine learning to determine the residual life cycle of boiler water wall tubes in thermal power plants. The proposed solution was successful in saving a huge amount of money by using the ...

What's new in OpenCV 4.0?

So, we come to the last section of the very first chapter. I will keep it short and to the point since you as a reader can safely skip it. The topic of our discussion is **OpenCV 4.0**.

OpenCV 4.0 is a result of three and a half years of hard work and bug fixes by OpenCV and was finally released in November 2018. In this section, we will look at some of the major changes and new features in OpenCV 4.0:

- With the OpenCV 4.0 release, OpenCV has officially become a C++11 library. This means that you have to make sure that a C++11 compliant compiler is present in your system when you are trying to compile OpenCV 4.0.
- In continuation of the previous point, a lot of C APIs have been removed. Some of the modules that have been affected include Video IO module (`videoio`), Object Detection module (`objdetect`), and others. File IO for XML, YAML, and JSON have also removed the C API.
- OpenCV 4.0 also has a lot of improvements in the DNN module (the deep learning module). ONNX support has been added. **Intel OpenVINO** also marks its presence in the new OpenCV version. We will be looking into this in some more detail in our later chapters.

- OpenCL acceleration has been fixed on AMD and NVIDIA GPUs.
- OpenCV Graph API has also been added, which is a highly efficient engine for image processing and other operations.
- As in every OpenCV release, there have been a lot of changes with the purpose of improving the performance. Some new features such as QR Code Detection and Decoding have also been added.

In short, there have been a lot of changes in OpenCV 4.0 and they have their own uses. For example, ONNX support helps in the portability of models across various languages and frameworks, OpenCL reduces runtime for computer vision applications, Graph API helps in increasing the efficiency of the applications, and the OpenVINO toolkit uses Intel's processors and a model zoo to provide highly efficient deep learning models. We will be focusing primarily on OpenVINO toolkit and DLDT as well as accelerating computer vision applications in later chapters. But, I should also point out here that both OpenCV 3.4.4 and OpenCV 4.0.0 are being modified at a high speed to fix bugs. So, if you are going to use either of them in any application, be prepared to modify your code and installation to incorporate the changes made. On a similar note, OpenCV 4.0.1 and OpenCV 3.4.5 are also out within a few months of their predecessors.

Summary

In this chapter, we talked about machine learning at a high abstraction level: what it is, why it is important, and what kinds of problems it can solve. We learned that machine learning problems come in three flavors: supervised learning, unsupervised learning, and reinforcement learning. We talked about the prominence of supervised learning, and that this field can be further divided into two subfields: classification and regression. Classification models allow us to categorize objects into known classes (such as animals into cats and dogs), whereas regression analysis can be used to predict continuous outcomes of target variables (such as the sales price of used cars).

We also learned how to set up a data science environment using ...

Working with Data in OpenCV

Now that we have whetted our appetite for machine learning, it is time to delve a little deeper into the different parts that make up a typical machine learning system.

Far too often, you hear someone throw around the phrase, *Just apply machine learning to your data!*, as if that will instantly solve all of your problems. You can imagine that the reality of this is much more intricate, although, I will admit that nowadays, it is incredibly easy to build your own machine learning system simply by cutting and pasting a few lines of code from the internet. However, to build a system that is truly powerful and effective, it is essential to have a firm grasp of the underlying concepts and an intimate knowledge of the strengths and weaknesses of each method. So, don't worry if you don't consider yourself a machine learning expert just yet. Good things take time.

Earlier, I described machine learning as a subfield of artificial intelligence. This might be true—mainly for historical reasons—but most often, machine learning is simply about making sense of data. Therefore, it might be more suitable to think of machine learning as a subfield of data science, where we build mathematical models to help us to understand data.

Hence, this chapter is all about data. We want to learn how data fits in with machine learning and how to work with data using the tools of our choice: OpenCV and Python.

In this chapter, we will cover the following topics:

- Understanding the machine learning workflow
- Understanding training data and test data
- Learning how to load, store, edit, and visualize data with OpenCV and Python

Technical requirements

You can refer to the code for this chapter from the following link: <https://github.com/PacktPublishing/Machine-Learning-for-OpenCV-Second-Edition/tree/master/Chapter02>.

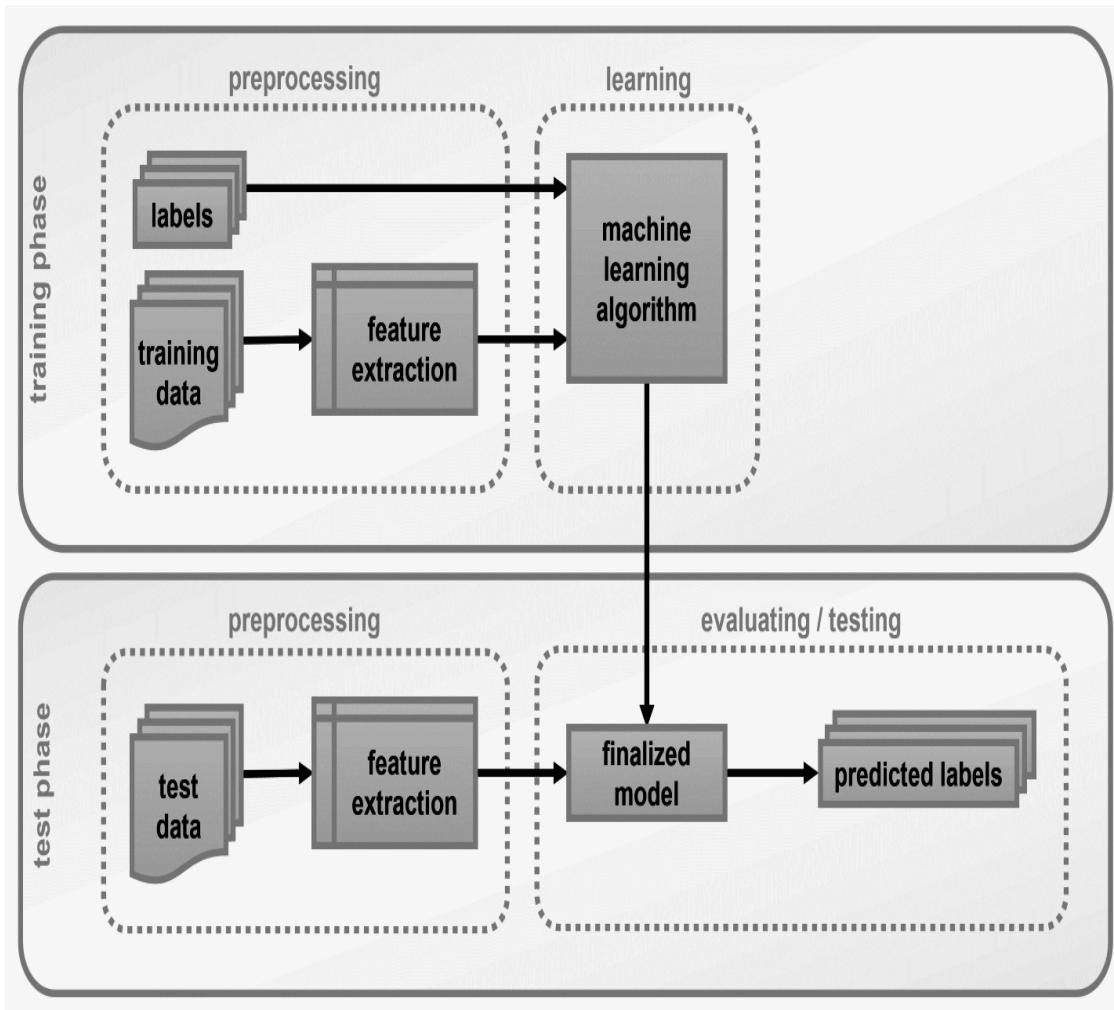
Here is a summary of the software and hardware requirements:

- You will need OpenCV version 4.1.x (4.1.0 or 4.1.1 will both work just fine).
- You will need Python version 3.6 (any Python version 3.x will be fine).
- You will need Anaconda Python 3 for installing Python and the required modules.
- You can use any OS—macOS, Windows, and Linux-based OSes along with this book. We recommend you have at least 4 GB RAM in your system.
- You don't need to have a GPU to run the code provided along with this book.

Understanding the machine learning workflow

As mentioned earlier, machine learning is all about building mathematical models to understand data. The learning aspect enters this process when we give a machine learning model the capability to adjust its **internal parameters**; we can tweak these parameters so that the model explains the data better. In a sense, this can be understood as the model learning from the data. Once the model has learned enough—whatever that means—we can ask it to explain newly observed data.

A typical classification process is illustrated in the following diagram:



Let's break it down step by step.

The first thing to notice is that machine learning problems are always split into (at least) two distinct phases:

- A **training phase**, during which we aim to train a machine learning model on a set of data that we call the **training dataset**
- A test phase, during which we evaluate the learned (or finalized) machine learning model on a new set of never-before-seen data that we call the **test dataset**

The importance of splitting our data into a training set and test set cannot be understated. We always evaluate our models on an independent test set because we are interested in knowing how well our models **generalize to new data**. In the end, isn't this what learning is all about—be it machine learning or human learning? Think back to school, when you were a learner yourself: the problems you had to solve as part of your homework would never show up in exactly the same form in the final exam. The same scrutiny should be applied to a machine learning model; we are not so much interested in how well our models can memorize a set of data points (such as a homework problem), but we want to know how our models will use what they have learned to solve new problems (such as the ones that show up in a final exam) and explain new data points.

*The workflow of an advanced machine learning problem will typically include a third set of data termed a **validation dataset**. For now, this distinction is not important. A validation set is typically formed by further partitioning the training dataset. It is used in advanced concepts such as model selection, which we will talk about in [Chapter 11, Selecting the Right Model with Hyperparameter Tuning](#), when we have become proficient in building machine learning systems.*

The next thing to notice is that machine learning is really all about the **data**. Data enters the previously described workflow diagram in its raw form—whatever that means—and is used in both training and test phases. Data can be anything from images and movies to text documents and audio files. Therefore, in its raw form, data might be made of pixels, letters, words, or even worse: pure bits. It is easy to see that data in such a raw form might not be very convenient to work with. Instead, we have to find ways to **preprocess** the data to bring it into a form that is easy to **parse or use the data**.

Data preprocessing comes in two stages:

- **Feature selection:** This is the process of identifying important attributes (or features) in the data. Possible features of an image might be the location of edges,

corners, or ridges. You might already be familiar with some more advanced feature descriptors that OpenCV provides, such as **Speeded Up Robust Features (SURF)** or the **Histogram of Oriented Gradients (HOG)**. Although these features can be applied to any image, they might not be that important (or work that well) for our specific task. For example, if our task was to distinguish between clean and dirty water, the most important feature might turn out to be the color of the water, and the use of SURF or HOG features might not help us much.

- **Feature extraction:** This is the actual process of transforming the raw data into the desired **feature space**. An example would be the **Harris operator**, which allows us to extract corners (that is, a selected feature) in an image.

A more advanced topic is the process of inventing informative features, which is known as **feature engineering**. After all, before people could select from popular features, someone had to invent them first. This is often more important for the success of our algorithm than the choice of the algorithm itself. We will talk about feature engineering extensively in [chapter 4, *Representing Data and Engineering Features*](#).

Don't let naming conventions confuse you! Sometimes, feature selection and feature extraction are hard to distinguish, mainly because of how things are named. For example, SURF stands for both the feature extractor as well as the actual name of the features. The same is true for the Scale-Invariant Feature Transform (SIFT), which is a feature extractor that yields what is known as SIFT features. Unfortunately, both the algorithms are patented and cannot be used for commercial purposes. We won't be sharing any code about either algorithms.

The last point to be made is that, in supervised learning, every data point must have a **label**. A label identifies a data point either of belonging to a certain class of things (such as cat or dog) or of having a certain value (such as the price of a house). At the end of the day, the goal of a supervised machine learning system is to predict the label of all data points in the test set (as shown in the previous diagram). We do this by learning regularities in the training data, using the labels that come with it, and then testing our performance on the test set.

Therefore, to build a functioning machine learning system, we first have to cover how to load, store, and manipulate data. How do you even do that in OpenCV with Python?

Dealing with data using OpenCV and Python

The world of data is full of various kinds of data types. This, at times, makes it very difficult for users to distinguish between the data type to use for a particular value. Here, we will try to keep it simple by treating everything as an array, except the scalar values, which will retain their standard data types. So, images will become 2D arrays because they have width and height. A 1D array could be a sound clip with intensity varying over time.

If you have mostly been using OpenCV's **C++ Application Programming Interface (API)** and plan on continuing to do so, you might find that dealing with data in C++ can be a bit of a pain. Not only will you have to deal with the syntactic overhead of the ...

Starting a new IPython or Jupyter session

Before we can get our hands on NumPy, we need to open an IPython shell or start a Jupyter Notebook:

1. Open a Terminal as we did in the previous chapter and navigate to the `OpenCV-ML` directory:

```
$ cd Desktop/OpenCV-ML
```

2. Activate the `conda` environment we created in the previous chapter:

```
$ source activate OpenCV-ML # Mac OS X / Linux  
$ activate OpenCV-ML       # Windows
```

3. Start a new IPython or Jupyter session:

```
$ ipython      # for an IPython session  
$ jupyter notebook # for a Jupyter session
```

If you chose to start an IPython session, the program should have greeted you with a welcome message like the following:

```
$ ipython
Python 3.6.0 | packaged by conda-forge | (default, Feb 9 2017,
14:36:55)
Type 'copyright', 'credits' or 'license' for more information
IPython 7.2.0 -- An enhanced Interactive Python. Type '?' for help.

In [1]:
```

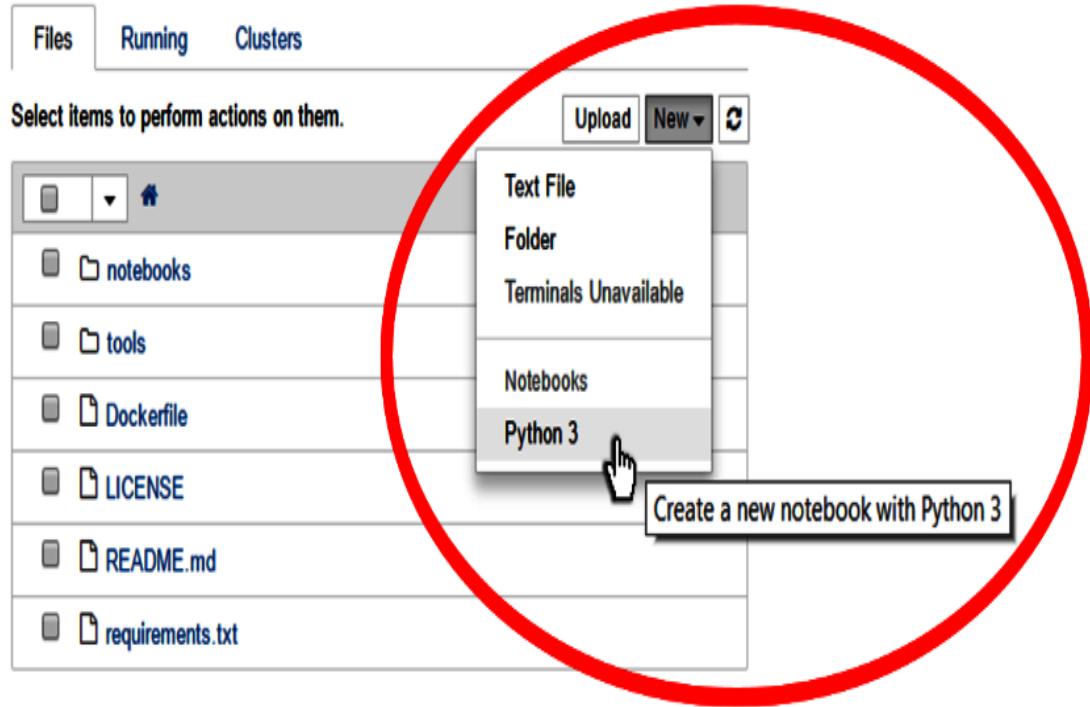
The line starting with `In [1]` is where you type in your regular Python commands. In addition, you can also use the *Tab* key while typing the names of variables and functions to have IPython automatically complete them.

A limited number of Unix and macOS system shell commands work too—such as `ls` and `pwd`. You can run any shell command by prefixing it with `!`, such as `!ping www.github.com`. For more information, check out the official IPython reference at <https://ipython.org/ipython-doc/3/interactive/tutorial.html>.

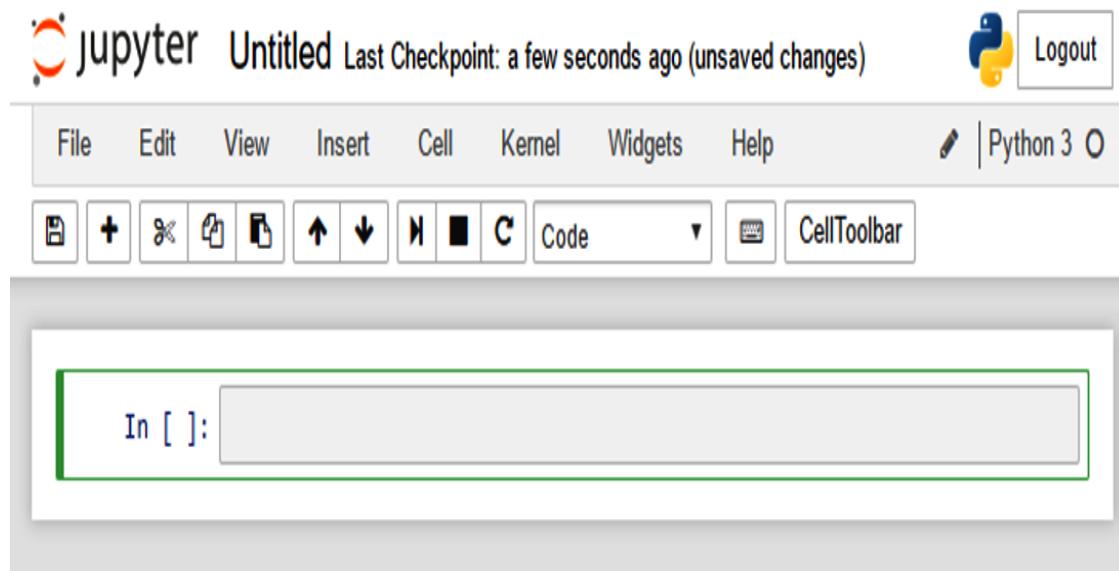
If you chose to start a Jupyter session, a new window should have opened in your web browser that is pointing to `http://localhost:8888`. You want to create a new notebook by clicking on New in the top-right corner and selecting Notebooks (Python 3):



Logout



This will open a new window that looks like this:



The cell (which looks like the preceding textbox) labeled with `In []` is the same as the command line in an IPython session. Now you

can start typing your Python code!

Dealing with data using Python's NumPy package

I am assuming that you already have NumPy installed in your virtual environment if you have installed Anaconda. If you have used Python's standard distribution or any other distribution, you can go to <http://www.numpy.org> and follow the installation instructions provided there.

As mentioned previously, it's okay if you aren't a Python expert yet. Who knows, perhaps you're just now switching from OpenCV's C++ API. This is all fine. I wanted to give you a quick overview of how to get started with NumPy. If you are a more advanced Python user, you can simply skip this section.

Once you are familiar with NumPy, you will find that most scientific computing tools in the Python world are built around ...

Importing NumPy

Once you start a new IPython or Jupyter session, you can import the NumPy module and verify its version as follows:

```
In [1]: import numpy  
In [2]: numpy.__version__  
Out[2]: '1.15.4'
```

Recall that in the Jupyter Notebook, you can hit `Ctrl + Enter` to execute a cell once you have typed the command. Alternatively, `Shift + Enter` executes the cell and automatically inserts or selects the cell below it. Check out all of the keyboard shortcuts by clicking on `Help | Keyboard Shortcut` or take a quick tour by clicking on `Help | User Interface Tour`.

For the parts of the package discussed here, I would recommend using NumPy version 1.8 or later. By convention, you'll find that most people in the scientific Python world will import NumPy using `np` as an alias:

```
In [3]: import numpy as np  
In [4]: np.__version__  
Out[4]: '1.15.4'
```

Throughout this chapter and the rest of this book, we will stick to the same convention.

Understanding NumPy arrays

You might already know that Python is a **weakly typed language**. This means that you do not have to specify a data type whenever you create a new variable. For example, the following will automatically be represented as an integer:

```
In [5]: a = 5
```

You can double-check this by typing the following:

```
In [6]: type(a)Out[6]: int
```

As the standard Python implementation is written in C, every Python object is basically a C structure in disguise. This is true even for integers in Python, which are actually pointers to compound C structures that contain more than just the raw integer value. Therefore, the default C data type used to represent Python integers will depend on your system architecture (that is, whether it is a 32-bit ...)

Accessing single array elements by indexing

If you have used Python's standard list indexing before, then you won't find many issues with indexing in NumPy. In a 1D array, the i^{th} value (counting from zero) can be accessed by specifying the desired index in square brackets, just as with Python lists:

```
In [13]: int_arr
Out[13]: array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
In [14]: int_arr[0]
Out[14]: 0
In [15]: int_arr[3]
Out[15]: 3
```

To index from the end of the array, you can use negative indices:

```
In [16]: int_arr[-1]
Out[16]: 9
In [17]: int_arr[-2]
Out[17]: 8
```

There are a few other cool tricks for **slicing arrays**, as follows:

```
In [18]: int_arr[2:5]    # from index 2 up to index 5 - 1
Out[18]: array([2, 3, 4])
In [19]: int_arr[:5]     # from the beginning up to index 5 - 1
Out[19]: array([0, 1, 2, 3, 4])
In [20]: int_arr[5:]     # from index 5 up to the end of the array
Out[20]: array([5, 6, 7, 8, 9])
In [21]: int_arr[::-2]   # every other element
Out[21]: array([0, 2, 4, 6, 8])
In [22]: int_arr[::-1]   # the entire array in reverse order
Out[22]: array([9, 8, 7, 6, 5, 4, 3, 2, 1, 0])
```

I encourage you to play around with these arrays yourself!

The general form of slicing arrays in NumPy is the same as it is for standard Python lists. To access a slice of an array, `x`, use `x[start:stop:step]`. If any of these are unspecified, they default to the `start=0`, `stop=size of dimension`, `step=1` values.

Creating multidimensional arrays

Arrays need not be limited to lists. In fact, they can have an arbitrary number of dimensions. In machine learning, we will often deal with at least 2D arrays, where the column index stands for the values of a particular feature and the rows contain the actual feature values.

With NumPy, it is easy to create multidimensional arrays from scratch. Let's say that we want to create an array with three rows and five columns, with all of the elements initialized to zero. If we don't specify a data type, NumPy will default to using floats:

```
In [23]: arr_2d = np.zeros((3, 5))...      arr_2dOut[23]: array([[0.,  
0., 0., 0., 0.],  
[0., 0., 0., 0., 0.],  
[0., 0., 0., 0., 0.]])
```

As you probably know from your OpenCV days, this ...

Loading external datasets in Python

Thanks to the SciPy community, there are many resources out there for getting our hands on some data.

A particularly useful resource comes in the form of the `sklearn.datasets` package of **scikit-learn**. This package comes preinstalled with some small datasets that do not require us to download any files from external websites. These datasets include the following:

- `load_boston`: The Boston dataset contains housing prices in different suburbs of Boston along with several interesting features such as per capita crime rate by town, proportion of residential land, and number of non-retail business
- `load_iris`: The Iris dataset contains three different types of Iris flowers (Setosa, Versicolor, and Virginica), along with four features describing the width and length of the sepals and petals
- `load_diabetes`: The diabetes dataset lets us classify patients as having diabetes or not, based on features such as patient age, sex, body mass index, average blood pressure, and six blood serum measurements

- `load_digits`: The digits dataset contains 8 x 8 pixel images of the digits 0-9
- `load_linnerud`: The Linnerud dataset contains 3 physiological variables and 3 exercise variables measured on 20 middle-aged men in a fitness club

Also, scikit-learn allows us to download datasets directly from external repositories, such as the following:

- `fetch_olivetti_faces`: The Olivetti faces dataset contains 10 different images each of 40 distinct subjects
- `fetch_20newsgroups`: The 20 newsgroup dataset contains around 18,000 newsgroup posts on 20 topics

Even better, it is possible to download datasets directly from the machine learning database at <http://openml.org>. For example, to download the Iris flower dataset, simply type the following:

```
In [1]: from sklearn import datasets
In [2]: iris = datasets.fetch_openml('iris', version=1)
In [3]: iris_data = iris['data']
In [4]: iris_target = iris['target']
```

The Iris flower database contains a total of 150 samples with 4 features—sepal length, sepal width, petal length, and petal width. The data is divided into three classes—Iris Setosa, Iris Versicolour, and Iris Virginica. Data and labels are delivered in two separate containers, which we can inspect as follows:

```
In [5]: iris_data.shape
Out[5]: (150, 4)
```

```
In [6]: iris_target.shape  
Out[6]: (150, )
```

Here, we can see that `iris_data` contains 150 samples, each with 4 features (and that's why the number 4 is in the shape). Labels are stored in `iris_target`, where there is only one label per sample.

We can further inspect the values of all targets, but we don't just want to print them all. Instead, we are interested to see all distinct target values, which is easy to do with NumPy:

```
In [7]: import numpy as np  
In [8]: np.unique(iris_target) # Find all unique elements in array  
Out[8]: array(['Iris-setosa', 'Iris-versicolor', 'Iris-virginica'],  
dtype=object)
```

*Another Python library for data analysis that you should have heard about is **pandas** (<http://pandas.pydata.org>). pandas implements several powerful data operations for both databases and spreadsheets. However great the library, at this point, pandas is a bit too advanced for our purposes.*

Visualizing the data using Matplotlib

Knowing how to load data is of limited use if we don't know how to look at the data. Thankfully, there is **Matplotlib**!

Matplotlib is a multiplatform data visualization library built on NumPy arrays—see, I promised you NumPy would show up again. It was conceived by John Hunter in 2002, originally designed as a patch to IPython to enable interactive MATLAB-style plotting from the command line. In more recent years, newer and shinier tools have popped up to eventually replace Matplotlib (such as `ggplot` and `ggvis` in the R language), but Matplotlib remains essential as a well-tested, cross-platform graphics engine.

Importing Matplotlib

You might be in luck again: if you followed the advice outlined in the previous chapter and installed the Python Anaconda stack, you already have Matplotlib installed and are ready to go. Otherwise, you might want to visit <http://matplotlib.org> for installation instructions.

Just as we used the `np` shorthand for NumPy, we will use some standard shorthand for the Matplotlib imports:

```
In [1]: import matplotlib as mpl  
In [2]: import matplotlib.pyplot as plt
```

The `plt` interface is what we will use most often, as we shall see throughout this book.

Producing a simple plot

Without further ado, let's create our first plot.

Let's say that we want to produce a simple line plot of the sine function, `sin(x)`. We want the function to be evaluated at all points on the x axis where $0 \leq x \leq 10$. We will use NumPy's `linspace` function to create a linear spacing on the x axis, from x values 0 to 10 , and a total of 100 sampling points:

```
In [3]: import numpy as npIn [4]: x = np.linspace(0, 10, 100)
```

We can evaluate the `sin` function at all points, x , using NumPy's `sin` function, and visualize the result by calling the `plot` function of `plt`:

```
In [5]: plt.plot(x, np.sin(x))
```

Did you try it yourself? What happened? Did anything show up?

The thing is, depending on where you are running this script, you might not ...

Visualizing data from an external dataset

As a final test for this chapter, let's visualize some data from an external dataset, such as the `digits` dataset from scikit-learn.

Specifically, we will need three tools for visualization:

- scikit-learn for the actual data
- NumPy for data processing
- Matplotlib

So, let's start by importing all of these:

```
In [1]: import numpy as np  
...      from sklearn import datasets  
...      import matplotlib.pyplot as plt  
...      %matplotlib inline
```

The first step is to actually load the data:

```
In [2]: digits = datasets.load_digits()
```

If we remember correctly, `digits` is supposed to have two different fields: a `data` field containing the actual image data and a `target` field containing the image labels. Rather than trusting our memory, we should simply investigate the `digits` object. We do this by typing out its name, adding a period, and then hitting the *Tab* key: `digits.<TAB>`.

This will reveal that the `digits` object also contains some other fields, such as one called `images`. The two fields, `images` and `data`, seem to simply differ by shape:

```
In [3]: print(digits.data.shape)
...     print(digits.images.shape)
Out[3]: (1797, 64)
(1797, 8, 8)
```

In both cases, the first dimension corresponds to the number of images in the dataset. However, `data` has all of the pixels lined up in one big vector, whereas `images` preserves the 8×8 spatial arrangement of each image.

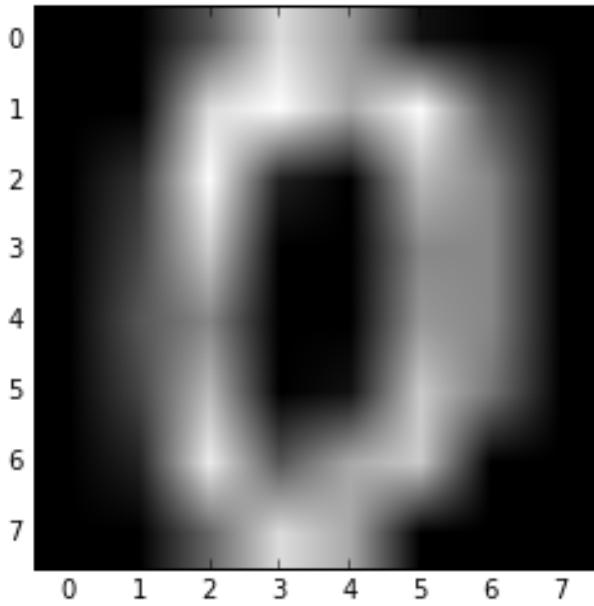
Hence, if we wanted to plot a single image, the `images` field would be more appropriate. First, we grab a single image from the dataset using NumPy's array slicing:

```
In [4]: img = digits.images[0, :, :]
```

Here, we are saying that we want to grab the first row in the 1,797-item-long array and all of the corresponding $8 \times 8 = 64$ pixels. We can then plot the image using the `imshow` function of `plt`:

```
In [5]: plt.imshow(img, cmap='gray')
...     plt.savefig('figures/02.04-digit0.png')
Out[5]: <matplotlib.image.AxesImage at 0x7efcd27f30f0>
```

The preceding command gives the following output. Note that the image is blurred because we have resized it to a larger size. The original image's size is just 8×8 :

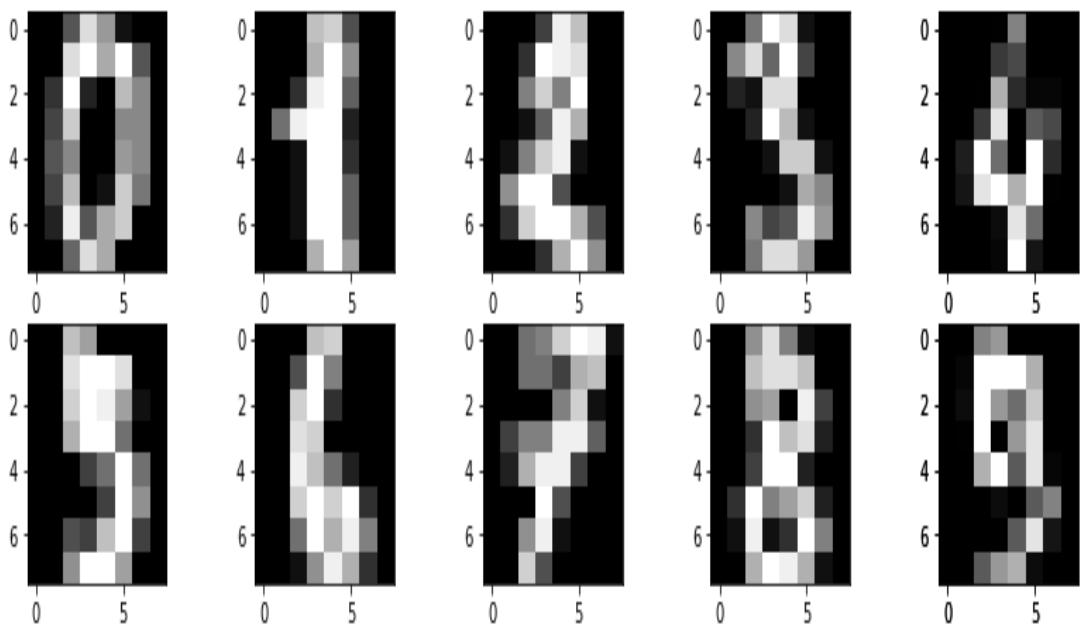


In addition, I also specified a colormap with the `cmap` argument. By default, Matplotlib uses MATLAB's default colormap **jet**. However, in the case of grayscale images, the **gray** colormap makes more sense.

Finally, we can plot a whole number of digit samples using the `subplot` function of `plt`. The `subplot` function is the same as in MATLAB, where we specify the number of rows, number of columns, and current subplot index (starts counting at 1). We will use a `for` loop to iterate over the first 10 images in the dataset and every image gets assigned its own subplot:

```
In [6]: plt.figure(figsize=(14, 4))
...
...     for image_index in range(10):
...         # images are 0-indexed, but subplots are 1-indexed
...         subplot_index = image_index + 1
...         plt.subplot(2, 5, subplot_index)
...         plt.imshow(digits.images[image_index, :, :], cmap='gray')
```

This leads to the following output:



Another great resource for all sorts of datasets is the machine learning repository of my alma mater, the University of California, Irvine: <http://archive.ics.uci.edu/ml/index.php>.

Dealing with data using OpenCV's TrainData container in C++

For the sake of completeness and for those who insist on using the C++ API of OpenCV, let's do a quick detour on OpenCV's `TrainData` container, which allows us to load numerical data from `.csv` files.

Among other things, in C++, the `m1` module contains a class called `TrainData`, which provides a container to work with data in C++. Its functionality is limited to reading (preferably) numerical data from `.csv` files (containing comma-separated values). Hence, if the data that you want to work with comes in a neatly organized `.csv` file, this class will save you a lot of time. If your data comes from a different source, I'm afraid your best option might be to create a `.csv` file by hand, using ...

Summary

In this chapter, we talked about a typical workflow to deal with machine learning problems: how we can extract informative features from raw data, how we can use data and labels to train a machine learning model, and how we can use the finalized model to predict new data labels. We learned that it is essential to split data into a training set and test set, as this is the only way to know how well a model will generalize to new data points.

On the software side of things, we significantly improved our Python skills. We learned how to use NumPy arrays to store and manipulate data and how to use Matplotlib for data visualization. We talked about scikit-learn and its many useful data resources. Finally, we also addressed OpenCV's own `TrainData` container, which provides some relief for users of OpenCV's C++ API.

With these tools in hand, we are now ready to implement our first real machine learning model! In the next chapter, we will focus on supervised learning and its two main problem categories, classification and regression.

First Steps in Supervised Learning

This is the moment you've been waiting for, isn't it?

We have covered all of the bases—we have a functioning Python environment, we have OpenCV installed, and we know how to handle data in Python. Now, it's time to build our first machine learning system! And what better way to start off than to focus on one of the most common and successful types of machine learning: **supervised learning**?

From the previous chapter, we already know that supervised learning is all about learning regularities in training data by using the labels that come with it so that we can predict the labels of some new, never-seen-before test data. In this chapter, we want to dig a little deeper and learn how to turn our theoretical knowledge ...

Technical requirements

You can refer to the code for this chapter at the following link: <https://github.com/PacktPublishing/Machine-Learning-for-OpenCV-Second-Edition/tree/master/Chapter03>.

The following is globally a summary of software and hardware requirements:

- You will need OpenCV version 4.1.x (4.1.0 or 4.1.1 will both work just fine).
- You will need Python version 3.6 (any Python 3.x version will be fine).
- You will need Anaconda Python 3 to install Python and the required modules.
- You can use any OS—macOS, Windows, and Linux-based OSes—with this book. We recommend you have at least 4 GB RAM in your system.
- You don't need to have a GPU to run the code provided with this book.

Understanding supervised learning

We have previously established that the goal of supervised learning is always to predict labels (or target values) for data. However, depending on the nature of these labels, supervised learning can come in two distinct forms:

- **Classification:** Supervised learning is called **classification** whenever we use the data to predict categories. A good example of this is when we try to predict whether an image contains a cat or a dog. Here, the labels of the data are categorical, either one or the other, but never a mixture of categories. For example, a picture contains either a cat or a dog, never 50% cat and 50% dog (before you ask, no, here we do not consider pictures of the cartoon character, CatDog), and our job ...

Having a look at supervised learning in OpenCV

Just knowing how supervised learning works is not going to be of any use if we can't put it into practice. Thankfully, OpenCV provides a pretty straightforward interface for all its statistical learning models, which includes all supervised learning models.

In OpenCV, every machine learning model derives from the `cv::ml::StatModel` base class. This is fancy talk for saying that if we want to use a machine learning model in OpenCV, we have to provide all of the functionality that `StatModel` tells us to. This includes a method to train the model (called `train`) and a method to measure the performance of the model (called `calcError`).

*In Object-Oriented Programming (OOP), we deal primarily with objects or classes. An object consists of several functions, called **methods**, as well as variables, called **members** or **attributes**. You can learn more about OOP in Python at <https://docs.python.org/3/tutorial/classes.html>.*

Thanks to this organization of the software, setting up a machine learning model in OpenCV always follows the same logic, as we will see later:

- **Initialization:** We call the model by name to create an empty instance of the model.
- **Set parameters:** If the model needs some parameters, we can set them via setter methods, which can be different for every model. For example, for a k-NN algorithm to work, we need to specify its open parameter, k (as we will find out later).

- **Train the model:** Every model must provide a method called `train`, used to fit the model to some data.
- **Predict new labels:** Every model must provide a method called `predict`, used to predict the labels of new data.
- **Score the model:** Every model must provide a method called `calcError`, used to measure performance. This calculation might be different for every model.

Because OpenCV is a vast and community-driven project, not every algorithm follows these rules to the extent that we as users might expect. For example, the k-NN algorithm does most of its work in a `findNearest` method, although `predict` still works. We will make sure to point out these discrepancies as we work through different examples.

As we will make occasional use of scikit-learn to implement some machine learning algorithms that OpenCV does not provide, it is worth pointing out that learning algorithms in scikit-learn follow an almost identical logic. The most notable difference is that scikit-learn sets all of the required model parameters in the initialization step. Also, it calls the training function, `fit`, instead of `train` and the scoring function `score` instead of `calcError`.

Measuring model performance with scoring functions

One of the most important parts of building a machine learning system is to find a way to measure the quality of model predictions. In real-life scenarios, a model will rarely get everything right. From earlier chapters, we know that we are supposed to use data from the test set to evaluate our model. But how exactly does that work?

The short, but not very helpful, answer is that it depends on the model. People have come up with all sorts of scoring functions that can be used to evaluate the trained model in all possible scenarios. The good news is that a lot of them are actually part of scikit-learn's `metrics` module.

Let's have a quick look at some of the most important scoring functions. ...

Scoring classifiers using accuracy, precision, and recall

In a binary classification task where there are only two different class labels, there are several different ways to measure classification performance. Some common metrics are as follows:

- `accuracy_score`: Accuracy counts the number of data points in the test set that have been predicted correctly and returns that number as a fraction of the test set size. Sticking to the example of classifying pictures as cats or dogs, accuracy indicates the fraction of pictures that have been correctly classified as containing either a cat or a dog. This is the most basic scoring function for classifiers.
- `precision_score`: Precision describes the ability of a classifier not to label as a cat a picture that contains a dog. In other words, out of all of the pictures in the test set that the classifier thinks contain a cat, precision is the fraction of pictures that actually contain a cat.
- `recall_score`: Recall (or sensitivity) describes the ability of a classifier to retrieve all of the pictures that contain a cat. In other words, out of all of the pictures of cats in the test set,

recall is the fraction of pictures that has been correctly identified as pictures of cats.

Let's say, we have some `ground truth` (correct according to the dataset we have) class labels that are either zeros or ones. We can generate them at random using NumPy's random number generator. Obviously, this means that, whenever we rerun our code, new data points will be generated at random. However, for the purpose of this book, this is not very helpful, as I want you to be able to run the code and always get the same result as me. A nice trick to achieve that is to fix the seed of the random number generator. This will make sure the generator is initialized the same way every time you run the script:

1. We can fix the seed of the random number generator using the following code:

```
In [1]: import numpy as np  
In [2]: np.random.seed(42)
```

2. Then we can generate five random labels that are either zeros or ones by picking random integers in the range, $(0, 2)$:

```
In [3]: y_true = np.random.randint(0, 2, size=5)  
...  
y_true  
Out[3]: array([0, 1, 0, 0, 0])
```

*In the literature, these two classes are sometimes also called **positives** (all data points with the class label, 1) and **negatives** (all other data points).*

Let's assume we have a classifier that tries to predict the class labels mentioned earlier. For the sake of argument, let's say the classifier is

not very smart and always predicts the label, `1`. We can mock this behavior by hardcoding the prediction labels:

```
In [4]: y_pred = np.ones(5, dtype=np.int32)
...
Out[4]: array([1, 1, 1, 1, 1], dtype=int32)
```

What is the accuracy of our prediction?

As mentioned earlier, accuracy counts the number of data points in the test set that have been predicted correctly and returns that number as a fraction of the test set size. We correctly predicted only the second data point (where the true label is `1`). In all other cases, the true label was `0`, yet we predicted `1`. Hence, our accuracy should be $1/5$ or 0.2 .

A naive implementation of an accuracy metric might sum up all occurrences where the predicted class label matched the true class label:

```
In [5]: test_set_size = len(y_true)
In [6]: predict_correct = np.sum(y_true == y_pred)
In [7]: predict_correct / test_set_size
Out[7]: 0.2
```

A smarter, and more convenient, implementation is provided by scikit-learn's `metrics` module:

```
In [8]: from sklearn import metrics
In [9]: metrics.accuracy_score(y_true, y_pred)
Out[9]: 0.2
```

That wasn't too hard, was it? However, to understand precision and recall, we need a general understanding of type I and type II errors. Let's recall that data points with the class label, `1`, are often called positives, and data points with the class label, `0` (or `-1`) are often

called negatives. Then, classifying a specific data point can have one of four possible outcomes, as illustrated by the following confusion matrix:

	Is truly positive	Is truly negative
Predicted positive	True Positive	False Positive
Predicted negative	False Negative	True Negative

Let's break this down. If a data point was truly positive, and we predicted a positive, we got it all right! In this case, the outcome is called a **true positive**. If we thought the data point was a positive, but it was really a negative, we falsely predicted a positive (hence the term, **false positive**). Analogously, if we thought the data point was negative, but it was really a positive, we falsely predicted a negative (false negative). Finally, if we predicted a negative and the data point was truly a negative, we found a true negative.

*In statistical hypothesis testing, false positives are also known as **type I errors** and false negatives are also known as **type II errors**.*

Let's quickly calculate these four metrics on our mock-up data. We have a true positive, where the true label is 1 and we predicted 1:

```
In [10]: truly_a_positive = (y_true == 1)
In [11]: predicted_a_positive = (y_pred == 1)
In [12]: true_positive = np.sum(predicted_a_positive * truly_a_positive
)
...      true_positive
Out[12]: 1
```

Similarly, a false positive is where we predicted 1 but ground truth was really 0:

```
In [13]: false_positive = np.sum((y_pred == 1) * (y_true == 0))
...
false_positive
Out[13]: 4
```

I'm sure by now you've got the hang of it. But do we even have to do the math to know about predicted negatives? Our not-so-smart classifier never predicted 0, so `(y_pred == 0)` should never be true:

```
In [14]: false_negative = np.sum((y_pred == 0) * (y_true == 1))
...
false_negative
Out[14]: 0
In [15]: true_negative = np.sum((y_pred == 0) * (y_true == 0))
...
true_negative
Out[15]: 0
```

Let's also draw the confusion matrix:

	Is truly positive	Is truly negative
Predicted positive	1	4
Predicted negative	0	0

To make sure we did everything right, let's calculate the accuracy one more time. Accuracy should be the number of true positives plus the number of true negatives (that is, everything we got right) divided by the total number of data points:

```
In [16]: accuracy = (true_positive + true_negative) / test_set_size
...
accuracy
```

```
| Out[16]: 0.2
```

Success! Precision is then given as the number of true positives divided by the number of all true predictions:

```
| In [17]: precision = true_positive / (true_positive + false_positive)
| ...      precision
| Out[17]: 0.2
```

It turns out that precision is no better than accuracy in our case. Let's check our math with scikit-learn:

```
| In [18]: metrics.precision_score(y_true, y_pred)
| Out[18]: 0.2
```

Finally, `recall` is given as the fraction of all positives that we correctly classified as positives:

```
| In [19]: recall = true_positive / (true_positive + false_negative)
| ...      recall
| Out[19]: 1.0
| In [20]: metrics.recall_score(y_true, y_pred)
| Out[20]: 1.0
```

Perfect recall! But, going back to our mock-up data, it should be clear that this excellent recall score was mere luck. Since there was only a single `1` label in our mock-up dataset, and we happened to correctly classify it, we got a perfect recall score. Does that mean our classifier is perfect? Not really! But we have found three useful metrics that seem to measure complementary aspects of our classification performance.

Scoring regressors using mean squared error, explained variance, and R squared

When it comes to regression models, our metrics, as shown earlier, don't work anymore. After all, we are now predicting continuous output values, not distinct classification labels. Fortunately, scikit-learn provides some other useful scoring functions:

- `mean_squared_error`: The most commonly used error metric for regression problems is to measure the squared error between the predicted and the true target value for every data point in the training set, averaged across all of the data points.
- `explained_variance_score`: A more sophisticated metric is to measure to what degree a model can explain the variation or dispersion of the test data. Often, the amount of explained ...

Using classification models to predict class labels

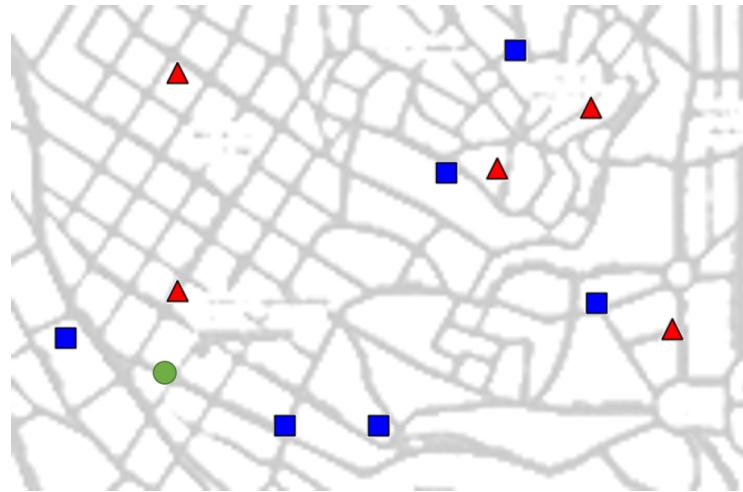
With these tools in hand, we can now take on our first real classification example.

Consider the small town of Randomville, where people are crazy about their two sports teams, the Randomville Reds and the Randomville Blues. The Reds had been around for a long time, and people loved them. But then, some out-of-town millionaire came along and bought the Reds' top scorer and started a new team, the Blues. To the discontent of most Reds fans, the top scorer would go on to win the championship title with the Blues. Years later, he would return to the Reds, despite some backlash from fans who could never forgive him for his earlier career choices. But anyway, you can see why fans of the Reds don't necessarily get along with fans of the Blues. In fact, these two fan bases are so divided that they never even live next to each other. I've even heard stories where the Red fans deliberately moved away once Blues fans moved in next door. True story!

Anyway, we are new in town and are trying to sell some Blues merchandise to people by going from door to door. However, every now and then we come across a bleeding-heart Reds fan who yells at us for selling Blues stuff and chases us off their lawn. Not nice! It would be much less stressful, and a better use of our time, to avoid these houses altogether and just visit the Blues fans instead.

Confident that we can learn to predict where the Reds fans live, we start keeping track of our encounters. If we come by a Reds fan's house, we draw a red triangle on our handy town map; otherwise, we

draw a blue square. After a while, we get a pretty good idea of where everyone lives:



However, now, we approach the house that is marked as a green circle in the preceding map. Should we knock on their door? We try to find some clue as to what team they prefer (perhaps a team flag hanging from the back porch), but we can't see any. How can we know if it is safe to knock on their door?

What this silly example illustrates is exactly the kind of problem a supervised learning algorithm can solve. We have a bunch of observations (houses, their locations, and their colors) that make up our training data. We can use this data to learn from experience so that, when we face the task of predicting the color of a new house, we can make a well-informed estimate.

As we mentioned earlier, fans of the Reds are really passionate about their team, so they would never move next to a Blues fan. Couldn't we use this information and look at all of the neighboring houses, to find out what kind of fan lives in the new house?

This is exactly what the k-NN algorithm would do.

Understanding the k-NN algorithm

The k-NN algorithm is arguably one of the simplest machine learning algorithms. The reason for this is that we basically only need to store the training dataset. Then, to predict a new data point, we only need to find the closest data point in the training dataset: its nearest neighbor.

In a nutshell, the k-NN algorithm argues that a data point probably belongs to the same class as its neighbors. Think about it: if our neighbor is a Reds fan, we're probably Reds fans, too; otherwise, we would have moved away a long time ago. The same can be said for the Blues.

Of course, some neighborhoods might be a little more complicated. In this case, we would not just consider our closest neighbor (where $k=1$), but instead ...

Implementing k-NN in OpenCV

Using OpenCV, we can easily create a k-NN model via the `cv2.ml.KNearest_create()` function. Building the model then involves the following steps:

1. Generate some training data.
2. Create a k-NN object for a given number, k .
3. Find the k nearest neighbors of a new data point that we want to classify.
4. Assign the class label of the new data point by majority vote.
5. Plot the result.

We first import all of the necessary modules: OpenCV for the k-NN algorithm, NumPy for data processing, and Matplotlib for plotting. If you are working in a Jupyter Notebook, don't forget to call the `%matplotlib inline` magic:

```
In [1]: import numpy as np  
...      import cv2  
...      import matplotlib.pyplot as plt  
...      %matplotlib inline  
In [2]: plt.style.use('ggplot')
```

Generating the training data

The first step is to generate some training data. For this, we will use NumPy's random number generator. As discussed in the previous section, we will fix the seed of the random number generator, so that re-running the script will always generate the same values:

```
In [3]: np.random.seed(42)
```

Alright, now let's get to it. What should our training data look like exactly?

In the previous example, each data point is a house on the town map. Every data point has two features (that is, the x and y coordinates of its location on the town map) and a class label (that is, a blue square if a Blues fan lives there and a red triangle if a Reds fan lives there).

The features of a single data point can, therefore, be represented ...

Training the classifier

As with all other machine learning functions, the k-NN classifier is part of the OpenCV 3.1 `ml` module. We can create a new classifier using the following command:

```
In [15]: knn = cv2.ml.KNearest_create()
```

In older versions of OpenCV, this function might be called `cv2.KNearest()` instead.

We then pass our training data to the `train` method:

```
In [16]: knn.train(train_data, cv2.ml.ROW_SAMPLE, labels)
Out[16]: True
```

Here, we have to tell `knn` that our data is an $N \times 2$ array (that is, every row is a data point). Upon success, the function returns `True`.

Predicting the label of a new data point

The other really helpful method that `knn` provides is called `findNearest`. It can be used to predict the label of a new data point based on its nearest neighbors.

Thanks to our `generate_data` function, it is actually really easy to generate a new data point! We can think of a new data point as a dataset of size 1:

```
In [17]: newcomer, _ = generate_data(1)...      newcomerOut[17]:  
array([[91., 59.]], dtype=float32)
```

Our function also returns a random label, but we are not interested in that. Instead, we want to predict it using our trained classifier! We can tell Python to ignore an output value with an underscore (_).

Let's have a look at our town map again. We will plot the training set as we did earlier, but ...

Using regression models to predict continuous outcomes

Now, let's turn our attention to a regression problem. As I'm sure you can recite in your sleep by now, regression is all about predicting continuous outcomes rather than predicting discrete class labels.

Understanding linear regression

The easiest regression model is called **linear regression**. The idea behind linear regression is to describe a target variable (such as Boston house pricing—recall the various datasets we studied in [Chapter 1, A Taste of Machine Learning](#)) with a linear combination of features.

To keep things simple, let's just focus on two features. Let's say we want to predict tomorrow's stock prices using two features: today's stock price and yesterday's stock price. We will denote today's stock price as the first feature, f_1 , and yesterday's stock price as f_2 . Then, the goal of linear regression would be to learn two weight coefficients, w_1 and w_2 , so that we can predict tomorrow's stock price as follows:

Here, is the

Linear regression in OpenCV

Before trying out linear regression on a real-life dataset, let's understand how we can use the `cv2.fitLine` function to fit a line to a 2D or 3D point set:

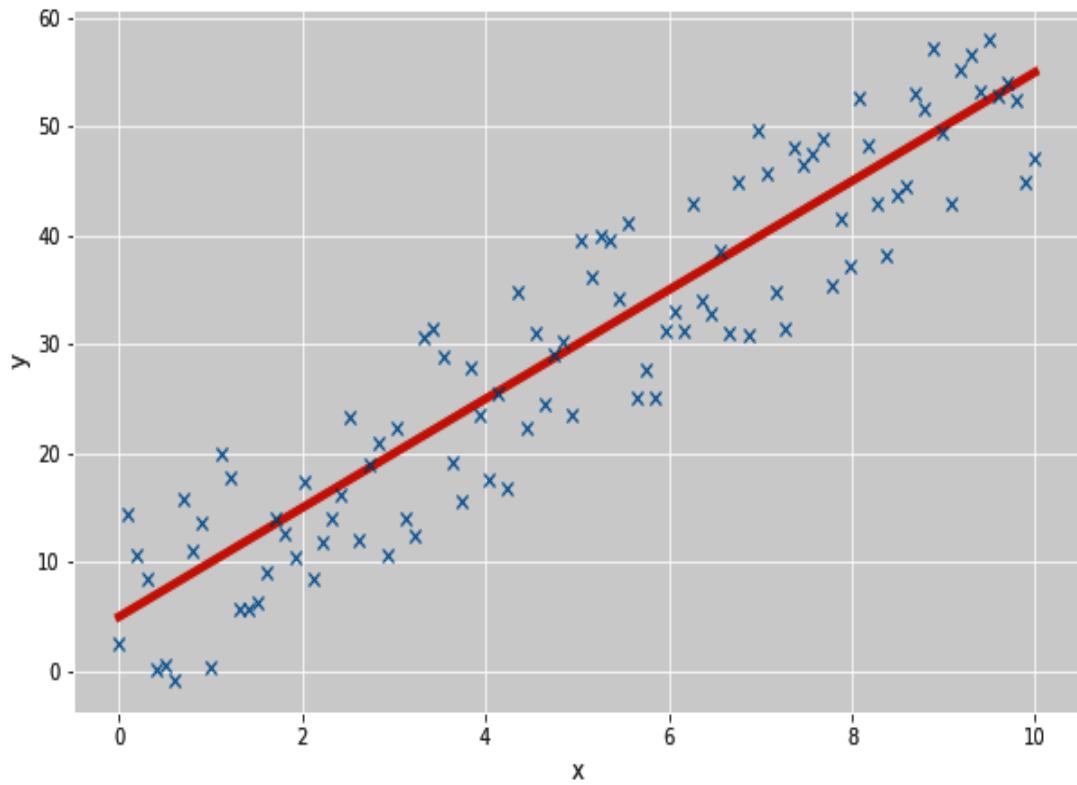
1. Let's start by generating some points. We will generate them by adding noise to the points lying on the line $y = 5x + 5$.

```
In [1]: import cv2
...
import numpy as np
...
import matplotlib.pyplot as plt
...
from sklearn import linear_model
...
from sklearn.model_selection import train_test_split
...
plt.style.use('ggplot')
...
%matplotlib inline
In [2]: x = np.linspace(0,10,100)
...
y_hat = x*5+5
...
np.random.seed(42)
...
y = x*5 + 20*(np.random.rand(x.size) - 0.5)+5
```

2. We can also visualize these points using the following code:

```
In [3]: plt.figure(figsize=(10, 6))
...
plt.plot(x, y_hat, linewidth=4)
...
plt.plot(x,y, 'x')
...
plt.xlabel('x')
...
plt.ylabel('y')
```

This gives us the following diagram, where the red line is the true function:



3. Next, we will split the points into training and testing sets. Here, we will split the data into a 70:30 ratio, meaning, 70% of the points will be used for training and 30% for testing:

```
In [4]: x_train, x_test, y_train, y_test =  
train_test_split(x,y,test_size=0.3,random_state=42)
```

4. Now, let's use `cv2.fitLine` to fit a line to this 2D point set. This function takes in the following arguments:

1. `points`: This is the set of points to which a line has to be fit.

2. `distType`: This is the distance used by the M-estimator.
3. `param`: This is the numerical parameter (C), which is used in some types of distances. We will keep it at 0 so that an optimal value can be chosen.
4. `reps`: This is the accuracy of the distance between the origin and the line. `0.01` is a good default value for `reps`.
5. `aeps`: This is the accuracy of the angle. `0.01` is a good default value for `aeps`.

For more information, have a look at the [documentation](#).

5. Let's see what kinds of result we get using different distance type options:

```
In [5]: distTypeOptions = [cv2.DIST_L2,\n...                      cv2.DIST_L1,\n...                      cv2.DIST_L12,\n...                      cv2.DIST_FAIR,\n...                      cv2.DIST_WELSCH,\n...                      cv2.DIST_HUBER]\n\nIn [6]: distTypeLabels = ['DIST_L2',\n...                      'DIST_L1',\n...                      'DIST_L12',\n...                      'DIST_FAIR',\n...                      'DIST_WELSCH',\n...                      'DIST_HUBER']\n\nIn [7]: colors = ['g', 'c', 'm', 'y', 'k', 'b']\nIn [8]: points = np.array([(xi,yi) for xi,yi in\nzip(x_train,y_train)])
```

6. We will also use scikit-learn's `LinearRegression` to fit the training points and then use the `predict` function to predict the `y`-values for them:

```
In [9]: linreg = linear_model.LinearRegression()
In [10]:
linreg.fit(x_train.reshape(-1,1),y_train.reshape(-1,1))
Out[10]:LinearRegression(copy_X=True, fit_intercept=True,
n_jobs=None,normalize=False)
In [11]: y_sklearn = linreg.predict(x.reshape(-1,1))
In [12]: y_sklearn = list(y_sklearn.reshape(1,-1)[0])
```

7. We use `reshape(-1,1)` and `reshape(1,-1)` to convert the NumPy arrays into a column vector and then back into a row vector:

```
In [13]: plt.figure(figsize=(10, 6))
...     plt.plot(x, y_hat,linewidth=2,label='Ideal')
...     plt.plot(x,y,'x',label='Data')

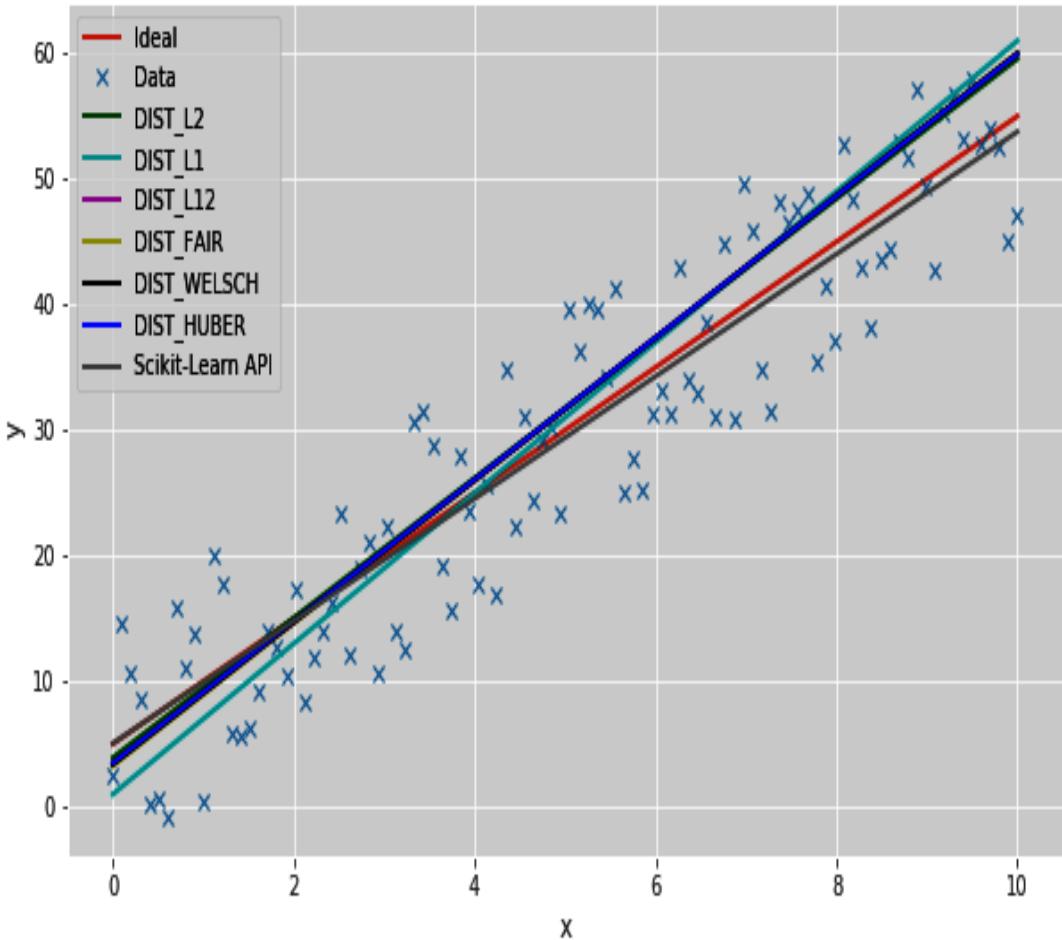
...     for i in range(len(colors)):
...         distType = distTypeOptions[i]
...         distTypeLabel = distTypeLabels[i]
...         c = colors[i]

...         [vxl, vyl, xl, yl] = cv2.fitLine(np.array(points,
dtype=np.int32), distType, 0, 0.01, 0.01)
...         y_cv = [vyl[0]/vxl[0] * (xi - xl[0]) + yl[0] for
xi in x]
...
plt.plot(x,y_cv,c=c,linewidth=2,label=distTypeLabel)

...     plt.plot(x,list(y_sklearn),c='0.5',\
linewidth=2,label='Scikit-Learn API')
...     plt.xlabel('x')
...     plt.ylabel('y')
...     plt.legend(loc='upper left')
```

The only purpose of this preceding (and lengthy) code was to create a plot that could be used to compare the results obtained using different distance measures.

Let's have a look at the plot:



As we can clearly see, scikit-learn's `LinearRegression` model performs much better than OpenCV's `fitLine` function. Now, let's use scikit-learn's API to predict Boston housing prices.

Using linear regression to predict Boston housing prices

To get a better understanding of linear regression, we want to build a simple model that can be applied to one of the most famous machine learning datasets: the **Boston housing prices dataset**. Here, the goal is to predict the value of homes in several Boston neighborhoods in the 1970s, using information such as crime rate, property tax rate, distance to employment centers, and highway accessibility.

Loading the dataset

We can again thank scikit-learn for easy access to the dataset. We first import all of the necessary modules, as we did earlier:

```
In [14]: from sklearn import datasets  
...       from sklearn import metrics
```

Then loading the dataset is a one-liner:

```
In [15]: boston = datasets.load_boston()
```

The structure of the `boston` object is identical to the `iris` object, as discussed in the preceding command. We can get more information about the dataset in '`DESCR`' and find all data in '`data`', all feature names in '`feature_names`', the physical location of the Boston CSV dataset in '`filename`', and all target values in '`target`':

```
In [16]: dir(boston)  
Out[16]: ['DESCR', 'data', 'feature_names', 'filename', 'target']
```

The dataset contains a total of 506 data points, each of which has 13 features:

```
In [17]: boston.data.shape  
Out[17]: (506, 13)
```

Of course, we have only a single target value, which is the housing price:

```
In [18]: boston.target.shape  
Out[18]: (506, )
```

Training the model

Now let's create a `LinearRegression` model that we will then train on the training set:

```
In [19]: linreg = linear_model.LinearRegression()
```

In the preceding command, we want to split the data into training and test sets. We are free to make the split as we see fit, but usually it is a good idea to reserve between 10 percent and 30 percent for testing. Here, we choose 10 percent, using the `test_size` argument:

```
In [20]: X_train, X_test, y_train, y_test = train_test_split(...  
boston.data, boston.target, test_size=0.1,...  
random_state=42...      )
```

In scikit-learn, the `train` function is called `fit` but otherwise behaves exactly the same as in OpenCV:

```
In [21]: linreg.fit(X_train, y_train)Out[21]:  
LinearRegression(copy_X=True, fit_intercept=True, ...
```

Testing the model

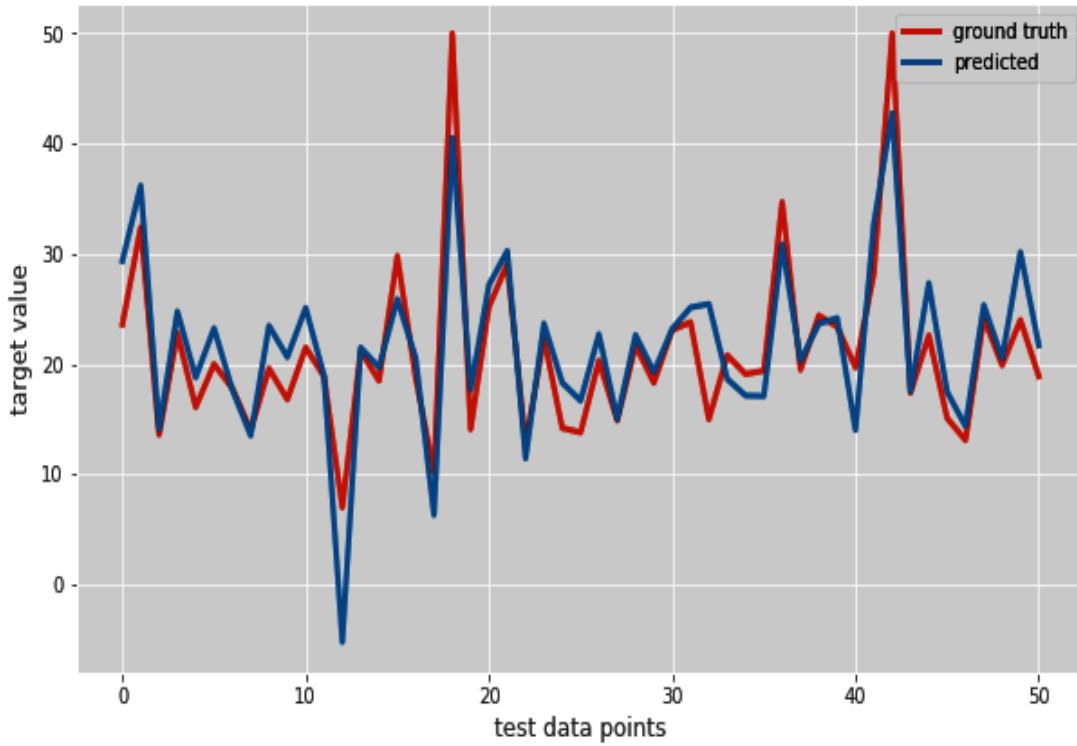
To test the generalization performance of the model, we calculate the mean squared error on the test data:

```
In [24]: y_pred = linreg.predict(X_test)
In [25]: metrics.mean_squared_error(y_test, y_pred)
Out[25]: 14.995852876582541
```

We note that the mean squared error is a little lower on the test set than the training set. This is good news, as we care mostly about the test error. However, from these numbers it is really hard to understand how good the model really is. Perhaps it's better to plot the data:

```
In [26]: plt.figure(figsize=(10, 6))
...     plt.plot(y_test, linewidth=3, label='ground truth')
...     plt.plot(y_pred, linewidth=3, label='predicted')
...     plt.legend(loc='best')
...     plt.xlabel('test data points')
...     plt.ylabel('target value')
Out[26]: <matplotlib.text.Text at 0x7ff46783c7b8>
```

This produces the following diagram:



This makes more sense! Here, we see the `ground truth` housing prices for all test samples in red and our predicted housing prices in blue. It's pretty close, if you ask me. It is interesting to note, though, that the model tends to be off the most for really high or really low housing prices, such as the peak values of data points **12**, **18**, and **42**. We can formalize the amount of variance in the data that we were able to explain by calculating R squared:

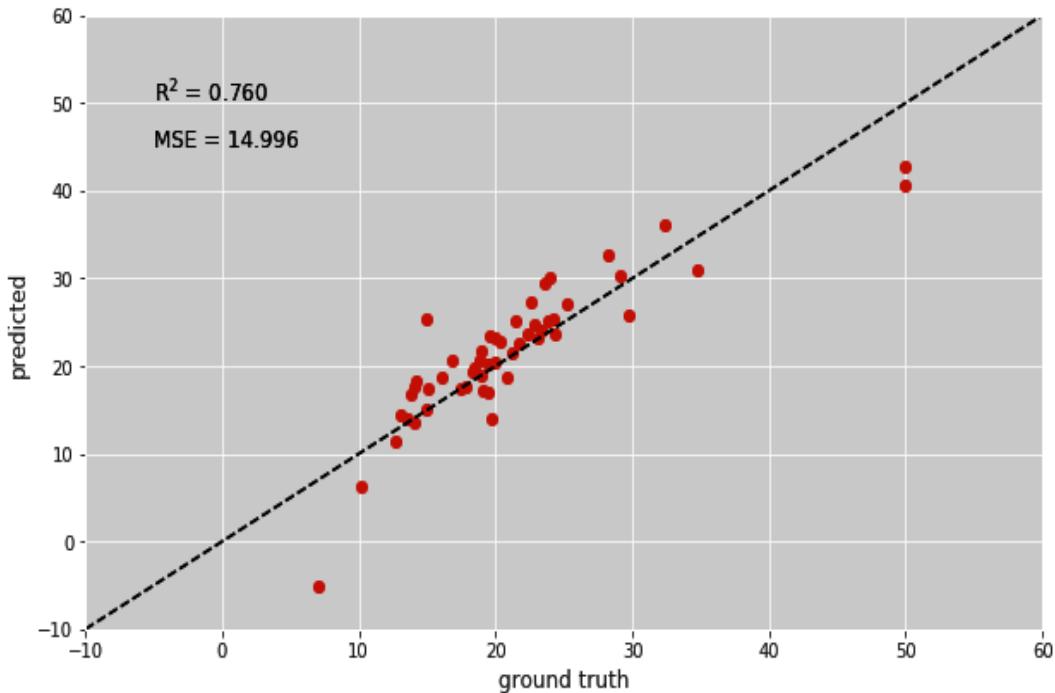
```
In [27]: plt.figure(figsize=(10, 6))
...
    plt.plot(y_test, y_pred, 'o')
...
    plt.plot([-10, 60], [-10, 60], 'k--')
...
    plt.axis([-10, 60, -10, 60])
...
    plt.xlabel('ground truth')
...
    plt.ylabel('predicted')
```

This will plot the `ground truth` prices, `y_test`, on the x axis and our predictions, `y_pred`, on the y axis. We also plot a diagonal line for reference (using a black dashed line, `'k--'`), as we will see soon. But

we also want to display the R^2 score and mean squared error in a textbox:

```
...     scorestr = r'R$^2$ = %.3f' % linreg.score(X_test, y_test)
...     errstr = 'MSE = %.3f' % metrics.mean_squared_error(y_test,
y_pred)
...     plt.text(-5, 50, scorestr, fontsize=12)
...     plt.text(-5, 45, errstr, fontsize=12)
Out[27]: <matplotlib.text.Text at 0x7ff4642d0400>
```

This will produce the following diagram and is a professional way of plotting a model fit:



If our model was perfect, then all data points would lie on the dashed diagonal, since y_{pred} would always be equal to y_{true} . Deviations from the diagonal indicate that the model made some errors, or that there is some variance in the data that the model was not able to explain. Indeed, R^2 indicates that we were able to explain 76% of the scatter in the data, with a mean squared error of 14.996. These are

some performance measures we can use to compare the linear regression model to some more complicated ones.

Applying Lasso and ridge regression

A common problem in machine learning is that an algorithm might work really well on the training set but, when applied to unseen data, it makes a lot of mistakes. You can see how this is problematic since, often, we are most interested in how a model generalizes to new data. Some algorithms (such as decision trees) are more susceptible to this phenomenon than others, but even linear regression can be affected.

*This phenomenon is also known as **overfitting**, and we will talk about it extensively in [chapter 5](#), *Using Decision Trees to Make a Medical Diagnosis*, and [chapter 11](#), *Selecting the Right Model with Hyperparameter Tuning*.*

A common technique for reducing overfitting is called **regularization**, which involves ...

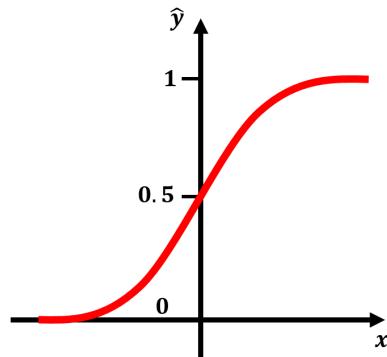
Classifying iris species using logistic regression

Another famous dataset in the world of machine learning is called the Iris dataset. The Iris dataset contains measurements of 150 iris flowers from three different species: Setosa, Versicolor, and Virginica. These measurements include the length and width of the petals and the length and width of the sepals, all measured in centimeters.

Our goal is to build a machine learning model that can learn the measurements of these iris flowers, the species of which are known, so that we can predict the species for a new iris flower.

Understanding logistic regression

Before we start off with this section, let me issue warning—logistic regression, despite its name, is actually a model for classification, specifically when you have two classes. It derives its name from the logistic function (or sigmoid) it uses to convert any real-valued input x into a predicted output value \hat{y} that takes values between **0** and **1**, as shown in the following diagram:



Rounding \hat{y} to the nearest integer effectively classifies the input as belonging either to class **0** or **1**.

Of course, most often, our problems have more than one input or feature value, x . For example, the Iris dataset provides a total ...

Loading the training data

The Iris dataset is included with scikit-learn. We first load all of the necessary modules, as we did in our earlier examples:

```
In [1]: import numpy as np  
...     import cv2  
...     from sklearn import datasets  
...     from sklearn import model_selection  
...     from sklearn import metrics  
...     import matplotlib.pyplot as plt  
...     %matplotlib inline  
In [2]: plt.style.use('ggplot')
```

Then loading the dataset is a one-liner:

```
In [3]: iris = datasets.load_iris()
```

This function returns a dictionary we call `iris`, which contains a bunch of different fields:

```
In [4]: dir(iris)  
Out[4]: ['DESCR', 'data', 'feature_names', 'filename', 'target',  
'target_names']
```

Here, all of the data points are contained in '`data`'. There are 150 data points, each of which has 4 feature values:

```
In [5]: iris.data.shape  
Out[5]: (150, 4)
```

These four features correspond to the sepal and petal dimensions mentioned earlier:

```
In [6]: iris.feature_names  
Out[6]: ['sepal length (cm)',  
         'sepal width (cm)',  
         'petal length (cm)',  
         'petal width (cm)']
```

For every data point, we have a class label stored in `target`:

```
In [7]: iris.target.shape  
Out[7]: (150,)
```

We can also inspect the class labels and find that there is a total of three classes:

```
In [8]: np.unique(iris.target)  
Out[8]: array([0, 1, 2])
```

Making it a binary classification problem

For the sake of simplicity, we want to focus on a binary classification problem for now, where we only have two classes. The easiest way to do this is to discard all data points belonging to a certain class, such as class label 2, by selecting all of the rows that do not belong to class 2:

```
In [9]: idx = iris.target != 2...      data =
iris.data[idx].astype(np.float32)...    target =
iris.target[idx].astype(np.float32)
```

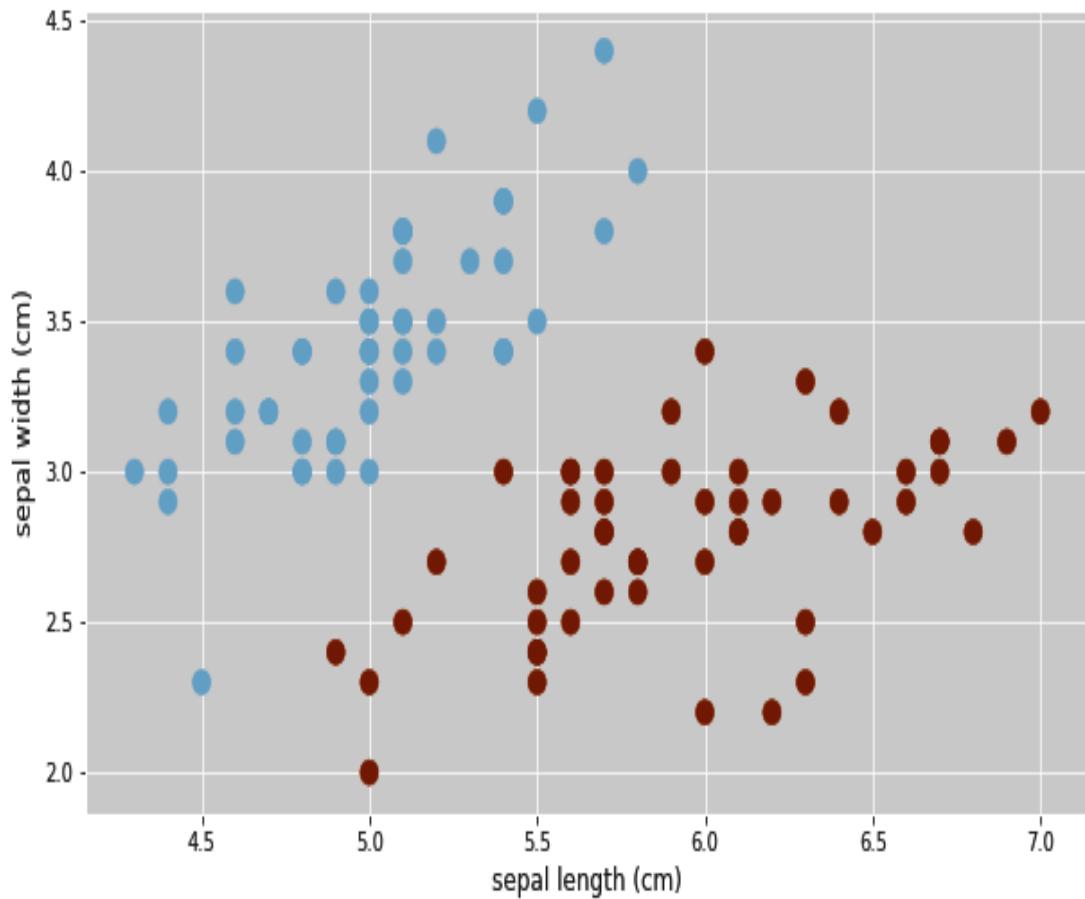
Next, let's inspect the data.

Inspecting the data

Before you get started with setting up a model, it is always a good idea to have a look at the data. We did this earlier for the town map example, so let's repeat it here too. Using Matplotlib, we create a **scatter plot** where the color of each data point corresponds to the class label:

```
In [10]: plt.scatter(data[:, 0], data[:, 1], c=target,
                     cmap=plt.cm.Paired, s=100)
...
...     plt.xlabel(iris.feature_names[0])
...     plt.ylabel(iris.feature_names[1])
Out[10]: <matplotlib.text.Text at 0x23bb5e03eb8>
```

To make plotting easier, we limit ourselves to the first two features (`iris.feature_names[0]` being the sepal length and `iris.feature_names[1]` being the sepal width). We can see a nice separation of classes in the following diagram:



The preceding image shows the plot for the first two features of the Iris dataset.

Splitting data into training and test sets

We learned, in the previous chapter, that it is essential to keep training and test data separate. We can easily split the data using one of scikit-learn's many helper functions:

```
In [11]: X_train, X_test, y_train, y_test =  
model_selection.train_test_split(...           data, target,  
test_size=0.1, random_state=42...       )
```

Here, we want to split the data into 90% training data and 10% test data, which we specify with `test_size=0.1`. By inspecting the return arguments, we note that we ended up with exactly 90 training data points and 10 test data points:

```
In [12]: X_train.shape, y_train.shapeOut[12]: ((90, 4), (90,))In [13]:  
X_test.shape, y_test.shapeOut[13]: ((10, 4), (10,))
```

Training the classifier

Creating a logistic regression classifier involves pretty much the same steps as setting up k-NN:

```
In [14]: lr = cv2.ml.LogisticRegression_create()
```

We then have to specify the desired training method. Here, we can choose `cv2.ml.LogisticRegression_BATCH` OR `cv2.ml.LogisticRegression_MINI_BATCH`. For now, all we need to know is that we want to update the model after every data point, which can be achieved with the following code:

```
In [15]: lr.setTrainMethod(cv2.ml.LogisticRegression_MINI_BATCH)
...           lr.setMiniBatchSize(1)
```

We also want to specify the number of iterations the algorithm should run before it terminates:

```
In [16]: lr.setIterations(100)
```

We can then call the `train` method of the object (in the exact same way as we did earlier), which will return `True` upon success:

```
In [17]: lr.train(X_train, cv2.ml.ROW_SAMPLE, y_train)
Out[17]: True
```

As we just saw, the goal of the training phase is to find a set of weights that best transform the feature values into an output label. A single data point is given by its four feature values (f_0, f_1, f_2 , and f_3).

Since we have four features, we should also get four weights, so that $x = w_0 f_0 + w_1 f_1 + w_2 f_2 + w_3 f_3$, and $\hat{y}=\sigma(x)$. However, as discussed previously, the algorithm adds an extra weight that acts as an offset or bias, so that $x = w_0 f_0 + w_1 f_1 + w_2 f_2 + w_3 f_3 + w_4$. We can retrieve these weights as follows:

```
In [18]: lr.get_learnt_thetas()
Out[18]: array([-0.04090132, -0.01910266, -0.16340332, 0.28743777,
0.11909772], dtype=float32)
```

This means that the input to the logistic function is $x = -0.0409 f_0 - 0.0191 f_1 - 0.163 f_2 + 0.287 f_3 + 0.119$. Then, when we feed in a new data point (f_0, f_1, f_2, f_3) that belongs to class 1, the output $\hat{y}=\sigma(x)$ should be close to 1. But how well does that actually work?

Testing the classifier

Let's see for ourselves by calculating the accuracy score on the training set:

```
In [19]: ret, y_pred = lr.predict(X_train)  
In [20]:  
metrics.accuracy_score(y_train, y_pred)  
Out[20]: 1.0
```

Perfect score! However, this only means that the model was able to perfectly **memorize** the training dataset. This does not mean that the model would be able to classify a new, unseen data point. For this, we need to check the test dataset:

```
In [21]: ret, y_pred = lr.predict(X_test)...  
metrics.accuracy_score(y_test, y_pred)  
Out[21]: 1.0
```

Luckily, we get another perfect score! Now we can be sure that the model we built is truly awesome.

Summary

In this chapter, we covered quite a lot of ground, didn't we?

In short, we learned a lot about different supervised learning algorithms, how to apply them to real datasets, and how to implement everything in OpenCV. We introduced classification algorithms such as k-NN and logistic regression and discussed how they could be used to predict labels as two or more discrete categories. We introduced various variants of linear regression (such as Lasso regression and ridge regression) and discussed how they could be used to predict continuous variables. Last but not least, we got acquainted with the Iris and Boston datasets, two classics in the history of machine learning.

In the following chapters, we will go into much greater depth within these topics and explore some more interesting examples of where these concepts can be useful.

But first, we need to talk about another essential topic in machine learning, feature engineering. Often, data does not come in nicely formatted datasets, and it is our responsibility to represent the data in a meaningful way. Therefore, the next chapter will talk about representing features and engineering data.

Representing Data and Engineering Features

In the last chapter, we built our very first supervised learning models and applied them to some classic datasets, such as the **Iris** and the **Boston** datasets. However, in the real world, data rarely comes in a neat `<n_samples x n_features>` **feature matrix** that is part of a pre-packaged database. Instead, it is our responsibility to find a way to represent the data in a meaningful way. The process of finding the best way to represent our data is known as **feature engineering**, and it is one of the main tasks of data scientists and machine learning practitioners trying to solve real-world problems.

I know you would rather jump right to the end and build the deepest neural network mankind has ever seen. ...

Technical requirements

You can refer to the code for this chapter from the following link: <https://github.com/PacktPublishing/Machine-Learning-for-OpenCV-Second-Edition/tree/master/Chapter04>.

Here is a summary of the software and hardware requirements:

- You will need OpenCV version 4.1.x (4.1.0 or 4.1.1 will both work just fine).
- You will need Python version 3.6 (any Python version 3.x will be fine).
- You will need Anaconda Python 3 for installing Python and the required modules.
- You can use any operating system—macOS, Windows, and Linux-based OSes, along with this book. We recommend you have at least 4 GB RAM in your system.
- You don't need to have a GPU to run the code provided along with this book.

Understanding feature engineering

Believe it or not, how well a machine learning system can learn is mainly determined by the quality of the training data. Although every learning algorithm has its strengths and weaknesses, differences in performance often come down to the way the data is prepared or represented. Feature engineering can hence be understood as a tool for **data representation**. Machine learning algorithms try to learn a solution to a problem from sample data, and feature engineering asks: what is the best representation of the sample data to use for learning a solution to the problem?

Remember, a couple of chapters ago, we talked about the whole machine learning pipeline. There, we already mentioned feature extraction, but we ...

Preprocessing data

The more disciplined we are in handling our data, the better results we are likely to achieve in the end. The first step in this procedure is known as **data preprocessing**, and it comes in (at least) three different flavors:

- **Data formatting:** The data may not be in a format that is suitable for us to work with; for example, the data might be provided in a proprietary file format, which our favorite machine learning algorithm does not understand.
- **Data cleaning:** The data may contain invalid or missing entries, which need to be cleaned up or removed.
- **Data sampling:** The data may be far too large for our specific purpose, forcing us to sample the data intelligently.

Once the data has been preprocessed, we are ready for the actual feature engineering: to transform the preprocessed data to fit our specific machine learning algorithm. This step usually involves one or more of three possible procedures:

- **Scaling:** Certain machine learning algorithms often require the data to be within a common range, such as to have zero

mean and unit variance. Scaling is the process of bringing all features (which might have different physical units) into a common range of values.

- **Decomposition:** Datasets often have many more features than we could possibly process. Feature decomposition is the process of compressing data into a smaller number of highly informative data components.
- **Aggregation:** Sometimes, it is possible to group multiple features into a single, more meaningful one. For example, a database might contain the date and time for each user who logged into a web-based system. Depending on the task, this data might be better represented by simply counting the number of logins per user.

Let's look at some of these processes in more detail.

Standardizing features

Standardization refers to the process of scaling the data to have zero mean and unit variance. This is a common requirement for a wide range of machine learning algorithms, which might behave badly if individual features do not fulfill this requirement. We could manually standardize our data by subtracting from every data point the mean value (μ) of all of the data and dividing that by the variance (σ) of the data; that is, for every feature x , we would compute $(x - \mu) / \sigma$.

Alternatively, scikit-learn offers a straightforward implementation of this process in its `preprocessing` module.

Let's consider a 3×3 data matrix, x , standing for three data points (rows) with three arbitrarily chosen feature values each (columns): ...

Normalizing features

Similar to standardization, normalization is the process of scaling individual samples to have a unit norm. I'm sure you know that the norm stands for the length of a vector and can be defined in different ways. We discussed two of them in the previous chapter: the L1 norm (or Manhattan distance), and the L2 norm (or Euclidean distance).

In scikit-learn, our data matrix, `x`, can be normalized using the `normalize` function, and the `l1` norm is specified by the `norm` keyword:

```
In [5]: X_normalized_l1 = preprocessing.normalize(X, norm='l1')
...
      X_normalized_l1
Out[5]: array([[ 0.2, -0.4,  0.4],
               [ 1. ,  0. ,  0. ],
               [ 0. ,  0.5, -0.5]])
```

Similarly, the L2 norm can be computed by specifying `norm='l2'`:

```
In [6]: X_normalized_l2 = preprocessing.normalize(X, norm='l2')
...
      X_normalized_l2
Out[6]: array([[ 0.33333333, -0.66666667,  0.66666667],
               [ 1. ,  0. ,  0. ],
               [ 0. ,  0.70710678, -0.70710678]])
```

Scaling features to a range

An alternative to scaling features to zero mean and unit variance is to get features to lie between a given minimum and maximum value. Often, these values are zero and one, so that the maximum absolute value of each feature is scaled to unit size. In scikit-learn, this can be achieved using `MinMaxScaler`:

```
In [7]: min_max_scaler = preprocessing.MinMaxScaler()...      X_min_max
= min_max_scaler.fit_transform(X)...      X_min_maxOut[7]: array([[0.33333333, 0. , 1. ], [ 1. , 0.66666667, 0.33333333], [ 0. , 1. , 0. ]])
```

By default, the data will be scaled to fall within 0 and 1. We can specify different ranges by passing a keyword argument, `feature_range`, to the `MinMaxScaler` constructor:

```
In [8]: min_max_scaler = preprocessing.MinMaxScaler(feature_range ...
```

Binarizing features

Finally, we might find ourselves not caring too much about the exact feature values of the data. Instead, we might just want to know whether a feature is present or absent. Binarizing the data can be achieved by thresholding the feature values. Let's quickly remind ourselves of our feature matrix, `x`:

```
In [9]: X
Out[9]: array([[ 1., -2.,  2.],
               [ 3.,  0.,  0.],
               [ 0.,  1., -1.]])
```

Let's assume that these numbers represent the thousands of dollars in our bank accounts. If there are more than 0.5 thousand dollars in the account, we consider the person rich, which we represent with a 1. Otherwise, we put a 0. This is akin to thresholding the data with `threshold=0.5`:

```
In [10]: binarizer = preprocessing.Binarizer(threshold=0.5)
...
...     x_binarized = binarizer.transform(X)
...
Out[10]: array([[ 1.,  0.,  1.],
               [ 1.,  0.,  0.],
               [ 0.,  1.,  0.]])
```

The result is a matrix made up entirely of ones and zeros.

Handling the missing data

Another common requirement in feature engineering is the handling of missing data. For example, we might have a dataset that looks like this:

```
In [11]: from numpy import nan...      x = np.array([[ nan,  0,   3
],...
[ 2,    9,   -8 ],...
[ 1,
nan,  1 ],...
[ 5,    2,   4 ],...
[ 7,    6,   -3 ]])
```

Most machine learning algorithms cannot handle the **Not a Number (NAN)** values (`nan` in Python). Instead, we first have to replace all of the `nan` values with some appropriate fill values. This is known as the **imputation** of missing values.

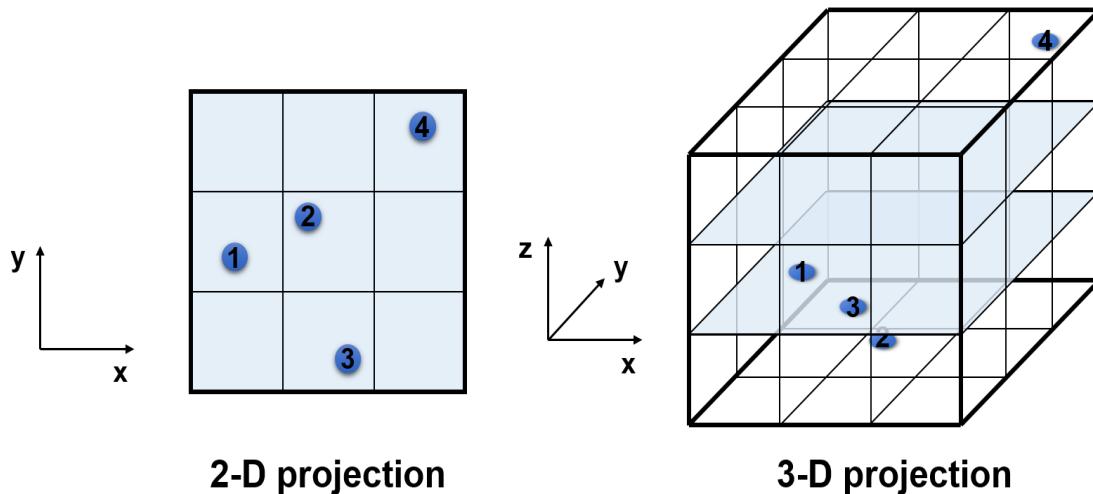
Three different strategies to impute missing values are offered by scikit-learn:

- `mean`: Replaces all of the `nan` values with a mean value along a specified axis of the matrix (default: `axis = 0`)
- `median`: Replaces all ...

Understanding dimensionality reduction

Datasets often have many more features than we could possibly process. For example, let's say our job was to predict a country's poverty rate. We would probably start by matching a country's name with its poverty rate, but that would not help us to predict the poverty rate of a new country. So, we start thinking about the possible causes of poverty. But how many possible causes of poverty are there? Factors might include a country's economy, lack of education, high divorce rate, overpopulation, and so on. If each one of these causes was a feature used to help to predict the poverty rate, we would end up with a countless number of features. If you're a mathematician, you might think of these features as axes in a high-dimensional space, and every country's poverty rate is then a single point in this high-dimensional space.

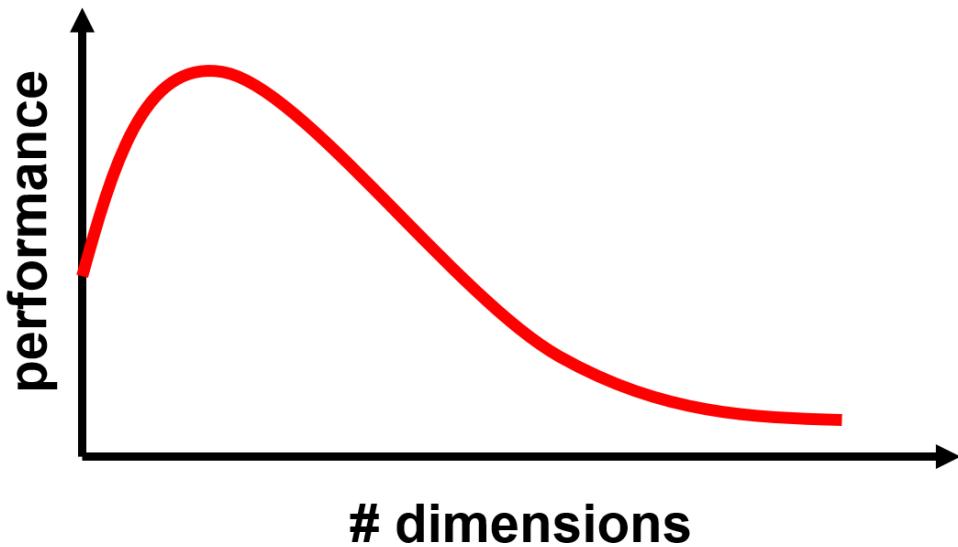
If you're not a mathematician, it might help to start small. Let's say, we first look at only two features: a country's **Gross Domestic Product (GDP)** and the number of citizens. We interpret the GDP as the *x axis*, and the number of citizens as the *y axis*, in a 2D space. Then, we look at the first country. It has a small GDP and an average number of citizens. We draw a point in the *x-y* plane that represents this country. We add a second, third, and fourth country. The fourth country just happens to have both a high GDP and a large number of citizens. Hence, our four data points might be spread across the *x-y* plane much like in the following screenshot:



However, what happens if we start adding a third feature, such as the country's divorce rate, to our analysis? This would add a third axis to our plot (*z axis*). Suddenly, we find that the data no longer spreads very nicely across the *x-y-z* cube, as most of the cube remains empty. While in two dimensions, it seemed like we had most of the *x-y* square covered, in three dimensions, we would need many more data points to fill the void between the data points 1 to 3 and the lonely data point 4 in the upper-right corner.

*This problem is also known as the **curse of dimensionality**: the number of data points needed to fill the available space grows exponentially with the number of dimensions (or plot axes). If a classifier is not fed with data points that span the entire feature space (such as shown in the preceding cube example), the classifier will not know what to do once a new data point is presented that lies far away from all of the previously encountered data points.*

The curse of dimensionality means that, after a certain number of features (or dimensions), the performance of the classifier will start degrading. Let's try to understand this. More features essentially mean that more variations in the dataset can be accounted for. But taking into account more than the required features will cause the classifier to even take into account any outliers or to overfit the dataset. Hence, the performance of the classifier will start degrading rather than improving:



But, how do we find this seemingly optimal number of dimensions for our dataset?

This is where dimensionality reduction comes in to play. These are a family of techniques that allow us to find a compact representation of our high-dimensional data, without losing too much information.

Implementing Principal Component Analysis (PCA) in OpenCV

One of the most common dimensionality reduction techniques is called **PCA**.

Similar to the 2D and 3D examples shown earlier, we can think of an image as a point in a high-dimensional space. If we flatten a 2D grayscale image of height m and width n by stacking all of the columns, we get a (feature) vector of length $m \times n \times 1$. The value of the i^{th} element in this vector is the grayscale value of the i^{th} pixel in the image. Now, imagine we would like to represent every possible 2D grayscale image with these exact dimensions. How many images would that give?

Since grayscale pixels usually take values between 0 and 255, there are a total of 256 raised to the power of $m \times n$ images. Chances ...

Implementing independent component analysis (ICA)

Other useful dimensionality reduction techniques that are closely related to PCA are provided by scikit-learn, but not OpenCV. We mention them here for the sake of completeness. ICA performs the same mathematical steps as PCA, but it chooses the components of the decomposition to be as independent as possible from each other, rather than per predictor as in PCA.

In scikit-learn, ICA is available from the `decomposition` module:

```
In [9]: from sklearn import decomposition  
In [10]: ica = decomposition.FastICA(tol=0.005)
```

Why do we use `tol=0.005`? Because we want the FastICA to converge to some particular value. There are two methods to do that—increase the number of iterations (the default value is 200) or decrease the tolerance (the default value is 0.0001). I tried to increase the iterations but, unfortunately, it didn't work, so I went ahead with the other option. Can you figure out why it didn't converge?

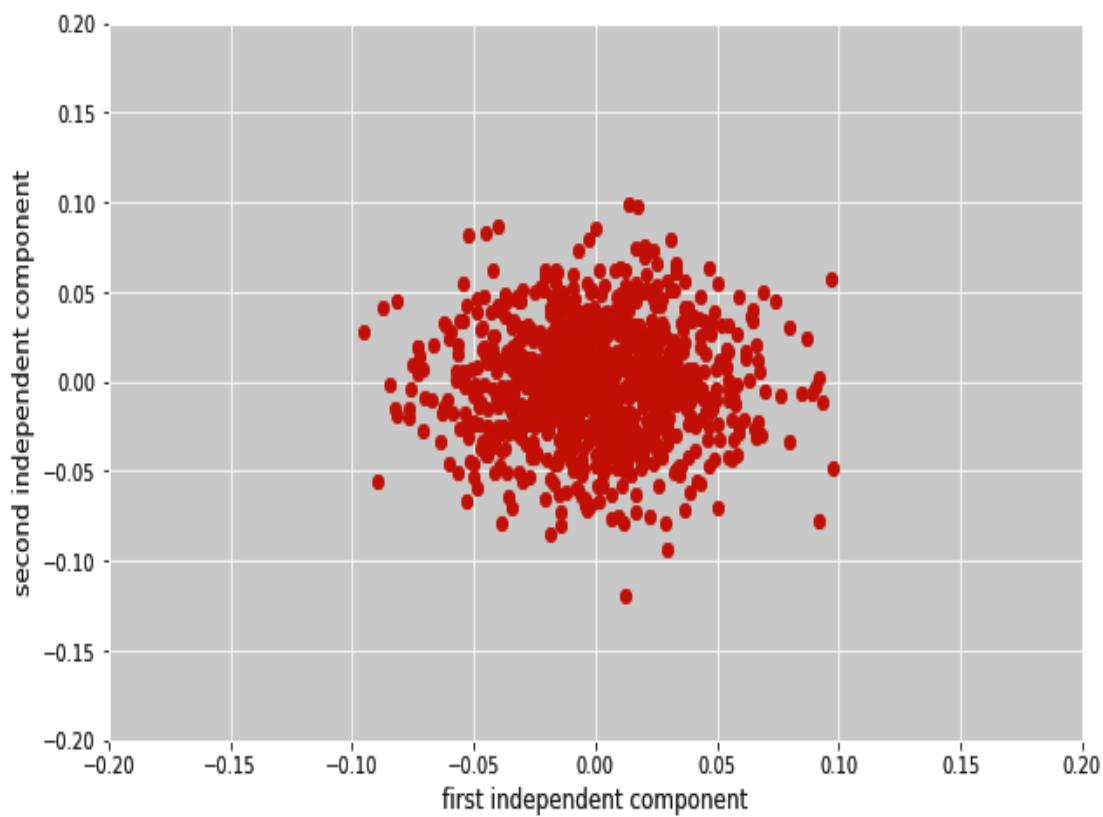
As seen before, the data transformation happens in the `fit_transform` function:

```
In [11]: X2 = ica.fit_transform(X)
```

In our case, plotting the rotated data leads to a similar result to that achieved with PCA earlier, as can be verified in the diagram that follows this code block.

```
In [12]: plt.figure(figsize=(10, 6))
...
...     plt.plot(X2[:, 0], X2[:, 1], 'o')
...     plt.xlabel('first independent component')
...     plt.ylabel('second independent component')
...     plt.axis([-0.2, 0.2, -0.2, 0.2])
Out[12]: [-0.2, 0.2, -0.2, 0.2]
```

This can be seen in the following diagram:



Implementing non-negative matrix factorization (NMF)

Another useful dimensionality reduction technique is called **NMF**. It again implements the same basic mathematical operations as PCA and ICA, but it has the additional constraint that it **only operates on non-negative data**. In other words, we cannot have negative values in our feature matrix if we want to use NMF; the resulting components of the decomposition will all have non-negative values as well.

In scikit-learn, NMF works exactly like ICA:

```
In [13]: nmf = decomposition.NMF()In [14]: X2 = nmf.fit_transform(X)In [15]: plt.plot(X2[:, 0], X2[:, 1], 'o')...      plt.xlabel('first non-negative component')...      plt.ylabel('second non-negative component')...      plt.axis([-5, 20, -5, 10])Out[15]: [-5, 20, ...]
```

Visualizing the dimensionality reduction using t-Distributed Stochastic Neighbor Embedding (t-SNE)

t-SNE is a technique for dimensionality reduction that is best suited to the visualization of high-dimensional data.

In this section, we will see an example of how to visualize high-dimensional datasets using t-SNE. Let's use the digits dataset in this case, which has handwritten images of digits from 0 to 9. It's a publicly available dataset, commonly referred to as the MNIST dataset. We will see how we can visualize the dimensionality reduction on this dataset using t-SNE:

1. First, let's load the dataset:

```
In [1]: import numpy as np
In [2]: from sklearn.datasets import load_digits
In [3]: digits = load_digits()
In [4]: X, y = digits.data/255.0, digits.target
In [5]: print(X.shape, y.shape)
Out[5]: (1797, 64) (1797, )
```

2. You should first apply a dimensional reduction technique such as PCA to reduce the high number of dimensions to a

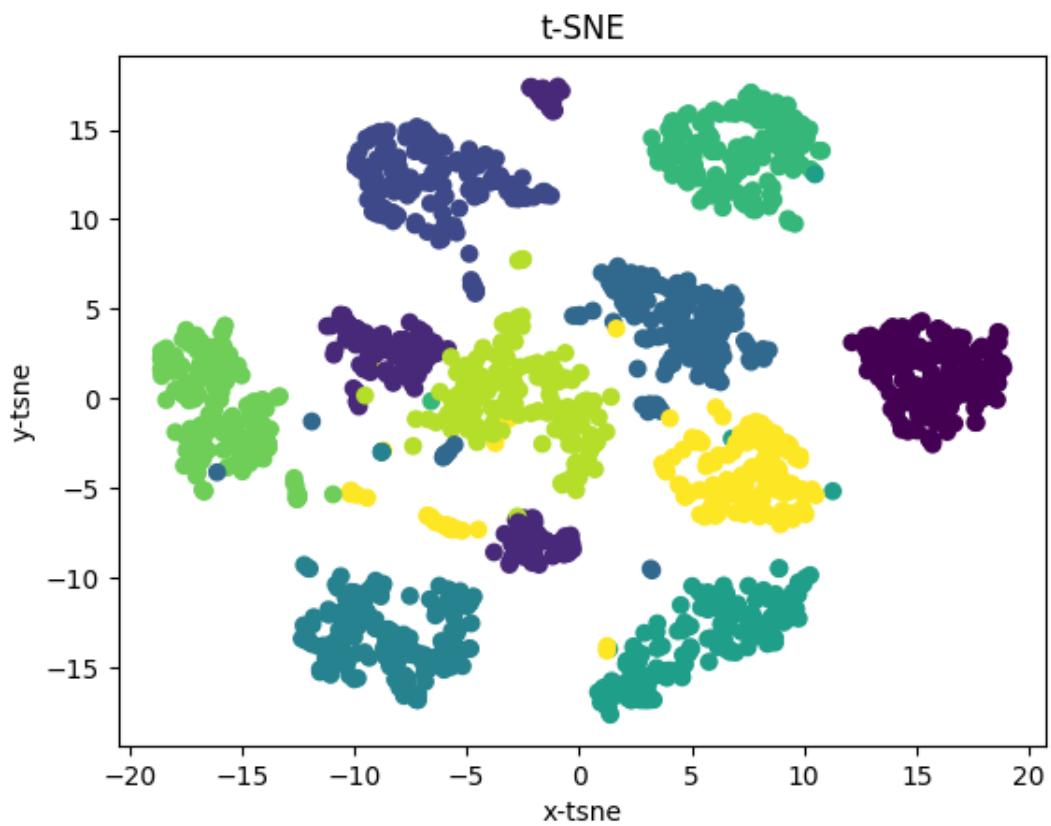
lower number and then use a technique such as t-SNE to visualize the data. But, in this case, let's use all of the dimensions and use t-SNE directly:

```
In [6]: from sklearn.manifold import TSNE
In [7]: tsne = TSNE(n_components=2, verbose=1, perplexity=40,
n_iter=300)
In [8]: tsne_results =
tsne.fit_transform(df.loc[:,features].values)
Out[8]: [t-SNE] Computing 121 nearest neighbors...
... [t-SNE] Indexed 1797 samples in 0.009s...
... [t-SNE] Computed neighbors for 1797 samples in 0.395s...
... [t-SNE] Computed conditional probabilities for sample 1000
/ 1797
... [t-SNE] Computed conditional probabilities for sample 1797
/ 1797
... [t-SNE] Mean sigma: 0.048776
... [t-SNE] KL divergence after 250 iterations with early
exaggeration: 61.094833
... [t-SNE] KL divergence after 300 iterations: 0.926492
```

3. Finally, let's visualize the two dimensions that we have extracted using t-SNE with the help of a scatterplot:

```
In [9]: import matplotlib.pyplot as plt
In [10]:
plt.scatter(tsne_results[:,0],tsne_results[:,1],c=y/10.0)
...     plt.xlabel('x-tsne')
...     plt.ylabel('y-tsne')
...     plt.title('t-SNE')
In [11]: plt.show()
```

And we get the following output:



Now, let's discuss how we can represent categorical variables in the next section.

Representing categorical variables

One of the most common data types we might encounter while building a machine learning system is **categorical features** (also known as **discrete features**), such as the color of a fruit or the name of a company. The challenge with categorical features is that they don't change in a continuous way, which makes it hard to represent them with numbers.

For example, a banana is either green or yellow, but not both. A product belongs either in the clothing department or in the books department, but rarely in both, and so on.

How would you go about representing such features?

For example, let's assume we are trying to encode a dataset consisting of a list of forefathers of machine learning and artificial intelligence:

...

Representing text features

Similar to categorical features, scikit-learn offers an easy way to encode another common feature type—text features. When working with text features, it is often convenient to encode individual words or phrases as numerical values.

Let's consider a dataset that contains a small corpus of text phrases:

```
In [1]: sample = [  
...     'feature engineering',  
...     'feature selection',  
...     'feature extraction'  
... ]
```

One of the simplest methods of encoding such data is by word count; for each phrase, we simply count the occurrences of each word within it. In scikit-learn, this is easily done using `CountVectorizer`, which functions akin to `DictVectorizer`:

```
In [2]: from sklearn.feature_extraction.text import CountVectorizer  
...     vec = CountVectorizer()  
...     X = vec.fit_transform(sample)  
...     X  
Out[2]: <3x4 sparse matrix of type '<class 'numpy.int64'>'  
           with 6 stored elements in Compressed Sparse Row format>
```

By default, this will store our feature matrix, `X`, as a sparse matrix. If we want to manually inspect it, we need to convert it into a regular array:

```
In [3]: X.toarray()  
Out[3]: array([[1, 0, 1, 0],
```

```
[0, 0, 1, 1],  
[0, 1, 1, 0]], dtype=int64)
```

To understand what these numbers mean, we have to look at the feature names:

```
In [4]: vec.get_feature_names()  
Out[4]: ['engineering', 'extraction', 'feature', 'selection']
```

Now, it becomes clear what the integers in `x` mean. If we look at the phrase that is represented in the top row of `x`, we see that it contains one occurrence of the word, `engineering`, and one occurrence of the word, `feature`. On the other hand, it does not contain the words `extraction` or `selection`. Does this make sense? A quick glance at our original data `sample` reveals that the phrase was indeed `feature engineering`.

Looking only at the `x` array (no cheating!), can you guess what the last phrase in `sample` was?

One possible shortcoming of this approach is that we might put too much weight on words that appear very frequently. One approach to fixing this is known as **Term Frequency-Inverse Document Frequency (TF-IDF)**. What TF-IDF does might be easier to understand than its name, which is basically to weigh the word counts by a measure of how often they appear in the entire dataset.

The syntax for TF-IDF is pretty much similar to the previous command:

```
In [5]: from sklearn.feature_extraction.text import TfidfVectorizer  
...     vec = TfidfVectorizer()  
...     X = vec.fit_transform(sample)  
...     X.toarray()  
Out[5]: array([[ 0.861037 , 0. , 0.50854232, 0. ],  
               [ 0. , 0. , 0.50854232, 0.861037 ],  
               [ 0. , 0.861037 , 0.50854232, 0. ]])
```

We note that the numbers are now smaller than before, with the third column taking the biggest hit. This makes sense, as the third column corresponds to the most frequent word across all three phrases, `feature`:

```
In [6]: vec.get_feature_names()  
Out[6]: ['engineering', 'extraction', 'feature', 'selection']
```

If you're interested in the math behind TF-IDF, you can start with this paper: [ht
tp://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.121.1424&rep=rep1&type=pdf](http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.121.1424&rep=rep1&type=pdf). For more information about its specific implementation in scikit-learn, have a look at the API documentation at [http://scikit-learn.org/stable/modules/feature_extraction.html#tfidf-te
rm-weighting](http://scikit-learn.org/stable/modules/feature_extraction.html#tfidf-te
rm-weighting).

Representing text features will become important in [Chapter 7, Implementing a Spam Filter with Bayesian Learning](#).

Representing images

One of the most common and important data types for computer vision is, of course, images. The most straightforward way to represent images is probably by using the grayscale value of each pixel in the image. Usually, grayscale values are not very indicative of the data they describe. For example, if we saw a single pixel with a grayscale value of 128, could we tell what object this pixel belonged to? Probably not. Therefore, grayscale values are not very effective image features.

Using color spaces

Alternatively, we might find that colors contain some information that raw grayscale values cannot capture. Most often, images come in the conventional RGB color space, where every pixel in the image gets an intensity value for its apparent **Redness (R)**, **Greenness (G)**, and **Blueness (B)**. However, OpenCV offers a whole range of other color spaces, such as **Hue Saturation Value (HSV)**, **Hue Saturation Lightness (HSL)**, and the Lab color space. Let's have a quick look at them.

Encoding images in the RGB space

I am sure that you are already familiar with the RGB color space, which uses additive mixing of different shades of red, green, and blue to produce different composite colors. The RGB color space is useful in everyday life because it covers a large part of the color space that the human eye can see. This is why color television sets or color computer monitors only need to care about producing mixtures of red, green, and blue light.

In OpenCV, RGB images are supported straight out of the box. All you need to know, or need to be reminded of, is that color images are actually stored as BGR images in OpenCV; that is, the order of color channels is blue-green-red instead of red-green-blue. The reasons for this ...

Encoding images in the HSV and HLS space

However, ever since the RGB color space was created, people have realized that it is actually quite a poor representation of human vision. Therefore, researchers have developed many alternative representations. One of them is called **HSV** (short for **Hue, Saturation, and Value**) and the other one is called **HLS (Hue, Lightness, and Saturation)**. You might have seen these color spaces in color pickers and common image editing software. In these color spaces, the hue of the color is captured by a single hue channel, the colorfulness is captured by a saturation channel, and the lightness or brightness is captured by a lightness or value channel.

In OpenCV, an RGB image can easily be converted into the HSV color space using `cv2.cvtColor`:

```
In [5]: img_hsv = cv2.cvtColor(img_bgr, cv2.COLOR_BGR2HSV)
```

The same is true for the HLS color space. In fact, OpenCV provides a whole range of additional color spaces, which are available via `cv2.cvtColor`. All we need to do is to replace the color flag with one of the following:

- HLS using `cv2.COLOR_BGR2HLS`
- LAB (lightness, green-red, and blue-yellow) using `cv2.COLOR_BGR2LAB`

- YUV (overall luminance, blue luminance, and red luminance) using `cv2.COLOR_BGR2YUV`

Detecting corners in images

One of the most straightforward traditional features to find in an image are probably corners (locations where several edges meet). OpenCV provides at least two different algorithms to find corners in an image:

- **Harris Corner Detection:** Knowing that edges are areas with high-intensity changes in all directions, Harris and Stephens came up with a fast way of finding such locations. This algorithm is implemented as `cv2.cornerHarris` in OpenCV.
- **Shi-Tomasi Corner Detection:** Shi and Tomasi have their own idea of what constitute good features to track, and they usually do better than Harris corner detection by finding the N strongest corners. This algorithm is implemented as `cv2.goodFeaturesToTrack` in OpenCV.

Harris ...

Using the star detector and BRIEF descriptor

However, corner detection is not sufficient when the scale of an image changes. Multiple papers have been published describing different algorithms for feature detection and description. We will look at a combination of the **Speeded Up Robust Features (SURF)** detector (for more information, see https://en.wikipedia.org/wiki/Speeded_up_robust_features) and the **Binary Robust Independent Elementary Features (BRIEF)** descriptor. The feature detector identifies the keypoints in the image and the feature descriptor calculates the actual feature value for all of the keypoints.

The details of these algorithms are beyond the scope of this book. Advanced users can refer to the papers describing these algorithms in detail.

For more details, you can refer to the following links:

- For SURF: <https://www.vision.ee.ethz.ch/~surf/eccv06.pdf>
- For BRIEF: https://www.cs.ubc.ca/~lowe/525/papers/calonder_eccv10.pdf

The entire process starts with reading the image, converting it into grayscale, using the star feature detector to find the interesting points, and finally, using the BRIEF descriptor to calculate the feature value.

1. Let's first read the image and convert it to grayscale:

```
In [23]: img = cv2.imread('data/rubic-cube.png')
In [24]: gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
```

2. Now, we will create the feature detector and descriptor:

```
In [25]: star = cv2.xfeatures2d.StarDetector_create()
In [26]: brief =
cv2.xfeatures2d.BriefDescriptorExtractor_create()
```

3. Next, it's time to use the star detector to get the keypoints and pass them to the BRIEF descriptor:

```
In [27]: keyPoints = star.detect(gray, None)
In [28]: keyPoints, descriptors = brief.compute(img, keyPoints)
```

There's a catch here. At the time of writing this book, the OpenCV version 4.0 didn't have the resolved version of the `cv2.drawKeypoints` function. So, I have written a similar function that we can use to draw the keypoints. You don't need to worry about the steps involved in the function—it's just for your reference. If you have installed the OpenCV version specified in this book (OpenCV 4.1.0 or OpenCV 4.1.1), you can use the `cv2.drawKeypoints` function directly:

```
In [29]: def drawKeypoint (img, keypoint, color):
...
        draw_shift_bits = 4
...
        draw_multiplier = 1 << draw_shift_bits

...
        center =
(int(round(keypoint.pt[0])),int(round(keypoint.pt[1])))

...
        radius = int(round(keypoint.size/2.0))

...
        # draw the circles around keypoints with the
```

```

keypoints size
...
cv2.circle(img, center, radius, color, 1,
cv2.LINE_AA)

...
# draw orientation of the keypoint, if it is
applicable
...
if keypoint.angle != -1:

...
srcAngleRad = keypoint.angle * np.pi/180.0

...
orient =
(int(round(np.cos(srcAngleRad)*radius)), \
int(round(np.sin(srcAngleRad)*radius)))

...
cv2.line(img, center, (center[0]+orient[0],\
center[1]+orient[1]),\
color, 1, cv2.LINE_AA)
...
else:
...
# draw center with R=1
...
radius = 1 * draw_multiplier
...
cv2.circle(img, center, radius,\ 
color, 1, cv2.LINE_AA)

...
return img
In [30]: from random import randint
...
def drawKeypoints(image, keypoints):
...
for keypoint in keypoints:
...
color =
(randint(0,256), randint(0,256), randint(0,256))
...
image = drawKeypoint(image, keypoint, color)
...
return image

```

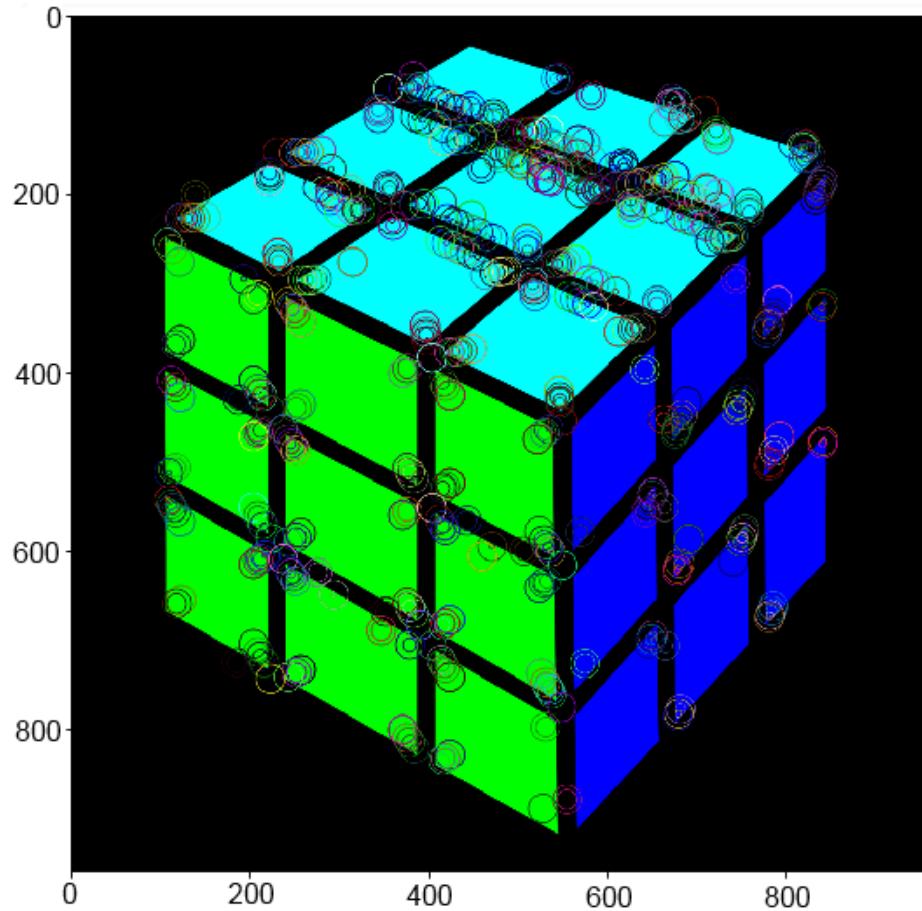
4. Let's now use this function to draw the detected keypoints:

```

In [31]: result = drawKeypoints(img, keyPoints)
In [32]: print("Number of keypoints =
{}".format(len(keyPoints)))
Out[32]: Number of keypoints = 453
In [33]: plt.figure(figsize=(18,9))
...
plt.imshow(result)

```

And we get the following output:



It's pretty awesome, right?

As easy and fast as BRIEF is, it doesn't work well with the rotation of an image. You can try it out by rotating the image (more information at <https://www.pyimagesearch.com/2017/01/02/rotate-images-correctly-with-opencv-and-python/>) and then running BRIEF. Let's see how ORB helps us to resolve this.

Using Oriented FAST and Rotated BRIEF (ORB)

Personally speaking, I am a huge fan of ORB. It's free and a good alternative to SIFT and SURF, which are both protected by patent laws. ORB actually works better than SURF. It's also interesting to note that Gary Bradski was one of the authors of the paper entitled *ORB: An Efficient Alternative to SIFT and SURF*. Can you figure out why that's interesting? Google Gary Bradski and OpenCV and you will get your answer.

The entire process will more or less stay the same, so let's quickly go through the code:

```
In [34]: img = cv2.imread('data/rubic-cube.png')...      gray =
cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)In [35]: orb = cv2.ORB_create()In
[36]: keyPoints = orb.detect(gray,None)In [37]: keyPoints, descriptors
...
...
```

Summary

In this chapter, we went deep down the rabbit hole and looked at several common feature engineering techniques, focusing on both feature selection and feature extraction. We successfully formatted, cleaned, and transformed data so that it could be understood by common machine learning algorithms. We learned about the curse of dimensionality and dabbled a bit in dimensionality reduction by implementing PCA in OpenCV. Finally, we took a short tour of common feature extraction techniques that OpenCV provides for image data.

With these skills under our belt, we are now ready to take on any data, be it numerical, categorical, text, or image data. We know exactly what to do when we encounter missing data, and we know how to transfer our data to make it fit our preferred machine learning algorithm.

In the next chapter, we will take the next step and talk about a specific use case, which is how to use our newly acquired knowledge to make medical diagnoses using decision trees.

Section 2: Operations with OpenCV

This section focuses on advanced machine learning concepts and how they can be implemented using OpenCV and scikit-learn. We will cover some advanced machine learning concepts, such as decision trees, support vector machines, and Bayesian learning, and then finally move on to discuss the second type of machine learning problems—unsupervised learning.

This section includes the following chapters:

- [Chapter 5, Using Decision Trees to Make a Medical Diagnosis](#)
- [Chapter 6, Detecting Pedestrians with Support Vector Machines](#)
- [Chapter 7, Implementing a Spam Filter with Bayesian Learning](#)
- [Chapter 8, Discovering Hidden Structures with Unsupervised Learning](#)

Using Decision Trees to Make a Medical Diagnosis

Now that we know how to handle data in all shapes and forms, be it numerical, categorical, text, or image data, it is time to put our newly gained knowledge to good use.

In this chapter, we will learn how to build a machine learning system that can make a medical diagnosis. We aren't all doctors, but we've probably all been to one at some point in our lives. Typically, a doctor would gain as much information as possible about a patient's history and symptoms to make an informed diagnosis. We will mimic a doctor's decision-making process with the help of what is known as **decision trees**. We will also cover the Gini coefficient, information gain, and variance reduction, along with overfitting and pruning.

A decision tree is a simple yet powerful supervised learning algorithm that resembles a flow chart; we will talk more about this in just a minute. Other than in medicine, decision trees are commonly used in fields such as astronomy (for example, for filtering noise from the Hubble Space Telescope images or to classify star-galaxy clusters), manufacturing and production (for example, by Boeing to discover flaws in the manufacturing process), and object recognition (for example, for recognizing 3D objects).

Specifically, we want to learn about the following in this chapter:

- Building simple decision trees from data and using them for either classification or regression

- Deciding which decision to make next using the Gini coefficient, information gain, and variance reduction
- Pruning a decision tree and its benefits

But first, let's talk about what decision trees actually are.

Technical requirements

You can refer to the code for this chapter from the following link: <https://github.com/PacktPublishing/Machine-Learning-for-OpenCV-Second-Edition/tree/master/Chapter05>.

Here is a summary of the software and hardware requirements:

- You will need OpenCV version 4.1.x (4.1.0 or 4.1.1 will both work just fine).
- You will need Python version 3.6 (any Python version 3.x will be fine).
- You will need Anaconda Python 3 for installing Python and the required modules.
- You can use any OS—macOS, Windows, and Linux-based OSes, along with this book. We recommend you have at least 4 GB RAM in your system.
- You don't need to have a GPU to run the code provided along with this book.

Understanding decision trees

A decision tree is a simple yet powerful model for supervised learning problems. As the name suggests, we can think of it as a tree in which information flows along different branches—starting at the trunk and going all of the way to the individual leaves, making decisions about which branch to take at each junction.

This is basically a decision tree! Here is a simple example of a decision tree:



A decision tree is made of a hierarchy of questions or tests about the data (also known as **decision nodes**) and their possible consequences.

One of the true difficulties with building decision trees is how to pull out suitable features from the data. To make this clear, let's use a concrete example. Let's say we have a dataset consisting of a single email:

```
In [1]: data = [
...     'I am Mohammed Abacha, the son of the late Nigerian Head of '
...     'State who died on the 8th of June 1998. Since i have been '
...     'unsuccessful in locating the relatives for over 2 years now
'
...
...     'I seek your consent to present you as the next of kin so '
...     'that the proceeds of this account valued at US$15.5 Million
'
...
...     'Dollars can be paid to you. If you are capable and willing '
...     'to assist, contact me at once via email with following '
...     'details: 1. Your full name, address, and telephone number. '
...     '2. Your Bank Name, Address. 3.Your Bank Account Number and '
...     'Beneficiary Name - You must be the signatory.'
...
]
```

This email can be vectorized in much the same way as we did in the previous chapter, using scikit-learn's CountVectorizer:

```
In [2]: from sklearn.feature_extraction.text import CountVectorizer
... vec = CountVectorizer()
... X = vec.fit_transform(data)
```

From the previous chapter, we know that we can have a look at the feature names in `x` using the following function:

```
In [3]: function:vec.get_feature_names()[:5]
Out[3]: ['15', '1998', '8th', 'abacha', 'account']
```

For the sake of clarity, we focus on the first five words only, which are sorted alphabetically. Then, the corresponding number of

occurrences can be found as follows:

```
In [4]: x.toarray()[0, :5]
Out[4]: array([1, 1, 1, 1, 2], dtype=int64)
```

This tells us that four out of five words show up just once in the email, but the word, `account` (the last one listed in `out[3]`), actually shows up twice. In the last chapter, we typed `x.toarray()` to convert the sparse array, `x`, into a human-readable array. The result is a 2D array, where rows correspond to the data samples, and columns correspond to the feature names described in the preceding command. Since there is only one sample in the dataset, we limit ourselves to row 0 of the array (that is, the first data sample) and the first five columns in the array (that is, the first five words).

So, how do we check whether the email is from a Nigerian prince?

One way to do this is to look at whether the email contains both the words, `nigerian` and `prince`:

```
In [5]: 'nigerian' in vec.get_feature_names()
Out[5]: True
In [6]: 'prince' in vec.get_feature_names()
Out[6]: False
```

What do we find to our surprise? The word `prince` does not occur in the email.

Does this mean the message is legit?

No, of course not. Instead of '`prince`', the email went with the words, `head of state`, effectively circumventing our all-too-simple spam detector.

Similarly, how do we even start to model the second decision in the tree: *wants me to send him money*? There is no straightforward feature in the text that answers this question. Hence, this is a problem of feature engineering, of combining the words that actually occur in the message in such a way that allows us to answer this question. Sure, a good sign would be to look for strings such as `us$` and `money`, but then we still wouldn't know the context in which these words were mentioned. For all we know, perhaps they were part of the sentence: *Don't worry, I don't want you to send me any money*.

To make matters worse, it turns out that the order in which we ask these questions can actually influence the final outcome. For example, what if we asked the last question first: *do I actually know a Nigerian prince*? Suppose we had a Nigerian prince for an uncle, then finding the words *Nigerian prince* in the email might no longer be suspicious.

As you can see, this seemingly simple example got quickly out of hand.

Luckily, the theoretical framework behind decision trees helps us to find both the right decision rules as well as which decisions to tackle next.

However, to understand these concepts, we have to dig a little deeper.

Building our first decision tree

I think we are ready for a more complex example. As promised earlier, let's now move into the medical domain.

Let's consider an example where several patients have suffered from the same illness, such as a rare form of basorexia. Let's further assume that the true causes of the disease remain unknown to this day and that all of the information that is available to us consists of a bunch of physiological measurements. For example, we might have access to the following information:

- A patient's blood pressure ($_{BP}$)
- A patient's cholesterol level ($_{cholesterol}$)
- A patient's gender ($_{sex}$)
- A patient's age ($_{age}$)
- A patient's blood sodium concentration ($_{Na}$)
- A patient's blood potassium concentration ($_{K}$)

Based on all of ...

Generating new data

Before proceeding with the further steps, let's quickly understand one very crucial step for every machine learning engineer—data generation. We know that all machine learning and deep learning techniques require a huge amount of data—in simple terms: the bigger, the better. But what if you don't have enough data? Well, you can end up with a model that doesn't have enough accuracy. The common technique employed (if you are not able to generate any new data) is to use the majority of the data for training. The major downside of this is that you have a model that is not generalized or, in other terms, suffers from overfitting.

One solution to deal with the preceding issue is to generate new data or, as it is commonly referred to, synthetic data. The key point to note here is that the synthetic data should have similar features to your real data. The more similar they are to the real data, the better it is for you as an ML engineer. This technique is referred to as **data augmentation**, where we use various techniques such as rotation and mirror images to generate new data that is based upon the existing data.

Since we are dealing with a hypothetical case here, we can write simple Python code to generate random data—since there are no set features for us here. In a real-world case, you would use data augmentation to generate realistic-looking new data samples. Let's see how to approach this for our case.

Here, the dataset is actually a list of dictionaries, where every dictionary constitutes a single data point that contains a patient's blood work, age, and gender, as well as the drug that was prescribed.

So, we know that we want to create new dictionaries and we know the keys to use in this dictionary. The next thing to focus on is the data type of the values in the dictionaries.

We start with `age`, which is an integer, then we have the `gender`, which is either `M` or `F`. Similarly, for other values, we can infer the data types and, in some cases, using common sense, we can even infer the range of the values to use.

It is very important to note that common sense and deep learning don't go well together most of the time. This is because you want your model to understand when something is an outlier. For example, we know that it's highly unlikely for someone to have an age of 130 but a generalized model should understand that this value is an outlier and should not be taken into account. This is why you should always have a small portion of data with such illogical values.

Let's see how we can generate some synthetic data for our case:

```
import random

def generateBasorexiaData(num_entries):
    # We will save our new entries in this list
    list_entries = []
    for entry_count in range(num_entries):
        new_entry = {}
        new_entry['age'] = random.randint(20,100)
        new_entry['sex'] = random.choice(['M','F'])
        new_entry['BP'] = random.choice(['low','high','normal'])
        new_entry['cholesterol'] =
random.choice(['low','high','normal'])
        new_entry['Na'] = random.random()
        new_entry['K'] = random.random()
        new_entry['drug'] = random.choice(['A','B','C','D'])
        list_entries.append(new_entry)
    return list_entries
```

We can call the preceding function using `entries = generateBasorexiaData(5)` if we want to generate five new entries.

Now that we know how to generate the data, let's have a look at what we can do with this data. Can we figure out the doctor's reasoning for

prescribing drugs A, B, C, or D? Can we see a relationship between a patient's blood values and the drug that the doctor prescribed?

Chances are, it is as difficult a question to answer for you as it is for me. Although the dataset might look random at first glance, I have, in fact, put in some clear relationships between a patient's blood values and the prescribed drug. Let's see whether a decision tree can uncover these hidden relationships.

Understanding the task by understanding the data

What is always the first step in tackling a new machine learning problem?

You are absolutely right: getting a sense of the data. The better we understand the data, the better we understand the problem we are trying to solve. In our future endeavors, this will also help us to choose an appropriate machine learning algorithm.

The first thing to realize is that the `drug` column is actually not a feature value like all of the other columns. Since it is our goal to predict which drug will be prescribed based on a patient's blood values, the `drug` column effectively becomes the target label. In other words, the inputs to our machine learning algorithm will be the blood values, age, and gender of a ...

Preprocessing the data

For our data to be understood by the decision tree algorithm, we need to convert all categorical features (`sex`, `BP`, and `cholesterol`) into numerical features. What is the best way to do that?

Exactly: we use scikit-learn's `DictVectorizer`. Like we did in the previous chapter, we feed the dataset that we want to convert to the `fit_transform` method:

```
In [10]: from sklearn.feature_extraction import DictVectorizer
...
...     vec = DictVectorizer(sparse=False)
...     data_pre = vec.fit_transform(data)
```

Then, `data_pre` contains the preprocessed data. If we want to look at the first data point (that is, the first row of `data_pre`), we match the feature names with the corresponding feature values:

```
In [12]: vec.get_feature_names()
Out[12]: ['BP=high', 'BP=low', 'BP=normal', 'K', 'Na', 'age',
...
...     'cholesterol=high', 'cholesterol=normal',
...
...     'sex=F', 'sex=M']
In [13]: data_pre[0]
Out[13]: array([ 1. ,  0. ,  0. ,  0.06,  0.66,  33. ,  1. ,  0. ,
...
...     1. ,  0. ])
```

From this, we can see that the three categorical variables—blood pressure (`BP`), cholesterol level (`cholesterol`), and gender (`sex`)—have been encoded using one-hot coding.

To make sure that our data variables are compatible with OpenCV, we need to convert everything into floating point values:

```
In [14]: import numpy as np  
...      data_pre = np.array(data_pre, dtype=np.float32)  
...      target = np.array(target, dtype=np.float32)
```

Then, all that's left to do is to split the data into training and tests sets, like we did in [Chapter 3, First Steps in Supervised Learning](#).

Remember that we always want to keep the training and test sets separate. Since we only have 20 data points to work with in this example, we should probably reserve more than 10 percent of the data for testing. A 15-5 split seems appropriate here. We can be explicit and order the `split` function to yield exactly five test samples:

```
In [15]: import sklearn.model_selection as ms  
...      X_train, X_test, y_train, y_test =  
...      ms.train_test_split(data_pre, target, test_size=5,  
...                          random_state=42)
```

Constructing the tree

Building the decision tree with OpenCV works in much the same way as in [Chapter 3, First Steps in Supervised Learning](#). Recall that all of the machine learning functions reside in OpenCV 3.1's `ml` module:

1. We can create an empty decision tree using the following code:

```
In [16]: import cv2...          dtree = cv2.ml.dtree_create()
```

2. To train the decision tree on the training data, we use the `train` method. That's why we converted the data to float before—so that we could use it in the `train` method:

```
In [17]: dtree.train(X_train, cv2.ml.ROW_SAMPLE, y_train)
```

Here, we have to specify whether the data samples in `x_train` occupy the rows (using `cv2.ml.ROW_SAMPLE`) or the columns (`cv2.ml.COL_SAMPLE`).

3. Then, we can predict the labels of the ...

Visualizing a trained decision tree

OpenCV's implementation of decision trees is good enough if you are just starting out and don't care too much about what's going on under the hood. However, in the following sections, we will switch to scikit-learn. Its implementation allows us to customize the algorithm and makes it a lot easier to investigate the inner workings of the tree. Its usage is also much better documented.

In scikit-learn, decision trees can be used for both classification and regression. They reside in the `tree` module:

1. Let's first import the `tree` module from `sklearn`:

```
In [21]: from sklearn import tree
```

2. Similar to OpenCV, we then create an empty decision tree using the `DecisionTreeClassifier` constructor:

```
In [22]: dtc = tree.DecisionTreeClassifier()
```

3. The tree can then be trained using the `fit` method:

```
In [23]: dtc.fit(X_train, y_train)
Out[23]: DecisionTreeClassifier(class_weight=None,
criterion='gini',
max_depth=None, max_features=None,
max_leaf_nodes=None,
min_impurity_split=1e-07, min_samples_leaf=1,
min_samples_split=2, min_weight_fraction_leaf=0.0,
presort=False, random_state=None, splitter='best')
```

4. We can then compute the accuracy score on both the training and the test sets using the `score` method:

```
In [24]: dtc.score(X_train, y_train)
Out[24]: 1.0
In [25]: dtc.score(X_test, y_test)
Out[25]: 0.4000000000000002
```

Now, here's the cool thing: if you want to know what the tree looks like, you can do so using GraphViz to create a PDF file (or any other supported file type) from the tree structure. For this to work, you need to install GraphViz first. Worry not as it is already present in the environment we created at the beginning of this book.

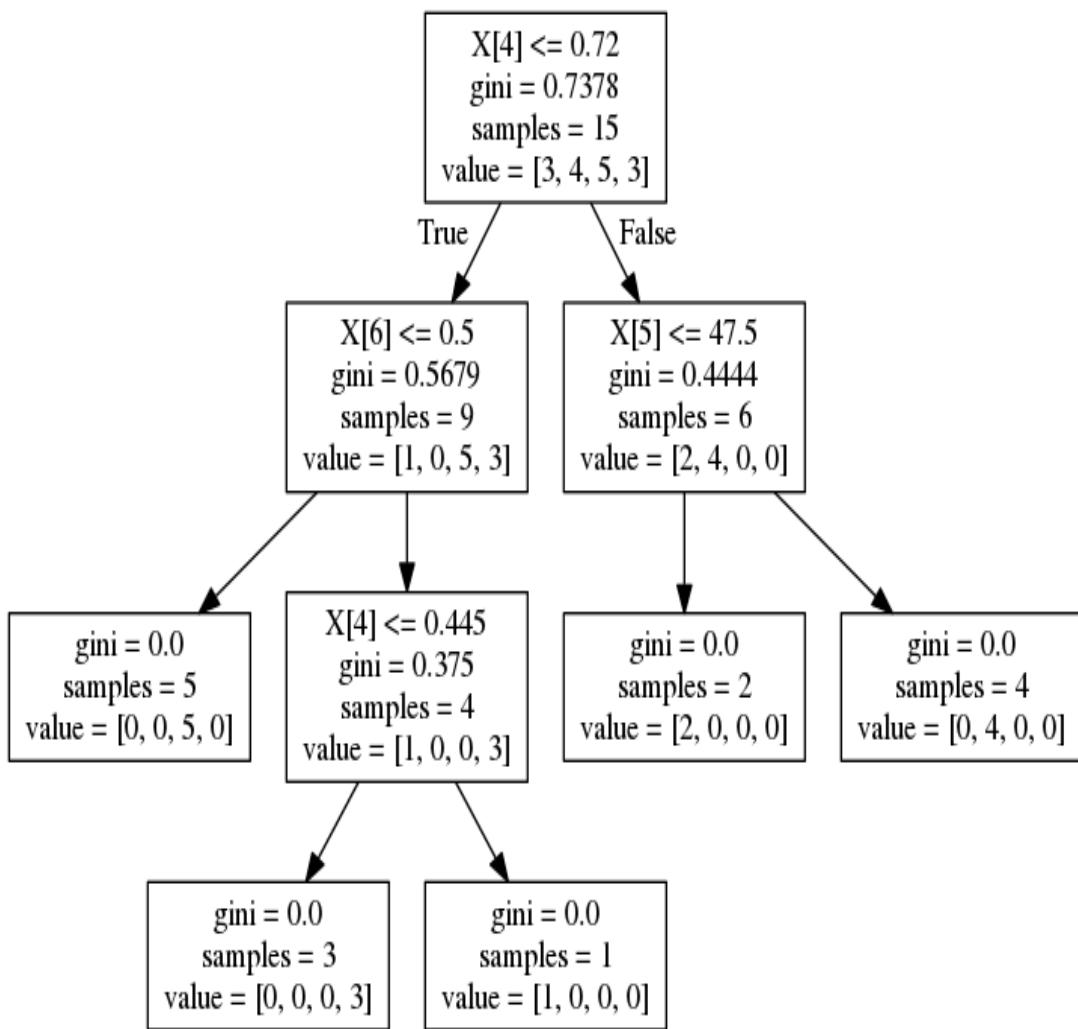
5. Then, back in Python, you can export the tree in GraphViz format to a file, `tree.dot`, using scikit-learn's `export_graphviz` exporter:

```
In [26]: with open("tree.dot", 'w') as f:
... tree.export_graphviz(clf, out_file=f)
```

6. Then, back on the command line, you can use GraphViz to turn `tree.dot` into (for example) a PNG file:

```
$ dot -Tpng tree.dot -o tree.png
```

Alternatively, you can also specify `-Tpdf` or any other supported image format. The result for the preceding tree looks like this:



What does this all mean? Let's break the diagram down step by step.

Investigating the inner workings of a decision tree

We established earlier that a decision tree is basically a flow chart that makes a series of decisions about the data. The process starts at the root node (which is the node at the very top), where we split the data into two groups (only for binary trees), based on some decision rule. Then, the process is repeated until all remaining samples have the same target label, at which point we have reached a leaf node.

In the spam filter example earlier, decisions were made by asking true/false questions. For example, we asked whether an email contained a certain word. If it did, we followed the edge labeled true and asked the next question. However, this works not just for categorical features, ...

Rating the importance of features

What I haven't told you yet is how you pick the features along which to split the data. The preceding root node split the data according to $Na \leq 0.72$, but who told the tree to focus on sodium first? Also, where does the number 0.72 come from anyway?

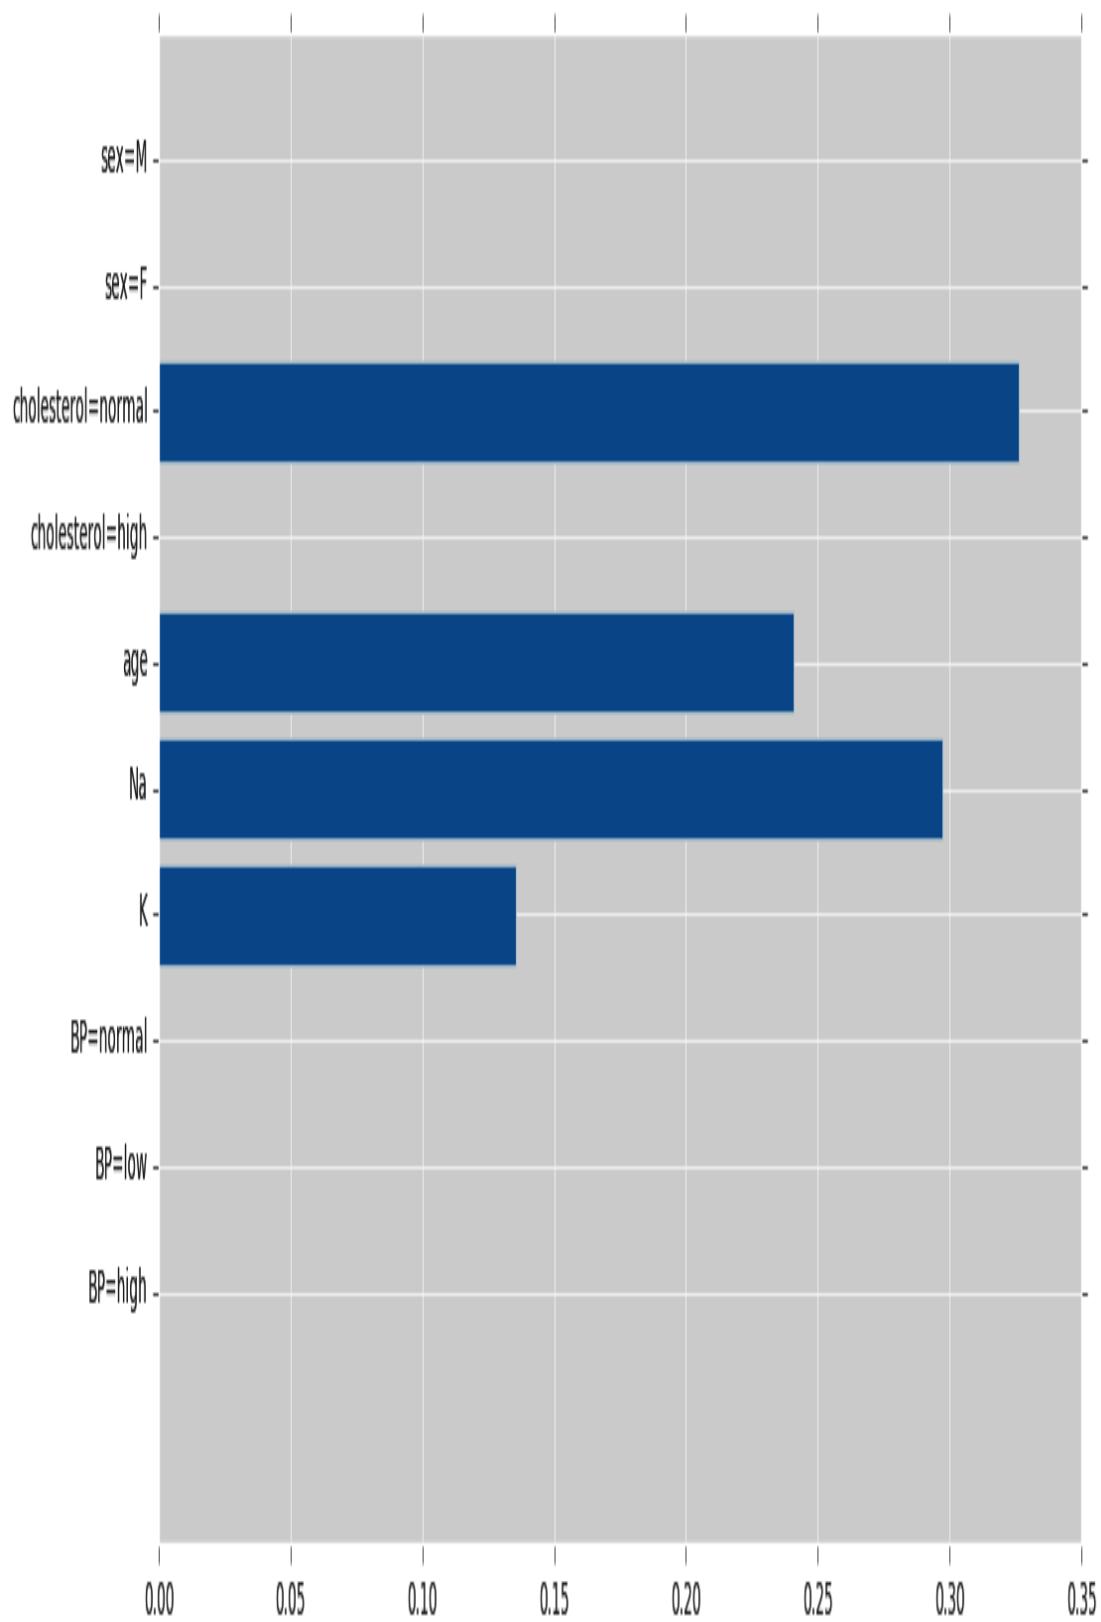
Apparently, some features might be more important than others. In fact, scikit-learn provides a function to rate feature importance, which is a number between 0 and 1 for each feature, where 0 means *not used at all in any decisions made* and 1 means *perfectly predicts the target*. The feature importances are normalized so that they all sum to 1:

```
In [27]: dtc.feature_importances_
Out[27]: array([ 0.        , 0.        , 0.        , 0.13554217, 0.29718876,
   0.24096386, 0.        , 0.32630522, 0.        , 0.        ])
```

If we remind ourselves of the feature names, it will become clear which feature seems to be the most important. A plot might be most informative:

```
In [28]: plt.barh(range(10), dtc.feature_importances_, align='center',
...           tick_label=vec.get_feature_names())
```

This will result in the following bar plot:



Now, it becomes evident that the most important feature for knowing which drug to administer to patients was actually whether the patient had a normal cholesterol level. Age, sodium levels, and potassium levels were also important. On the other hand, gender and blood pressure did not seem to make any difference at all. However, this does not mean that gender or blood pressure are uninformative. It only means that these features were not picked by the decision tree, likely because another feature would have led to the same splits.

But, hold on. If cholesterol level is so important, why was it not picked as the first feature in the tree (that is, in the root node)? Why would you choose to split on the sodium level first? This is where I need to tell you about that ominous `gini` label in the diagram earlier.

Feature importances tell us which features are important for classification, but not which class label they are indicative of. For example, we only know that the cholesterol level is important, but we don't know how that led to different drugs being prescribed. In fact, there might not be a simple relationship between features and classes.

Understanding the decision rules

To build the perfect tree, you would want to split the tree at the most informative feature, resulting in the purest daughter nodes. However, this simple idea leads to some practical challenges:

- It's not actually clear what most informative means. We need a concrete value, a score function, or a mathematical equation that can describe how informative a feature is.
- To find the best split, we have to search over all of the possibilities at every decision node.

Fortunately, the decision tree algorithm actually does these two steps for you. Two of the most commonly used criteria that scikit-learn supports are the following:

- `criterion='gini'`: The Gini impurity is a measure of misclassification, with the aim ...

Controlling the complexity of decision trees

If you continue to grow a tree until all leaves are pure, you will typically arrive at a tree that is too complex to interpret. The presence of pure leaves means that the tree is 100 percent correct on the training data, as was the case with our tree shown earlier. As a result, the tree is likely to perform very poorly on the test dataset, as was the case with our tree shown earlier. We say the tree overfits the training data.

There are two common ways to avoid overfitting:

- **Pre-pruning:** This is the process of stopping the creation of the tree early.
- **Post-pruning (or just pruning):** This is the process of first building the tree but then removing or collapsing nodes that contain only a little information.

There are several ways to pre-prune a tree, all of which can be achieved by passing optional arguments to the `DecisionTreeClassifier` constructor:

- Limiting the maximum depth of the tree via the `max_depth` parameter
- Limiting the maximum number of leaf nodes via `max_leaf_nodes`

- Requiring a minimum number of points in a node to keep splitting it via `min_samples_split`

Often pre-pruning is sufficient to control overfitting.

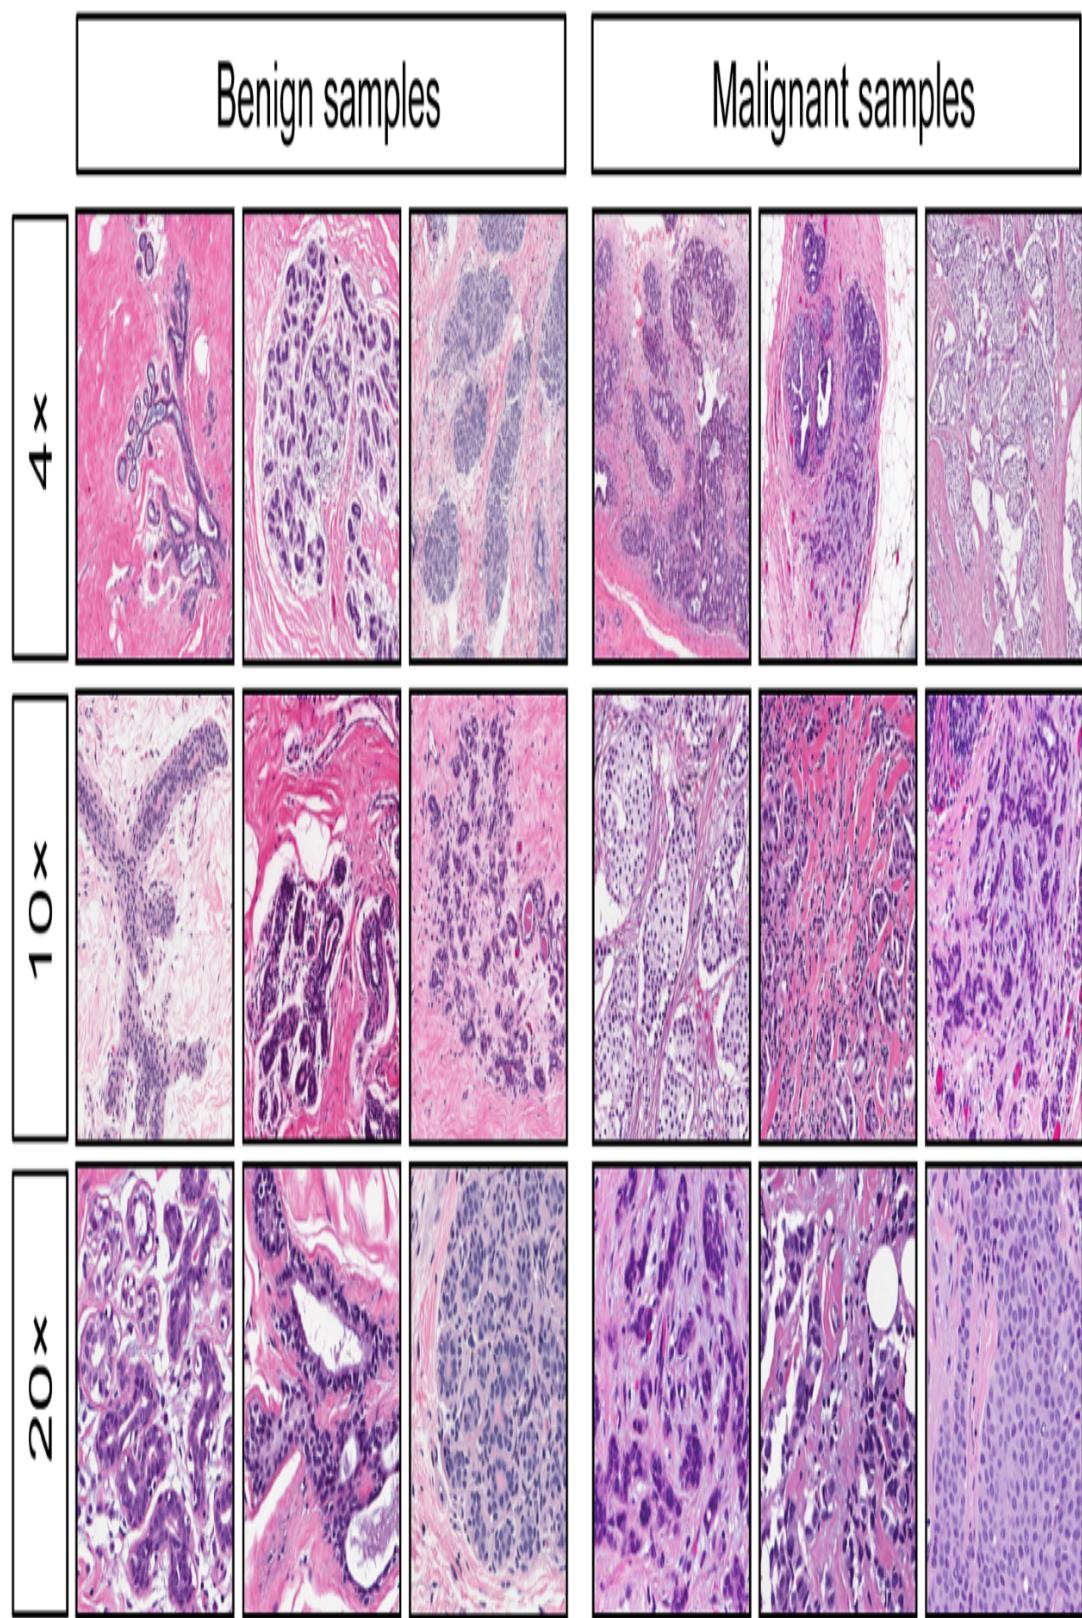
Try it out on our toy dataset! Can you get the score on the test set to improve at all? How does the tree layout change as you start playing with the earlier parameters?

*In more complicated real-world scenarios, pre-pruning is no longer sufficient to control overfitting. In such cases, we want to combine multiple decision trees into what is known as a **random forest**. We will talk about this in [chapter 10](#), *Ensemble Methods for Classification*.*

Using decision trees to diagnose breast cancer

Now that we have built our first decision tree, it's time to turn our attention to a real dataset: the Breast Cancer Wisconsin dataset ([https://archive.ics.uci.edu/ml/datasets/Breast+Cancer+Wisconsin+\(Diagnostic\)](https://archive.ics.uci.edu/ml/datasets/Breast+Cancer+Wisconsin+(Diagnostic))).

This dataset is a direct result of medical imaging research and is considered a classic today. The dataset was created from digitized images of healthy (benign) and cancerous (malignant) tissues. Unfortunately, I wasn't able to find any public-domain examples from the original study, but the images look similar to the following screenshot:



The goal of the research was to classify tissue ...

Loading the dataset

The full dataset is part of scikit-learn's example datasets. We can import it using the following commands:

1. First, let's load the dataset using the `load_breast_cancer` function:

```
In [1]: from sklearn import datasets  
...      data = datasets.load_breast_cancer()
```

2. As in the previous examples, all data is contained in a 2D feature matrix, `data.data`, where the rows represent data samples and the columns are the feature values:

```
In [2]: data.data.shape  
Out[2]: (569, 30)
```

3. With a look at the provided feature names, we recognize some that we mentioned earlier:

```
In [3]: data.feature_names  
Out[3]: array(['mean radius', 'mean texture', 'mean perimeter',  
           'mean area', 'mean smoothness', 'mean  
           compactness',  
           'mean concavity', 'mean concave points',  
           'mean symmetry', 'mean fractal dimension',
```

```
'radius error', 'texture error', 'perimeter  
error',  
        'area error', 'smoothness error',  
        'compactness error', 'concavity error',  
        'concave points error', 'symmetry error',  
        'fractal dimension error', 'worst radius',  
        'worst texture', 'worst perimeter', 'worst  
area',  
        'worst smoothness', 'worst compactness',  
        'worst concavity', 'worst concave points',  
        'worst symmetry', 'worst fractal dimension'],  
        dtype='<U23')
```

4. Since this is a binary classification task, we expect to find exactly two target names:

```
In [4]: data.target_names  
Out[4]: array(['malignant', 'benign'], dtype='<U9')
```

5. Let's reserve some 20 percent of all data samples for testing:

```
In [5]: import sklearn.model_selection as ms  
...      X_train, X_test, y_train, y_test =  
...      ms.train_test_split(data_pre, target, test_size=0.2,  
...      random_state=42)
```

6. You could certainly choose a different ratio, but most commonly people use something like 70-30, 80-20, or 90-10. It all depends a bit on the dataset size but, in the end, should not make too much of a difference. Splitting the data 80-20 should result in the following set sizes:

```
In [6]: X_train.shape, X_test.shape  
Out[6]: ((455, 30), (114, 30))
```

Building the decision tree

As shown earlier, we can create a decision tree using scikit-learn's `tree` module. For now, let's not specify any optional arguments:

1. We will start off by creating a decision tree:

```
In [5]: from sklearn import tree...      dtc =  
tree.DecisionTreeClassifier()
```

2. Do you remember how to train the decision tree? We will use the `fit` function for that:

```
In [6]: dtc.fit(X_train, y_train)Out[6]:  
DecisionTreeClassifier(class_weight=None, criterion='gini',  
max_depth=None, max_features=None,  
max_leaf_nodes=None,  
min_impurity_split=1e-07,  
min_samples_leaf=1,  
min_samples_split=2,  
min_weight_fraction_leaf=0.0,  
presort=False, random_state=None,  
splitter='best')
```

3. Since we did not specify any pre-pruning parameters, we would expect this ...

Using decision trees for regression

Although we have so far focused on using decision trees in classification tasks, you can also use them for regression. But you will need to use scikit-learn again, as OpenCV does not provide this flexibility. We will therefore only briefly review its functionality here:

1. Let's say we wanted to use a decision tree to fit a sin wave.

To make things interesting, we will also add some noise to the data points using NumPy's random number generator:

```
In [1]: import numpy as np  
...      rng = np.random.RandomState(42)
```

2. We then create 100 randomly spaced x values between 0 and 5 and calculate the corresponding sin values:

```
In [2]: X = np.sort(5 * rng.rand(100, 1), axis=0)  
...      y = np.sin(X).ravel()
```

3. We then add noise to every other data point in y (using $y[::2]$), scaled by 0.5 so we don't introduce too much jitter:

```
In [3]: y[::2] += 0.5 * (0.5 - rng.rand(50))
```

4. You can then create a regression tree like any other tree before.

A small difference is that the `gini` and `entropy` split criteria do not apply to regression tasks. Instead, scikit-learn provides two different split criteria:

- **mse** (also known as variance reduction):
This criterion calculates the **Mean Squared Error (MSE)** between ground truth and prediction and splits the node that leads to the smallest MSE.
 - **mae**: This criterion calculates the **Mean Absolute Error (MAE)** between ground truth and prediction and splits the node that leads to the smallest MAE.

5. Using the MSE criterion, we will build two trees. Let's first build a tree with depth 2:

```
presort=False, random_state=42,
splitter='best')
```

6. Next, we will build a decision tree with a maximum depth of 5:

```
In [6]: regr2 = tree.DecisionTreeRegressor(max_depth=5,
...     random_state=42)
...     regr2.fit(X, y)
Out[6]: DecisionTreeRegressor(criterion='mse', max_depth=5,
                               max_features=None,
                               max_leaf_nodes=None,
                               min_impurity_split=1e-07,
                               min_samples_leaf=1,
                               min_samples_split=2,
                               min_weight_fraction_leaf=0.0,
                               presort=False, random_state=42,
                               splitter='best')
```

We can then use the decision tree like a linear regressor from [Chapter 3](#), *First Steps in Supervised Learning*.

7. For this, we create a test set with x values densely sampled in the whole range from 0 through 5:

```
In [7]: X_test = np.arange(0.0, 5.0, 0.01)[:, np.newaxis]
```

8. The predicted y values can then be obtained with the `predict` method:

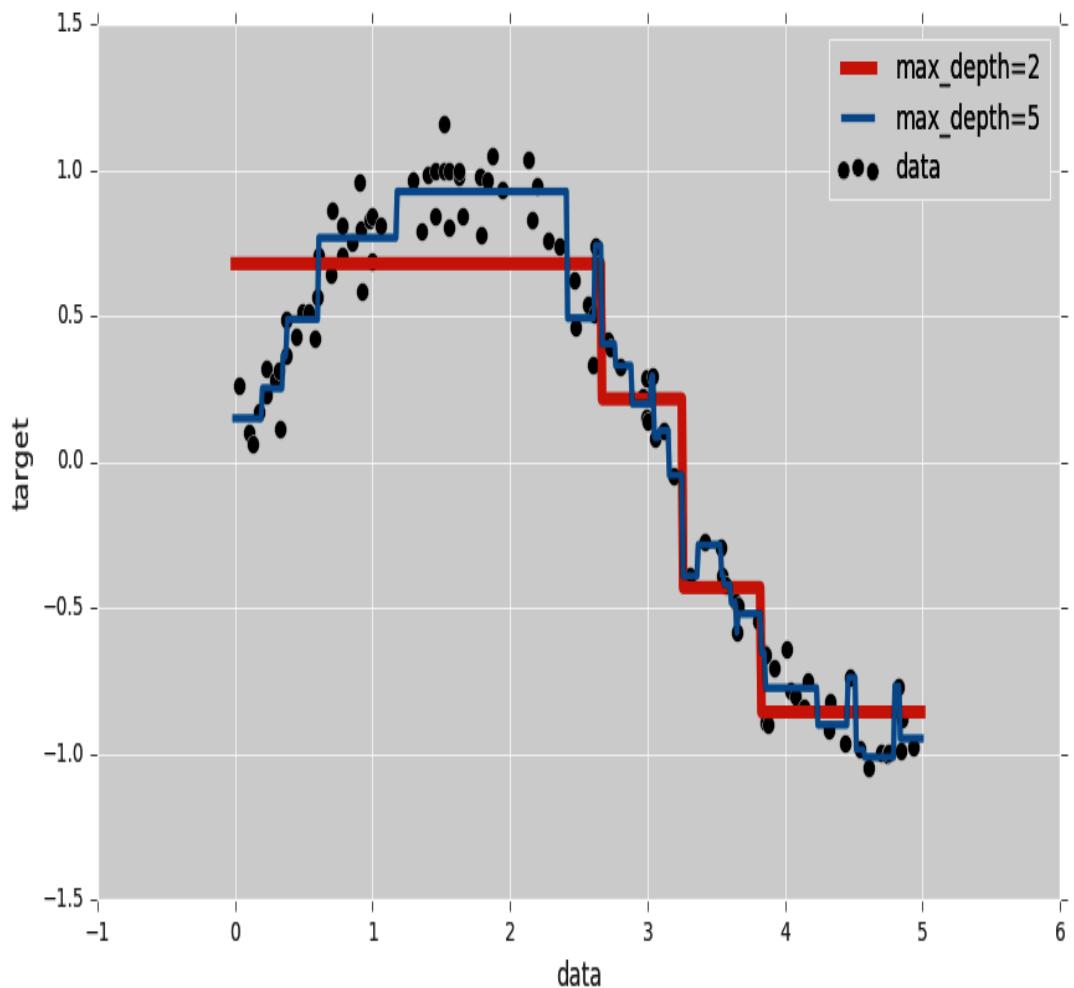
```
In [8]: y_1 = regr1.predict(X_test)
...     y_2 = regr2.predict(X_test)
```

9. If we plot all of these together, we can see how the decision trees differ:

```
In [9]: import matplotlib.pyplot as plt
... %matplotlib inline
... plt.style.use('ggplot')

... plt.scatter(X, y, c='k', s=50, label='data')
... plt.plot(X_test, y_1, label="max_depth=2", linewidth=5)
... plt.plot(X_test, y_2, label="max_depth=5", linewidth=3)
... plt.xlabel("data")
... plt.ylabel("target")
... plt.legend()
Out[9]: <matplotlib.legend.Legend at 0x12d2ee345f8>
```

This will produce the following plot:



Here, the thick red line represents the regression tree with depth 2. You can see how the tree tries to approximate the data using these crude steps. The thinner blue line belongs to the regression tree with depth 5; the added depth has allowed the tree to make many finer-grained approximations. Therefore, this tree can approximate the data even better. However, because of this added power, the tree is also more susceptible to fitting noisy values, as can be seen especially from the spikes on the right-hand side of the plot.

Summary

In this chapter, we learned all about decision trees and how to apply them to both classification and regression tasks. We talked a bit about data generation, overfitting, and ways to avoid this phenomenon by tweaking pre-pruning and post-pruning settings. We also learned about how to rate the quality of a node split using metrics such as the Gini impurity and information gain. Finally, we applied decision trees to medical data to detect cancerous tissues. We will come back to decision trees towards the end of this book when we will combine multiple trees into what is known as a random forest. But for now, let's move on to a new topic.

In the next chapter, we will introduce another staple of the machine learning world: support vector ...

Detecting Pedestrians with Support Vector Machines

In the previous chapter, we talked about how to use decision trees for classification and regression. In this chapter, we want to direct our attention to another well-established supervised learner in the machine learning world: **support vector machines (SVMs)**. Soon after their introduction in early 1990, SVMs quickly became popular in the machine learning community, largely because of their success in early handwritten digit classification. They remain relevant to this day, especially in application domains, such as computer vision.

The goal of this chapter is to apply SVMs to a popular problem in computer vision: pedestrian detection. In contrast to a recognition task (where we name the category of an object), the goal of a detection task is to say whether a particular object (or in our case, a pedestrian) is present in an image or not. You might already know that OpenCV can do this in two to three lines of code. But, we won't learn anything if we do it that way. So instead, we'll build the whole pipeline from scratch! We will obtain a real-world dataset, perform feature extraction using the **histogram of oriented gradients (HOG)**, and apply an SVM to it.

In this chapter, we will implement SVMs in OpenCV with Python. We will learn to deal with nonlinear decision boundaries and understand the kernel trick. By the end of the chapter, we will learn to detect pedestrians in the wild.

Along the way, we will cover the following topics:

- Implementing SVMs in OpenCV with Python
- Dealing with nonlinear decision boundaries
- Understanding the kernel trick
- Detecting pedestrians in the wild

Excited? Then let's go!

Technical requirement

You can refer to the codes for this chapter from the following link: <https://github.com/PacktPublishing/Machine-Learning-for-OpenCV-Second-Edition/tree/master/Chapter06>.

Here is a short summary of the software and hardware requirements:

- OpenCV version 4.1.x (4.1.0 or 4.1.1 will both work just fine).
- Python version 3.6 (any Python version 3.x will be fine).
- Anaconda Python 3 for installing Python and the required modules.
- You can use any operating system—macOS, Windows, and Linux-based OS—with this book. We recommend you to have at least 4 GB RAM in your system.
- You don't need to have a GPU to run the code provided with the book.

Understanding linear SVMs

In order to understand how SVMs work, we have to think about decision boundaries. When we used linear classifiers or decision trees in earlier chapters, our goal was always to minimize the classification error. We did this by assessing the accuracy using mean squared error. An SVM tries to achieve low classification errors too, but it does so only implicitly. An SVM's explicit objective is to maximize the margins between data points of

Learning optimal decision boundaries

Let's look at a simple example. Consider some training samples with only two features (x and y values) and a corresponding target label (positive (+) or negative (-)). Since the labels are categorical, we know that this is a classification task. Moreover, because we only have two distinct classes (+ and -), it's a binary classification task.

In a binary classification task, a decision boundary is a line that partitions the training set into two subsets, one for each class. An **optimal decision boundary** partitions the data so that all data samples from one class (say, +) are to the left of the decision boundary, and all other data samples (say, -) are to the right of it.

An SVM updates its choice of a decision ...

Implementing our first SVM

But enough with the theory. Let's do some coding!

It might be a good idea to pace ourselves. For our very first SVM, we should probably focus on a simple dataset, perhaps a binary classification task.

A cool trick about scikit-learn's `datasets` module that I haven't told you about is that you can generate random datasets of controlled size and complexity. A few notable ones are as follows:

- `datasets.make_classification([n_samples, ...])`: This function generates a random n -class classification problem, where we can specify the number of samples, the number of features, and the number of target labels
- `datasets.make_regression([n_samples, ...])`: This function generates a random regression problem
- `datasets.make_blobs([n_samples, n_features, ...])`: This function generates a number of Gaussian blobs we can use for clustering

This means that we can use `make_classification` to build a custom dataset for a binary classification task.

Generating the dataset

As we can now recite in our sleep, a binary classification problem has exactly two distinct target labels (`n_classes=2`). For the sake of simplicity, let's limit ourselves to only two feature values (`n_features=2`; for example, an x and a y value). Let's say we want to create 100 data samples:

```
In [1]: from sklearn import datasets...      X, y =
datasets.make_classification(n_samples=100, n_features=2, ...
n_redundant=0, n_classes=2, ...
random_state=7816)
```

We expect X to have 100 rows (data samples) and 2 columns (features), whereas the y vector should have a single column that contains all the target labels:

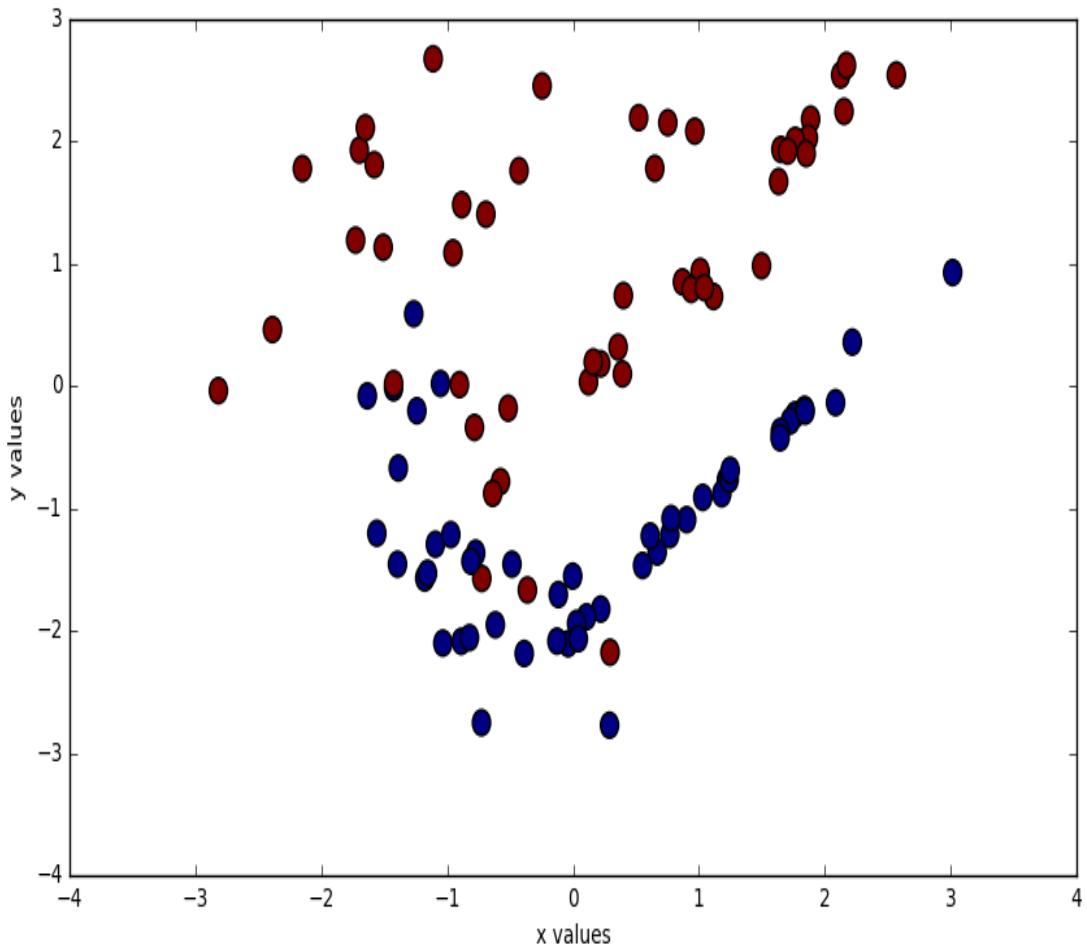
```
In [2]: X.shape, y.shape
Out[2]: ((100, 2), (100,))
```

Visualizing the dataset

We can plot these data points in a scatter plot using Matplotlib. Here, the idea is to plot the x values (found in the first column of x , $x[:, 0]$) against the y values (found in the second column of x , $x[:, 1]$). A neat trick is to pass the target labels as color values ($c=y$):

```
In [3]: import matplotlib.pyplot as plt
...
...     %matplotlib inline
...     plt.scatter(X[:, 0], X[:, 1], c=y, s=100)
...     plt.xlabel('x values')
...     plt.ylabel('y values')
Out[3]: <matplotlib.text.Text at 0x24f7ffb00f0>
```

This will produce the following output:



The preceding output shows the randomly generated data for a binary classification problem. You can see that, for the most part, data points of the two classes are clearly separated. However, there are a few regions (particularly near the left and bottom of the plot) where the data points of both classes intermingle. These will be hard to classify correctly, as we will see in just a second.

Preprocessing the dataset

The next step is to split the data points into training and test sets, as we have done before. But, before we do that, we have to prepare the data for OpenCV as follows:

- All feature values in x must be 32-bit floating point numbers
- Target labels must be either -1 or +1

We can achieve this with the following code:

```
In [4]: import numpy as np...      x = x.astype(np.float32)...      y = y
* 2 - 1
```

Now we can pass the data to scikit-learn's `train_test_split` function as we did in the earlier chapters:

```
In [5]: from sklearn import model_selection as ms...      x_train,
x_test, y_train, y_test = ms.train_test_split(...      x, y,
test_size=0.2, random_state=42...      )
```

Here, I chose to reserve 20 percent of all data points for the test set, but ...

Building the support vector machine

In OpenCV, SVMs are built, trained, and scored the same exact way as every other learning algorithm we have encountered so far, using the following four steps:

1. Call the `create` method to construct a new SVM:

```
In [6]: import cv2  
...      svm = cv2.ml.SVM_create()
```

As shown in the following command, there are different *modes* in which we can operate an SVM. For now, all we care about is the case we discussed in the previous example: an SVM that tries to partition the data with a straight line. This can be specified with the `setKernel` method:

```
In [7]: svm.setKernel(cv2.ml.SVM_LINEAR)
```

2. Call the classifier's `train` method to find the optimal decision boundary:

```
In [8]: svm.train(X_train, cv2.ml.ROW_SAMPLE, y_train)  
Out[8]: True
```

3. Call the classifier's `predict` method to predict the target labels of all data samples in the test set:

```
In [9]: _, y_pred = svm.predict(X_test)
```

4. Use scikit-learn's `metrics` module to score the classifier:

```
In [10]: from sklearn import metrics  
...         metrics.accuracy_score(y_test, y_pred)  
Out[10]: 0.8000000000000004
```

Congratulations, we got 80 percent correctly classified test samples!

Of course, so far we have no idea what happened under the hood. For all we know, we might as well have got these commands off a web search and typed them into the Terminal, without really knowing what we're doing. But this is not who we want to be. Getting a system to work is one thing and understanding it is another. Let's get to that!

Visualizing the decision boundary

What was true in trying to understand our data is true for trying to understand our classifier: visualization is the first step in understanding a system. We know the SVM somehow came up with a decision boundary that allowed us to correctly classify 80 percent of the test samples. But how can we find out what that decision boundary actually looks like?

For this, we will borrow a trick from the guys behind scikit-learn. The idea is to generate a fine grid of x and y coordinates and run that through the SVM's `predict` method. This will allow us to know, for every (x, y) point, what target label the classifier would have predicted.

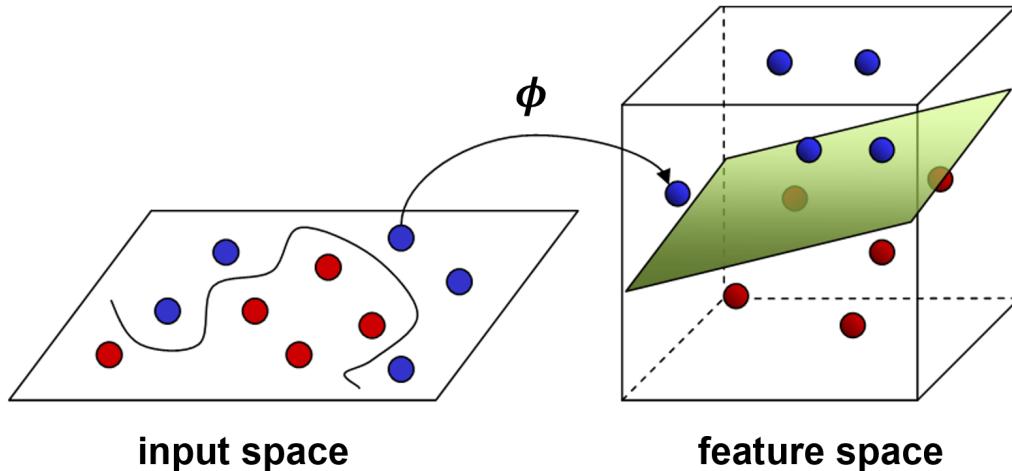
We will do this in a dedicated function, which we call `plot_decision_boundary` ...

Dealing with nonlinear decision boundaries

What if the data cannot be optimally partitioned using a linear decision boundary? In such a case, we say the data is not linearly separable.

The basic idea to deal with data that is not linearly separable is to create nonlinear combinations of the original features. This is the same as saying we want to project our data to a higher-dimensional space (for example, from 2D to 3D), in which the data suddenly becomes linearly separable.

This concept is illustrated in the following diagram:



The preceding diagram shows how to find linear hyperplanes in higher-dimensional spaces. If data in its original input space (left) cannot be linearly separated, we can apply a mapping function $\phi(.)$ that projects the data from 2D into a 3D (or a high-

dimensional) space. In this higher-dimensional space, we may find that there is now a linear decision boundary (which, in 3D, is a plane) that can separate the data.

*A linear decision boundary in an n-dimensional space is called a **hyperplane**. For example, a decision boundary in 6D feature space is a 5D hyperplane; in 3D feature space, it's a regular 2D plane; and in 2D space, it's a straight line.*

However, one problem with this mapping approach is that it is impractical in large dimensions because it adds a lot of extra terms to do the mathematical projections between the dimensions. This is where the so-called **kernel trick** comes into play.

Understanding the kernel trick

Granted, we won't have time to develop all the mathematics needed to truly understand the kernel trick. A more realistic section title would have been *Acknowledging that something called the kernel trick exists and accepting that it works*, but that would have been a bit wordy.

Here's the kernel trick in a nutshell.

In order to figure out the slope and orientation of the decision hyperplane in the high-dimensional space, we have to multiply all the feature values with appropriate weight values and sum them all up. The more dimensions our feature space has, the more work we have to do.

However, mathematicians smarter than us have long realized that an SVM has no need to explicitly work in the higher-dimensional ...

Knowing our kernels

OpenCV provides a whole range of SVM kernels to experiment with. Some of the most commonly used ones include the following:

- `cv2.ml.SVM_LINEAR`: This is the kernel we used previously. It provides a linear decision boundary in the original feature space (the x and y values).
- `cv2.ml.SVM_POLY`: This kernel provides a decision boundary that is a polynomial function in the original feature space. In order to use this kernel, we also have to specify a coefficient via `svm.setCoef0` (usually set to 0) and the degree of the polynomial via `svm.setDegree`.
- `cv2.ml.SVM_RBF`: This kernel implements the kind of Gaussian function we discussed earlier.
- `cv2.ml.SVM_SIGMOID`: This kernel implements a sigmoid function, similar to the one we encountered when talking about logistic regression in [Chapter 3, First Steps in Supervised Learning](#).
- `cv2.ml.SVM_INTER`: This kernel is a new addition to OpenCV 3. It separates classes based on the similarity of their histograms.

Implementing nonlinear SVMs

In order to test some of the SVM kernels we just talked about, we will return to our code sample mentioned earlier. We want to repeat the process of building and training the SVM on the dataset generated earlier, but this time, we want to use a whole range of different kernels:

```
In [13]: kernels = [cv2.ml.SVM_LINEAR, cv2.ml.SVM_INTER, ...
cv2.ml.SVM_SIGMOID, cv2.ml.SVM_RBF]
```

Do you remember what all of these stand for?

Setting a different SVM kernel is relatively simple. We take an entry from the `kernels` list and pass it to the `setKernels` method of the SVM class. That's all.

The laziest way to repeat things is to use a `for` loop as shown here:

```
In [14]: for idx, kernel in enumerate(kernels):
```

Then the steps are as follows: ...

Detecting pedestrians in the wild

We briefly talked about the difference between detection and recognition. While recognition is concerned with classifying objects (for example, as pedestrians, cars, bicycles, and so on), detection is basically answering the question: is there a pedestrian present in this image?

The core idea behind most detection algorithms is to split up an image into many small patches, and then classify each image patch as either containing a pedestrian or not. This is exactly what we are going to do in this section. In order to arrive at our own pedestrian detection algorithm, we need to perform the following steps:

1. Build a database of images containing pedestrians. These will be our positive data samples.
2. Build a database of images not containing pedestrians. These will be our negative data samples.
3. Train an SVM on the dataset.
4. Apply the SVM to every possible patch of a test image in order to decide whether the overall image contains a pedestrian.

Obtaining the dataset

For the purpose of this section, we will work with the MIT People dataset, which we are free to use for non-commercial purposes. So make sure not to use this in your groundbreaking autonomous start-up company before obtaining a corresponding software license.

However, if you followed our installation instructions from earlier and checked out the code on GitHub, you already have the dataset and are ready to go! The file can be found at <https://github.com/PacktPublishing/Machine-Learning-for-OpenCV-Second-Edition/blob/master/data/chapter6/pedestrians128x64.tar.gz>.

By referring to the following steps, you will learn to detect pedestrians in the wild:

1. Since we are supposed to run this code from a Jupyter Notebook in the ...

Taking a glimpse at the histogram of oriented gradients (HOG)

The HOG might just provide the help we're looking for in order to get this project done. The HOG is a feature descriptor for images, much like the ones we discussed in [Chapter 4, Representing Data and Engineering Features](#). It has been successfully applied to many different tasks in computer vision but seems to work especially well for classifying people.

The essential idea behind HOG features is that the local shapes and appearance of objects within an image can be described by the distribution of edge directions. The image is divided into small connected regions, within which a histogram of gradient directions (or edge directions) is compiled. Then, the descriptor is assembled by concatenating the different histograms. For improved performance, the local histograms can also be contrast-normalized, which results in better invariance to changes in illumination and shadowing.

The HOG descriptor is fairly accessible in OpenCV by means of `cv2.HOGDescriptor`, which takes a bunch of input arguments, such as the detection window size (minimum size of the object to be detected, 48 x 96), the block size (how large each box is, 16 x 16), the cell size (8 x 8), and the cell stride (how many pixels to move from one cell to the next, 8 x 8). For each of these cells, the HOG descriptor then calculates a histogram of oriented gradients using nine bins:

```
In [7]: win_size = (48, 96)
...     block_size = (16, 16)
...     block_stride = (8, 8)
...     cell_size = (8, 8)
...     num_bins = 9
...     hog = cv2.HOGDescriptor(win_size, block_size, block_stride,
...                         cell_size, num_bins)
```

Although this function call looks fairly complicated, these are actually the only values for which the HOG descriptor is implemented. The argument that matters the most is the window size (`win_size`).

All that's left to do is call `hog.compute` on our data samples. For this, we build a dataset of positive samples (`x_pos`) by randomly picking pedestrian images from our data directory. In the following code snippet, we randomly select 400 pictures from the over 900 available, and apply the HOG descriptor to them:

```
In [8]: import numpy as np
...     import random
...     random.seed(42)
...     X_pos = []
...     for i in random.sample(range(900), 400):
...         filename = "%s/per%05d.ppm" % (extractdir, i)
...         img = cv2.imread(filename)
...         if img is None:
...             print('Could not find image %s' % filename)
...             continue
...         X_pos.append(hog.compute(img, (64, 64)))
```

We should also remember that OpenCV wants the feature matrix to contain 32-bit floating point numbers, and the target labels to be 32-bit integers. We don't mind, since converting to NumPy arrays will allow us to easily investigate the sizes of the matrices we created:

```
In [9]: X_pos = np.array(X_pos, dtype=np.float32)
...     y_pos = np.ones(X_pos.shape[0], dtype=np.int32)
```

```
| ...      X_pos.shape, y_pos.shape  
| Out[9]: ((399, 1980, 1), (399,))
```

It looks like we picked a total of 399 training samples, each of which has 1,980 feature values (which are the HOG feature values).

Generating negatives

The real challenge, however, is to come up with the perfect example of a non-pedestrian. After all, it's easy to think of example images of pedestrians. But what is the opposite of a pedestrian?

This is actually a common problem when trying to solve new machine learning problems. Both research labs and companies spend a lot of time creating and annotating new datasets that fit their specific purpose.

If you're stumped, let me give you a hint on how to approach this. A good first approximation to finding the opposite of a pedestrian is to assemble a dataset of images that look like the images of the positive class but do not contain pedestrians. These images could contain anything like cars, bicycles, streets, houses, ...

Implementing the SVM

We already know how to build an SVM in OpenCV, so there's nothing much to see here. Planning ahead, we wrap the training procedure into a function, so that it's easier to repeat the procedure in the future:

```
In [15]: def train_svm(X_train, y_train):
...         svm = cv2.ml.SVM_create()
...         svm.train(X_train, cv2.ml.ROW_SAMPLE, y_train)
...         return svm
```

The same can be done for the scoring function. Here we pass a feature matrix, x and a label vector, y , but we do not specify whether we're talking about the training or the test set. In fact, from the viewpoint of the function, it doesn't matter what set the data samples belong to, as long as they have the right format:

```
In [16]: def score_svm(svm, X, y):
...         from sklearn import metrics
...         _, y_pred = svm.predict(X)
...         return metrics.accuracy_score(y, y_pred)
```

Then we can train and score the SVM with two short function calls:

```
In [17]: svm = train_svm(X_train, y_train)
In [18]: score_svm(svm, X_train, y_train)
Out[18]: 1.0
In [19]: score_svm(svm, X_test, y_test)
Out[19]: 0.64615384615384619
```

Thanks to the HOG feature descriptor, we make no mistake on the training set. However, our generalization performance is quite abysmal (64.6 percent), as it is much less than the training performance (100 percent). This is an indication that the model is overfitting the data. The fact that it is performing way better on the training set than the test set means that the model has resorted to memorizing the training samples, rather than trying to abstract it into a meaningful decision rule. What can we do to improve the model performance?

Bootstrapping the model

An interesting way to improve the performance of our model is to use bootstrapping. This idea was actually applied in one of the first papers on using SVMs in combination with HOG features for pedestrian detection. So let's pay a little tribute to the pioneers and try to understand what they did.

Their idea was quite simple. After training the SVM on the training set, they scored the model and found that the model produced some false positives. Remember that false positive means that the model predicted a positive (+) for a sample that was really a negative (-). In our context, this would mean the SVM falsely believed an image to contain a pedestrian. If this happens for a particular image in the dataset, this example ...

Detecting pedestrians in a larger image

What's left to do is to connect the SVM classification procedure with the process of detection. The way to do this is to repeat our classification for every possible patch in the image. This is similar to what we did earlier when we visualized decision boundaries; we created a fine grid and classified every point on that grid. The same idea applies here. We divide the image into patches and classify every patch as either containing a pedestrian or not.

By following these steps, you will be able to detect a pedestrian in an image:

1. We first have to loop over all possible patches in an image as follows, each time shifting our region of interest by a small number of `stride` pixels:

```
In [23]: stride = 16
...
...     found = []
...     for ystart in np.arange(0, img_test.shape[0], stride):
...         for xstart in np.arange(0, img_test.shape[1],
... stride):
```

2. We want to make sure that we do not go beyond the image boundaries:

```
...           if ystart + hroi > img_test.shape[0]:  
...               continue  
...           if xstart + wroi > img_test.shape[1]:  
...               continue
```

3. Then we cut out the ROI, preprocess it, and classify it:

```
...           roi = img_test[ystart:ystart + hroi,  
...                           xstart:xstart + wroi, :]  
...           feat = np.array([hog.compute(roi, (64, 64))])  
..., ypred = svm.predict(feat)
```

4. If that particular patch happens to be classified as a pedestrian, we add it to the list of successes:

```
...           if np.allclose(ypred, 1):  
...               found.append((ystart, xstart, hroi, wroi))
```

5. Because pedestrians could appear not just at various locations but also in various sizes, we would have to rescale the image and repeat the whole process. Thankfully, OpenCV has a convenience function for this multi-scale detection task in the form of the `detectMultiScale` function. This is a bit of a hack, but we can pass all SVM parameters to the `hog` object:

```
In [24]: rho, _, _ = svm.getDecisionFunction(0)  
...         sv = svm.getSupportVectors()  
...         hog.setSVMClassifier(np.append(sv.ravel(), rho))
```

6. Then it's possible to call the detection function:

```
In [25]: found = hog.detectMultiScale(img_test)
```

The function will return a list of bounding boxes that contain detected pedestrians.

This seems to work only for linear SVM classifiers. The OpenCV documentation is terribly inconsistent across versions in this regard, so I'm not sure at which version this started or stopped working. Be careful!

7. In practice, when people are faced with a standard task such as pedestrian detection, they often rely on pre-scanned SVM classifiers that are built into OpenCV. This is the method that I hinted at in the very beginning of this chapter. By loading either `cv2.HOGDescriptor_getDaimlerPeopleDetector()` or `cv2.HOGDescriptor_getDefaultPeopleDetector()`, we can get started with only a few lines of code:

```
In [26]: hogdef = cv2.HOGDescriptor()  
...          pdetect = cv2.HOGDescriptor_getDefaultPeopleDetector()  
In [27]: hogdef.setSVMClassifier(pdetect)  
In [28]: found, _ = hogdef.detectMultiScale(img_test)
```

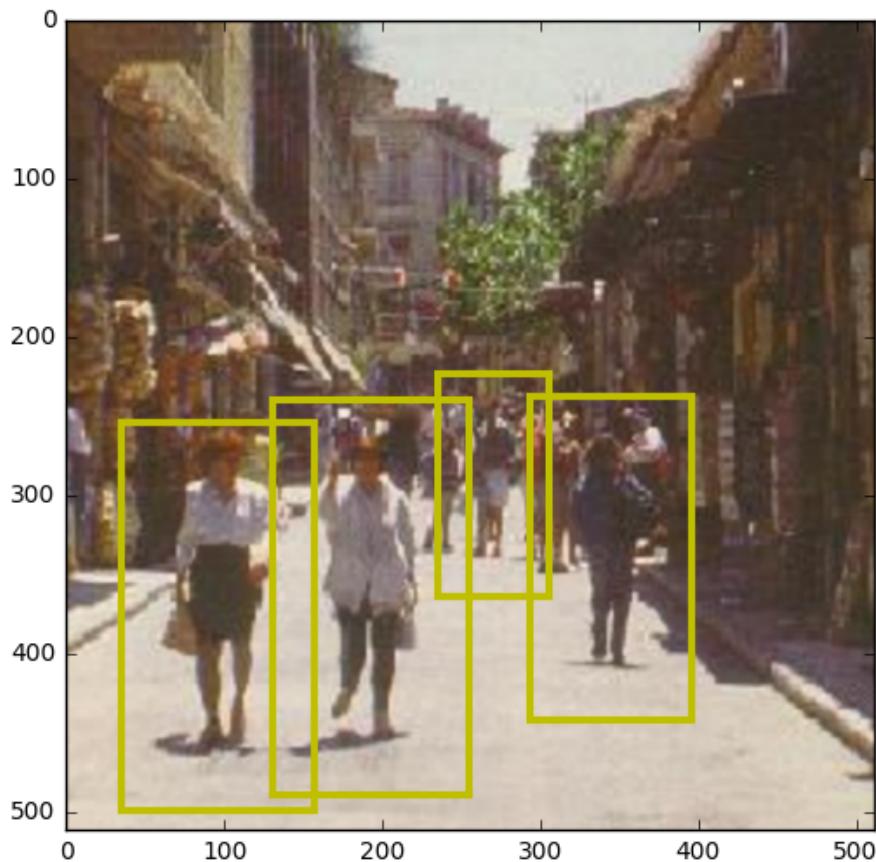
8. It's easy to plot the test image with matplotlib as shown here:

```
In [29]: from matplotlib import patches  
...          fig = plt.figure()  
...          ax = fig.add_subplot(111)  
...          ax.imshow(cv2.cvtColor(img_test, cv2.COLOR_BGR2RGB))
```

9. Then we can mark the detected pedestrians in the image by looping over the bounding boxes in `found`:

```
...     for f in found:  
...         ax.add_patch(patches.Rectangle((f[0], f[1]), f[2],  
f[3],  
...                                         color='y',  
linewidth=3,  
...                                         fill=False))
```

The result looks like this:



The preceding screenshot shows detected pedestrians in a test image.

Further improving the model

Although the RBF kernel makes for a good default kernel, it is not always the one that works best for our problem. The only real way to know which kernel works best on our data is to try them all and compare the classification performance across models. There are strategic ways to perform this so-called **hyperparameter tuning**, which we'll talk about in detail in [chapter 11, *Selecting the Right Model with Hyperparameter Tuning*](#).

What if we don't know how to do hyperparameter tuning properly yet?

Well, I'm sure you remember the first step in data understanding, *visualize the data*. Visualizing the data could help us understand if a linear SVM was powerful enough to classify the data, in which case there would be no ...

Multiclass classification using SVMs

SVMs are inherently two-class classifiers. In particular, the most prevalent method of multi-class classification in practice has been to create $|C|$ one-versus-rest classifiers (commonly referred to as **one-versus-all (OVA)** classification) where $|C|$ is the number of classes and to choose the class that classifies the test datum with the highest margin. Another approach is to develop a set of one-versus-one classifiers and to select the class that is chosen by the most classifiers. While this involves building $|C|(|C| - 1)/2$ classifiers, the time for training classifiers may decrease, since the training data set for each classifier is much smaller.

Now let's quickly jump onto how you can apply multi-class classification using SVMs with the help of a real-life dataset.

For the purpose of this section, we will work with the UCI Human Activity Recognition using smartphones dataset, which we are free to use for non-commercial purposes. So make sure not to use this in your groundbreaking autonomous start-up company before obtaining a corresponding software license.

The dataset can be obtained from the Kaggle website, <https://www.kaggle.com/uciml/human-activity-recognition-with-smartphones>. There you should find a Download button that leads you to a file called <https://www.kaggle.com/uciml/human-activity-recognition-with-smartphones/downloads/human-activity-recognition-with-smartphones.zip/1>.

However, if you followed our installation instructions from earlier and checked out the code on GitHub, you already have the dataset and are ready to go! The file can be found at `notebooks/data/multiclass`.

About the data

A group of 30 volunteers were chosen within an age group of 19-48 years and the experiments were carried out on them. Six activities were performed by each person, namely Walking, Walking_Upsairs, Walking_Downstairs, Sitting, Standing, and Laying with the help of a smartphone fastened around the waist. Mainly three-axial linear acceleration and three-axial angular velocity at a constant rate of 50 Hz were captured using embedded accelerometer and gyroscope. To label the data, the experiments have been video-recorded. The dataset has been randomly split into two sets, where 70% of the volunteers were selected for generating the training data and 30% the test data.

Attribute information

For each entry in the dataset, the following is provided:

- Triaxial acceleration from the accelerometer and the approximate acceleration of the body
- Triaxial angular velocity from the gyroscope
- Time and frequency domain variables with 561-feature vector
- Various labels of activity
- An identifier of the subject who was observed

By referring to the following steps, you will learn how to build a multi-class classification using SVMs:

1. Let's quickly import all the necessary libraries that you will need in order to implement an SVM with multi-class classification:

```
In [1]: import numpy as np
...
...     import pandas as pd
...     import matplotlib.pyplot as plt
...     %matplotlib inline
...     from sklearn.utils import shuffle
...     from sklearn.svm import SVC
...     from sklearn.model_selection import cross_val_score,
GridSearchCV
```

2. Next, you will be loading the dataset. Since we are supposed to run this code from a Jupyter Notebook in the `notebooks/` directory, the relative path to the data directory is simply `data/`:

```
In [2]: datadir = "data"  
...     dataset = "multiclass"  
...     train = shuffle(pd.read_csv("data/dataset/train.csv"))  
...     test = shuffle(pd.read_csv("data/dataset/test.csv"))
```

3. Let's check whether there are any missing values in the training and testing dataset; if there are any, then we will simply drop them from the dataset:

```
In [3]: train.isnull().values.any()  
Out[3]: False  
In [4]: test.isnull().values.any()  
Out[4]: False
```

4. Next, we will find the frequency distribution of the classes in the data, which means that we will check how many samples belong to each of the six classes:

```
In [5]: train_outcome = pd.crosstab(index=train["Activity"], #  
Make a crosstab  
columns="count") # Name the count column  
... train_outcome
```

From the following screenshot, you can observe that the `LAYING` class has the most samples, but overall, the data is

approximately equally distributed and there are no major signs of class imbalance:

Activity	col_0	count
LAYING	1407	
SITTING	1286	
STANDING	1374	
WALKING	1226	
WALKING_DOWNSTAIRS	986	
WALKING_UPSTAIRS	1073	

5. Next, we will separate out the predictors (input values) and outcome values (class labels) from the train and test datasets:

```
In [6]: X_train =  
pd.DataFrame(train.drop(['Activity','subject'],axis=1))  
...     Y_train_label = train.Activity.values.astype(object)  
...     X_test =  
pd.DataFrame(test.drop(['Activity','subject'],axis=1))  
...     Y_test_label = test.Activity.values.astype(object)
```

6. Since the SVM expects numerical input and labels, you will now transform the non-numerical labels into numerical labels. But first, we will have to import a `preprocessing` module from the `sklearn` library:

```
In [7]: from sklearn import preprocessing  
... encoder = preprocessing.LabelEncoder()
```

7. Now, we will encode the train and test labels into numerical values:

```
In [8]: encoder.fit(Y_train_label)
...
...     Y_train = encoder.transform(Y_train_label)
...
...     encoder.fit(Y_test_label)
...
...     Y_test = encoder.transform(Y_test_label)
```

8. Next, we will scale (normalise) the train and test feature set and for this, you will import `StandardScaler` from `sklearn`:

```
In [9]: from sklearn.preprocessing import StandardScaler
...
...     scaler = StandardScaler()
...
...     X_train_scaled = scaler.fit_transform(X_train)
...
...     X_test_scaled = scaler.transform(X_test)
```

9. Once the data is scaled and the labels are in a correct format, now is the time when we will fit the data. But before that, we will define a dictionary with the different parameter settings that the SVM will use while training itself, and this technique is called `GridSearchCV`. The parameter grid will be based on the results of a random search:

```
In [10]: params_grid = [ {'kernel': ['rbf'], 'gamma': [1e-3, 1e-4],
...
...             'C': [1, 10, 100, 1000]}, {
'kernel': ['linear'], 'C': [1, 10, 100,
1000]}]
```

10. Finally, we will call `GridSearchCV` on the data using the preceding parameters for the best SVM fit:

```
In [11]: svm_model = GridSearchCV(SVC(), params_grid, cv=5)
...         svm_model.fit(X_train_scaled, Y_train)
```

11. It's time to check how well the SVM model was trained on the data; in short, we will find the accuracy. Not only that, but we will also check what the parameter settings were for which SVM performed the best:

```
In [12]: print('Best score for training data:',
svm_model.best_score_, "\n")
...     print('Best C:', svm_model.best_estimator_.C, "\n")
...     print('Best
Kernel:', svm_model.best_estimator_.kernel, "\n")
...     print('Best
Gamma:', svm_model.best_estimator_.gamma, "\n")
Out[12]: Best score for training data: 0.986
...
...     Best C: 100
...     Best Kerne: rbf
...     Best Gamma: 0.001
```

Voila! As we can see, the SVM achieved 98.6% accuracy on the training data on a multi-class classification problem. But hold down your horses until we find the accuracy on the test data. So, let's quickly check that:

```
In [13]: final_model = svm_model.best_estimator_
... print("Training set score for SVM: %f" %
final_model.score(X_train_scaled , Y_train))
... print("Testing set score for SVM: %f" %
final_model.score(X_test_scaled , Y_test ))
Out[13]: Training set score for SVM: 1.00
... Testing set score for SVM: 0.9586
```

Wow! Isn't that amazing? We were able to achieve 95.86% accuracy on the testing set; that's the power of SVMs.

Summary

In this chapter, we learned about SVMs in all their forms and flavors. We now know how to draw decision boundaries in 2D and hyperplanes in high-dimensional spaces. We learned about different SVM kernels and look at how to implement them in OpenCV.

In addition, we also applied our newly gained knowledge to the practical example of pedestrian detection. For this, we had to learn about the HOG feature descriptor, and how to collect suitable data for the task. We used bootstrapping to improve the performance of our classifier and combined the classifier with OpenCV's multi-scale detection mechanism.

Not only was that a lot to digest in a single chapter, but you have also made it through half of the book. Congrats!

In the next chapter, ...

Implementing a Spam Filter with Bayesian Learning

Before we get to grips with advanced topics, such as cluster analysis, deep learning, and ensemble models, let's turn our attention to a much simpler model that we have overlooked so far: the Naive Bayes classifier.

Naive Bayes classifiers have their roots in Bayesian inference, named after the famed statistician and philosopher Thomas Bayes (1701-1761). Bayes' theorem famously describes the probability of an event based on prior knowledge of conditions that might lead to the event. We can use Bayes' theorem to build a statistical model that not only can classify data but can also provide us with an estimate of how likely it is that our classification is correct. In our case, we can use Bayesian inference to dismiss an email as spam with high confidence and to determine the probability of a woman having breast cancer, given a positive screening test.

We have now gained enough experience with the mechanics of implementing machine learning methods, and so we should no longer be afraid to try and understand the theory behind them. Don't worry, we won't write a book on it, but we need some understanding of the theory to appreciate a model's inner workings. After that, I am sure you will find that Bayesian classifiers are easy to implement, are computationally efficient, and tend to perform quite well on relatively small datasets. In this chapter, we will understand the Naive Bayes classifier and then implement our first Bayesian classifier. We will then classify emails using the Naive Bayes classifier.

In this chapter, we will cover the following topics:

- Understanding the Naive Bayes classifier
- Implementing your first Bayesian classifier
- Classifying emails using the Naive Bayes classifier

Technical requirements

You can refer to the code for this chapter from the following link: <https://github.com/PacktPublishing/Machine-Learning-for-OpenCV-Second-Edition/tree/master/Chapter07>.

Here is a summary of the software and hardware requirements:

- You will need OpenCV version 4.1.x (4.1.0 or 4.1.1 will both work just fine).
- You will need Python version 3.6 (any Python version 3.x will be fine).
- You will need Anaconda Python 3 for installing Python and the required modules.
- You can use any OS—macOS, Windows, and Linux-based OSes—with this book. We recommend you have at least 4 GB RAM in your system.
- You don't need to have a GPU to run the code provided along with this book.

Understanding Bayesian inference

Although Bayesian classifiers are relatively simple to implement, the theory behind them can be quite counter-intuitive at first, especially if you are not too familiar with probability theory yet. However, the beauty of Bayesian classifiers is that they understand the underlying data better than all of the classifiers we have encountered so far. For example, standard classifiers, such as the k -nearest neighbor algorithm or decision trees, might be able to tell us the target label of a never-before-seen data point. However, these algorithms have no concept of how likely it is for their predictions to be right or wrong. We call them discriminative models. Bayesian models, on the other hand, have an understanding of the underlying probability distribution that caused the data. We call them generative models because they don't just put labels on existing data points—they can also generate new data points with the same statistics.

If this last paragraph was a bit over your head, you might enjoy the following brief on probability theory. It will be important for the upcoming sections.

Taking a short detour through probability theory

In order to appreciate Bayes' theorem, we need to get a hold of the following technical terms:

- **Random variable:** This is a variable whose value depends on chance. A good example is the act of flipping a coin, which might turn up heads or tails. If a random variable can take on only a limited number of values, we call it discrete (such as a coin flip or a dice roll); otherwise, we call it a continuous random variable (such as the temperature on a given day). Random variables are often typeset as capital letters.
- **Probability:** This is a measure of how likely it is for an event to occur. We denote the probability of an event, e , happening as $p(e)$, which must be a number between 0 and 1 (or between ...

Understanding Bayes' theorem

There are quite a few scenarios where it would be really good to know how likely it is for our classifier to make a mistake. For example, in chapter 5, *Using Decision Trees to Make a Medical Diagnosis*, we trained a decision tree to diagnose women with breast cancer based on some medical tests. You can imagine that, in this case, we would want to avoid a misdiagnosis at all costs; diagnosing a healthy woman with breast cancer (a false positive) would be both soul-crushing and lead to unnecessary, expensive medical procedures, whereas missing a woman's breast cancer (a false negative) might eventually cost the woman her life.

It's good to know we have Bayesian models to count on. Let's walk through a specific (and quite famous) example from <http://yudkowsky.net/rational/bayes>:

"1% of women at age forty who participate in routine screening have breast cancer. 80% of women with breast cancer will get positive mammographies. 9.6% of women without breast cancer will also get positive mammographies. A woman in this age group had a positive mammography in a routine screening. What is the probability that she actually has breast cancer?"

What do you think the answer is?

Well, given that her mammography was positive, you might reason that the probability of her having cancer is quite high (somewhere near 80%). It seems much less likely that the woman would belong to 9.6% with false positives, so the real probability is probably somewhere between 70% and 80%.

I'm afraid that's not correct.

Here's one way to think about this problem. For the sake of simplicity, let's assume we're looking at some concrete number of patients, say 10,000. Before the mammography screening, the 10,000 women can be divided into two groups:

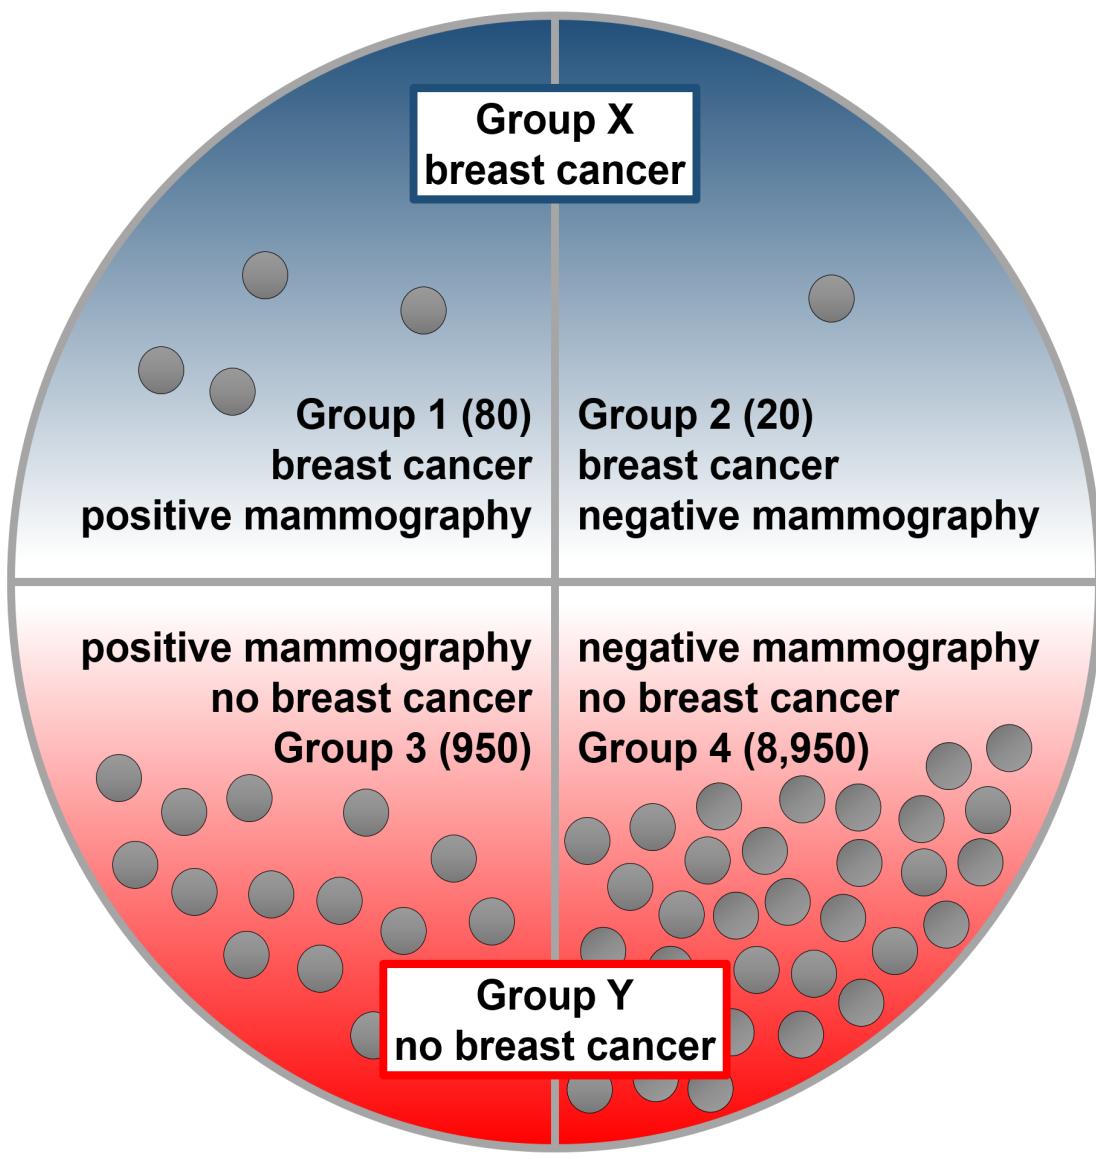
- **Group X:** 100 women *with* breast cancer
- **Group Y:** 9,900 women *without* breast cancer

So far, so good. If we sum up the numbers in the two groups, we get a total of 10,000 patients, confirming that nobody has been lost in the math. After the mammography screening, we can divide the 10,000 women into four groups:

- **Group 1:** 80 women with breast cancer and a positive mammography
- **Group 2:** 20 women with breast cancer and a negative mammography
- **Group 3:** Around 950 women without breast cancer and with a positive mammography
- **Group 4:** Approx. 8,950 women without breast cancer and with a negative mammography

From the preceding analysis, you can see that the sum of all the four groups is 10,000. The sum of **Group 1** and **Group 2** (with breast cancer) corresponds to **Group X**, and the sum of **Group 3** and **Group 4** (without breast cancer) corresponds to **Group Y**.

This might become clearer when we draw it out:



In this diagram, the top half corresponds to **Group X**, and the bottom half corresponds to **Group Y**. Analogously, the left half corresponds to all women with positive mammographies, and the right half corresponds to all women with negative mammographies.

Now, it is easier to see that what we are looking for concerns only the left half of the diagram. The proportion of cancer patients with

positive results within the group of all patients with positive results is the proportion of Group 1 within Groups 1 and 3:

$$80 / (80 + 950) = 80 / 1,030 = 7.8\%$$

In other words, if you offer a mammography to 10,000 patients, then out of the 1,030 with a positive mammographies, there would be 80 patients with positive mammography having cancer. The answer a doctor should give a positive mammography patient if she asks about the chance she has breast cancer: roughly 1 out of 13 will have cancer, given that 13 patients ask this question.

What we just calculated is called a **conditional probability**: what is our **degree of belief** that a woman has breast cancer **under the condition** of (we also say **given**) a positive mammography? As in the last subsection, we denote this with $p(\text{cancer}|\text{mammography})$, or $p(C|M)$ for short. Using capital letters once again enforces the idea that both health and the mammography can have several outcomes, depending on several underlying (and possibly unknown) causes. Therefore, they are random variables.

Then, we can express $P(C|M)$ with the following formula:

$$p(C|M) = \frac{p(C, M)}{p(C, M) + p(\sim C, M)}$$

Here, $p(C, M)$ denotes the probability that both C and M are true (meaning the probability that a woman both has cancer and a positive mammography). This is equivalent to the probability of a woman belonging to Group 1, as shown earlier.

The comma (,) means logical *and*, and the tilde (\sim) stands for logical *not*. Hence, $p(\sim C, M)$ denotes the probability that C is not true and M is true (meaning the probability that a woman does not have cancer but has a positive mammography). This is equivalent to the probability of a woman belonging to Group 3. So, the denominator

basically adds up women in Group 1 ($p(C, M)$) and Group 3 ($p(\sim C, M)$).

But wait! Those two groups together simply denote the probability of a woman having a positive mammography, $p(M)$. Therefore, we can simplify the preceding equation:

$$p(C|M) = \frac{p(C, M)}{p(M)}$$

The Bayesian version is to re-interpret what $p(C, M)$ means. We can express $p(C, M)$ as follows:

$$p(C, M) = p(M|C)p(C)$$

Now it gets a bit confusing. Here, $p(C)$ is simply the probability that a woman has cancer (corresponding to the aforementioned Group X). Given that a woman has cancer, what is the probability that her mammography will be positive? From the problem question, we know it is 80%. This is $p(M|C)$, the probability of M given C .

Replacing $p(C, M)$ in the first equation with this new formula, we get the following equation:

$$p(C|M) = \frac{p(C)p(M|C)}{p(M)}$$

In the Bayesian world, these terms all have their specific names:

- $p(C|M)$ is called the **posterior**, which is always the thing we want to compute. In our example, this corresponds to the degree of belief that a woman has breast cancer, given a positive mammography.

- $p(C)$ is called the **prior** as it corresponds to our initial knowledge about how common breast cancer is. We also call this our initial degree of belief in C .
- $p(M|C)$ is called the **likelihood**.
- $p(M)$ is called **evidence**.

So, you can rewrite the equation one more time, as follows:

$$posterior = \frac{prior \times likelihood}{evidence}$$

Mostly, there is interest only in the numerator of that fraction because the denominator does not depend on C , so the denominator is constant and can be neglected.

Understanding the Naive Bayes classifier

So far, we have only talked about one piece of evidence. However, in most real-world scenarios, we have to predict an outcome (such as a random variable, Y) given multiple pieces of evidence (such as random variables X_1 and X_2). So, instead of calculating $p(Y|X)$, we would often have to calculate $p(Y|X_1, X_2, \dots, X_n)$. Unfortunately, this makes the math very complicated. For two random variables, X_1 and X_2 , the joint probability would be computed like this:

$$p(C, X_1, X_2) = p(X_1|X_2, C)p(X_2|C)p(C)$$

The ugly part is the term $p(X_1|X_2, C)$, which says that the conditional probability of X_1 depends on all other variables, including C . This gets even ...

Implementing your first Bayesian classifier

But enough with the math, let's do some coding!

In the previous chapter, we learned how to generate a number of Gaussian blobs using scikit-learn. Do you remember how that is done?

Creating a toy dataset

The function I'm referring to resides within scikit-learn's `datasets` module. Let's create 100 data points, each belonging to one of two possible classes, and group them into two Gaussian blobs. To make the experiment reproducible, we specify an integer to pick a seed for `random_state`. You can again pick whatever number you prefer. Here, I went with Thomas Bayes' year of birth (just for kicks):

```
In [1]: from sklearn import datasets...      x, y =
datasets.make_blobs(100, 2, centers=2,          random_state=1701,
cluster_std=2)
```

Let's have a look at the dataset we just created using our trusty friend, Matplotlib:

```
In [2]: import matplotlib.pyplot as plt...
plt.style.use('ggplot')...    %matplotlib inlineIn [3]:
plt.scatter(X[:, 0], X[:, ...]
```

Classifying the data with a normal Bayes classifier

We will then use the same procedure as in earlier chapters to train a **normal Bayes classifier**. Wait, why not a Naive Bayes classifier? Well, it turns out OpenCV doesn't really provide a true Naive Bayes classifier. Instead, it comes with a Bayesian classifier that doesn't necessarily expect features to be independent, but rather expects the data to be clustered into Gaussian blobs. This is exactly the kind of dataset we created earlier!

By following these steps, you will learn how to build a classifier with a normal Bayes classifier:

1. We can create a new classifier using the following function:

```
In [5]: import cv2  
...      model_norm = cv2.ml.NormalBayesClassifier_create()
```

2. Then, training is done via the `train` method:

```
In [6]: model_norm.train(X_train, cv2.ml.ROW_SAMPLE, y_train)  
Out[6]: True
```

3. Once the classifier has been trained successfully, it will return `True`. We go through the motions of predicting and scoring the

classifier, just like we have done a million times before:

```
In [7]: _, y_pred = model_norm.predict(X_test)
In [8]: from sklearn import metrics
...      metrics.accuracy_score(y_test, y_pred)
Out[8]: 1.0
```

4. Even better—we can reuse the plotting function from the last chapter to inspect the decision boundary! If you recall, the idea was to create a mesh grid that would encompass all data points and then classify every point on the grid. The mesh grid is created via the NumPy function of the same name:

```
In [9]: def plot_decision_boundary(model, X_test, y_test):
...
...      # create a mesh to plot in
...      h = 0.02 # step size in mesh
...      x_min, x_max = X_test[:, 0].min() - 1, X_test[:, 0].max() +
...                      1
...      y_min, y_max = X_test[:, 1].min() - 1, X_test[:, 1].max() +
...                      1
...      xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
...                          np.arange(y_min, y_max, h))
```

5. The `meshgrid` function will return two floating-point matrices, `xx` and `yy`, that contain the x and y coordinates of every coordinate point on the grid. We can flatten these matrices into column vectors using the `ravel` function and stack them to form a new matrix, `x_hypo`:

```
...           X_hypo =
np.column_stack((xx.ravel().astype(np.float32),
...
yy.ravel().astype(np.float32)))
```

6. `x_hypo` now contains all `x` values in `x_hypo[:, 0]` and all `y` values in `x_hypo[:, 1]`. This is a format that the `predict` function can understand:

```
...           ret = model.predict(X_hypo)
```

7. However, we want to be able to use models from both OpenCV and scikit-learn. The difference between the two is that OpenCV returns multiple variables (a Boolean flag indicating success/failure and the predicted target labels), whereas scikit-learn returns only the predicted target labels. Hence, we can check whether the `ret` output is a tuple, in which case, we know we're dealing with OpenCV. In this case, we store the second element of the tuple (`ret[1]`). Otherwise, we are dealing with scikit-learn and don't need to index into `ret`:

```
...           if isinstance(ret, tuple):
...               zz = ret[1]
...           else:
...               zz = ret
...           zz = zz.reshape(xx.shape)
```

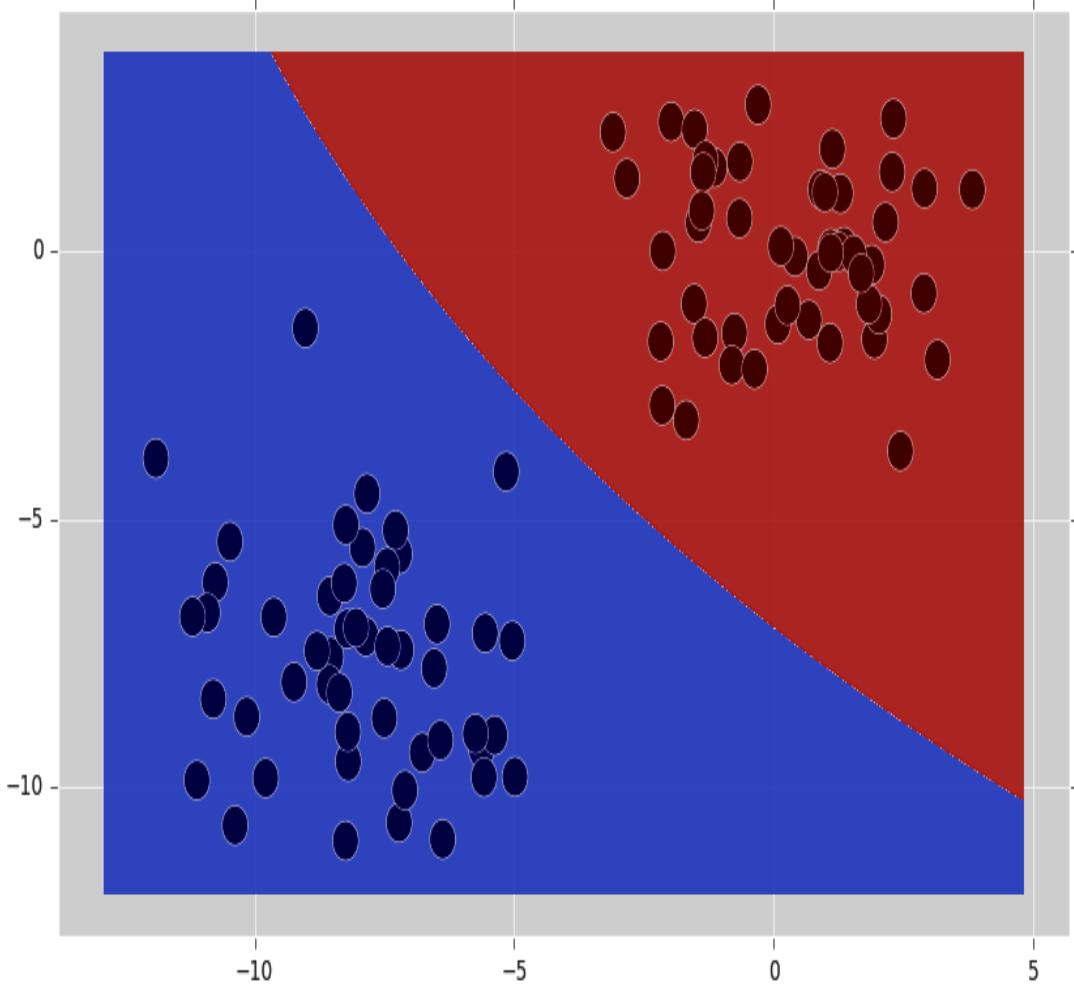
8. All that's left to do is to create a contour plot where `zz` indicates the color of every point on the grid. On top of that, we plot the data points using our trusty scatter plot:

```
...         plt.contourf(xx, yy, zz, cmap=plt.cm.coolwarm,
alpha=0.8)
...
...         plt.scatter(X_test[:, 0], X_test[:, 1], c=y_test,
s=200)
```

9. We call the function by passing a model (`model_norm`), a feature matrix (`x`), and a target label vector (`y`):

```
In [10]: plot_decision_boundary(model_norm, X, y)
```

The output looks like this:



So far, so good. The interesting part is that a Bayesian classifier also returns the probability with which each data point has been classified:

```
In [11]: ret, y_pred, y_proba = model_norm.predictProb(X_test)
```

The function returns a Boolean flag (`True` for success and `False` for failure), the predicted target labels (`y_pred`), and the conditional probabilities (`y_proba`). Here, `y_proba` is an $N \times 2$ matrix that indicates, for every one of the N data points, the probability with which it was classified as either class 0 or class 1:

```
In [12]: y_proba.round(2)
Out[12]: array([[ 0.15000001,  0.05        ],
               [ 0.08        ,  0.          ],
               [ 0.          ,  0.27000001],
               [ 0.          ,  0.13        ],
               [ 0.          ,  0.          ],
               [ 0.18000001,  1.88        ],
               [ 0.          ,  0.          ],
               [ 0.          ,  1.88        ],
               [ 0.          ,  0.          ],
               [ 0.          ,  0.          ]], dtype=float32)
```

This means that, for the first data point (top row), the probability of it belonging to class 0 (that is, $p(C_0|X)$) is 0.15 (or 15%). Similarly, the probability of belonging to class 1 is $p(C_1|X) = 0.05$.

*The reason why some of the rows show values greater than 1 is that OpenCV does not really return probability values. Probability values are always between 0 and 1, and each row in the preceding matrix should add up to 1. Instead, what is being reported is a **likelihood**, which is basically the numerator of the conditional probability equation, $p(C) p(M|C)$. The denominator, $p(M)$, does not need to be computed. All we need to know is that $0.15 > 0.05$ (top row). Hence, the data point most likely belongs to class 0.*

Classifying the data with a Naive Bayes classifier

The following steps will help you build a Naive Bayes classifier:

1. We can compare the result to a true Naive Bayes classifier by asking scikit-learn for help:

```
In [13]: from sklearn import naive_bayes...           model_naive =  
naive_bayes.GaussianNB()
```

2. As usual, training the classifier is done via the `fit` method:

```
In [14]: model_naive.fit(X_train, y_train)Out[14]:  
GaussianNB(priors=None)
```

3. Scoring the classifier is built in:

```
In [15]: model_naive.score(X_test, y_test)Out[15]: 1.0
```

4. Again a perfect score! However, in contrast to OpenCV, this classifier's `predict_proba` method returns true probability values, because all values are between 0 and 1 and because all rows add up to 1:

```
In [16]: yprob = model_naive.predict_proba(X_test) ...
```

Visualizing conditional probabilities

By referring to the following steps, you will be able to visualize conditional probabilities:

1. For this, we will slightly modify the plot function from the previous example. We start out by creating a mesh grid between (x_{\min} , x_{\max}) and (y_{\min} , y_{\max}):

```
In [18]: def plot_proba(model, X_test, y_test):
    ...
    # create a mesh to plot in
    ...
    h = 0.02 # step size in mesh
    ...
    x_min, x_max = X_test[:, 0].min() - 1, X_test[:, 0].max() + 1
    ...
    y_min, y_max = X_test[:, 1].min() - 1, X_test[:, 1].max() + 1
    ...
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                         np.arange(y_min, y_max, h))
```

2. Then, we flatten xx and yy and add them column-wise to the feature matrix, X_{hypo} :

```
...
X_hypo =
np.column_stack((xx.ravel().astype(np.float32),
...
yy.ravel().astype(np.float32)))
```

3. If we want to make this function work with both OpenCV and scikit-learn, we need to implement a switch for `predictProb` (in the case of OpenCV) and `predict_proba` (in the case of scikit-learn). For this, we check whether `model` has a method called `predictProb`. If the method exists, we can call it; otherwise, we assume we're dealing with a model from scikit-learn:

```
...         if hasattr(model, 'predictProb'):
...             _, _, y_proba = model.predictProb(X_hypo)
...         else:
...             y_proba = model.predict_proba(X_hypo)
```

4. Like in `In [16]`, which we saw earlier, `y_proba` will be a 2D matrix containing, for each data point, the probability of the data belonging to class 0 (in `y_proba[:, 0]`) and to class 1 (in `y_proba[:, 1]`). An easy way to convert these two values into a color that the contour function can understand is to simply take the difference of the two probability values:

```
...         zz = y_proba[:, 1] - y_proba[:, 0]
...         zz = zz.reshape(xx.shape)
```

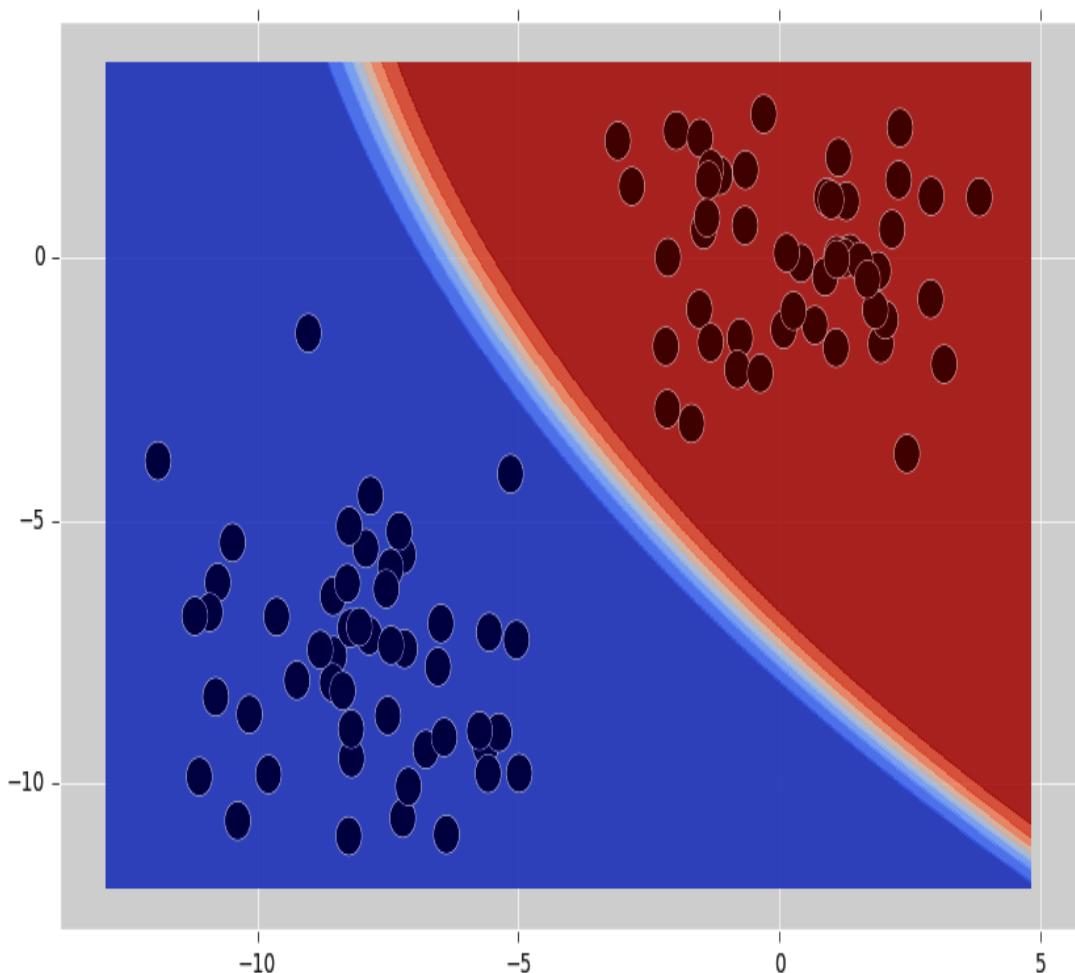
5. The last step is to plot `x_test` as a scatter plot on top of the colored mesh grid:

```
... plt.contourf(xx, yy, zz, cmap=plt.cm.coolwarm, alpha=0.8)
... plt.scatter(X_test[:, 0], X_test[:, 1], c=y_test, s=200)
```

6. Now, we are ready to call the function:

```
In [19]: plot_proba(model_naive, X, y)
```

The result looks like this:



The preceding screenshot shows conditional probabilities of a Naive Bayes classifier.

Classifying emails using the Naive Bayes classifier

The final task of this chapter will be to apply our newly gained skills to a real spam filter! This task deals with solving a binary-class (spam/ham) classification problem using the Naive Bayes algorithm.

Naive Bayes classifiers are actually a very popular model for email filtering. Their naivety lends itself nicely to the analysis of text data, where each feature is a word (or a **bag of words**), and it would not be feasible to model the dependence of every word on every other word.

There are a bunch of good email datasets out there, such as the following:

- The Hewlett-Packard spam database: <https://archive.ics.uci.edu/ml/machine-learning-databases/spambase>
- The Enron-Spam dataset: <http://www.aueb.gr/users/ion/data/enron-spam> ...

Loading the dataset

You can refer to these steps to load the dataset:

1. If you downloaded the latest code from GitHub, you will find several `.zip` files in the `notebooks/data/chapter7` directory. These files contain raw email data (with fields for `To:`, `Cc:`, and text body) that are either classified as spam (with the `SPAM = 1` class label) or not (also known as ham, the `HAM = 0` class label).
2. We build a variable called `sources`, which contains all of the raw data files:

```
In [1]: HAM = 0
...
SPAM = 1
...
datadir = 'data/chapter7'
...
sources = [
...
    ('beck-s.tar.gz', HAM),
...
    ('farmer-d.tar.gz', HAM),
...
    ('kaminski-v.tar.gz', HAM),
...
    ('kitchen-l.tar.gz', HAM),
...
    ('lokay-m.tar.gz', HAM),
...
    ('williams-w3.tar.gz', HAM),
...
    ('BG.tar.gz', SPAM),
...
    ('GP.tar.gz', SPAM),
...
    ('SH.tar.gz', SPAM)
...
]
```

3. The first step is to extract these files into subdirectories. For this, we can use the `extract_tar` function we wrote in the previous chapter:

```
In [2]: def extract_tar(datafile, extractdir):
...
    try:
...
        import tarfile
...
    except ImportError:
...
        raise ImportError("You do not have tarfile
installed. ")
...
...
    "Try unzipping the file
outside of "
...
...
    "Python.")
...
    tar = tarfile.open(datafile)
...
    tar.extractall(path=extractdir)
...
    tar.close()
...
    print("%s successfully extracted to %s" %
(datafile,
...
extractdir))
```

4. To apply the function to all data files in the sources, we need to run a loop. The `extract_tar` function expects a path to the `.tar.gz` file—which we build from `datadir` and an entry in `sources`—and a directory to extract the files to (`datadir`). This will extract all emails in, for example, `data/chapter7/beck-s.tar.gz` to the `data/chapter7/beck-s/` subdirectory:

```
In [3]: for source, _ in sources:
...
    datafile = '%s/%s' % (datadir, source)
...
    extract_tar(datafile, datadir)
Out[3]: data/chapter7/beck-s.tar.gz successfully extracted to
data/chapter7
data/chapter7/farmer-d.tar.gz successfully extracted to
data/chapter7
data/chapter7/kaminski-v.tar.gz successfully extracted
to
data/chapter7
data/chapter7/kitchen-l.tar.gz successfully extracted
to
data/chapter7
data/chapter7/lokay-m.tar.gz successfully extracted to
data/chapter7
```

```
          data/chapter7/williams-w3.tar.gz successfully extracted
          to
          data/chapter7
          data/chapter7/BG.tar.gz successfully extracted to
data/chapter7
          data/chapter7/GP.tar.gz successfully extracted to
data/chapter7
          data/chapter7/SH.tar.gz successfully extracted to
data/chapter7
```

Now here's the tricky bit. Every one of these subdirectories contains many other directories, wherein the text files reside. So, we need to write two functions:

- `read_single_file(filename)`: This is a function that extracts the relevant content from a single file called `filename`.
- `read_files(path)`: This is a function that extracts the relevant content from all files in a particular directory called `path`.

To extract the relevant content from a single file, we need to be aware of how each file is structured. The only thing we know is that the header section of the email (From:, To:, and Cc:) and the main body of text are separated by a newline character, '`\n`'. So, what we can do is iterate over every line in the text file and keep only those lines that belong to the main text body, which will be stored in the variable `lines`. We also want to keep a Boolean flag, `past_header`, around, which is initially set to `False` but will be flipped to `True` once we are past the header section:

1. We start by initializing those two variables:

```
In [4]: import os
...
...     def read_single_file(filename):
...             past_header, lines = False, []
```

2. Then, we check whether a file with the name `filename` exists.

If it does, we start looping over it line by line:

```
...     if os.path.isfile(filename):
...         f = open(filename, encoding="latin-1")
...         for line in f:
```

You may have noticed the `encoding="latin-1"` part. Since some of the emails are not in Unicode, this is an attempt to decode the files correctly.

We do not want to keep the header information, so we keep looping until we encounter the '`\n`' character, at which point we flip `past_header` from `False` to `True`.

3. At this point, the first condition of the following `if-else` clause is met, and we append all remaining lines in the text file to the `lines` variable:

```
...             if past_header:
...                 lines.append(line)
...             elif line == '\n':
...                 past_header = True
...             f.close()
```

4. In the end, we concatenate all lines into a single string, separated by the newline character, and return both the full path to the file and the actual content of the file:

```
...             content = '\n'.join(lines)
...             return filename, content
```

5. The job of the second function will be to loop over all files in a folder and call `read_single_file` on them:

```
In [5]: def read_files(path):
...     for root, dirnames, filenames in os.walk(path):
...         for filename in filenames:
...             filepath = os.path.join(root, filename)
...             yield read_single_file(filepath)
```

Here, `yield` is a keyword that is similar to `return`. The difference is that `yield` returns a generator instead of the actual values, which is desirable if you expect to have a large number of items to iterate over.

Building a data matrix using pandas

Now, it's time to introduce another essential data science tool that comes preinstalled with Python Anaconda: **pandas**. pandas is built on NumPy and provides several useful tools and methods to deal with data structures in Python. Just as we generally import NumPy under the alias, `np`, it is common to import pandas under the `pd` alias:

```
In [6]: import pandas as pd
```

pandas provide a useful data structure called a DataFrame, which can be understood as a generalization of a 2D NumPy array, as shown here:

```
In [7]: pd.DataFrame({...           'model': [...           'Normal  
Bayes',...           'Multinomial Bayes',...           'Bernoulli  
Bayes'...           ],...           'class': [...           'cv2.ml.NormalBayesClassifier_create()',...           'sklearn.naive_bayes.MultinomialNB()',...           'sklearn.naive_bayes.BernoulliNB()' ...
```

Preprocessing the data

Scikit-learn offers several options when it comes to encoding text features, which we discussed in [chapter 4, *Representing Data and Engineering Features*](#). One of the simplest methods of encoding text data, as you may recall, is by **word count**; for each phrase, you count the number of occurrences of each word within it. In scikit-learn, this is easily done using `CountVectorizer`:

```
In [10]: from sklearn import feature_extraction  
...      counts = feature_extraction.text.CountVectorizer()  
...      X = counts.fit_transform(data['text'].values)  
...      X.shape  
Out[10]: (52076, 643270)
```

The result is a giant matrix, which tells us that we harvested a total of 52,076 emails that collectively contain 643,270 different words. However, scikit-learn is smart and saved the data in a sparse matrix:

```
In [11]: X  
Out[11]: <52076x643270 sparse matrix of type '<class 'numpy.int64'>'  
           with 8607632 stored elements in Compressed Sparse Row  
           format>
```

To build the vector of target labels (`y`), we need to access data in the pandas DataFrame. This can be done by treating the DataFrame like a dictionary, where the `values` attribute will give us access to the underlying NumPy array:

```
In [12]: y = data['class'].values
```

Training a normal Bayes classifier

From here on out, things are (almost) like they always were. We can use scikit-learn to split the data into training and test sets (let's reserve 20% of all data points for testing):

```
In [13]: from sklearn import model_selection as ms...      x_train,
x_test, y_train, y_test = ms.train_test_split(...           x, y,
test_size=0.2, random_state=42...)
```

We can instantiate a new normal Bayes classifier with OpenCV:

```
In [14]: import cv2...      model_norm =
cv2.ml.NormalBayesClassifier_create()
```

However, OpenCV does not know about sparse matrices (at least its Python interface does not). If we were to pass `x_train` and `y_train` to the `train` function as we did earlier, OpenCV would complain that the data matrix is not a NumPy array. ...

Training on the full dataset

However, if you want to classify the full dataset, we need a more sophisticated approach. We turn to scikit-learn's Naive Bayes classifier, as it understands how to handle sparse matrices. In fact, if you didn't pay attention and treated `x_train` like every NumPy array before, you might not even notice that anything is different:

```
In [17]: from sklearn import naive_bayes  
...      model_naive = naive_bayes.MultinomialNB()  
...      model_naive.fit(X_train, y_train)  
Out[17]: MultinomialNB(alpha=1.0, class_prior=None, fit_prior=True)
```

Here, we used `MultinomialNB` from the `naive_bayes` module, which is the version of Naive Bayes classifier that is best suited to handling categorical data, such as word counts.

The classifier is trained almost instantly and returns the scores for both the training and test sets:

```
In [18]: model_naive.score(X_train, y_train)  
Out[18]: 0.95086413826212191  
In [19]: model_naive.score(X_test, y_test)  
Out[19]: 0.94422043010752688
```

And there we have it: 94.4% accuracy on the test set! Pretty good for not doing much other than using the default values, isn't it?

However, what if we were super critical of our own work and wanted to improve the result even further? There are a couple of things we could do.

Using n-grams to improve the result

One thing to do is to use n -gram counts instead of plain word counts. So far, we have relied on what is known as a bag of words: we simply threw every word of an email into a bag and counted the number of its occurrences. However, in real emails, the order in which words appear can carry a great deal of information!

This is exactly what n -gram counts are trying to convey. You can think of an n -gram as a phrase that is n words long. For example, the phrase *Statistics has its moments* contains the following 1-grams: *Statistics*, *has*, *its*, and *moments*. It also has the following 2-grams: *Statistics has*, *has its*, and *its moments*. It also has two 3-grams (*Statistics has its* and *has its moments*) and only a single ...

Discovering Hidden Structures with Unsupervised Learning

So far, we have focused our attention exclusively on supervised learning problems, where every data point in the dataset had a known label or target value. However, what do we do when there is no known output or no teacher to supervise the learning algorithm?

This is what **unsupervised learning** is all about. In unsupervised learning, the learning process is shown only in the input data and is asked to extract knowledge from this data without further instruction. We have already talked about one of the many forms that unsupervised learning comes in—**dimensionality reduction**. Another popular domain is **cluster analysis**, which aims to partition data into distinct groups of similar items.

Some of the problems where clustering techniques can be useful are document analysis, image retrieval, finding spam emails, identifying fake news, identifying criminal activities, and so on.

In this chapter, we want to understand how different clustering algorithms can be used to extract hidden structures in simple, unlabeled datasets. These hidden structures have many benefits, whether they are used in feature extraction, image processing, or even as a preprocessing step for supervised learning tasks. As a concrete example, we will learn how to apply clustering to images to reduce their color spaces to 16 bits.

More specifically, we will cover the following topics:

- **k-means clustering** and **expectation-maximization** and implementing these in OpenCV
- Arranging clustering algorithms in hierarchical trees and what are the benefits that come from that
- Using unsupervised learning for preprocessing, image processing, and classification

Let's get started!

Using TF-IDF to improve the result

It was called the **Term Frequency-Inverse Document Frequency (TF-IDF)**, and we encountered it in [Chapter 4, Representing Data and Engineering Features](#). If you recall, what TF-IDF does is basically weigh the word count by a measure of how often the words appear in the entire dataset. A useful side effect of this method is the IDF part —the inverse frequency of words. This makes sure that frequent words, such as *and*, *the*, and *but*, carry only a small weight in the classification.

We apply TF-IDF to the feature matrix by calling `fit_transform` on our existing feature matrix, `x`:

```
In [24]: tfidf = feature_extraction.text.TfidfTransformer()  
In [25]:  
X_new = tfidf.fit_transform(X)
```

Don't forget to split the data; also, ...

Summary

In this chapter, we took our first look at probability theory, learning about random variables and conditional probabilities, which allowed us to get a glimpse of Bayes' theorem—the underpinning of a Naive Bayes classifier. We talked about the difference between discrete and continuous random variables, likelihoods and probabilities, priors and evidence, and normal and Naive Bayes classifiers.

Finally, our theoretical knowledge would be of no use if we didn't apply it to practical examples. We obtained a dataset of raw email messages, parsed it, and trained Bayesian classifiers on it to classify emails as either spam or ham (not spam) using a variety of feature extraction approaches.

In the next chapter, we will switch gears and, for once, discuss what to do if we have to deal with unlabeled data.

Technical requirements

You can refer to the code for this chapter from the following link: <https://github.com/PacktPublishing/Machine-Learning-for-OpenCV-Second-Edition/tree/master/Chapter08>.

Here is a summary of the software and hardware requirements:

- You will need OpenCV version 4.1.x (4.1.0 or 4.1.1 will both work just fine).
- You will need Python version 3.6 (any Python version 3.x will be fine).
- You will need Anaconda Python 3 for installing Python and the required modules.
- You can use any operating system—macOS, Windows, and Linux-based OSes—along with this book. We recommend you have at least 4 GB RAM in your system.
- You don't need to have a GPU to run the code provided along with this book.

Understanding unsupervised learning

Unsupervised learning might come in many shapes and forms, but the goal is always to convert original data into a richer, more meaningful representation, whether that means making it easier for humans to understand or easier for machine learning algorithms to parse.

Some common applications of unsupervised learning include the following:

- **Dimensionality reduction:** This takes a high-dimensional representation of data consisting of many features and tries to compress the data so that its main characteristics can be explained with a small number of highly informative features. For example, when applied to housing prices in the neighborhoods of Boston, dimensionality reduction might be able to tell us that the indicators we should pay most attention to are the property tax and the neighborhood's crime rate.
- **Factor analysis:** This tries to find the hidden causes or unobserved components that gave rise to the observed data. For example, when applied to all of the episodes of the 1970s TV show, *Scooby-Doo, Where Are You!*, factor analysis might be able to tell us that (spoiler alert!) every ghost or monster on the show is essentially some disgruntled count playing an elaborate hoax on the town.

- **Cluster analysis:** This tries to partition the data into distinct groups of similar items. This is the type of unsupervised learning we will focus on in this chapter. For example, when applied to all of the movies on Netflix, cluster analysis might be able to automatically group them into genres.

To make things more complicated, these analyses have to be performed on unlabeled data, where we do not know beforehand what the right answer should be. Consequently, a major challenge in unsupervised learning is to determine whether an algorithm did well or learned anything useful. Often, the only way to evaluate the result of an unsupervised learning algorithm is to inspect it manually and determine by hand whether the result makes sense.

That being said, unsupervised learning can be immensely helpful, for example, as a preprocessing or feature extraction step. You can think of unsupervised learning as a **data transformation**—a way to transform data from its original representation into a more informative form. Learning a new representation might give us deeper insights into our data, and sometimes, it might even improve the accuracy of supervised learning algorithms.

Understanding k-means clustering

The essential clustering algorithm that OpenCV provides is k -means clustering, which searches for a predestined number of k -clusters (or groups) from an unlabeled multi-dimensional data.

It achieves this by using two simple hypotheses about what optimal clustering should look like:

- The center of each cluster is basically the mean of all of the points belonging to that cluster, also known as the centroid.
- Each data point in that cluster is closer to its center than to all other cluster centers.

It's easiest to understand the algorithm by looking at a concrete example.

Implementing our first k-means example

First, let's generate a 2D dataset containing four distinct blobs. To emphasize that this is an unsupervised approach, we will leave the labels out of the visualization:

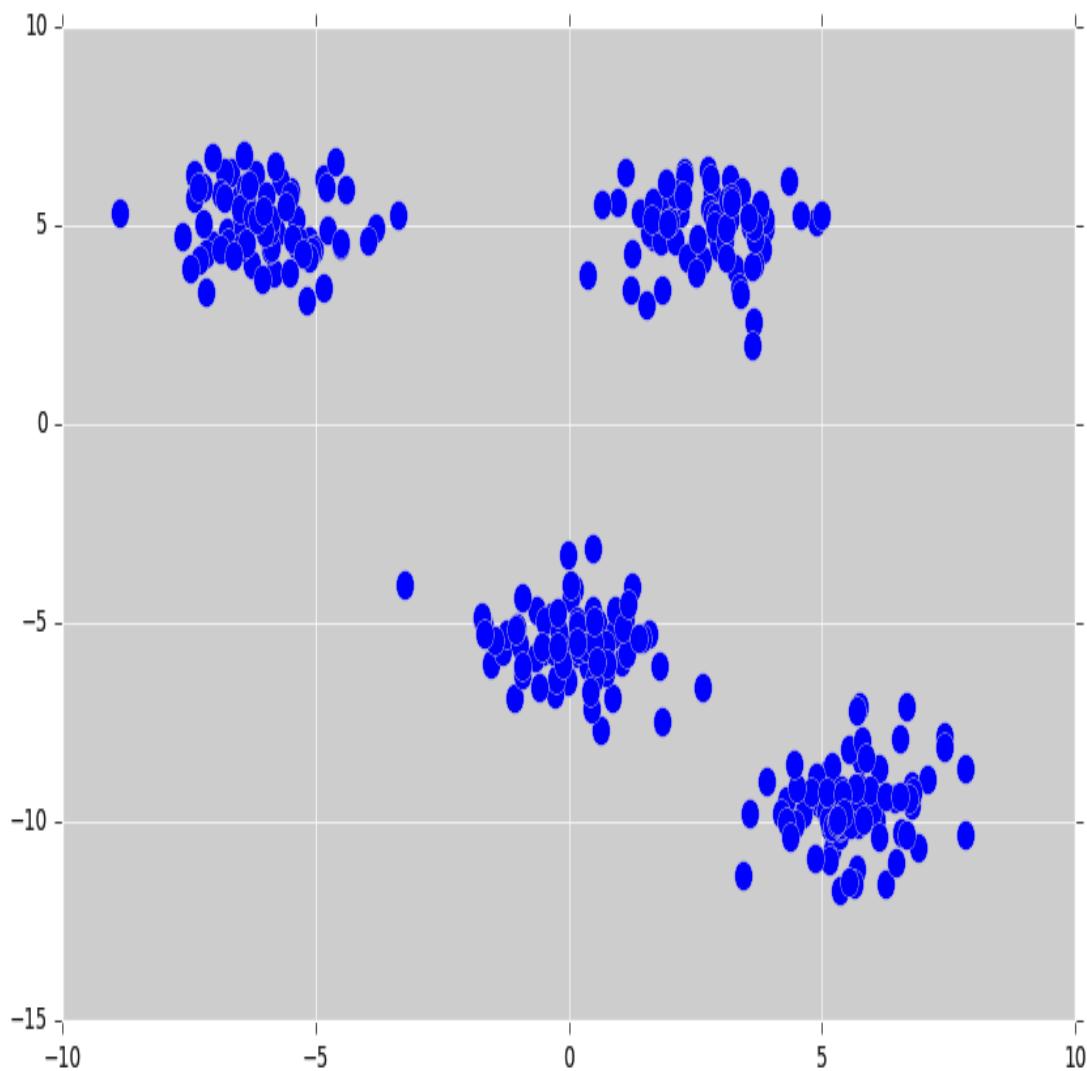
1. We will continue using `matplotlib` for all of our visualization purposes:

```
In [1]: import matplotlib.pyplot as plt  
...      %matplotlib inline  
...      plt.style.use('ggplot')
```

2. Following the same recipe from earlier chapters, we will create a total of 300 blobs (`n_samples=300`) belonging to four distinct clusters (`centers=4`):

```
In [2]: from sklearn.datasets.samples_generator import  
make_blobs  
...      X, y_true = make_blobs(n_samples=300, centers=4,  
...                           cluster_std=1.0,  
...                           random_state=10)  
...      plt.scatter(X[:, 0], X[:, 1], s=100);
```

This will generate the following diagram:



The preceding diagram shows an example dataset of 300 unlabeled points organized into four distinct clusters. Even without assigning target labels to the data, it is straightforward to pick out the four clusters by eye. The k -means algorithm can do this, too, without having any information about target labels or underlying data distributions.

3. Although k -means is, of course, a statistical model, in OpenCV, it does not come via the `ml` module and the common `train` and `predict` API calls. Instead, it is directly available as

`cv2.kmeans`. To use the model, we have to specify some arguments, such as the termination criteria and some initialization flags. Here, we will tell the algorithm to terminate whenever the error is smaller than 1.0 (`cv2.TERM_CRITERIA_EPS`) or when ten iterations have been executed (`cv2.TERM_CRITERIA_MAX_ITER`):

```
In [3]: import cv2
...
    criteria = (cv2.TERM_CRITERIA_EPS +
cv2.TERM_CRITERIA_MAX_ITER,
...
            10, 1.0)
...
    flags = cv2.KMEANS_RANDOM_CENTERS
```

4. Then, we can pass the preceding data matrix (`x`) to `cv2.means`. We also specify the number of clusters (`4`) and the number of attempts the algorithm should make with different random initial guesses (`10`), as shown in the following snippet:

```
In [4]: import numpy as np
...
    compactness, labels, centers =
cv2.kmeans(X.astype(np.float32),
...
           4, None,
criteria,
...
           10, flags)
```

Three different variables are returned.

5. The first one, `compactness`, returns the sum of squared distances from each point to their corresponding cluster centers. A high compactness score indicates that all points are close to their

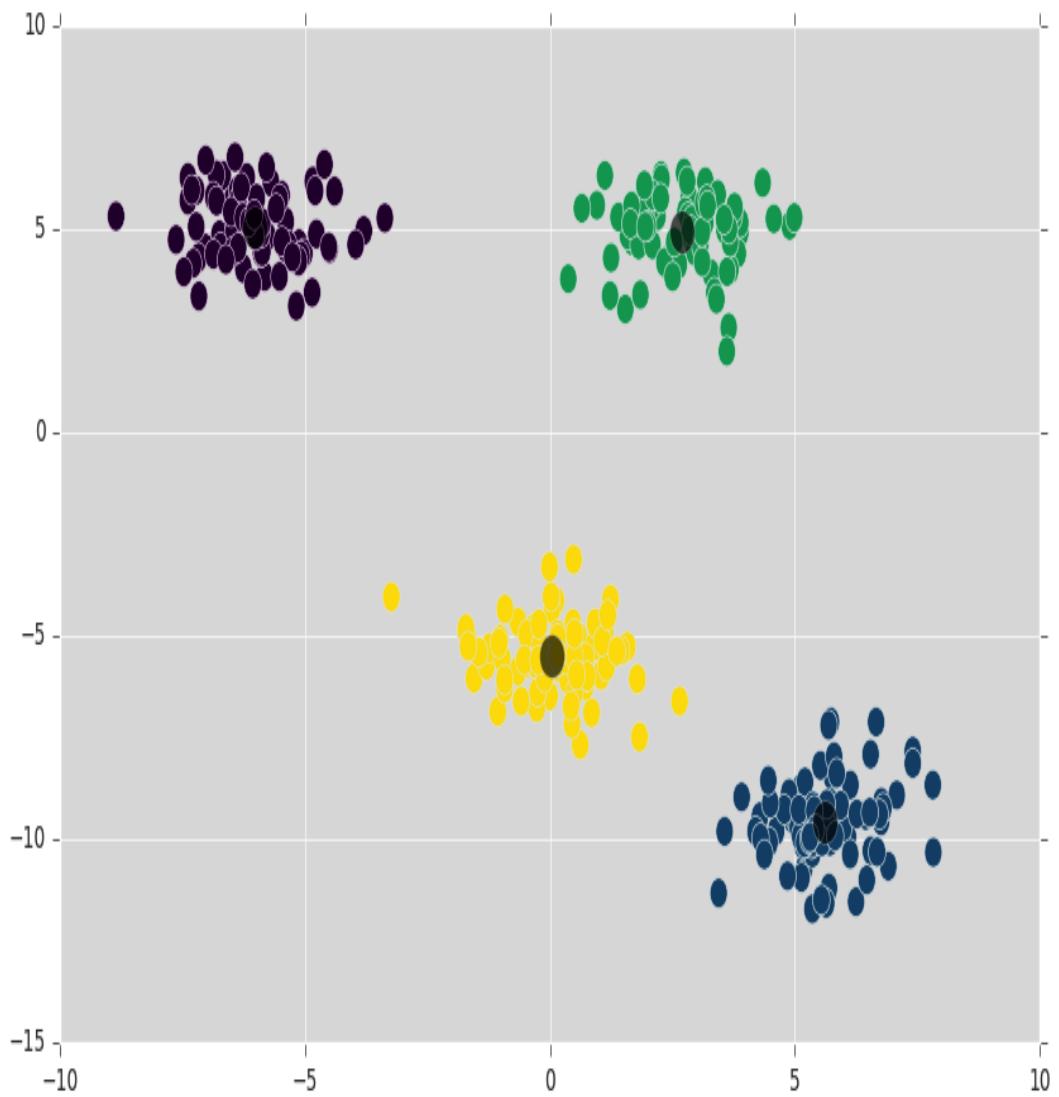
cluster centers, whereas a low compactness score indicates that the different clusters might not be well separated:

```
In [5]: compactness  
Out[5]: 527.01581170992
```

6. Of course, this number strongly depends on the actual values in x . If the distances between points were large, to begin with, we could not expect an arbitrarily small compactness score. Hence, it is more informative to plot the data points, colored to their assigned cluster labels:

```
In [6]: plt.scatter(x[:, 0], x[:, 1], c=labels, s=50,  
cmap='viridis')  
...     plt.scatter(centers[:, 0], centers[:, 1], c='black',  
s=200,  
...                         alpha=0.5);
```

7. This produces a scatter plot of all of the data points colored according to whichever cluster they belong to, with the corresponding cluster centers indicated with a blob of a darker shade in the center of each cluster:



The preceding diagram shows the result of k -means clustering for $k=4$. The good news here is that the k -means algorithm (at least, in this simple case) assigns the points to clusters very similarly to how we might have, had we done the job by eye. But how did the algorithm find these different clusters so quickly? After all, the number of possible combinations of cluster assignments is exponential to the number of data points! By hand, trying all possible combinations would have certainly taken forever.

Fortunately, an exhaustive search is not necessary. Instead, the typical approach that k -means takes is to use an iterative algorithm, also known as **expectation-maximization**.

Understanding expectation-maximization

k-means clustering is but one concrete application of a more general algorithm known as expectation-maximization. In short, the algorithm works as follows:

1. Start with some random cluster centers.
2. Repeat until convergence:
 1. **Expectation step:** Assign all data points to their nearest cluster center.
 2. **Maximization step:** Update the cluster centers by taking the mean of all of the points in the cluster.

Here, the expectation step is so named because it involves updating our expectation of which cluster each point in the dataset belongs to. The maximization step is so named because it involves maximizing a fitness function that defines the location of the cluster centers. In the case of *k*-means, maximization ...

Implementing our expectation-maximization solution

The expectation-maximization algorithm is simple enough for us to code it ourselves. To do so, we will define a function, `find_clusters(x, n_clusters, rseed=5)`, that takes as input a data matrix (`x`), the number of clusters we want to discover (`n_clusters`), and a random seed (optional, `rseed`). As will become clear in a second, scikit-learn's `pairwise_distances_argmin` function will come in handy:

```
In [7]: from sklearn.metrics import pairwise_distances_argmin
...     def find_clusters(X, n_clusters, rseed=5):
```

We can implement expectation-maximization for k -means in five essential steps:

1. **Initialization:** Randomly choose a number of cluster centers, `n_clusters`. We don't just pick any random number but instead pick actual data points to be the cluster centers. We do this by permuting `x` along its first axis and picking the first `n_clusters` points in this random permutation:

```
...         rng = np.random.RandomState(rseed)
...         i = rng.permutation(X.shape[0])[:n_clusters]
centers = X[i]
```

2. **while looping forever:** Assign labels based on the closest cluster centers. Here, scikit-learn's `pairwise_distance_argmin` function does exactly what we want. It computes, for each data point in `x`, the index of the closest cluster center in `centers`:

```
...           while True:  
...             labels = pairwise_distances_argmin(X, centers)
```

3. **Find the new cluster centers:** In this step, we have to take the arithmetic mean of all data points in `x` that belong to a specific cluster (`x[labels == i]`):

```
...           new_centers = np.array([X[labels ==  
i].mean(axis=0)
```

4. **Check for convergence and break the `while` loop if necessary:** This is the last step to make sure that we stop the execution of the algorithm once the job is done. We determine whether the job is done by checking whether all of the new cluster centers are equal to the old cluster centers. If this is true, we exit the loop; otherwise, we keep looping:

```
...           for i in range(n_clusters))  
...             if np.all(centers == new_centers):  
...               break  
...             centers = new_centers
```

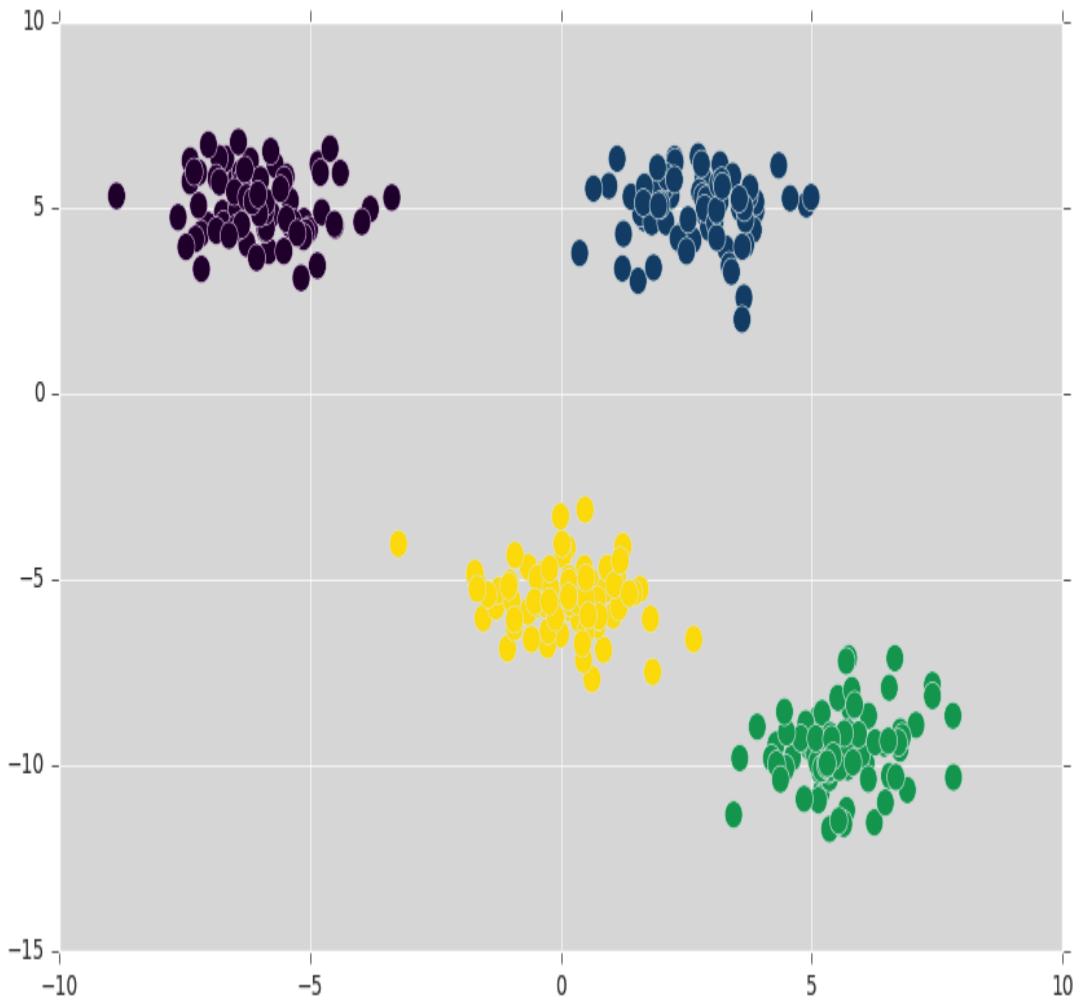
5. Exit the function and return the result:

```
...             return centers, labels
```

We can apply our function to the preceding data matrix, x , we created. Since we know what the data looks like, we know that we are looking for four clusters:

```
In [8]: centers, labels = find_clusters(X, 4)
...     plt.scatter(X[:, 0], X[:, 1], c=labels, s=100, cmap='viridis');
```

This will generate the following plot. The vital point to observe from the following diagram is that, before applying k -means clustering, all data points were categorized to the same one color; however, after using k -means clustering, each color is a different cluster (similar data points are clustered or grouped in one color) :



The preceding diagram shows the outcome of our home-made k -means using expectation-maximization. As we can see, our home-made algorithm got the job done! Granted, this particular clustering example was fairly easy, and most real-life implementations of k -means clustering will do a bit more under the hood. But for now, we are happy.

Knowing the limitations of expectation-maximization

For all its simplicity, expectation-maximization performs incredibly well in a range of scenarios. That being said, there are some potential limitations that we need to be aware of:

- Expectation-maximization does not guarantee that we will find the globally best solution.
- We must know the number of desired clusters beforehand.
- The decision boundaries of the algorithms are linear.
- The algorithm is slow for large datasets.

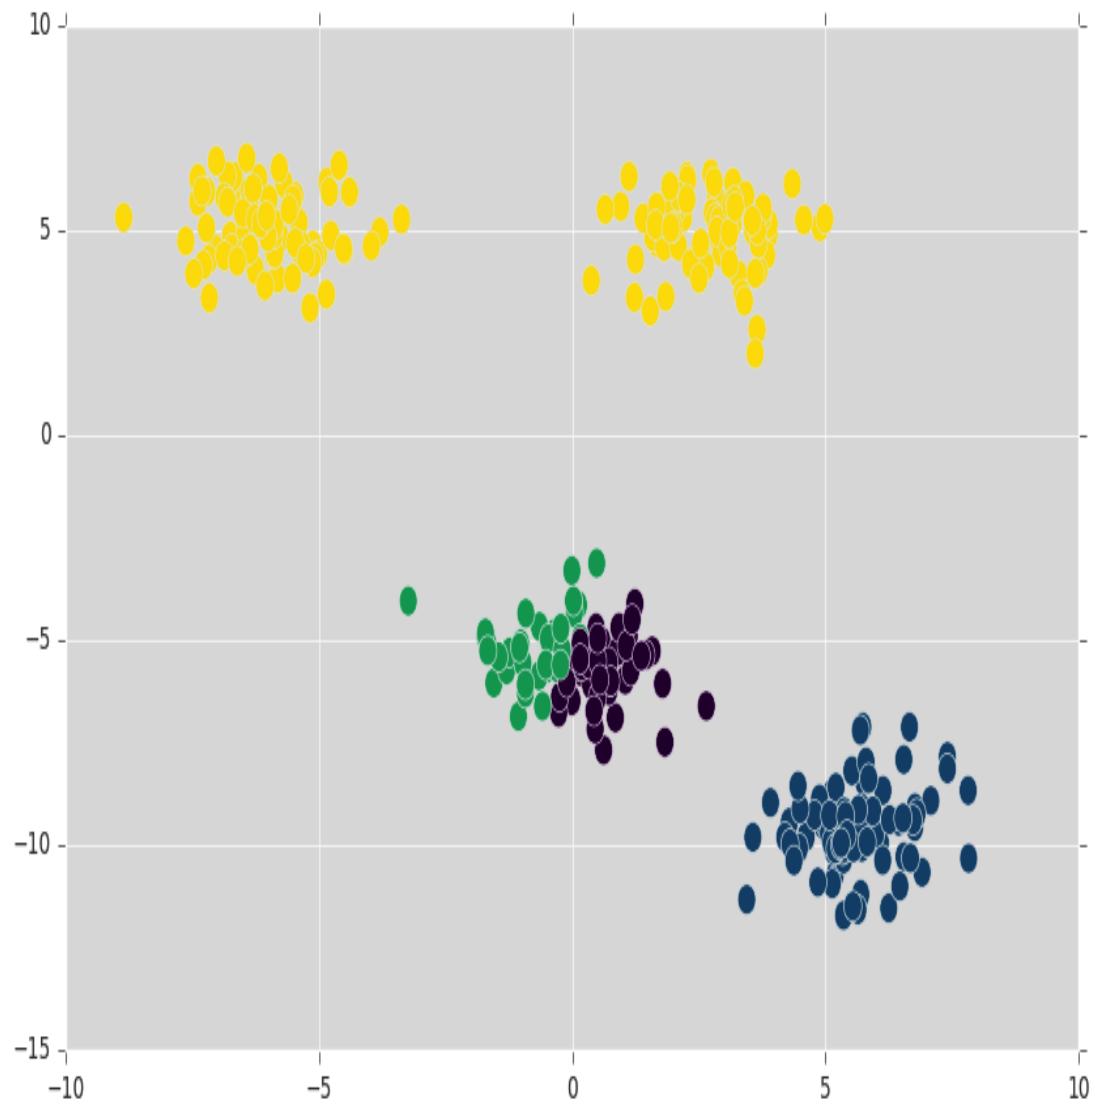
Let's quickly discuss these potential caveats in a little more detail.

The first caveat – no guarantee of finding the global optimum

Although mathematicians have proved that the expectation-maximization step improves the result in each step, there is still no guarantee that, in the end, we will find the global best solution. For example, if we use a different random seed in our simple example (such as using seed 10 instead of 5), we suddenly get very poor results:

```
In [9]: centers, labels = find_clusters(X, 4, rseed=10)
...     plt.scatter(X[:, 0], X[:, 1], c=labels, s=100, cmap='viridis');
```

This will generate the following diagram:



The preceding diagram shows an example of k -means missing the global optimum. What happened?

The short answer is that the random initialization of cluster centers was unfortunate. It led to the center of the yellow cluster migrating in-between the two top blobs, essentially combining them into one. As a result, the other clusters got confused because they suddenly had to split two visually distinct blobs into three clusters.

For this reason, it is common for the algorithm to be run for multiple initial states. Indeed, OpenCV does this by default (set by the optional `attempts` parameter).

The second caveat – we must select the number of clusters beforehand

Another potential limitation is that k -means cannot learn the number of clusters from the data. Instead, we must tell it how many clusters we expect beforehand. You can see how this could be problematic for complicated real-world data that you don't fully understand yet.

From the viewpoint of k -means, there is no wrong or nonsensical number of clusters. For example, if we ask the algorithm to identify six clusters in the dataset generated in the preceding section, it will happily proceed and find the best six clusters:

```
In [10]: criteria = (cv2.TERM_CRITERIA_EPS +
cv2.TERM_CRITERIA_MAX_ITER,...           10, 1.0)...      flags
= cv2.KMEANS_RANDOM_CENTERS... compactness, labels, centers ...
```

The third caveat – cluster boundaries are linear

The k -means algorithm is based on a simple assumption, which is that points will be closer to their own cluster center than to others. Consequently, k -means always assumes linear boundaries between clusters, meaning that it will fail whenever the geometry of the clusters is more complicated than that.

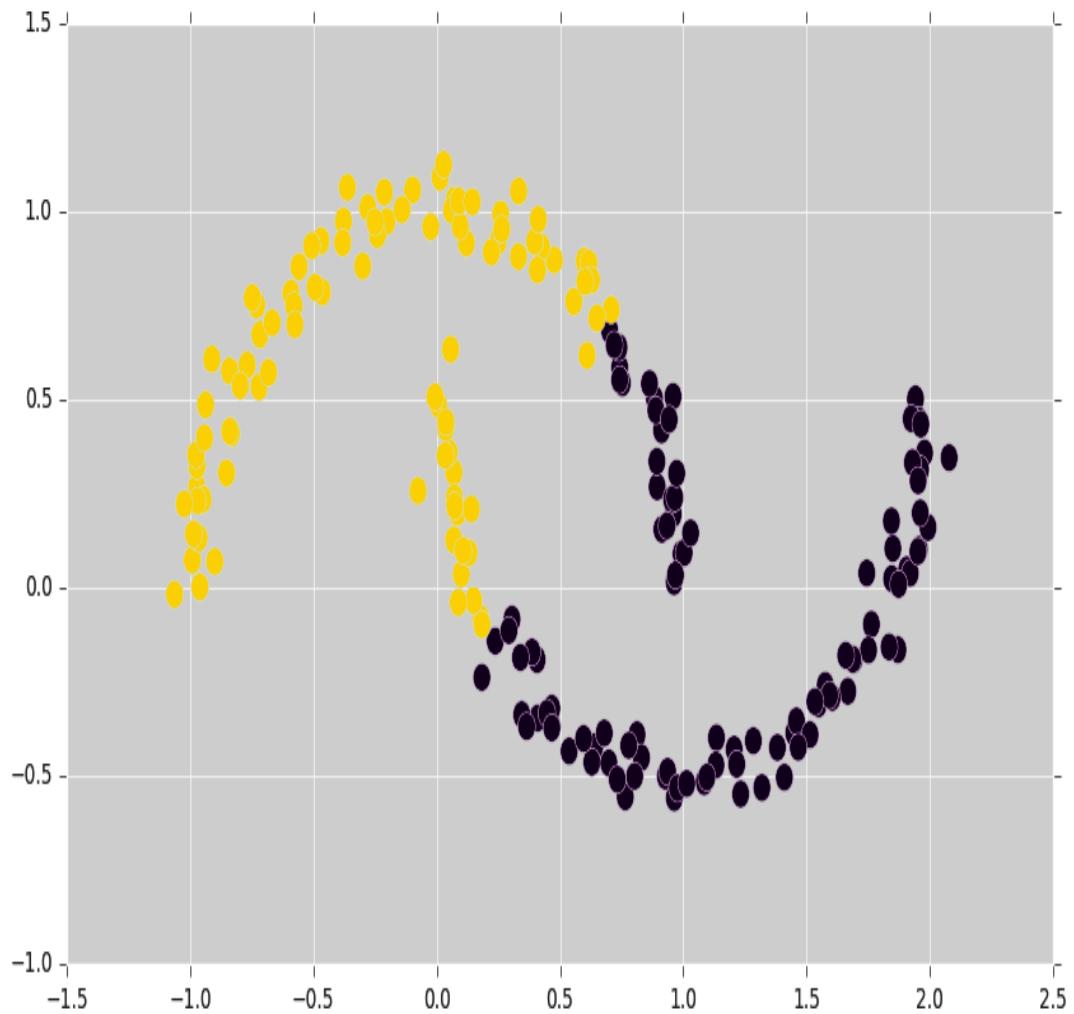
We see this limitation for ourselves by generating a slightly more complicated dataset. Instead of generating data points from Gaussian blobs, we want to organize the data into two overlapping half circles. We can do this using scikit-learn's `make_moons`. Here, we choose 200 data points belonging to two half circles, in combination with some Gaussian noise:

```
In [14]: from sklearn.datasets import make_moons  
...       X, y = make_moons(200, noise=.05, random_state=12)
```

This time, we tell k -means to look for two clusters:

```
In [15]: criteria = (cv2.TERM_CRITERIA_EPS +  
...           cv2.TERM_CRITERIA_MAX_ITER, 10, 1.0)  
...       flags = cv2.KMEANS_RANDOM_CENTERS  
...       compactness, labels, centers =  
cv2.kmeans(X.astype(np.float32),  
...           2, None, criteria,  
...           10, flags)  
...       plt.scatter(X[:, 0], X[:, 1], c=labels, s=100,  
cmap='viridis');
```

The resulting scatter plot looks like this diagram:



The preceding diagram shows an example of k -means finding linear boundaries in nonlinear data. As is evident from the plot, k -means failed to identify the two half circles and instead split the data with what looks like a diagonal straight line (from bottom-left to top-right).

This scenario should ring a bell. We had the same problem when we talked about linear SVMs in [Chapter 6, Detecting Pedestrians with Support Vector Machines](#). The idea there was to use the kernel trick to transform the data into a higher-dimensional feature space. Can we do the same here?

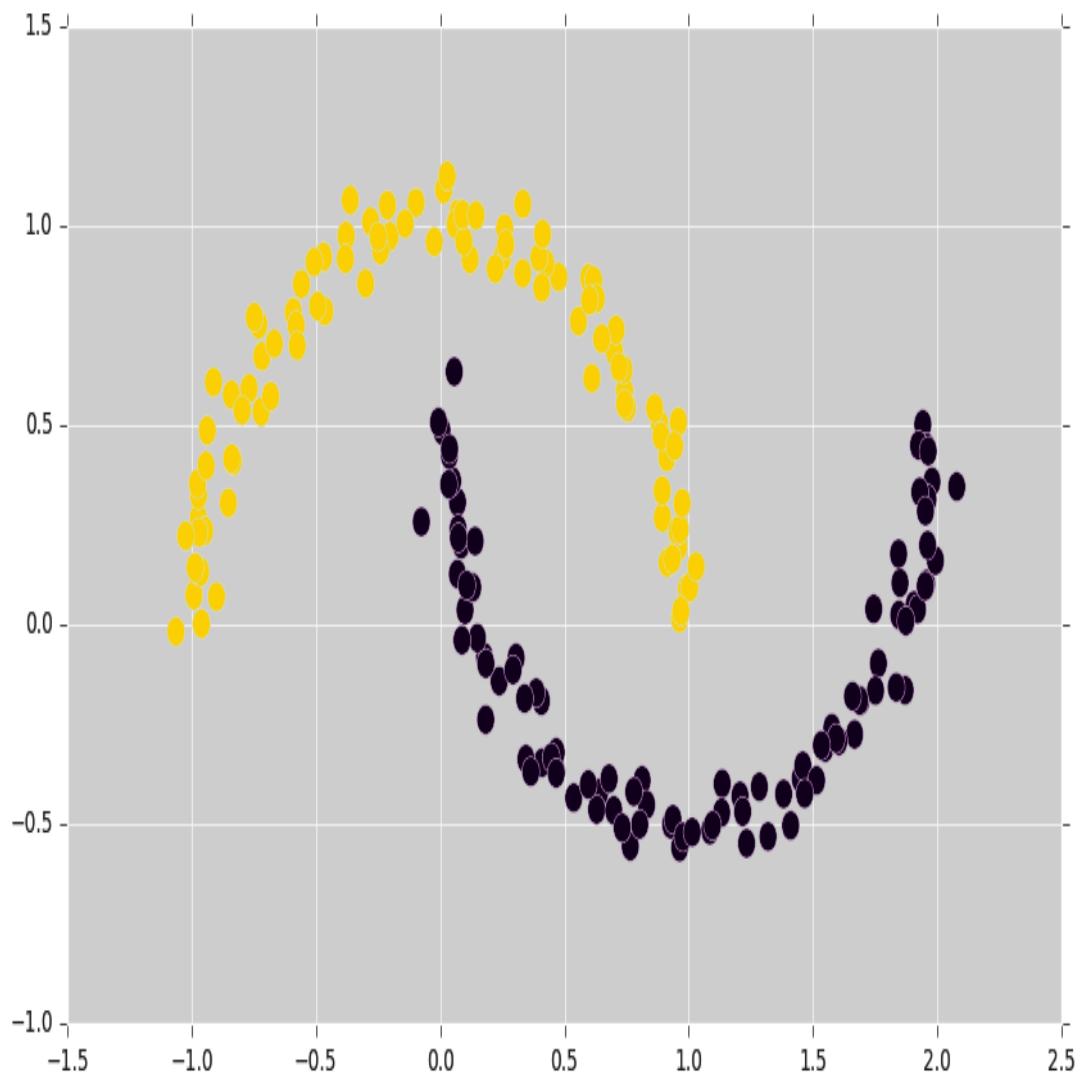
We most certainly can. There is a form of kernelized k -means that works akin to the kernel trick for SVMs, called **spectral clustering**. Unfortunately, OpenCV does not provide an implementation of spectral clustering. Fortunately, scikit-learn does:

```
In [16]: from sklearn.cluster import SpectralClustering
```

The algorithm uses the same API as all other statistical models: we set optional arguments in the constructor and then call `fit_predict` on the data. Here, we want to use the graph of the nearest neighbors to compute a higher-dimensional representation of the data and then assign labels using k -means:

```
In [17]: model = SpectralClustering(n_clusters=2,
...                                 affinity='nearest_neighbors',
...                                 assign_labels='kmeans')
...
...     labels = model.fit_predict(X)
...     plt.scatter(X[:, 0], X[:, 1], c=labels, s=100,
cmap='viridis');
```

The output of spectral clustering looks like this:



We see that spectral clustering gets the job done. Alternatively, we could have transformed the data into a more suitable representation ourselves and then applied OpenCV's linear k -means to it. The lesson of all of this is that, perhaps, again, feature engineering saved the day.

The fourth caveat – k-means is slow for a large number of samples

The final limitation of k -means is that it is relatively slow for large datasets. You can imagine that quite a lot of algorithms might suffer from this problem. However, k -means is affected especially badly: each iteration of k -means must access every single data point in the dataset and compare it to all of the cluster centers.

You might wonder whether the requirement to access all data points during each iteration is really necessary. For example, you might just use a subset of the data to update the cluster centers at each step. Indeed, this is the exact idea that underlies a variation of the algorithm called **batch-based k -means**. Unfortunately, this algorithm is not implemented ...

Compressing color spaces using k-means

One interesting use case of k -means is the compression of image color spaces. For example, a standard **color image** comes with 24-bit color depth, providing for a total of 16,777,216 color varieties. However, in most images, a large number of colors will be unused, and many of the pixels in the image will have similar values. The compressed image can then be sent over the internet at a faster speed, and at the receiver end, it can be decompressed to get back the original image. Hence, reducing the storage and transmission cost. However, the image color space compression will be **lossy** and you might not notice fine details in the image after compression.

Alternatively, we can use k -means to reduce the color palette too. The idea here is to think of the cluster centers as the decreased color palette. Then, k -means will organize the millions of colors in the original image into the appropriate number of colors.

Visualizing the true-color palette

By performing the following steps, you will be able to visualize the true-color palette of a color image:

1. Let's have a look at a particular image:

```
In [1]: import cv2...      import numpy as np...      lena =
cv2.imread('data/lena.jpg', cv2.IMREAD_COLOR)
```

2. By now, we know how to start Matplotlib in our sleep:

```
In [2]: import matplotlib.pyplot as plt...      %matplotlib
inline...      plt.style.use('ggplot')
```

3. However, this time, we want to disable the grid lines that the `ggplot` option typically displays over images:

```
In [3]: plt.rc('axes', **{'grid': False})
```

4. Then, we can visualize Lena with the following command (don't forget to switch the BGR ordering of the color channels to RGB):

```
In [4]: plt.imshow(cv2.cvtColor(lena, cv2.COLOR_BGR2RGB)) ...
```

Reducing the color palette using k-means

By referring to the following steps, you will be able to project a color image into a reduced color palette using *k*-means clustering:

1. Now, let's reduce the 16 million colors to a mere 16 by indicating *k*-means to cluster all of the 16 million color variations into 16 distinct clusters. We will use the previously mentioned procedure, but now define 16 as the number of clusters:

```
In [9]: criteria = (cv2.TERM_CRITERIA_EPS +
cv2.TERM_CRITERIA_MAX_ITER,
...                      10, 1.0)
...     flags = cv2.KMEANS_RANDOM_CENTERS
...     img_data = img_data.astype(np.float32)
...     compactness, labels, centers = cv2.kmeans(img_data,
...                                              16, None,
...     criteria,
...                                              10, flags)
```

2. The 16 different colors of your reduced color palette correspond to the resultant clusters. The output from the `centers` array reveals that all colors have three entries—`B`, `G`, and `R`—with values between 0 and 1:

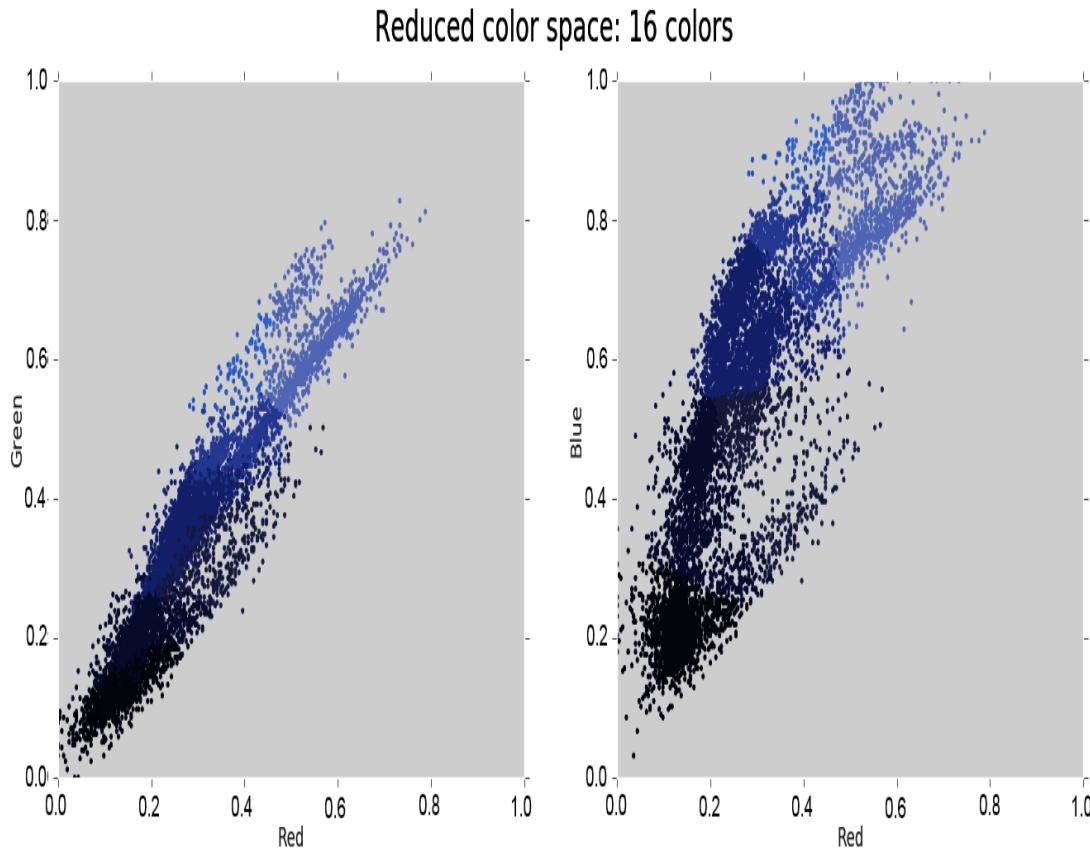
```
In [10]: centers
Out[10]: array([[ 0.29973754,  0.31500012,  0.48251548],
       [ 0.27192295,  0.35615689,  0.64276862],
       [ 0.17865284,  0.20933454,  0.41286203],
       [ 0.39422086,  0.62827665,  0.94220853],
       [ 0.34117648,  0.58823532,  0.90196079],
       [ 0.42996961,  0.62061119,  0.91163337],
       [ 0.06039202,  0.07102439,  0.1840712 ],
       [ 0.5589878 ,  0.6313886 ,  0.83993536],
       [ 0.37320262,  0.54575169,  0.88888896],
       [ 0.35686275,  0.57385623,  0.88954246],
       [ 0.47058824,  0.48235294,  0.59215689],
       [ 0.34346411,  0.57483661,  0.88627452],
       [ 0.13815609,  0.12984112,  0.21053818],
       [ 0.3752504 ,  0.47029912,  0.75687987],
       [ 0.31909946,  0.54829341,  0.87378371],
       [ 0.40409693,  0.58062142,  0.85475557 ]],  
dtype=float32)
```

3. The `labels` vector contains the 16 colors corresponding to the 16 cluster `labels`. So, all of the data points with the label 0 will be colored according to row 0 in the `centers` array; similarly, all data points with the label 1 will be colored according to row 1 in the `centers` array and so on. Hence, we want to use `labels` as an index in the `centers` array—these are our new colors:

```
In [11]: new_colors = centers[labels].reshape((-1, 3))
```

4. We can plot the data again, but this time, we will use `new_colors` to color the data points accordingly:

The result is the recoloring of the original pixels, where each pixel is assigned the color of its closest cluster center:



5. To observe the effect of recoloring, we have to plot

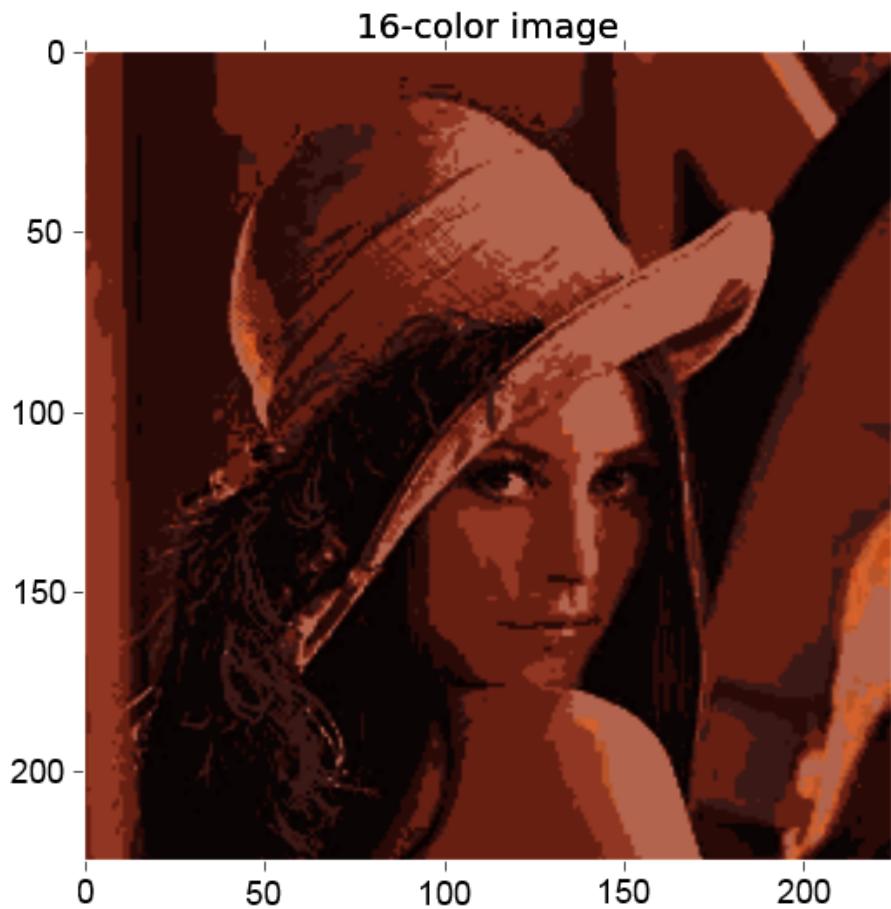
`new_colors` as an image. We flattened the earlier image to get from the image to the data matrix. To get back to the image now, we need to do the inverse, which is reshape `new_colors` according to the shape of the Lena image:

```
In [13]: lena_recolored = new_colors.reshape(lena.shape)
```

6. Then, we can visualize the recolored Lena image like any other image:

```
In [14]: plt.figure(figsize=(10, 6))
...
plt.imshow(cv2.cvtColor(lena_recolored,
cv2.COLOR_BGR2RGB));
...
plt.title('16-color image')
```

The result looks like this:



It's pretty awesome, right?

Overall, the preceding screenshot is quite clearly recognizable although some details are arguably lost. Given that you compressed the image by a factor of around 1 million, this is pretty remarkable.

You can repeat this procedure for any desired number of colors.

*Another way to reduce the color palette of images involves the use of **bilateral filters**. The resulting images often look like cartoon versions of the original image. You can find an example of this in the book, OpenCV with Python Blueprints, by M. Beyeler, Packt Publishing.*

Another potential application of k -means is something you might not expect: using it for image classification.

Classifying handwritten digits using k-means

Although the last application was a pretty creative use of k -means, we can do better still. We have previously discussed k -means in the context of unsupervised learning, where we tried to discover some hidden structure in the data.

However, doesn't the same concept apply to most classification tasks? Let's say our task was to classify handwritten digits. Don't most zeros look similar, if not the same? And don't all zeros look categorically different from all possible ones? Isn't this exactly the kind of *hidden structure* we set out to discover with unsupervised learning? Doesn't this mean we could use clustering for classification as well?

Let's find out together. In this section, we will attempt ...

Loading the dataset

From the earlier chapters, you might recall that scikit-learn provides a whole range of handwritten digits via its `load_digits` utility function. The dataset consists of 1,797 samples with 64 features each, where each of the features has the brightness of one pixel in an 8×8 image:

```
In [1]: from sklearn.datasets import load_digits  
...      digits = load_digits()  
...      digits.data.shape  
Out[1]: (1797, 64)
```

Running k-means

Setting up k -means works exactly the same as in the previous examples. We tell the algorithm to perform at most 10 iterations and stop the process if our prediction of the cluster centers does not improve within a distance of 1.0:

```
In [2]: import cv2...      criteria = (cv2.TERM_CRITERIA_EPS +
cv2.TERM_CRITERIA_MAX_ITER, ...                      10, 1.0)...      flags =
cv2.KMEANS_RANDOM_CENTERS
```

Then, we apply k -means to the data as we did before. Since there are 10 different digits (0-9), we tell the algorithm to look for 10 distinct clusters:

```
In [3]: import numpy as np...      digits.data =
digits.data.astype(np.float32)...      compactness, clusters, centers =
cv2.kmeans(digits.data, 10, None,... criteria, 10, flags)
```

And we're done!

Similar to the $N \times 3$ matrices ...

Organizing clusters as a hierarchical tree

An alternative to k -means is **hierarchical clustering**. One advantage of hierarchical clustering is that it allows us to organize the different clusters in a hierarchy (also known as a **dendrogram**), which can make it easier to interpret the results. Another useful advantage is that we do not need to specify the number of clusters upfront.

Understanding hierarchical clustering

There are two approaches to hierarchical clustering:

- In **agglomerative hierarchical clustering**, we start with each data point potentially being its own cluster, and we subsequently merge the closest pair of clusters until only one cluster remains.
- In **divisive hierarchical clustering**, it's the other way around; we start by assigning all of the data points to the same cluster, and we subsequently split the cluster into smaller clusters until each cluster only contains one sample.

Of course, we can specify the number of desired clusters if we wish to. In the following screenshot, we asked the algorithm to find a total of three clusters:

The preceding screenshot shows a step-by-step example of agglomerative ...

Implementing agglomerative hierarchical clustering

Although OpenCV does not provide an implementation of agglomerative hierarchical clustering, it is a popular algorithm that should, by all means, belong to our machine learning skill set:

1. We start by generating 10 random data points, just like in the previous screenshot:

```
In [1]: from sklearn.datasets import make_blobs  
...      X, y = make_blobs(random_state=100, n_samples=10)
```

2. Using the familiar statistical modeling API, we import the `AgglomerativeClustering` algorithm and specify the desired number of clusters:

```
In [2]: from sklearn import cluster  
...      agg = cluster.AgglomerativeClustering(n_clusters=3)
```

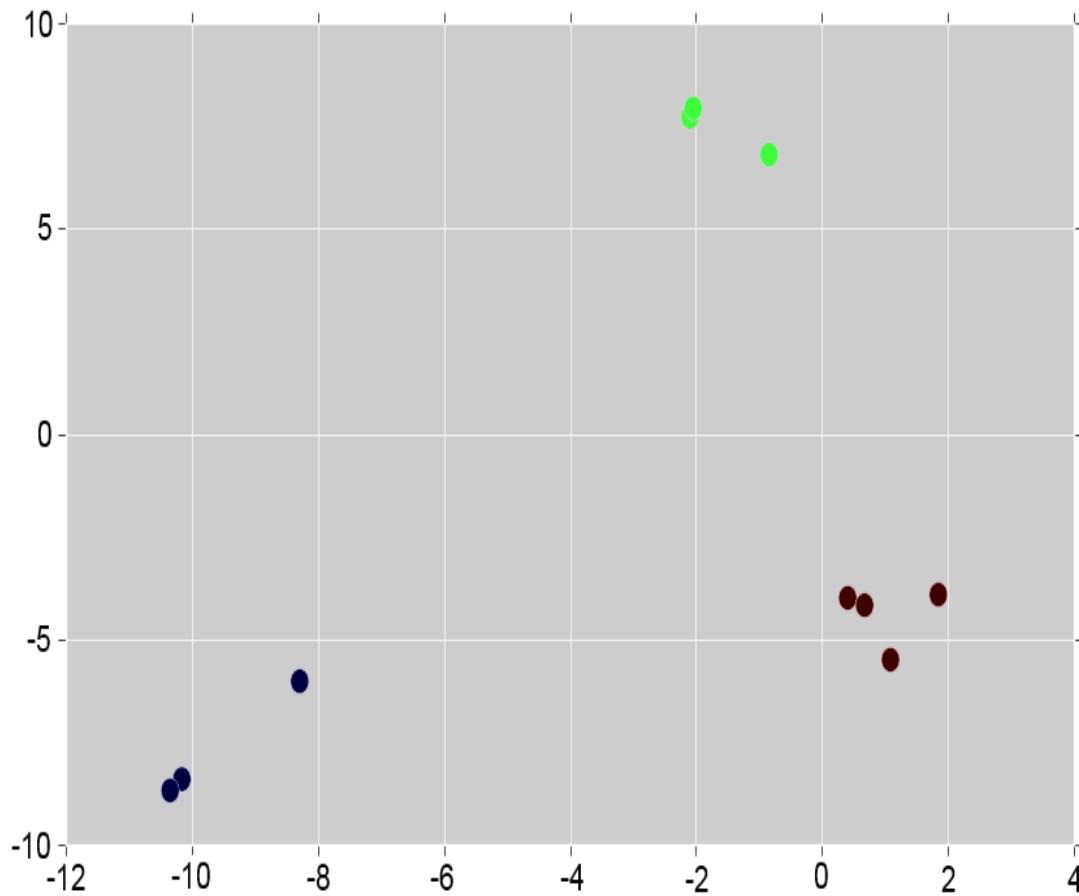
3. Fitting the model to the data works, as usual, via the `fit_predict` method:

```
In [3]: labels = agg.fit_predict(X)
```

4. We can generate a scatter plot where every data point is colored according to the predicted label:

```
In [4]: import matplotlib.pyplot as plt  
... %matplotlib inline  
... plt.style.use('ggplot')  
... plt.scatter(X[:, 0], X[:, 1], c=labels, s=100)
```

The resulting clustering is equivalent to the following diagram:



Finally, before we end this chapter, let's look at how to compare clustering algorithms and choose the correct clustering algorithm for the data you have!

Comparing clustering algorithms

There are around thirteen different clustering algorithms in the `sklearn` library. Having thirteen different sets of choices, the question is: what clustering algorithms should you use? The answer is your data. What type of data you have and which clustering you would like to apply on it is how you will select the algorithm. Having said that, there can be many possible algorithms that could be useful for the kind of problem and data you have. Each of the thirteen classes in `sklearn` is specialized for specific tasks (such as co-clustering and bi-clustering or clustering features instead of data points). An algorithm specializing in text clustering would be the right choice for clustering text data. Hence, if ...

Summary

In this chapter, we talked about some unsupervised learning algorithms, including k -means, spherical clustering, and agglomerative hierarchical clustering. We saw that k -means is just a specific application of the more general expectation-maximization algorithm, and we discussed its potential limitations. Furthermore, we applied k -means to two specific applications, which were reducing the color palette of images and classifying handwritten digits.

In the next chapter, we will move back into the world of supervised learning and talk about some of the most powerful current machine learning algorithms: neural networks and deep learning.

Section 3: Advanced Machine Learning with OpenCV

The final section of this book will cover important and advanced topics such as deep learning, ensemble machine learning methods, and hyperparameter tuning. We will also cover the latest addition to OpenCV—OpenVINO toolkit by Intel. We will briefly introduce OpenVINO, how to install it, and what are its various components, and then finally see how it can be used with OpenCV for image classification problems.

This section includes the following chapters:

- [Chapter 9, Using Deep Learning to Classify Handwritten Digits](#)
- [Chapter 10, Ensemble Methods for Classification](#)
- [Chapter 11, Selecting the Right Model with Hyperparameter Tuning](#)
- [Chapter 12, Using OpenVINO with OpenCV](#)
- [Chapter 13, Conclusion](#)

Using Deep Learning to Classify Handwritten Digits

Let's now return to supervised learning and discuss a family of algorithms known as **artificial neural networks**. Early studies of neural networks go back to the 1940s when Warren McCulloch and Walter Pitts first described how biological nerve cells (or neurons) in the brain might work. More recently, artificial neural networks have seen a revival under the buzzword deep learning, which powers state-of-the-art technologies such as Google's DeepMind and Facebook's DeepFace algorithms.

In this chapter, we want to wrap our heads around some simple versions of artificial neural networks, such as the McCulloch-Pitts neuron, the perceptron, and the multilayer perceptron. Once we have familiarized ourselves with the basics, we will be ready to implement a more sophisticated deep neural network to classify handwritten digits from the popular **MNIST database** (short for **Mixed National Institute of Standards and Technology database**). For this, we will be making use of Keras, a high-level neural network library, which is also frequently used by researchers and tech companies.

Along the way, we will address the following topics:

- Implementing perceptrons and multilayer perceptrons in OpenCV
- Differentiating stochastic and batch gradient descent, and how they fit in with backpropagation

- Finding the size of your neural network
- Using Keras to build sophisticated deep neural networks

Excited? Then let's go!

Technical requirements

You can refer the code for this chapter at the following link: <https://github.com/PacktPublishing/Machine-Learning-for-OpenCV-Second-Edition/tree/master/Chapter09>.

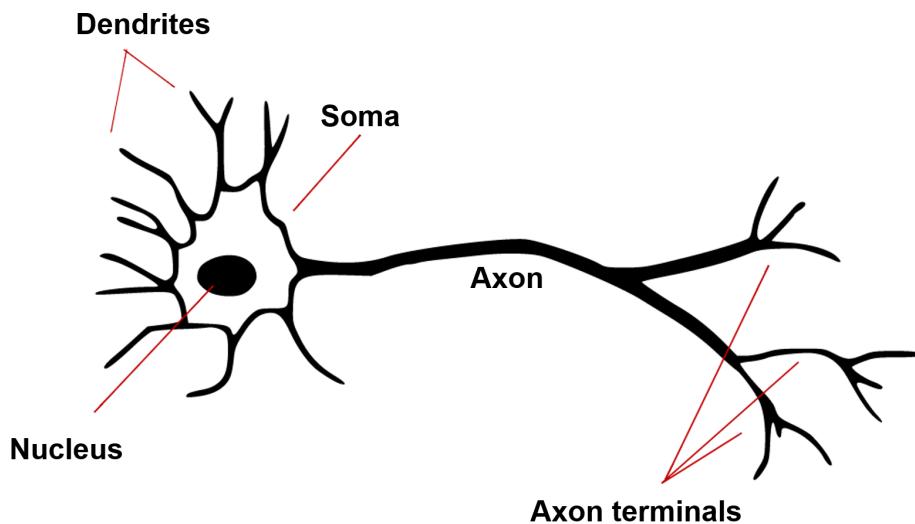
Here is a short summary of the software and hardware requirements:

- OpenCV version 4.1.x (4.1.0 or 4.1.1 will both work just fine).
- Python version 3.6 (any Python version 3.x will be fine).
- Anaconda Python 3 for installing Python and the required modules.
- You can use any operating system—macOS, Windows, or Linux-based—with this book. We recommend you have at least 4 GB RAM in your system.
- You don't need to have a GPU to run the code provided with the book.

Understanding the McCulloch-Pitts neuron

In 1943, Warren McCulloch and Walter Pitts published a mathematical description of neurons as they were believed to operate in the brain. A neuron receives input from other neurons through connections on its dendritic tree, which are integrated to produce an output at the cell body (or soma). The output is then communicated to other neurons via a long wire (or axon), which eventually branches out to make one or more connections (at axon terminals) on the dendritic tree of other neurons.

An example neuron is shown in the following diagram:



McCulloch and Pitts described the inner workings of such a neuron as a simple logic gate that would be either on or off, depending on the input it received on its dendritic tree. Specifically, the neuron would

sum up all of its inputs, and if the sum exceeded a certain threshold, an output signal would be generated and passed on by the axon.

However, today we know that real neurons are much more complicated than that. Biological neurons perform intricate nonlinear mathematical operations on thousands of inputs and can change their responsiveness dynamically depending on the context, importance, or novelty of the input signal. You can think of real neurons being as complex as computers and of the human brain being as complex as the internet.

Let's consider a simple artificial neuron that receives exactly two inputs, x_0 and x_1 . The job of the artificial neuron is to calculate a sum of the two inputs (usually in the form of a weighted sum), and if this sum exceeds a certain threshold (often zero), the neuron will be considered active and output a one; else it will be considered silent and output a minus one (or zero). In more mathematical terms, the output, y , of this McCulloch-Pitts neuron can be described as follows:

$$y = \begin{cases} 1 & \text{if } x_0w_0 + x_1w_1 \geq 0 \\ -1 & \text{otherwise} \end{cases}$$

In the preceding equation, w_0 and w_1 are weight coefficients, which, together with x_0 and x_1 , make up the weighted sum. In textbooks, the two different scenarios where the output, y , is either +1 and -1 would often be masked by an activation function, ϕ , which could take on two different values:

$$\begin{aligned} y &= \phi(x_0w_0 + x_1w_1) = \phi(z) \\ \phi(z) &= \begin{cases} 1 & \text{if } z \geq \theta \\ -1 & \text{otherwise} \end{cases} \end{aligned}$$

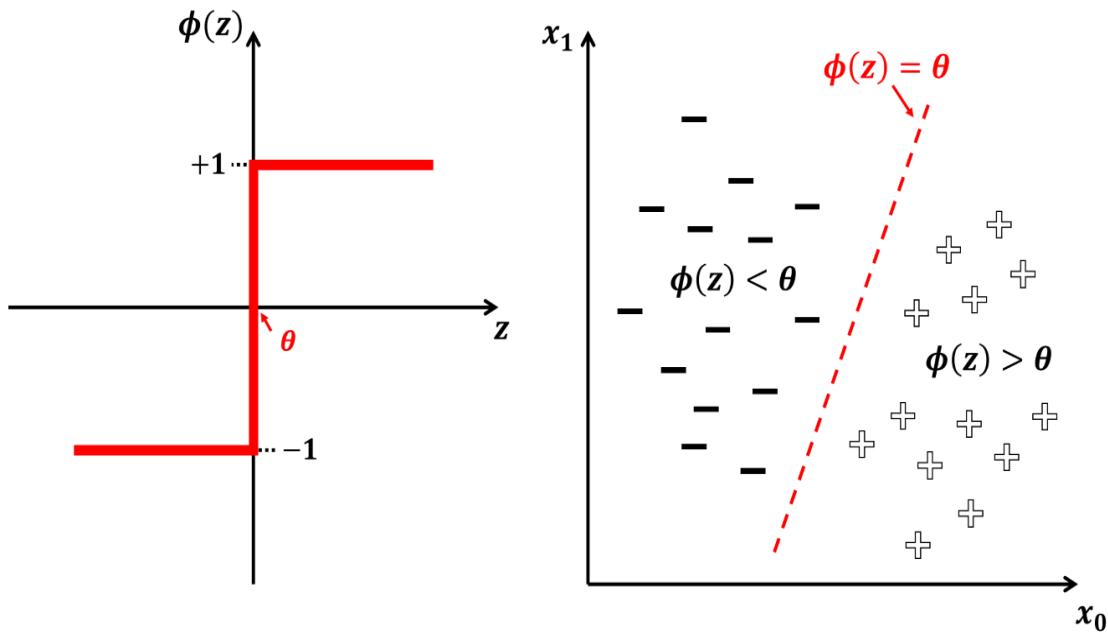
Here, we introduce a new variable, z (the so-called **network input**), which is equivalent to the weighted sum: $z = w_0x_0 + w_1x_1$. The weighted sum is then compared to a threshold, θ , to determine the value of ϕ and subsequently the value of y . Apart from that, these two equations say exactly the same thing as the preceding one.

If these equations look strangely familiar, you might be reminded of chapter 1, *A Taste of Machine Learning*, when we were talking about linear classifiers.

And you are right, a McCulloch-Pitts neuron is essentially a linear, binary classifier!

You can think of it this way: x_0 and x_1 are the input features, w_0 and w_1 are weights to be learned, and the classification is performed by the activation function, ϕ . If we do a good job of learning the weights, which we would do with the help of a suitable training set, we could classify data as positive or negative samples. In this scenario, $\phi(z)=\theta$ would act as the decision boundary.

This might all make more sense with the help of the following diagram:



On the left, you can see the neuron's activation function, ϕ , plotted against z . Remember that z is nothing more than the weighted sum of the two inputs x_0 and x_1 . The rule is that as long as the weighted sum

is below some threshold, θ , the output of the neuron is -1; above θ , the output is +1.

On the right, you can see the decision boundary denoted by $\phi(z)=\theta$, which splits the data into two regimes, $\phi(z)<\theta$ (where all data points are predicted to be negative samples) and $\phi(z)>\theta$ (where all data points are predicted to be positive samples).

The decision boundary does not need to be vertical or horizontal, it can be tilted as shown in the preceding diagram. But in the case of a single McCulloch-Pitts neuron, the decision boundary will always be a straight line.

Of course, the magic lies with learning the weight coefficients, w_0 and w_1 , such that the decision boundary comes to lie right between all positive and all negative data points.

To train a neural network, we generally need three things:

- **Training data:** It is no surprise to learn that we need some data samples with which the effectiveness of our classifier can be verified.
- **Cost function (also known as loss function):** A cost function provides a measure of how good the current weight coefficients are. There is a wide range of cost functions available, which we will talk about toward the end of this chapter. One solution is to count the number of misclassifications. Another one is to calculate the **sum of squared errors**.
- **Learning rule:** A learning rule specifies mathematically how we have to update the weight coefficients from one iteration to the next. This learning rule usually depends on the error

(measured by the cost function) we observed on the training data.

This is where the work of renowned researcher Frank Rosenblatt comes in.

Understanding the perceptron

In the 1950s, American psychologist and artificial intelligence researcher Frank Rosenblatt invented an algorithm that would automatically learn the optimal weight coefficients w_0 and w_1 needed to perform an accurate binary classification: the perceptron learning rule.

Rosenblatt's original perceptron algorithm can be summed up as follows:

1. Initialize the weights to zero or some small random numbers.
2. For each training sample, s_i , perform the following steps:
 1. Compute the predicted target value, \hat{y}_i .
 2. Compare \hat{y}_i to the ground truth, y_i , and update the weights accordingly:
 1. If the two are the same (correct prediction), skip ahead.
 2. If the two are different (wrong prediction), push the weight coefficients, w_0 and

Implementing your first perceptron

Perceptrons are easy enough to be implemented from scratch. We can mimic the typical OpenCV or scikit-learn implementation of a classifier by creating a perceptron object. This will allow us to initialize new perceptron objects that can learn from data via a `fit` method and make predictions via a separate `predict` method.

When we initialize a new perceptron object, we want to pass a learning rate (`lr`, or η in the previous section) and the number of iterations after which the algorithm should terminate (`n_iter`):

```
In [1]: import numpy as np
In [2]: class Perceptron(object):
...     def __init__(self, lr=0.01, n_iter=10):
...         self.lr = lr
...         self.n_iter = n_iter
... 
```

The `fit` method is where most of the work is done. This method should take as input some data samples (`x`) and their associated target labels (`y`). We will then create an array of weights (`self.weights`), one for each feature (`x.shape[1]`), initialized to zero. For convenience, we will keep the bias term (`self.bias`) separate from the weight vector and initialize it to zero as well. One of the reasons for initializing the bias to zero is because the small random numbers in the weights provide asymmetry breaking in the network:

```
...
    def fit(self, X, y):
        self.weights = np.zeros(X.shape[1])
        self.bias = 0.0
... 
```

The `predict` method should take in a number of data samples (x) and, for each of them, return a target label, either +1 or -1. In order to perform this classification, we need to implement $\phi(z) > \theta$. Here we will choose $\theta = 0$, and the weighted sum can be computed with NumPy's dot product:

```
...     def predict(self, X):
...         return np.where(np.dot(X, self.weights) + self.bias >=
0.0,
...                         1, -1)
```

Then we will calculate the Δw terms for every data sample (x_i, y_i) in the dataset and repeat this step for a number of iterations (`self.n_iter`). For this, we need to compare the ground-truth label (y_i) to the predicted label (aforementioned `self.predict(xi)`). The resulting delta term will be used to update both the weights and the bias term:

```
...         for _ in range(self.n_iter):
...             for xi, yi in zip(X, y):
...                 delta = self.lr * (yi - self.predict(xi))
...                 self.weights += delta * xi
...                 self.bias += delta
```

That's it!

Generating a toy dataset

In the following steps, you will learn how to create and plot a toy dataset:

1. To test our perceptron classifier, we need to create some mock data. Let's keep things simple for now and generate 100 data samples (`n_samples`) belonging to one of two blobs (`centers`), again relying on scikit-learn's `make_blobs` function:

```
In [3]: from sklearn.datasets.samples_generator import  
make_blobs...      x, y = make_blobs(n_samples=100,  
centers=2,...                      cluster_std=2.2,  
random_state=42)
```

2. One thing to keep in mind is that our perceptron classifier expects target labels to be either +1 or -1, whereas `make_blobs` returns 0 and 1. An easy way to adjust the labels is with the following equation:

```
In [4]: y = 2 * y - 1
```

4. In the following code, we ...

Fitting the perceptron to data

In the following steps, you will learn to fit a perceptron algorithm on the given data:

1. We can instantiate our perceptron object similar to other classifiers we encountered with OpenCV:

```
In [6]: p = Perceptron(lr=0.1, n_iter=10)
```

Here, we chose a learning rate of 0.1 and told the perceptron to terminate after 10 iterations. These values are chosen rather arbitrarily at this point, although we will come back to them in a little while.

Choosing an appropriate learning rate is critical, but it's not always clear what the most appropriate choice is. The learning rate determines how quickly or slowly we move toward the optimal weight coefficients. If the learning rate is too large, we might accidentally skip the optimal solution. If it is too small, we will need a large number of iterations to converge to the best values.

2. Once the perceptron is set up, we can call the `fit` method to optimize the weight coefficients:

```
In [7]: p.fit(X, y)
```

3. Did it work? Let's have a look at the learned weight values:

```
In [8]: p.weights  
Out[8]: array([ 2.20091094, -0.4798926 ])
```

4. And don't forget to have a peek at the bias term:

```
In [9]: p.bias  
Out[9]: 0.2000000000000001
```

If we plug these values into our equation for ϕ , it becomes clear that the perceptron learned a decision boundary of the form $2.2x_1 - 0.48x_2 + 0.2 \geq 0$.

Evaluating the perceptron classifier

In the following steps, you will be evaluating the trained perceptron on the test data:

1. In order to find out how good our perceptron performs, we can calculate the accuracy score on all data samples:

```
In [10]: from sklearn.metrics import accuracy_score...
accuracy_score(p.predict(X), y)Out[10]: 1.0
```

Perfect score!

2. Let's have a look at the decision landscape by bringing back our `plot_decision_boundary` from the earlier chapters:

```
In [10]: def plot_decision_boundary(classifier, X_test,
y_test):...           # create a mesh to plot in...          h =
0.02 # step size in mesh...           x_min, x_max = X_test[:, 0].min() - 1, X_test[:, 0].max() + 1...
y_min, y_max = X_test[:, 1].min() - 1, X_test[:, 1].max() + 1... xx, yy =
np.meshgrid(np.arange(x_min, ...
```

Applying the perceptron to data that is not linearly separable

In the following steps, you will learn to build a perceptron to separate a nonlinear data:

1. Since the perceptron is a linear classifier, you can imagine that it would have trouble trying to classify data that is not linearly separable. We can test this by increasing the spread (`cluster_std`) of the two blobs in our toy dataset so that the two blobs start overlapping:

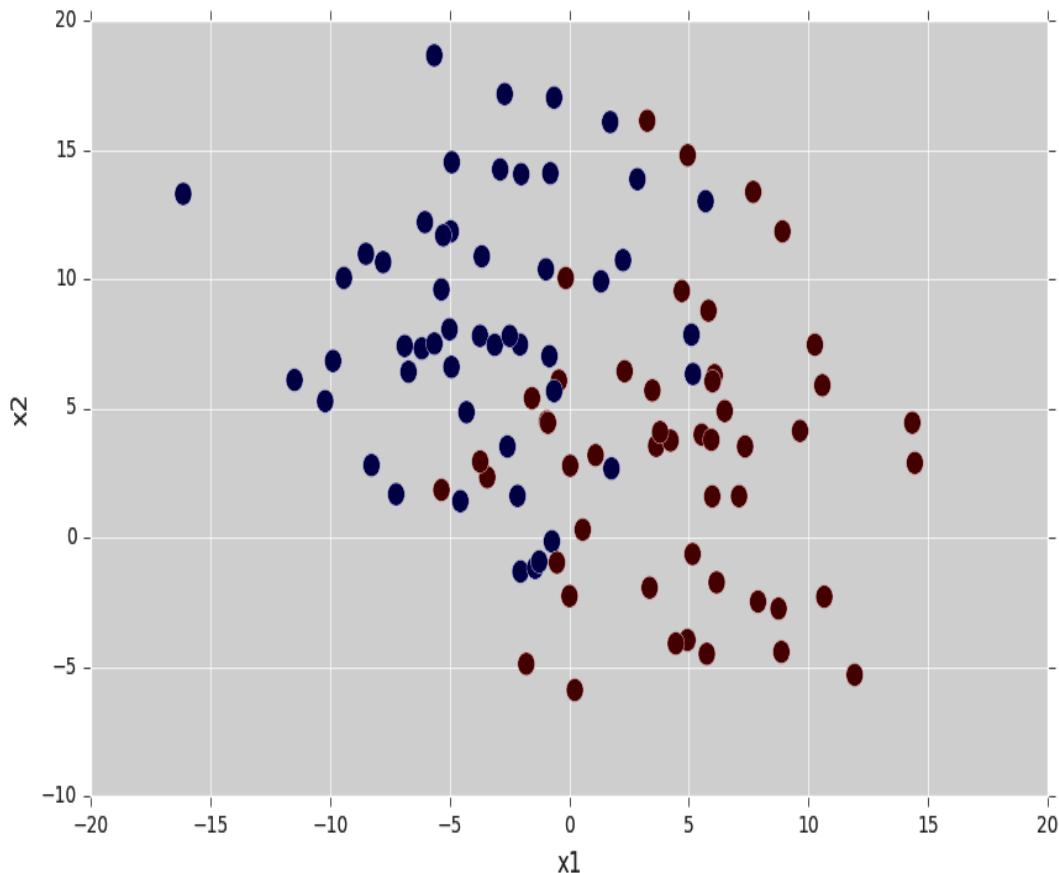
```
In [12]: X, y = make_blobs(n_samples=100, centers=2,
...                      cluster_std=5.2, random_state=42)
...          y = 2 * y - 1
```

2. We can plot the dataset again using matplotlib's `scatter` function:

```
In [13]: plt.scatter(X[:, 0], X[:, 1], s=100, c=y);
...         plt.xlabel('x1')
...         plt.ylabel('x2')
```

As is evident in the following screenshot, this data is no longer linearly separable because there is no straight line that

perfectly separates the two blobs:



The preceding screenshot shows an example of data that is not linearly separable. So what would happen if we applied the perceptron classifier to this dataset?

3. We can find an answer to this question by repeating the preceding steps:

```
In [14]: p = Perceptron(lr=0.1, n_iter=10)
...       p.fit(X, y)
```

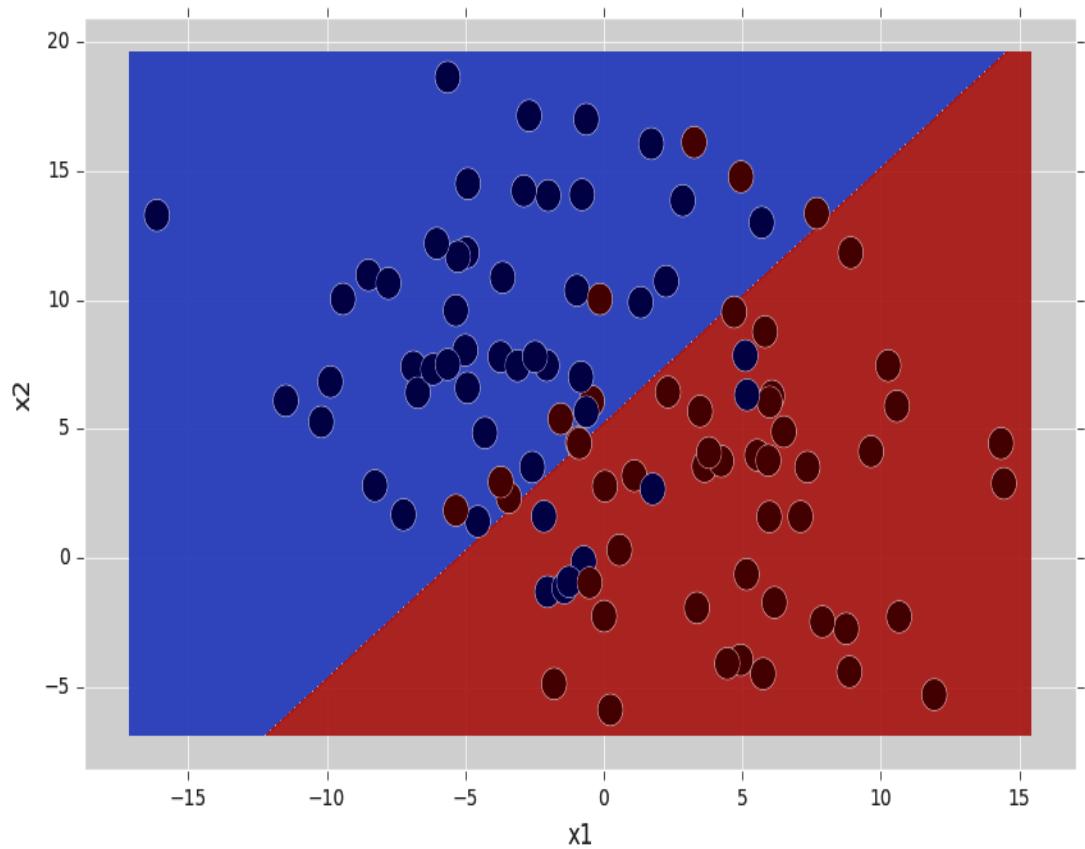
4. Then we find an accuracy score of 81%:

```
In [15]: accuracy_score(p.predict(X), y)
Out[15]: 0.8100000000000005
```

5. In order to find out which data points were misclassified, we can again visualize the decision landscape using our helper function:

```
In [16]: plot_decision_boundary(p, X, y)
...
    plt.xlabel('x1')
...
    plt.ylabel('x2')
```

The following graph makes the limitations of the perceptron classifier evident. Being a linear classifier, it tried to separate the data using a straight line but ultimately failed. The main reason it failed is because the data was not linearly separable even though we achieved 81% accuracy. However, from the following plot, it is clear that many of the red dots lie in the blue region and vice versa. So, unlike a perceptron, we need a nonlinear algorithm that can create not a straight but a nonlinear (circular) decision boundary:

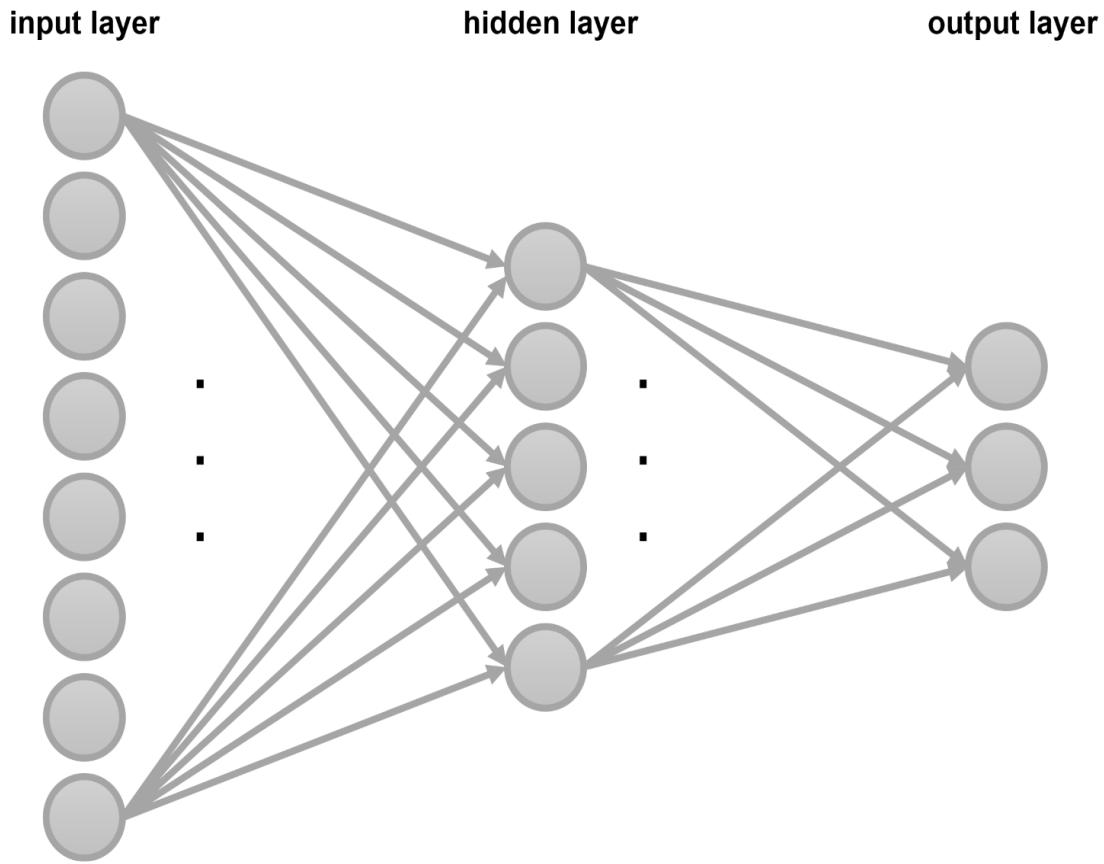


Fortunately, there are ways to make the perceptron more powerful and ultimately create nonlinear decision boundaries.

Understanding multilayer perceptrons

In order to create nonlinear decision boundaries, we can combine multiple perceptrons to form a larger network. This is also known as a **multilayer perceptron (MLP)**. MLPs usually consist of at least three layers, where the first layer has a node (or neuron) for every input feature of the dataset, and the last layer has a node for every class label. The layer in between is called the **hidden layer**.

An example of this feedforward neural network architecture is shown in the following diagram:



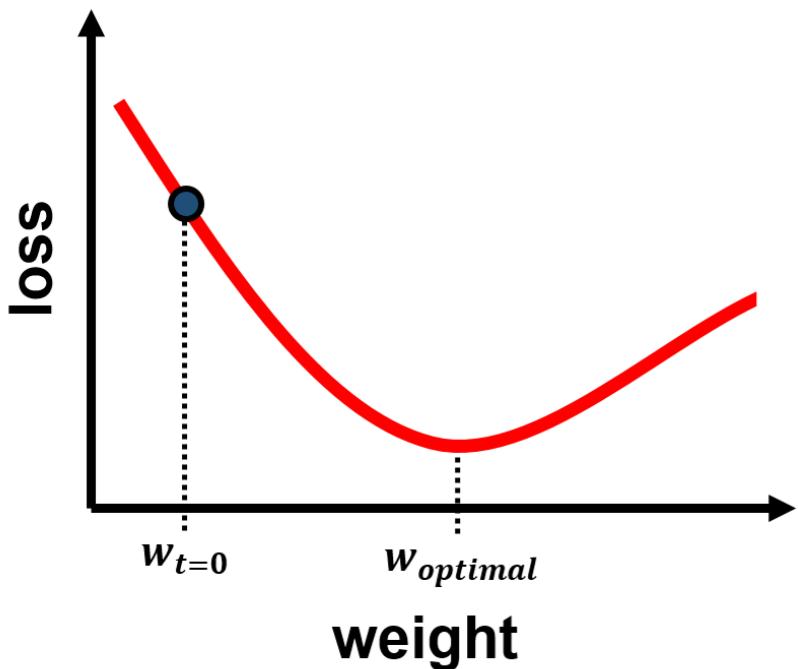
In this network, every circle is an artificial neuron (or, essentially, a perceptron), and the output of one artificial ...

Understanding gradient descent

When we talked about the perceptron earlier in this chapter, we identified three of the essential ingredients needed for training: training data, a cost function, and a learning rule. While the learning rule worked great for a single perceptron, unfortunately, it did not generalize to MLPs, so people had to come up with a more general rule.

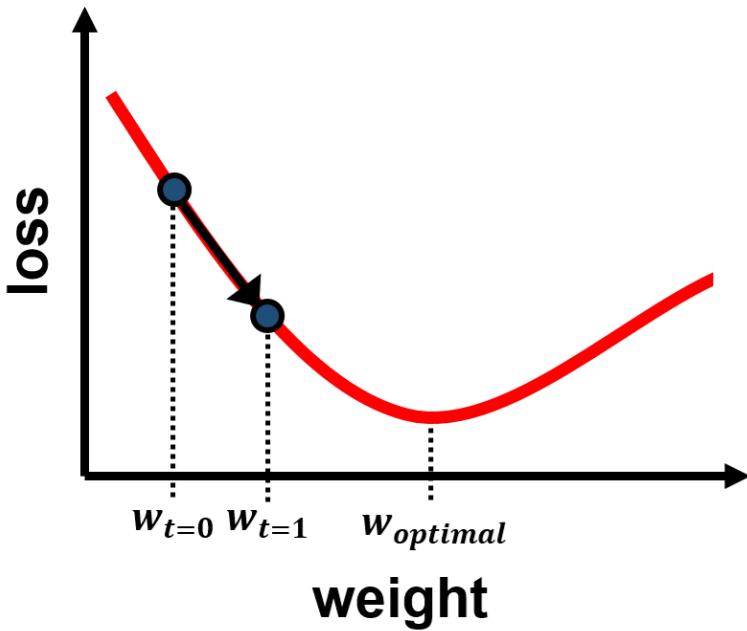
If you think about how we measure the success of a classifier, we usually do so with the help of a cost function. A typical example is the number of misclassifications of the network or the mean squared error. This function (also known as a **loss function**) usually depends on the parameters we are trying to tweak. In neural networks, these parameters are the weight coefficients.

Let's assume a simple neural network has a single weight to tweak, w . Then we can visualize the cost as a function of the weight:



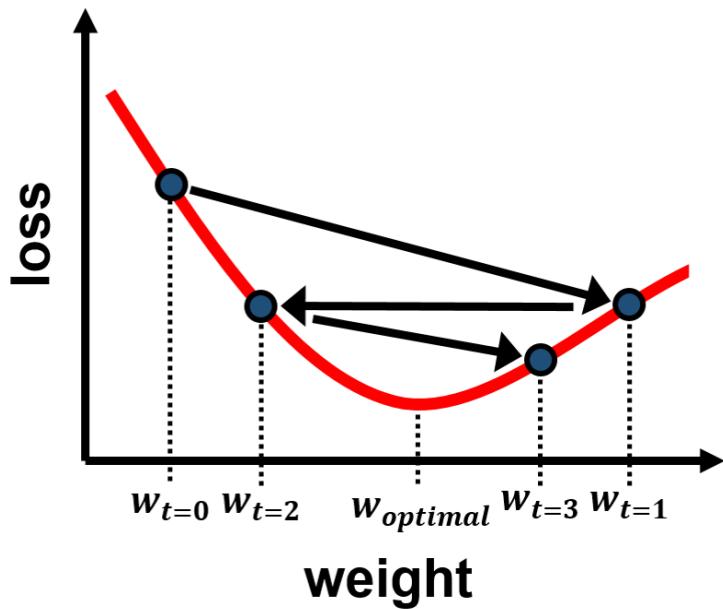
At the beginning of training, at time zero, we may start out way on the left of this graph ($w_{t=0}$). But from the graph, we know that there would be a better value for w , namely $w_{optimal}$, which would minimize the cost function. The smallest cost means the lowest error, so it should be our highest goal to reach $w_{optimal}$ through learning.

This is exactly what gradient descent does. You can think of the gradient as a vector that points up the hill. In gradient descent, we are trying to walk opposite of the gradient, effectively walking down the hill, from the peaks down to the valley:



Once you reach the valley, the gradient goes to zero, and that completes the training.

There are several ways to reach the valley—we could approach from the left, or we could approach from the right. The starting point of our descent is determined by the initial weight values. Furthermore, we have to be careful not to take too large a step, otherwise we might miss the valley:



Hence, in stochastic gradient descent (sometimes also called iterative or on-line gradient descent), the goal is to take small steps but to take them as often as possible. The effective step size is determined by the learning rate of the algorithm.

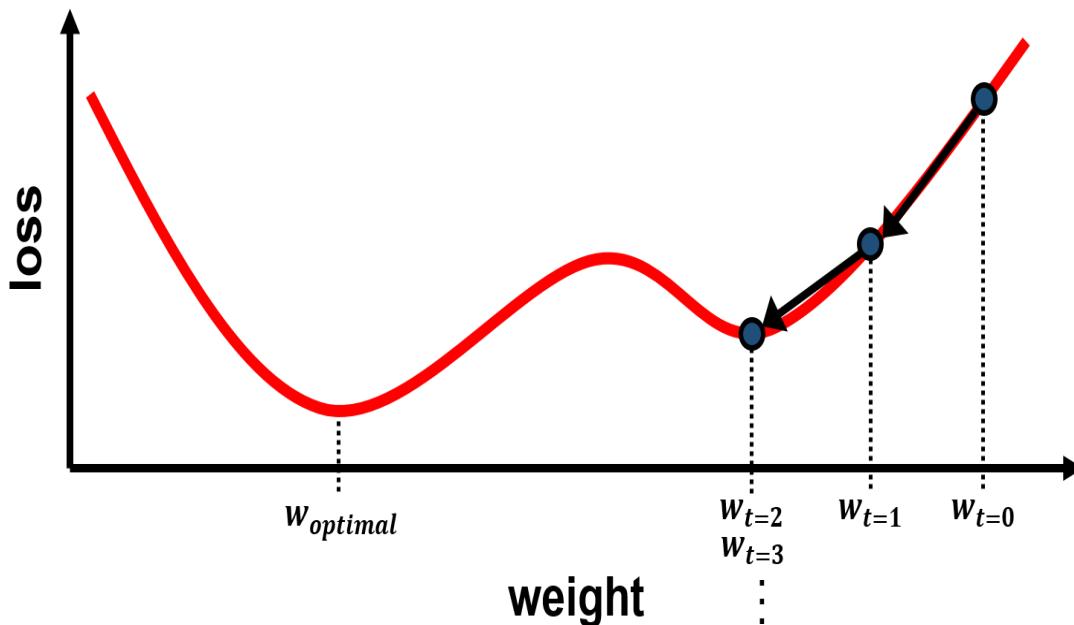
Specifically, we would perform the following procedure over and over:

1. Present a small number of training samples to the network (called the **batch size**).
2. On this small batch of data, calculate the gradient of the cost function.
3. Update the weight coefficients by taking a small step in the opposite direction of the gradient, toward the valley.
4. Repeat steps 1-3 until the weight cost no longer goes down.
This is an indication that we have reached the valley.

Some other ways to improve SGD are using the learning rate finder in the Keras framework, decreasing the step size (learning rate) in epochs, and, as discussed in the preceding point, using a batch size (or mini-batch), which will compute the weight update faster.

Can you think of an example where this procedure might fail?

One scenario that comes to mind is where the cost function has multiple valleys, some deeper than others, as shown in the following diagram:



If we start on the left, we should arrive at the same valley as before—no problem. But, if our starting point is all the way to the right, we might encounter another valley on the way. Gradient descent will lead us straight down to the valley, but it will not have any means to climb out of it.

*This is also known as **getting stuck in a local minimum**. Researchers have come up with different ways to try and avoid this issue, one of them being to add noise to the process.*

There is one piece left in the puzzle. Given our current weight coefficients, how do we know the slope of the cost function?

Training MLPs with backpropagation

This is where backpropagation comes in, which is an algorithm for estimating the gradient of the cost function in neural networks. Some might say that it is basically a fancy word for the chain rule, which is a means to calculate the partial derivative of functions that depend on more than one variable. Nonetheless, it is a method that helped bring the field of artificial neural networks back to life, so we should be thankful for that.

Understanding backpropagation involves quite a bit of calculus, so I will only give you a brief introduction here.

Let's remind ourselves that the cost function, and therefore its gradient, depends on the difference between the true output (y_i) and the current output (\hat{y}_i)

Implementing a MLP in OpenCV

Implementing an MLP in OpenCV uses the same syntax that we have seen at least a dozen times before. In order to see how an MLP compares to a single perceptron, we will operate on the same toy data as before:

Preprocessing the data

However, since we are working with OpenCV, this time we want to make sure the input matrix is made up of 32-bit floating point numbers, otherwise the code will break:

```
In [2]: import numpy as np... x = X_raw.astype(np.float32)
```

Furthermore, we need to think back to [chapter 4, *Representing Data and Engineering Features*](#), and remember how to represent categorical variables. We need to find a way to represent target labels, not as integers but with a one-hot encoding. The easiest way to achieve this is by using scikit-learn's `preprocessing` module:

```
In [3]: from sklearn.preprocessing import OneHotEncoder... enc =
OneHotEncoder(sparse=False, dtype=np.float32)... y =
enc.fit_transform(y_raw.reshape(-1, 1))
```

Creating an MLP classifier in OpenCV

The syntax to create an MLP in OpenCV is the same as for all the other classifiers:

```
In [4]: import cv2  
...     mlp = cv2.ml.ANN_MLP_create()
```

However, now we need to specify how many layers we want in the network and how many neurons there are per layer. We do this with a list of integers, which specify the number of neurons in each layer. Since the data matrix x has two features, the first layer should also have two neurons in it (n_{input}). Since the output has two different values, the last layer should also have two neurons in it (n_{output}).

In between these two layers, we can put as many hidden layers with as many neurons as we want. Let's choose a single hidden layer with an arbitrary number of 10 neurons in it (n_{hidden}):

```
In [5]: n_input = 2  
...     n_hidden = 10  
...     n_output = 2  
...     mlp.setLayerSizes(np.array([n_input, n_hidden, n_output]))
```

Customizing the MLP classifier

Before we move on to training the classifier, we can customize the MLP classifier via a number of optional settings:

- `mlp.setActivationFunction`: This defines the activation function to be used for every neuron in the network.
- `mlp.setTrainMethod`: This defines a suitable training method.
- `mlp.setTermCriteria`: This sets the termination criteria of the training phase.

Whereas our home-brewed perceptron classifier used a linear activation function, OpenCV provides two additional options:

- `cv2.ml.ANN_MLP_IDENTITY`: This is the linear activation function, $f(x) = x$.
- `cv2.ml.ANN_MLP_SIGMOID_SYM`: This is the symmetrical sigmoid function (also known as **hyperbolic tangent**), $f(x) = \beta (1 - \exp(-\alpha x)) / (1 + \exp(-\alpha x))$. Whereas ...

Training and testing the MLP classifier

This is the easy part. Training the MLP classifier is the same as with all other classifiers:

```
In [11]: mlp.train(X, cv2.ml.ROW_SAMPLE, y)
Out[11]: True
```

The same goes for predicting target labels:

```
In [12]: _, y_hat = mlp.predict(X)
```

The easiest way to measure accuracy is by using scikit-learn's helper function:

```
In [13]: from sklearn.metrics import accuracy_score
...     accuracy_score(y_hat.round(), y)
Out[13]: 0.88
```

It looks like we were able to increase our performance from 81% with a single perceptron to 88% with an MLP consisting of 10 hidden-layer neurons and 2 output neurons. In order to see what changed, we can look at the decision boundary one more time:

```
In [14]: def plot_decision_boundary(classifier, X_test, y_test):
... # create a mesh to plot in
... h = 0.02 # step size in mesh
... x_min, x_max = X_test[:, 0].min() - 1, X_test[:, 0].max() + 1
... y_min, y_max = X_test[:, 1].min() - 1, X_test[:, 1].max() + 1
... xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
```

```
... np.arange(y_min, y_max, h))
...
... X_hypo = np.c_[xx.ravel().astype(np.float32),
... yy.ravel().astype(np.float32)]
... _, zz = classifier.predict(X_hypo)
```

However, there is a problem right here, in that `zz` is now a one-hot encoded matrix. In order to transform the one-hot encoding into a number that corresponds to the class label (zero or one), we can use NumPy's `argmax` function:

```
...         zz = np.argmax(zz, axis=1)
```

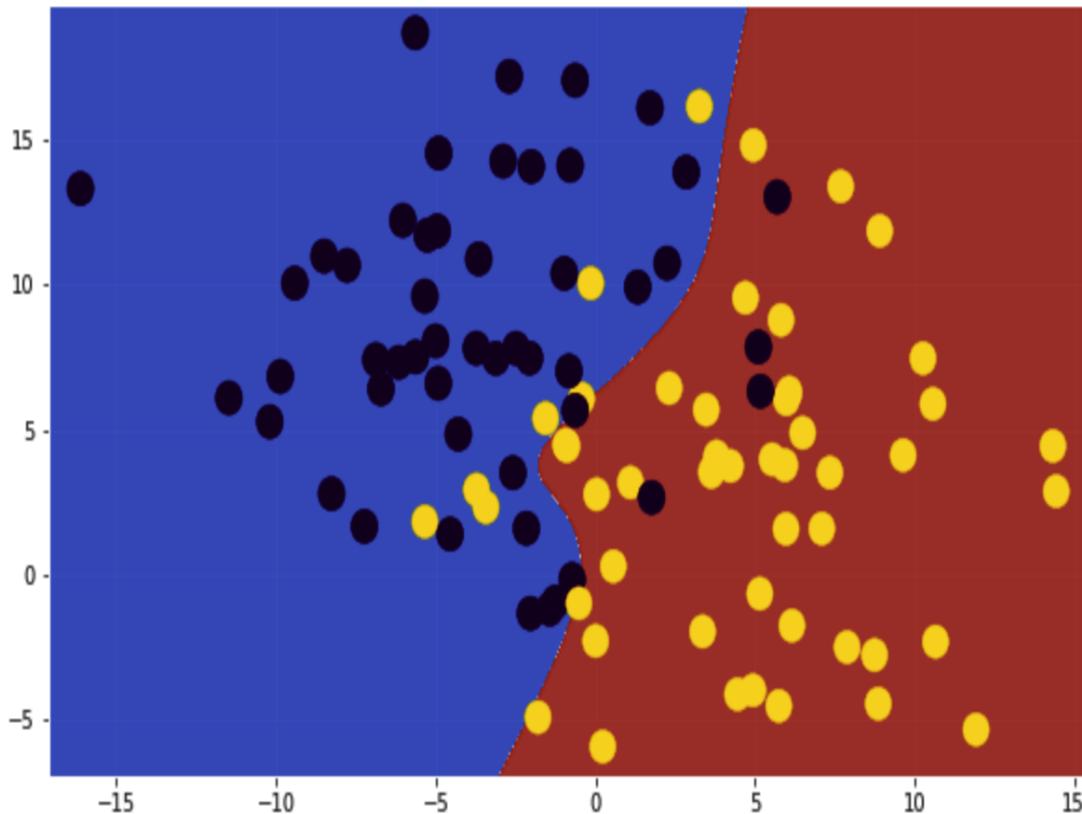
Then the rest stays the same:

```
...         zz = zz.reshape(xx.shape)
...         plt.contourf(xx, yy, zz, cmap=plt.cm.coolwarm, alpha=0.8)
...         plt.scatter(X_test[:, 0], X_test[:, 1], c=y_test, s=200)
```

Then we can call the function like this:

```
In [15]: plot_decision_boundary(mlp, X, y_raw)
```

The output looks like this:



The preceding output shows a decision boundary of an MLP with one hidden layer.

And voila! The decision boundary is no longer a straight line. That being said, you got a great performance increase and might have expected a more drastic performance increase. But nobody said we have to stop here!

There are at least two different things we can try from here on out:

- We can add more neurons to the hidden layer. You can do this by replacing `n_hidden` on line 6 with a larger value and running the code again. Generally speaking, the more neurons you put in the network, the more powerful the MLP will be.

- We can add more hidden layers. It turns out that this is where neural networks really get their power from.

Hence, this is where I should tell you about deep learning.

Getting acquainted with deep learning

Back when deep learning didn't have a fancy name yet, it was called artificial neural networks. So you already know a great deal about it!

Eventually, interest in neural networks was rekindled in 1986 when David Rumelhart, Geoffrey Hinton, and Ronald Williams were involved in the (re)discovery and popularization of the aforementioned backpropagation algorithm. However, it was not until recently that computers became powerful enough so they could actually execute the backpropagation algorithm on large-scale networks, leading to a surge in deep learning research.

You can find more information on the history and origin of deep learning in the following scientific article: Wang and Raj (2017), On the Origin ...

Getting acquainted with Keras

The core data structure of Keras is a model, which is similar to OpenCV's classifier object, except it focuses on neural networks only. The simplest type of model is the sequential model, which arranges the different layers of the neural network in a linear stack, just like we did for the MLP in OpenCV:

```
In [1]: from keras.models import Sequential  
...      model = Sequential()  
Out[1]: Using TensorFlow backend.
```

Then, different layers can be added to the model one by one. In Keras, layers do not just contain neurons, they also perform a function. Some core layer types include the following:

- **Dense:** This is a densely connected layer. This is exactly what we used when we designed our MLP: a layer of neurons that is connected to every neuron in the previous layer.
- **Activation:** This applies an activation function to an output. Keras provides a whole range of activation functions, including OpenCV's identify function (`linear`), the hyperbolic tangent (`tanh`), a sigmoidal squashing function (`sigmoid`), a softmax function (`softmax`), and many more.
- **Reshape:** This reshapes an output to a certain shape.

There are other layers that calculate arithmetic or geometric operations on their inputs:

- **Convolutional layers:** These layers allow you to specify a kernel with which the input layer is convolved. This allows you to perform operations such as a Sobel filter or apply a Gaussian kernel in 1D, 2D, or even 3D.
- **Pooling layers:** These layers perform a max pooling operation on their input, where the output neuron's activity is given by the maximally active input neuron.

Some other layers that are popular in deep learning are as follows:

- **Dropout:** This layer randomly sets a fraction of input units to zero at each update. This is a way to inject noise into the training process, making it more robust.
- **Embedding:** This layer encodes categorical data, similar to some functions from scikit-learn's `preprocessing` module.
- **Gaussian noise:** This layer applies additive zero-centered Gaussian noise. This is another way of injecting noise into the training process, making it more robust.

A perceptron similar to the preceding one could thus be implemented using a dense layer that has two inputs and one output. Staying true to our earlier example, we will initialize the weights to zero and use the hyperbolic tangent as an activation function:

```
In [2]: from keras.layers import Dense  
...      model.add(Dense(1, activation='tanh',  
input_dim=2,  
...                      kernel_initializer='zeros'))
```

Finally, we want to specify the training method. Keras provides a number of optimizers, including the following:

- **Stochastic gradient descent (SGD)**: This is what we discussed earlier.
- **Root mean square propagation (RMSprop)**: This is a method in which the learning rate is adapted for each of the parameters.
- **Adaptive moment estimation (Adam)**: This is an update to the root mean square propagation.

In addition, Keras also provides a number of different loss functions:

- **Mean squared error (mean_squared_error)**: This is what was discussed earlier.
- **Hinge loss (hinge)**: This is a maximum-margin classifier often used with SVM, as discussed in [chapter 6, Detecting Pedestrians with Support Vector Machines](#).

You can see that there's a plethora of parameters to be specified and methods to choose from. To stay true to our aforementioned perceptron implementation, we will choose SGD as an optimizer, the

mean squared error as a cost function, and accuracy as a scoring function:

```
In [3]: model.compile(optimizer='sgd',
...                     loss='mean_squared_error',
...                     metrics=['accuracy'])
```

In order to compare the performance of the Keras implementation to our home-brewed version, we will apply the classifier to the same dataset:

```
In [4]: from sklearn.datasets.samples_generator import make_blobs
...       X, y = make_blobs(n_samples=100, centers=2,
...                         cluster_std=2.2, random_state=42)
```

Finally, a Keras model is fit to the data with a very familiar syntax. Here, we can also choose how many iterations to train for (`epochs`), how many samples to present before we calculate the error gradient (`batch_size`), whether to shuffle the dataset (`shuffle`), and whether to output progress updates (`verbose`):

```
In [5]: model.fit(X, y, epochs=400, batch_size=100, shuffle=False,
...                  verbose=0)
```

After the training completes, we can evaluate the classifier as follows:

```
In [6]: model.evaluate(X, y)
Out[6]: 32/100 [=====>.....] - ETA: 0s
[0.040941802412271501, 1.0]
```

Here, the first reported value is the mean squared error, whereas the second value denotes accuracy. This means that the final mean squared error was 0.04, and we had 100% accuracy. Way better than our own implementation!

You can find more information on Keras, source code documentation, and a number of tutorials at <http://keras.io>.

With these tools in hand, we are now ready to approach a real-world dataset!

Classifying handwritten digits

In the previous section, we covered a lot of the theory around neural networks, which can be a little overwhelming if you are new to this topic. In this section, we will use the famous MNIST dataset, which contains 60,000 samples of handwritten digits along with their labels.

We will train two different networks on it:

- An MLP using OpenCV
- A deep neural network using Keras

Loading the MNIST dataset

The easiest way to obtain the MNIST dataset is by using Keras:

```
In [1]: from keras.datasets import mnist
...      (X_train, y_train), (X_test, y_test) = mnist.load_data()
Out[1]: Using TensorFlow backend.
        Downloading data from
        https://s3.amazonaws.com/img-datasets/mnist.npz
```

This will download the data from Amazon Cloud (might take a while depending on your internet connection) and automatically split the data into training and test sets.

MNIST provides its own predefined train-test split. This way, it is easier to compare the performance of different classifiers because they will all use the same data for training and the same data for testing.

This data comes in a format that we are already familiar with:

```
In [2]: X_train.shape, y_train.shape
Out[2]: ((60000, 28, 28), (60000,))
```

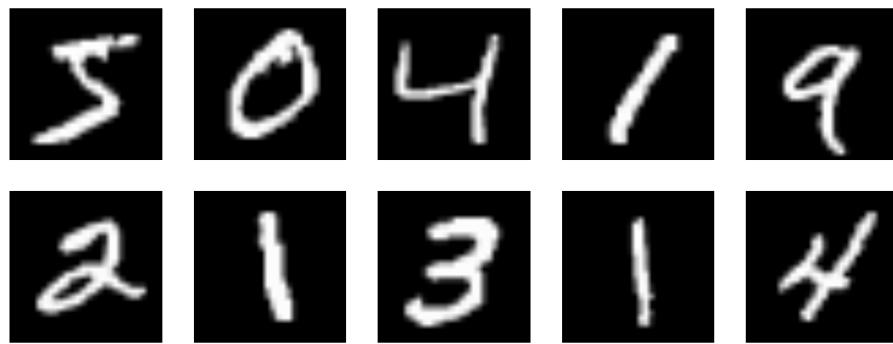
We should take note that the labels come as integer values between zero and nine (corresponding to the digits 0-9):

```
In [3]: import numpy as np
...      np.unique(y_train)
Out[3]: array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9], dtype=uint8)
```

We can have a look at some example digits:

```
In [4]: import matplotlib.pyplot as plt
...      %matplotlib inline
In [5]: for i in range(10):
...     plt.subplot(2, 5, i + 1)
...     plt.imshow(X_train[i, :, :], cmap='gray')
...     plt.axis('off')
```

The digits look like this:



In fact, the MNIST dataset is the successor to the NIST digits dataset provided by scikit-learn that we used before (`sklearn.datasets.load_digits`; refer to [chapter 2](#), *Working with Data in OpenCV*). Some notable differences are as follows:

- MNIST images are significantly larger (28 x 28 pixels) than NIST images (8 x 8 pixels), thus paying more attention to fine details such as distortions and individual differences between images of the same digit.
- The MNIST dataset is much larger than the NIST dataset, providing 60,000 training and 10,000 test samples (as compared to a total of 5,620 NIST images).

Preprocessing the MNIST dataset

As we learned in [Chapter 4, Representing Data and Engineering Features](#), there are a number of preprocessing steps we might like to apply here:

- **Centering:** It is important that all the digits are centered in the image. For example, take a look at all the example images of the digit 1 in the preceding diagram, which are all made of an almost-vertical strike. If the images were misaligned, the strike could lie anywhere in the image, making it hard for the neural network to find commonalities in the training samples. Fortunately, images in MNIST are already centered.
- **Scaling:** The same is true for scaling the digits so that they all have the same size. This way, the location of strikes, curves, and loops are important. ...

Training an MLP using OpenCV

We can set up and train an MLP in OpenCV with the following recipe:

1. Instantiate a new MLP object:

```
In [9]: import cv2  
...      mlp = cv2.ml.ANN_MLP_create()
```

2. Specify the size of every layer in the network. We are free to add as many layers as we want, but we need to make sure that the first layer has the same number of neurons as input features (784 in our case), and that the last layer has the same number of neurons as class labels (10 in our case), while there are two hidden layers each having 512 nodes:

```
In [10]: mlp.setLayerSizes(np.array([784, 512, 512, 10]))
```

3. Specify an activation function. Here we use the sigmoidal activation function from before:

```
In [11]: mlp.setActivationFunction(cv2.ml.ANN_MLP_SIGMOID_SYM,  
...                                2.5, 1.0)
```

4. Specify the training method. Here, we use the backpropagation algorithm described earlier. We also need to make sure that we choose a small enough learning rate. Since we have on the order of 10^5 training samples, it is a good idea to set the learning rate to at most 10^{-5} :

```
In [12]: mlp.setTrainMethod(cv2.ml.ANN_MLP_BACKPROP)  
...          mlp.setBackpropWeightScale(0.00001)
```

5. Specify the termination criteria. Here, we use the same criteria as before: to run training for 10 iterations (`term_max_iter`) or until the error does no longer increase significantly (`term_eps`):

```
In [13]: term_mode = (cv2.TERM_CRITERIA_MAX_ITER +  
...                  cv2.TERM_CRITERIA_EPS)  
...          term_max_iter = 10  
...          term_eps = 0.01  
...          mlp.setTermCriteria((term_mode, term_max_iter,  
...                                term_eps))
```

6. Train the network on the training set (`x_train_pre`):

```
In [14]: mlp.train(X_train_pre, cv2.ml.ROW_SAMPLE, y_train_pre)  
Out[14]: True
```

Before you call `mlp.train`, here is a word of caution: this might take several hours to run, depending on your computer setup! For comparison, it took just under an hour on my own laptop. We are now dealing with a real-world dataset of 60,000 samples: if we run 100 training epochs, we have to compute 6 million gradients! So beware.

When the training completes, we can calculate the accuracy score on the training set to see how far we got:

```
In [15]: _, y_hat_train = mlp.predict(X_train_pre)
In [16]: from sklearn.metrics import accuracy_score
...      accuracy_score(y_hat_train.round(), y_train_pre)
Out[16]: 0.9297666666666663
```

But, of course, what really counts is the accuracy score we get on the held-out test data which was not taken into account in training process:

```
In [17]: _, y_hat_test = mlp.predict(X_test_pre)
...      accuracy_score(y_hat_test.round(), y_test_pre)
Out[17]: 0.9169000000000005
```

91.7% accuracy is not bad at all if you ask me! The first thing you should try is to change the layer sizes in the preceding In [10] and see how the test score changes. As you add more neurons to the network, you should see the training score increase—and with it, hopefully, the test score. However, having N neurons in a single layer is not the same as having them spread out over several layers! Can you confirm this observation?

Training a deep neural network using Keras

Although we achieved a formidable score with the preceding MLP, our result does not hold up to state-of-the-art results. Currently, the best result has close to 99.8% accuracy—better than human performance! This is why, nowadays, the task of classifying handwritten digits is largely regarded as solved.

To get closer to the state-of-the-art results, we need to use state-of-the-art techniques. Thus, we return to Keras.

Preprocessing the MNIST dataset

In the following steps, you will learn to preprocess the data before it is fed to the neural network:

1. To make sure we get the same result every time we run the experiment, we will pick a random seed for NumPy's random number generator. This way, shuffling the training samples from the MNIST dataset will always result in the same order:

```
In [1]: import numpy as np  
...      np.random.seed(1337)
```

2. Keras provides a loading function similar to `train_test_split` from scikit-learn's `model_selection` module. Its syntax might look strangely familiar to you:

```
In [2]: from keras.datasets import mnist  
...      (x_train, y_train), (x_test, y_test) =  
mnist.load_data()
```

In contrast to other datasets we have encountered so far, MNIST comes with a predefined train-test split. This allows the dataset to be used as a benchmark, as the test score reported by different algorithms will always apply to the same test samples.

3. The neural networks in Keras act on the feature matrix slightly differently than the standard OpenCV and scikit-learn estimators. Whereas the rows of a feature matrix in Keras still correspond to the number of samples (`x_train.shape[0]` in the following code), we can preserve the two-dimensional nature of the input images by adding more dimensions to the feature matrix:

```
In [3]: img_rows, img_cols = 28, 28
...      X_train = X_train.reshape(X_train.shape[0], img_rows,
img_cols, 1)
...      X_test = X_test.reshape(X_test.shape[0], img_rows,
img_cols, 1)
...      input_shape = (img_rows, img_cols, 1)
```

4. Here, we have reshaped the feature matrix into a four-dimensional matrix with dimensions `n_features` x 28 x 28 x 1. We also need to make sure we operate on 32-bit floating point numbers between [0, 1], rather than unsigned integers in [0, 255]:

```
...      X_train = X_train.astype('float32') / 255.0
...      X_test = X_test.astype('float32') / 255.0
```

5. Then, we can one-hot encode the training labels as we did before. This will make sure each category of target labels can be assigned to a neuron in the output layer. We could do this with scikit-learn's `preprocessing`, but in this case, it is easier to use Keras' own utility function:

```
In [4]: from keras.utils import np_utils
...
...     n_classes = 10
...     Y_train = np_utils.to_categorical(y_train, n_classes)
...     Y_test = np_utils.to_categorical(y_test, n_classes)
```

Creating a convolutional neural network

In the following steps, you will create a neural network and train on the data you preprocessed earlier:

1. Once we have preprocessed the data, it is time to define the actual model. Here, we will once again rely on the `Sequential` model to define a feedforward neural network:

```
In [5]: from keras.models import Sequential... model =  
Sequential()
```

2. However, this time, we will be smarter about the individual layers. We will design our neural network around a convolutional layer, where the kernel is a 3 x 3-pixel two-dimensional convolution:

```
In [6]: from keras.layers import Convolution2D...      n_filters  
= 32...      kernel_size = (3, 3)...  
model.add(Convolution2D(n_filters, kernel_size[0],  
kernel_size[1],... border_mode='valid', ...
```

Model summary

You can also visualize the model's summary, which will list all the layers along with their respective dimensions and the number of weights each layer consists. It will also provide you with information about the total number of parameters (weights and biases) in your network:

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_4 (Conv2D)	(None, 26, 26, 32)	320
activation_5 (Activation)	(None, 26, 26, 32)	0
conv2d_5 (Conv2D)	(None, 24, 24, 32)	9248
activation_6 (Activation)	(None, 24, 24, 32)	0
max_pooling2d_3 (MaxPooling2D)	(None, 12, 12, 32)	0
dropout_3 (Dropout)	(None, 12, 12, 32)	0
flatten_3 (Flatten)	(None, 4608)	0
dense_3 (Dense)	(None, 128)	589952
activation_7 (Activation)	(None, 128)	0
dropout_4 (Dropout)	(None, 128)	0
dense_4 (Dense)	(None, 10)	1290
activation_8 (Activation)	(None, 10)	0
<hr/>		
Total params:	600,810	
Trainable params:	600,810	
Non-trainable params:	0	

We can see that there are in total 600,810 parameters that will be trained and will require a good amount of computation power! Please note that how we calculate the number of parameters in each layer is out of the scope of this book.

Fitting the model

We fit the model as we do with all other classifiers (caution, this might take a while):

```
In [12]: model.fit(X_train, Y_train, batch_size=128,  
nb_epoch=12, ...  
verbose=1, validation_data=  
(X_test, Y_test))
```

After training completes, we can evaluate the classifier:

```
In [13]: model.evaluate(X_test, Y_test, verbose=0)Out[13]: 0.99
```

And we achieved 99% accuracy! This is worlds apart from the MLP classifier we implemented before. And this is just one way to do things. As you can see, neural networks provide a plethora of tuning parameters, and it is not at all clear which ones will lead to the best performance.

Summary

In this chapter, we added a whole bunch of skills to our list as a machine learning practitioner. Not only did we cover the basics of artificial neural networks, including perceptrons and MLPs, we also got our hands on some advanced deep learning software. We learned how to build a simple perceptron from scratch and how to build state-of-the-art networks using Keras. Furthermore, we learned about all the details of neural nets: activation functions, loss functions, layer types, and training methods. All in all, this was probably the most intensive chapter yet.

Now that you know about most of the essential supervised learners, it is time to talk about how to combine different algorithms into a more powerful one. Thus, in the next chapter, we will talk about how to build ensemble classifiers.

Ensemble Methods for Classification

So far, we have looked at a number of interesting machine learning algorithms, from classic methods such as linear regression to more advanced techniques such as deep neural networks. At various points, we pointed out that every algorithm has its own strengths and weaknesses—and we took note of how to spot and overcome these weaknesses.

However, wouldn't it be great if we could simply stack together a bunch of average classifiers to form a much stronger **ensemble** of classifiers?

In this chapter, we will do just that. Ensemble methods are techniques that bind multiple different models together in order to solve a shared problem. Their use has become a common practice in competitive machine learning—making ...

Technical requirements

You can refer to the code for this chapter from the following link: <https://github.com/PacktPublishing/Machine-Learning-for-OpenCV-Second-Edition/tree/master/Chapter10>.

Here is a short summary of the software and hardware requirements:

- OpenCV version 4.1.x (4.1.0 or 4.1.1 will both work just fine).
- Python version 3.6 (any Python version 3.x will be fine).
- Anaconda Python 3 for installing Python and the required modules.
- You can use any OS—macOS, Windows, or Linux-based—with this book. We recommend you have at least 4 GB RAM in your system.
- You don't need to have a GPU to run the code provided with the book.

Understanding ensemble methods

The goal of ensemble methods is to combine the predictions of several individual estimators built with a given learning algorithm in order to solve a shared problem. Typically, an ensemble consists of two major components:

- A set of models
- A set of decision rules that govern how the results of these models are combined into a single output

The idea behind ensemble methods has much to do with the wisdom of the crowd concept. Rather than the opinion of a single expert, we consider the collective opinion of a group of individuals. In the context of machine learning, these individuals would be classifiers or regressors. The idea is that if we just ask a large enough number of classifiers, one of them ought to get ...

Understanding averaging ensembles

Averaging methods have a long history in machine learning and are commonly applied to fields such as molecular dynamics and audio signal processing. Such ensembles are typically seen as exact replicas of a given system.

An averaging ensemble is essentially a collection of models that train on the same dataset. Their results are then aggregated in a number of ways.

One common method involves creating multiple model configurations that take different parameter subsets as input. Techniques that take this approach are referred to collectively as bagging methods.

Bagging methods come in many different flavors. However, they typically only differ in the way they draw random subsets of the training set:

- Pasting methods draw random subsets of the samples without replacement of data samples.
- Bagging methods draw random subsets of the samples with replacement of data samples.
- Random subspace methods draw random subsets of the features but train on all data samples.

- Random patch methods draw random subsets of both samples and features.

Averaging ensembles can be used to reduce the variability of a model's performance.

In scikit-learn, bagging methods can be realized using the `BaggingClassifier` and `BaggingRegressor` meta-estimators. These are meta-estimators because they allow us to build an ensemble from any other base estimator.

Implementing a bagging classifier

We can, for instance, build an ensemble from a collection of 10 k -NN classifiers as follows:

```
In [1]: from sklearn.ensemble import BaggingClassifier...      from
sklearn.neighbors import KNeighborsClassifier...      bag_knn =
BaggingClassifier(KNeighborsClassifier(),...
n_estimators=10)
```

The `BaggingClassifier` class provides a number of options to customize the ensemble:

- `n_estimators`: As shown in the preceding code, this specifies the number of base estimators in the ensemble.
- `max_samples`: This denotes the number (or fraction) of samples to draw from the dataset to train each base estimator. We can set `bootstrap=True` to sample with replacement (effectively implementing bagging), or we can set `bootstrap=False` to implement ...

Implementing a bagging regressor

Similarly, we can use the `BaggingRegressor` class to form an ensemble of regressors.

For example, we could build an ensemble of decision trees to predict housing prices from the Boston dataset of [chapter 3, First Steps in Supervised Learning](#).

In the following steps, you will learn how to use a bagging regressor for forming an ensemble of regressors:

1. The syntax is almost identical to setting up a bagging classifier:

```
In [7]: from sklearn.ensemble import BaggingRegressor
...
...     from sklearn.tree import DecisionTreeRegressor
...     bag_tree = BaggingRegressor(DecisionTreeRegressor(),
...                                 max_features=0.5,
...                                 n_estimators=10,
...                                 random_state=3)
```

2. Of course, we need to load and split the dataset as we did for the breast cancer dataset:

```
In [8]: from sklearn.datasets import load_boston
...
...     dataset = load_boston()
```

```
...     X = dataset.data
...     y = dataset.target
In [9]: from sklearn.model_selection import train_test_split
...     X_train, X_test, y_train, y_test = train_test_split(
...             X, y, random_state=3
...     )
```

3. Then, we can fit the bagging regressor on `x_train` and score it
on `x_test`:

```
In [10]: bag_tree.fit(X_train, y_train)
...     bag_tree.score(X_test, y_test)
Out[10]: 0.82704756225081688
```

In the preceding example, we find a performance boost of roughly 5%, from 77.3% accuracy for a single decision tree to 82.7% accuracy.

Of course, we wouldn't just stop here. Nobody said the ensemble needs to consist of 10 individual estimators, so we are free to explore different-sized ensembles. On top of that, the `max_samples` and `max_features` parameters allow for a great deal of customization.

A more sophisticated version of bagged decision trees is called random forests, which we will talk about later in this chapter.

Understanding boosting ensembles

Another approach to building ensembles is through boosting. Boosting models use multiple individual learners in sequence to iteratively boost the performance of the ensemble.

Typically, the learners used in boosting are relatively simple. A good example is a decision tree with only a single node—a decision stump. Another example could be a simple linear regression model. The idea is not to have the strongest individual learners, but quite the opposite—we want the individuals to be weak learners so that we get a superior performance only when we consider a large number of individuals.

At each iteration of the procedure, the training set is adjusted so that the next classifier is applied to the data points that ...

Weak learners

Weak learners are classifiers that are only slightly correlated with the actual classification; they can be somewhat better than the random predictions. On the contrary, strong learners are arbitrarily well correlated with the correct classification.

The idea here is that you don't use just one but a broad set of weak learners, each one slightly better than random. Many instances of the weak learners can be pooled using boosting, bagging, and so on together to create a strong ensemble classifier. The benefit is that the final classifier will not lead to *overfitting* on your training data.

For example, AdaBoost fits a sequence of weak learners on different weighted training data. It starts by predicting the training dataset and gives equal weight to each observation/sample. If the first learner prediction is incorrect, then it gives higher weight to the observation/sample that has been mispredicted. Since it is an iterative process, it continues to add learners until a limit is reached in the number of models or accuracy.

Implementing a boosting classifier

For example, we can build a boosting classifier from a collection of 10 decision trees as follows:

```
In [11]: from sklearn.ensemble import GradientBoostingClassifier...
boost_class = GradientBoostingClassifier(n_estimators=10, ...
random_state=3)
```

These classifiers support both binary and multiclass classification.

Similar to the `BaggingClassifier` class, the `GradientBoostingClassifier` class provides a number of options to customize the ensemble:

- `n_estimators`: This denotes the number of base estimators in the ensemble. A large number of estimators typically results in better performance.
- `loss`: This denotes the loss function (or cost function) to be optimized. Setting `loss='deviance'` implements logistic regression ...

Implementing a boosting regressor

Implementing a boosted regressor follows the same syntax as the boosted classifier:

```
In [15]: from sklearn.ensemble import GradientBoostingRegressor  
...      boost_reg = GradientBoostingRegressor(n_estimators=10,  
...                                         random_state=3)
```

We have seen earlier that a single decision tree can achieve 79.3% accuracy on the Boston dataset. A bagged decision tree classifier made of 10 individual regression trees achieved 82.7% accuracy. But how does a boosted regressor compare?

Let's reload the Boston dataset and split it into training and test sets. We want to make sure we use the same value for `random_state` so that we end up training and testing on the same subsets of the data:

```
In [16]: dataset = load_boston()  
...      x = dataset.data  
...      y = dataset.target  
In [17]: x_train, x_test, y_train, y_test = train_test_split(  
...           x, y, random_state=3  
...           )
```

As it turns out, the boosted decision tree ensemble actually performs worse than the previous code:

```
In [18]: boost_reg.fit(x_train, y_train)  
...      boost_reg.score(x_test, y_test)  
Out[18]: 0.71991199075668488
```

This result might be confusing at first. After all, we used 10 times more classifiers than we did for the single decision tree. Why would our numbers get worse?

You can see this is a good example of an expert classifier being smarter than a group of weak learners. One possible solution is to make the ensemble larger. In fact, it is customary to use in the order of 100 weak learners in a boosted ensemble:

```
In [19]: boost_reg = GradientBoostingRegressor(n_estimators=100)
```

Then, when we retrain the ensemble on the Boston dataset, we get a test score of 89.8%:

```
In [20]: boost_reg.fit(X_train, y_train)
...       boost_reg.score(X_test, y_test)
Out[20]: 0.89984081091774459
```

What happens when you increase the number to `n_estimators=500`? There's a lot more we could do by playing with the optional parameters.

As you can see, boosting is a powerful procedure that allows you to get massive performance improvements by combining a large number of relatively simple learners.

A specific implementation of boosted decision trees is the AdaBoost algorithm, which we will talk about later in this chapter.

Understanding stacking ensembles

All the ensemble methods we have seen so far share a common design philosophy: to fit multiple individual classifiers to the data and incorporate their predictions with the help of some simple decision rules (such as averaging or boosting) into a final prediction.

Stacking ensembles, on the other hand, build ensembles with hierarchies. Here, individual learners are organized into multiple layers where the output of one layer of learners is used as training data for a model at the next layer. This way, it is possible to successfully blend hundreds of different models.

Unfortunately, discussing stacking ensembles in detail is beyond the scope of this book.

However, these models can be very powerful, as seen, ...

Combining decision trees into a random forest

A popular variation of bagged decision trees are the so-called random forests. These are essentially a collection of decision trees, where each tree is slightly different from the others. In contrast to bagged decision trees, each tree in a random forest is trained on a slightly different subset of data features.

Although a single tree of unlimited depth might do a relatively good job of predicting the data, it is also prone to overfitting. The idea behind random forests is to build a large number of trees, each of them trained on a random subset of data samples and features. Because of the randomness of the procedure, each tree in the forest will overfit the data in a slightly different way. The effect of overfitting can then be reduced by averaging the predictions of the individual trees.

Understanding the shortcomings of decision trees

The effect of overfitting the dataset, which a decision tree often falls victim to, is best demonstrated through a simple example.

For this, we will return to the `make_moons` function from scikit-learn's `datasets` module, which we previously used in [Chapter 8, *Discovering Hidden Structures with Unsupervised Learning*](#), to organize data into two interleaving half circles. Here, we choose to generate 100 data samples belonging to two half circles, in combination with some Gaussian noise with a standard deviation of `0.25`:

```
In [1]: from sklearn.datasets import make_moons...      x, y =
make_moons(n_samples=100, noise=0.25, ...
random_state=100)
```

We can visualize this data using matplotlib and the `scatter`

Implementing our first random forest

In OpenCV, random forests can be built using the `RTrees_create` function from the `m1` module:

```
In [7]: import cv2
...      rtree = cv2.m1.RTrees_create()
```

The tree object provides a number of options, the most important of which are the following:

- `setMaxDepth`: This sets the maximum possible depth of each tree in the ensemble. The actual obtained depth may be smaller if other termination criteria are met first.
- `setMinSampleCount`: This sets the minimum number of samples that a node can contain for it to get split.
- `setMaxCategories`: This sets the maximum number of categories allowed. Setting the number of categories to a smaller value than the actual number of classes in the data leads to subset estimation.
- `setTermCriteria`: This sets the termination criteria of the algorithm. This is also where you set the number of trees in the forest.

Although we might have hoped for a `setNumTrees` method to set the number of trees in the forest (kind of the most important parameter of them all, no?), we instead need to rely on the `setTermCriteria` method. Confusingly, the number of trees is conflated with `cv2.TERM_CRITERIA_MAX_ITER`, which is usually reserved for the number of iterations that an algorithm is run for, not for the number of estimators in an ensemble.

We can specify the number of trees in the forest by passing an integer, `n_trees`, to the `setTermCriteria` method. Here, we also want to tell the algorithm to quit once the score does not increase by at least `eps` from one iteration to the next:

```
In [8]: n_trees = 10
...
    eps = 0.01
...
    criteria = (cv2.TERM_CRITERIA_MAX_ITER + cv2.TERM_CRITERIA_EPS,
...
                n_trees, eps)
...
    rtree.setTermCriteria(criteria)
```

Then, we are ready to train the classifier on the data from the preceding code:

```
In [9]: rtree.train(X_train.astype(np.float32), cv2.ml.ROW_SAMPLE,
                  y_train);
```

The test labels can be predicted with the `predict` method:

```
In [10]: _, y_hat = rtree.predict(X_test.astype(np.float32))
```

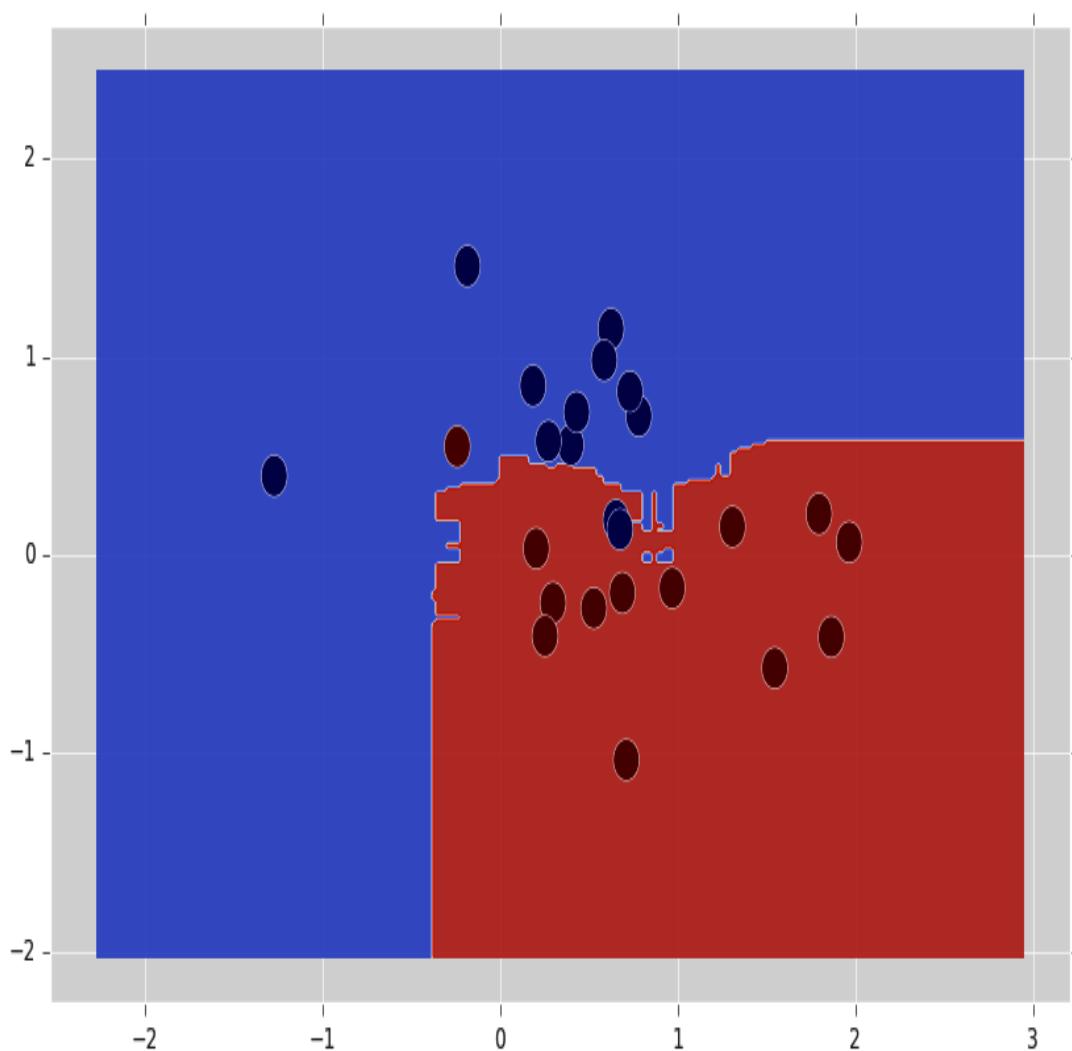
Using scikit-learn's `accuracy_score`, we can evaluate the model on the test set:

```
In [11]: from sklearn.metrics import accuracy_score
...
    accuracy_score(y_test, y_hat)
Out[11]: 0.8399999999999997
```

After training, we can pass the predicted labels to the `plot_decision_boundary` function:

```
In [12]: plot_decision_boundary(rtree, X_test, y_test)
```

This will produce the following plot:



The preceding image shows the decision landscape of a random forest classifier.

Implementing a random forest with scikit-learn

Alternatively, we can implement random forests using scikit-learn:

```
In [13]: from sklearn.ensemble import RandomForestClassifier...
forest = RandomForestClassifier(n_estimators=10, random_state=200)
```

Here, we have a number of options to customize the ensemble:

- `n_estimators`: This specifies the number of trees in the forest.
- `criterion`: This specifies the node-splitting criterion. Setting `criterion='gini'` implements the Gini impurity, whereas setting `criterion='entropy'` implements information gain.
- `max_features`: This specifies the number (or fraction) of features to consider at each node split.
- `max_depth`: This specifies the maximum depth of each tree.
- `min_samples`: This specifies the minimum number ...

Implementing extremely randomized trees

Random forests are already pretty arbitrary. But what if we wanted to take the randomness to its extreme?

In extremely randomized trees (see the `ExtraTreesClassifier` and `ExtraTreesRegressor` classes), the randomness is taken even further than in random forests. Remember how decision trees usually choose a threshold for every feature so that the purity of the node split is maximized? Extremely randomized trees, on the other hand, choose these thresholds at random. The best one of these randomly generated thresholds is then used as the splitting rule.

We can build an extremely randomized tree as follows:

```
In [16]: from sklearn.ensemble import ExtraTreesClassifier  
...      extra_tree = ExtraTreesClassifier(n_estimators=10,  
random_state=100)
```

To illustrate the difference between a single decision tree, a random forest, and extremely randomized trees, let's consider a simple dataset, such as the Iris dataset:

```
In [17]: from sklearn.datasets import load_iris  
...      iris = load_iris()  
...      X = iris.data[:, [0, 2]]  
...      y = iris.target  
In [18]: X_train, X_test, y_train, y_test = train_test_split(  
...          X, y, random_state=100  
...      )
```

We can then fit and score the tree object the same way we did before:

```
In [19]: extra_tree.fit(X_train, y_train)
...       extra_tree.score(X_test, y_test)
Out[19]: 0.92105263157894735
```

For comparison, using a random forest would have resulted in the same performance:

```
In [20]: forest = RandomForestClassifier(n_estimators=10,
                                         random_state=100)
...       forest.fit(X_train, y_train)
...       forest.score(X_test, y_test)
Out[20]: 0.92105263157894735
```

In fact, the same is true for a single tree:

```
In [21]: tree = DecisionTreeClassifier()
...       tree.fit(X_train, y_train)
...       tree.score(X_test, y_test)
Out[21]: 0.92105263157894735
```

So, what's the difference between them? To answer this question, we have to look at the decision boundaries. Fortunately, we have already imported our `plot_decision_boundary` helper function in the preceding section, so all we need to do is pass the different classifier objects to it.

We will build a list of classifiers, where each entry in the list is a tuple that contains an index, a name for the classifier, and the classifier object:

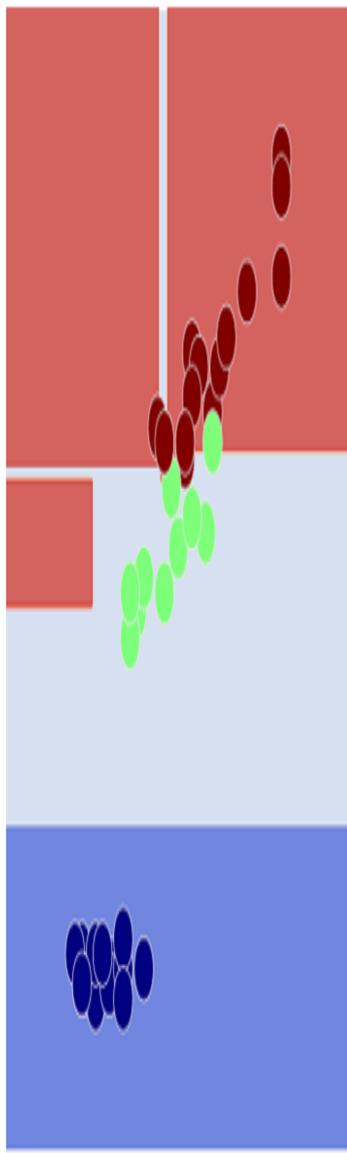
```
In [22]: classifiers = [
...       (1, 'decision tree', tree),
...       (2, 'random forest', forest),
...       (3, 'extremely randomized trees', extra_tree)
...     ]
```

Then, it's easy to pass the list of classifiers to our helper function such that the decision landscape of every classifier is drawn in its own subplot:

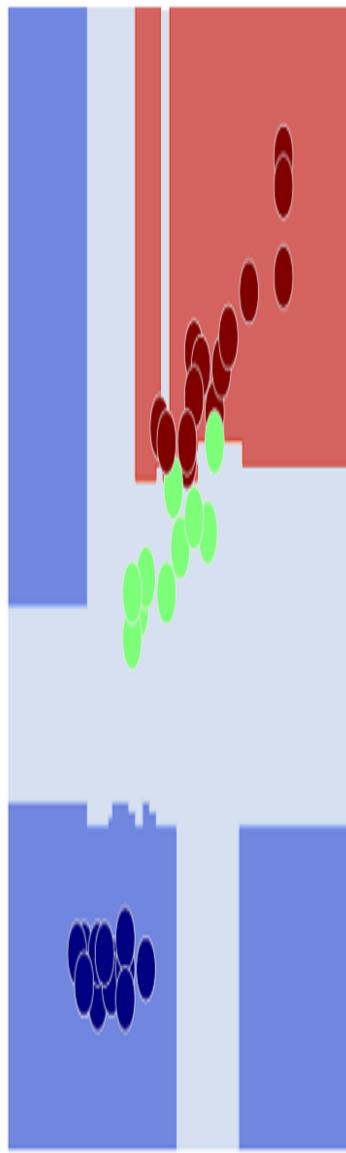
```
In [23]: for sp, name, model in classifiers:  
...     plt.subplot(1, 3, sp)  
...     plot_decision_boundary(model, X_test, y_test)  
...     plt.title(name)  
...     plt.axis('off')
```

The result looks like this:

decision tree



random forest



extremely randomized trees



Now the differences between the three classifiers become clearer. We see the single tree drawing by far the simplest decision boundaries, splitting the landscape using horizontal decision boundaries. The random forest is able to more clearly separate the cloud of data points in the lower-left of the decision landscape. However, only extremely

randomized trees were able to corner the cloud of data points toward the center of the landscape from all sides.

Now that we know about all the different variations of tree ensembles, let's move on to a real-world dataset.

Using random forests for face recognition

A popular dataset that we haven't talked much about yet is the Olivetti faces dataset.

The Olivetti face dataset was collected in 1990 by AT&T Laboratories Cambridge. The dataset comprises facial images of 40 distinct subjects, taken at different times and under different lighting conditions. In addition, subjects varied their facial expressions (open/closed eyes, smiling/not smiling) and their facial details (glasses/no glasses).

Images were then quantized into 256 grayscale levels and stored as unsigned 8-bit integers. Because there are 40 distinct subjects, the dataset comes with 40 distinct target labels. Recognizing faces thus constitutes an example of a multiclass classification task.

Loading the dataset

Like many other classic datasets, the Olivetti face dataset can be loaded using scikit-learn:

```
In [1]: from sklearn.datasets import fetch_olivetti_faces  
...      dataset = fetch_olivetti_faces()  
In [2]: X = dataset.data  
...      y = dataset.target
```

Although the original images consisted of 92 x 112 pixel images, the version available through scikit-learn contains images downscaled to 64 x 64 pixels.

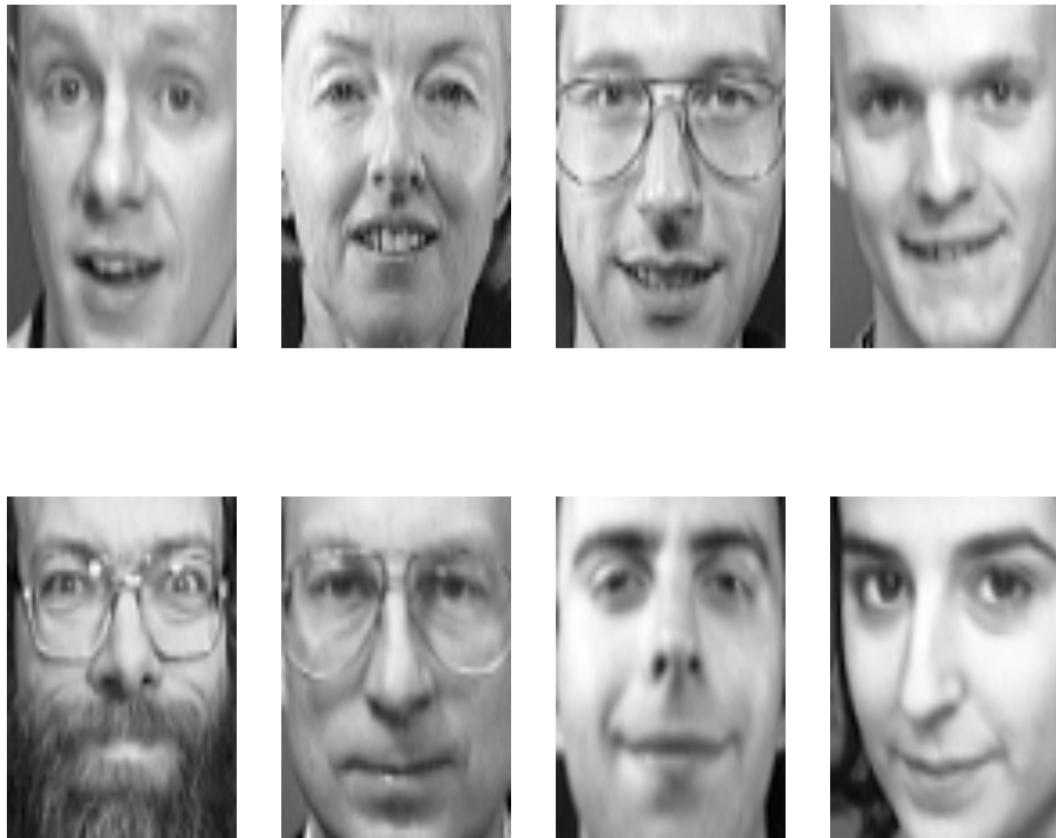
To get a sense of the dataset, we can plot some example images. Let's pick eight indices from the dataset in random order:

```
In [3]: import numpy as np  
...      np.random.seed(21)  
...      idx_rand = np.random.randint(len(X), size=8)
```

We can plot these example images using matplotlib, but we need to make sure we reshape the column vectors to 64 x 64-pixel images before plotting:

```
In [4]: import matplotlib.pyplot as plt  
...      %matplotlib inline  
...      for p, i in enumerate(idx_rand):  
...          plt.subplot(2, 4, p + 1)  
...          plt.imshow(X[i, :].reshape((64, 64)), cmap='gray')  
...          plt.axis('off')
```

The preceding code produces the following output:



You can see how all the faces are taken against a dark background and are portrait. The facial expressions vary drastically from image to image, making this an interesting classification problem. Try not to laugh at some of them!

Preprocessing the dataset

Before we can pass the dataset to the classifier, we need to preprocess it following the best practices from [Chapter 4, Representing Data and Engineering Features](#).

Specifically, we want to make sure that all example images have the same mean grayscale level:

```
In [5]: n_samples, n_features = x.shape[:2]...      x -= x.mean(axis=0)
```

We repeat this procedure for every image to make sure the feature values of every data point (that is, a row in x) are centered around zero:

```
In [6]: x -= x.mean(axis=1).reshape(n_samples, -1)
```

The preprocessed data can be visualized using the following code:

```
In [7]: for p, i in enumerate(idx_rand):...      plt.subplot(2, 4, p + 1)...      plt.imshow(x[i, :].reshape((64, 64)), cmap='gray')...      plt.axis('off') ...
```

Training and testing the random forest

We continue to follow our best practice to split the data into training and test sets:

```
In [8]: from sklearn.model_selection import train_test_split  
...     X_train, X_test, y_train, y_test = train_test_split(  
...         X, y, random_state=21  
...     )
```

Then, we are ready to apply a random forest to the data:

```
In [9]: import cv2  
...     rtree = cv2.ml.RTrees_create()
```

Here, we want to create an ensemble with 50 decision trees:

```
In [10]: n_trees = 50  
...     eps = 0.01  
...     criteria = (cv2.TERM_CRITERIA_MAX_ITER +  
cv2.TERM_CRITERIA_EPS,  
...                 n_trees, eps)  
...     rtree.setTermCriteria(criteria)
```

Because we have a large number of categories (that is, 40), we want to make sure the random forest is set up to handle them accordingly:

```
In [10]: rtree.setMaxCategories(len(np.unique(y)))
```

We can play with other optional arguments, such as the number of data points required in a node before it can be split:

```
In [11]: rtree.setMinSampleCount(2)
```

However, we might not want to limit the depth of each tree. This is again a parameter we will have to experiment with in the end. But for now, let's set it to a large integer value, making the depth effectively unconstrained:

```
In [12]: rtree.setMaxDepth(1000)
```

Then, we can fit the classifier to the training data:

```
In [13]: rtree.train(X_train, cv2.ml.ROW_SAMPLE, y_train);
```

We can check the resulting depth of the tree using the following function:

```
In [13]: rtree.getMaxDepth()  
Out[13]: 25
```

This means that although we allowed the tree to go up to depth 1,000, in the end, only 25 layers were needed.

The evaluation of the classifier is done once again by predicting the labels first (`y_hat`) and then passing them to the `accuracy_score` function:

```
In [14]: _, y_hat = tree.predict(X_test)  
In [15]: from sklearn.metrics import accuracy_score  
...       accuracy_score(y_test, y_hat)  
Out[15]: 0.87
```

We find 87% accuracy, which turns out to be much better than with a single decision tree:

```
In [16]: from sklearn.tree import DecisionTreeClassifier  
...      tree = DecisionTreeClassifier(random_state=21, max_depth=25)  
...      tree.fit(X_train, y_train)  
...      tree.score(X_test, y_test)  
Out[16]: 0.4699999999999997
```

Not bad! We can play with the optional parameters to see whether we get better. The most important one seems to be the number of trees in the forest. We can repeat the experiment with a forest made from 1,000 trees instead of 50 trees:

```
In [18]: num_trees = 1000  
... eps = 0.01  
... criteria = (cv2.TERM_CRITERIA_MAX_ITER + cv2.TERM_CRITERIA_EPS,  
... num_trees, eps)  
... rtree.setTermCriteria(criteria)  
... rtree.train(X_train, cv2.ml.ROW_SAMPLE, y_train);  
... _, y_hat = rtree.predict(X_test)  
... accuracy_score(y_test, y_hat)  
Out[18]: 0.94
```

With this configuration, we get 94% accuracy!

Here, we tried to improve the performance of our model through creative trial and error: we varied some of the parameters we deemed important and observed the resulting change in performance until we found a configuration that satisfied our expectations. We will learn more sophisticated techniques for improving a model in [Chapter 11](#), Selecting the Right Model with Hyperparameter Tuning.

Another interesting use case of decision tree ensembles is AdaBoost.

Implementing AdaBoost

When the trees in the forest are trees of depth 1 (also known as **decision stumps**) and we perform boosting instead of bagging, the resulting algorithm is called **AdaBoost**.

AdaBoost adjusts the dataset at each iteration by performing the following actions:

- Selecting a decision stump
- Increasing the weighting of cases that the decision stump labeled incorrectly while reducing the weighting of correctly labeled cases

This iterative weight adjustment causes each new classifier in the ensemble to prioritize training the incorrectly labeled cases. As a result, the model adjusts by targeting highly weighted data points.

Eventually, the stumps are combined to form a final classifier.

Implementing AdaBoost in OpenCV

Although OpenCV provides a very efficient implementation of AdaBoost, it is hidden under the Haar cascade classifier. Haar cascade classifiers are a very popular tool for face detection, which we can illustrate through the example of the Lena image:

```
In [1]: img_bgr = cv2.imread('data/lena.jpg', cv2.IMREAD_COLOR)
...      img_gray = cv2.cvtColor(img_bgr, cv2.COLOR_BGR2GRAY)
```

After loading the image in both color and grayscale, we load a pretrained Haar cascade:

```
In [2]: import cv2
...      filename = 'data/haarcascade_frontalface_default.xml'
...      face_cascade = cv2.CascadeClassifier(filename)
```

The classifier will then detect faces present in the image using the following function call:

```
In [3]: faces = face_cascade.detectMultiScale(img_gray, 1.1, 5)
```

Note that the algorithm operates only on grayscale images. That's why we saved two pictures of Lena, one to which we can apply the classifier (`img_gray`), and one on which we can draw the resulting bounding box (`img_bgr`):

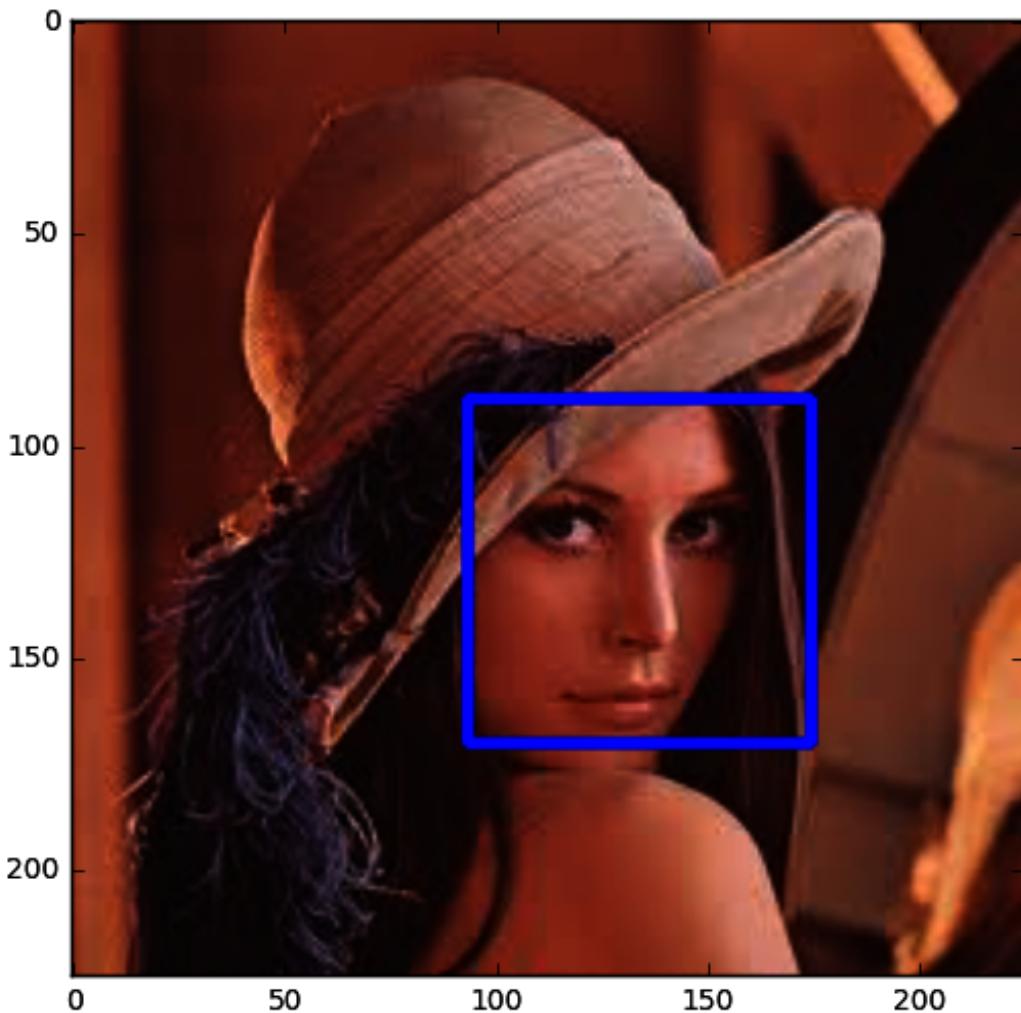
```
In [4]: color = (255, 0, 0)
...      thickness = 2
```

```
...     for (x, y, w, h) in faces:  
...         cv2.rectangle(img_bgr, (x, y), (x + w, y + h),  
...                         color, thickness)
```

Then, we can plot the image using the following code:

```
In [5]: import matplotlib.pyplot as plt  
...     %matplotlib inline  
...     plt.imshow(cv2.cvtColor(img_bgr, cv2.COLOR_BGR2RGB));
```

This results in the following output, with the location of the face indicated by a blue bounding box:



Obviously, this screenshot contains only a single face. However, the preceding code will work even on images where multiple faces could be detected. Try it out!

Implementing AdaBoost in scikit-learn

In scikit-learn, AdaBoost is just another ensemble estimator. We can create an ensemble from 50 decision stumps, as follows:

```
In [6]: from sklearn.ensemble import AdaBoostClassifier...      ada =
AdaBoostClassifier(n_estimators=50, ...
random_state=456)
```

We can load the breast cancer set once more and split it 75-25:

```
In [7]: from sklearn.datasets import load_breast_cancer...      cancer =
load_breast_cancer()...      X = cancer.data...      y = cancer.targetIn
[8]: from sklearn.model_selection import train_test_split...
X_train, X_test, y_train, y_test = train_test_split(...           X, y,
random_state=456...       )
```

Then, `fit` and `score` AdaBoost using the familiar procedure:

```
In [9]: ada.fit(X_train, y_train)...      ada.score(X_test, y_test)
```

Combining different models into a voting classifier

So far, we have seen how to combine different instances of the same classifier or regressor into an ensemble. In this chapter, we are going to take this idea a step further and combine conceptually different classifiers into what is known as a **voting classifier**.

The idea behind voting classifiers is that the individual learners in the ensemble don't necessarily need to be of the same type. After all, no matter how the individual classifiers arrived at their prediction, in the end, we are going to apply a decision rule that integrates all the votes of the individual classifiers. This is also known as a **voting scheme**.

Understanding different voting schemes

Two different voting schemes are common among voting classifiers:

- In **hard voting** (also known as **majority voting**), every individual classifier votes for a class, and the majority wins. In statistical terms, the predicted target label of the ensemble is the mode of the distribution of individually predicted labels.
- In **soft voting**, every individual classifier provides a probability value that a specific data point belongs to a particular target class. The predictions are weighted by the classifier's importance and summed up. Then, the target label with the greatest sum of weighted probabilities wins the vote.

For example, let's assume we have three different classifiers in the ensemble that perform a ...

Implementing a voting classifier

Let's look at a simple example of a voting classifier that combines three different algorithms:

- A logistic regression classifier from [chapter 3, *First Steps in Supervised Learning*](#)
- A Gaussian Naive Bayes classifier from [chapter 7, *Implementing a Spam Filter with Bayesian Learning*](#)
- A random forest classifier from this chapter

We can combine these three algorithms into a voting classifier and apply it to the breast cancer dataset with the following steps:

1. Load the dataset, and split it into training and test sets:

```
In [1]: from sklearn.datasets import load_breast_cancer  
...      cancer = load_breast_cancer()  
...      X = cancer.data  
...      y = cancer.target  
In [2]: from sklearn.model_selection import train_test_split  
...      X_train, X_test, y_train, y_test = train_test_split(X,  
y, random_state=13)
```

2. Instantiate the individual classifiers:

```
In [3]: from sklearn.linear_model import LogisticRegression  
...      model1 = LogisticRegression(random_state=13)  
In [4]: from sklearn.naive_bayes import GaussianNB  
...      model2 = GaussianNB()  
In [5]: from sklearn.ensemble import RandomForestClassifier  
...      model3 = RandomForestClassifier(random_state=13)
```

3. Assign the individual classifiers to the voting ensemble.

Here, we need to pass a list of tuples (`estimators`), where every tuple consists of the name of the classifier (a string of letters depicting a short name of each classifier) and the model object. The voting scheme can be either `voting='hard'` or `voting='soft'`. For now, we will choose `voting='hard'`:

```
In [6]: from sklearn.ensemble import VotingClassifier  
...      vote = VotingClassifier(estimators=[('lr', model1),  
...                                         ('gnb', model2), ('rfc',  
model3)], voting='hard')
```

4. Fit the ensemble to the training data and score it on the test data:

```
In [7]: vote.fit(X_train, y_train)  
...      vote.score(X_test, y_test)  
Out[7]: 0.95104895104895104
```

In order to convince us that 95.1% is a great accuracy score, we can compare the ensemble's performance to the theoretical performance of each individual classifier. We do this by fitting the individual classifiers to the data. Then, we will see that the logistic regression model achieves 94.4% accuracy on its own:

```
In [8]: model1.fit(X_train, y_train)
...      model1.score(X_test, y_test)
Out[8]: 0.94405594405594406
```

Similarly, the Naive Bayes classifier achieves 93.0% accuracy:

```
In [9]: model2.fit(X_train, y_train)
...      model2.score(X_test, y_test)
Out[9]: 0.93006993006993011
```

Last but not least, the random forest classifier also achieved 94.4% accuracy:

```
In [10]: model3.fit(X_train, y_train)
...      model3.score(X_test, y_test)
Out[10]: 0.94405594405594406
```

All in all, we were just able to gain a good percentage in performance by combining three unrelated classifiers into an ensemble. Each of these classifiers might have made different mistakes on the training set, but that's OK because, on average, we need just two out of three classifiers to be correct.

Plurality

In the previous sections, we discussed ensemble methods. What we didn't mention earlier was how the results are aggregated across the individual models prepared by the ensemble techniques. The concept that is used for this is called **plurality**, which is nothing but voting. The higher the vote a class gets, the higher the chances of it being the final class. Imagine if we had three models prepared during ensemble techniques and 10 possible classes (think of them as digits from 0 to 9). Each model would choose one class based on the highest probability it obtained. Finally, the class with the maximum number of votes would be selected. This is the concept of plurality. In practice, plurality tries to bring benefit to both k -NN and Naive ...

Summary

In this chapter, we talked about how to improve various classifiers by combining them into an ensemble. We discussed how to average the predictions of different classifiers using bagging and how to have different classifiers correct each other's mistakes using boosting. A lot of time was spent discussing all possible ways to combine decision trees, be it decision stumps (AdaBoost), random forests, or extremely randomized trees. Finally, we learned how to combine even different types of classifiers in an ensemble by building a voting classifier.

In the next chapter, we will talk more about how to compare the results of different classifiers by diving into the world of model selection and hyperparameter tuning.

Selecting the Right Model with Hyperparameter Tuning

Now that we have explored a wide variety of machine learning algorithms, I am sure you have realized that most of them come with a great number of settings to choose from. These settings or tuning knobs, the so-called **hyperparameters**, help us to control the behavior of the algorithm when we try to maximize performance.

For example, we might want to choose the depth or split criterion in a decision tree or tune the number of neurons in a neural network. Finding the values of important parameters of a model is a tricky task but necessary for almost all models and datasets.

In this chapter, we will dive deeper into **model evaluation** and **hyperparameter tuning**. Assume that we have two different ...

Technical requirements

You can refer to the code for this chapter from the following link: <https://github.com/PacktPublishing/Machine-Learning-for-OpenCV-Second-Edition/tree/master/Chapter11>.

Here is a summary of the software and hardware requirements:

- You will need OpenCV version 4.1.x (4.1.0 or 4.1.1 will both work just fine).
- You will need Python version 3.6 (any Python version 3.x will be fine).
- You will need Anaconda Python 3 for installing Python and the required modules.
- You can use any operating system—macOS, Windows, and Linux-based OSes along with this book. We recommend you have at least 4 GB RAM in your system.
- You don't need to have a GPU to run the code provided along with this book.

Evaluating a model

Model evaluation strategies come in many different forms and shapes. In the following sections, we will, therefore, highlight three of the most commonly used techniques to compare models against each other:

- k-fold cross-validation
- Bootstrapping
- McNemar's test

In principle, model evaluation is simple: after training a model on some data, we can estimate its effectiveness by comparing model predictions to some ground truth values. We learned early on that we should split the data into training and test sets, and we tried to follow this instruction whenever possible. But why exactly did we do that again?

Evaluating a model the wrong way

The reason we never evaluate a model on the training set is that, in principle, any dataset can be learned if we throw a strong enough model at it.

A quick demonstration of this can be given with the help of the Iris dataset, which we talked about extensively in [chapter 3, First Steps in Supervised Learning](#). There, the goal was to classify species of iris flowers based on their physical dimensions. We can load the Iris dataset using scikit-learn:

```
In [1]: from sklearn.datasets import load_iris  
...      iris = load_iris()
```

An innocent approach to this problem would be to store all data points in the matrix, x , and all class labels in the vector, y :

```
In [2]: import numpy as np  
...      X = iris.data.astype(np.float32)  
...      y = iris.target
```

Next, we choose a model and its hyperparameters. For example, let's use the algorithm from [chapter 3, First Steps in Supervised Learning](#), which provides only a single hyperparameter: the number of neighbors, k . With $k=1$, we get a very simple model that classifies the label of an unknown point as belonging to the same class as its closest neighbor.

In the following steps, you will learn how to build a **k-nearest neighbour (k-NN)** and compute its accuracy:

1. In OpenCV, kNN instantiates as follows:

```
In [3]: import cv2  
...      knn = cv2.ml.KNearest_create()  
...      knn.setDefaultK(1)
```

2. Then, we train the model and use it to predict labels for the data that we already know:

```
In [4]: knn.train(X, cv2.ml.ROW_SAMPLE, y)  
...      _, y_hat = knn.predict(X)
```

3. Finally, we compute the fraction of the correctly labeled points:

```
In [5]: from sklearn.metrics import accuracy_score  
...      accuracy_score(y, y_hat)  
Out[5]: 1.0
```

As we can see, the accuracy score is `1.0`, it indicates that 100% of points were correctly labeled by our model.

If a model gets 100% accuracy on the training set, we say the model memorized the data.

But is the expected accuracy truly being measured? Have we come up with a model that we expect to be correct 100% of the time?

As you may have gathered, the answer is no. This example shows that even a simple algorithm is capable of memorizing a real-world dataset. Imagine how easy this task would have been for a deep neural network! Usually, the more parameters a model has, the more powerful it is. We will come back to this shortly.

Evaluating a model in the right way

A better sense of a model's performance can be found using what's known as a test set, but you already knew this. When presented with data held out from the training procedure, we can check whether a model has learned some dependencies in the data that hold across the board or whether it just memorized the training set.

We can split the data into training and test sets using the familiar `train_test_split` from scikit-learn's `model_selection` module:

```
In [6]: from sklearn.model_selection import train_test_split
```

But how do we choose the right train-test ratio? Is there even such a thing as a right ratio? Or is this considered another hyperparameter of the model?

There are two competing concerns here:

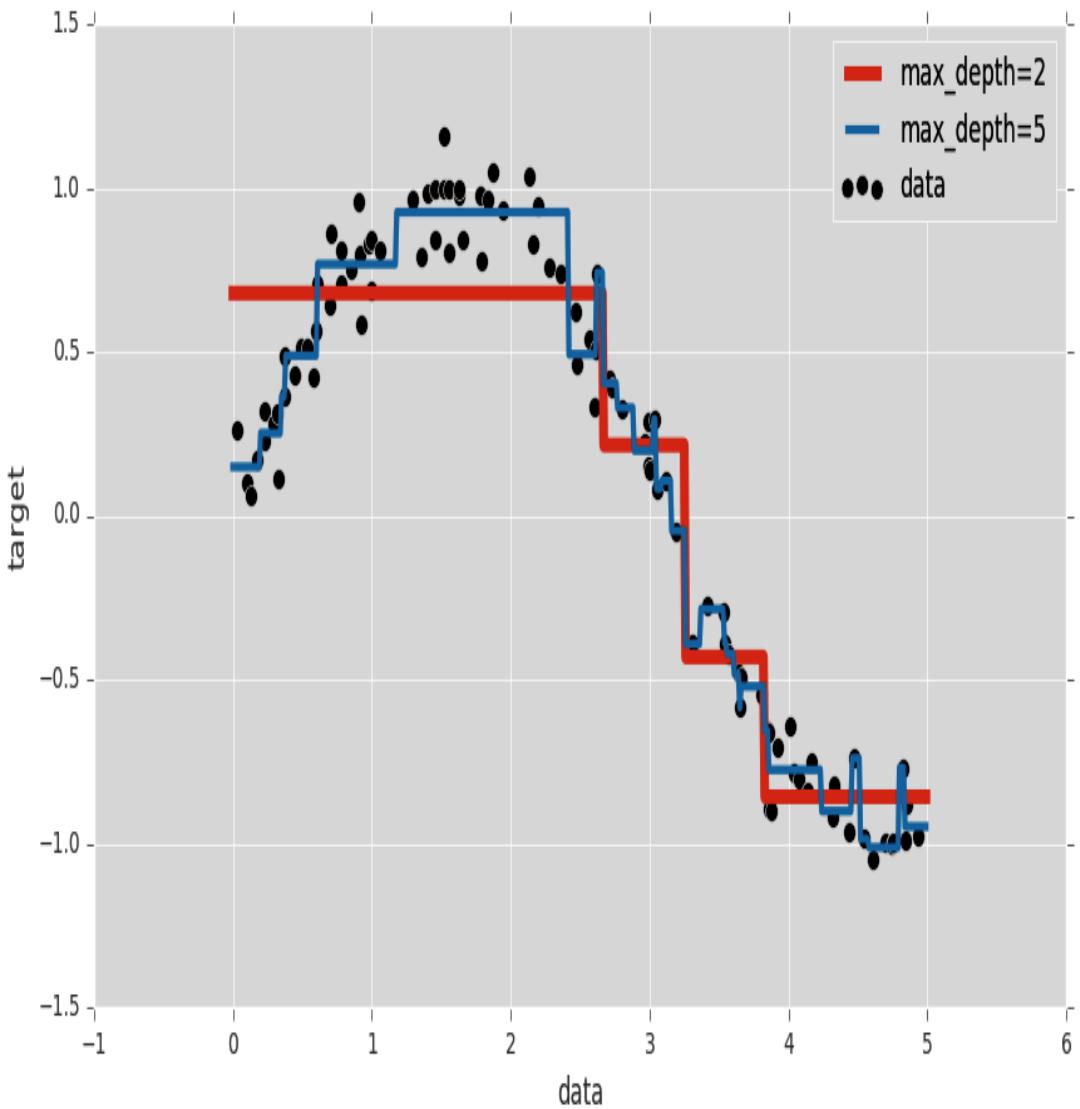
- If our ...

Selecting the best model

When a model is under-performing, it is often not clear how to make it better. Throughout this book, I have declared a rule of thumbs, for example, how to select the number of layers in a neural network. Even worse, the answer is often counter-intuitive! For example, adding another layer to the network might make the results worse, and adding more training data might not change performance at all.

You can see why these issues are some of the most important aspects of machine learning. At the end of the day, the ability to determine what steps will or will not improve our model is what separates the successful machine learning practitioner from all others.

Let's have a look at a specific example. Remember [Chapter 5, Using Decision Trees to Make a Medical Diagnosis](#), where we used decision trees in a regression task? We were fitting two different trees to a sin function—one with depth 2 and one with depth 5. As a reminder, the regression result looked like this:



It should be clear that neither of these fits are particularly good. However, the two decision trees fail in two different ways!

The decision tree with depth 2 (thick line in the preceding screenshot) attempts to fit four straight lines through the data. Because the data is intrinsically more complicated than a few straight lines, this model fails. We could train it as much as we wanted, on as many training samples as we could generate—it would never be able to describe this dataset well. Such a model is said to underfit the data. In other words,

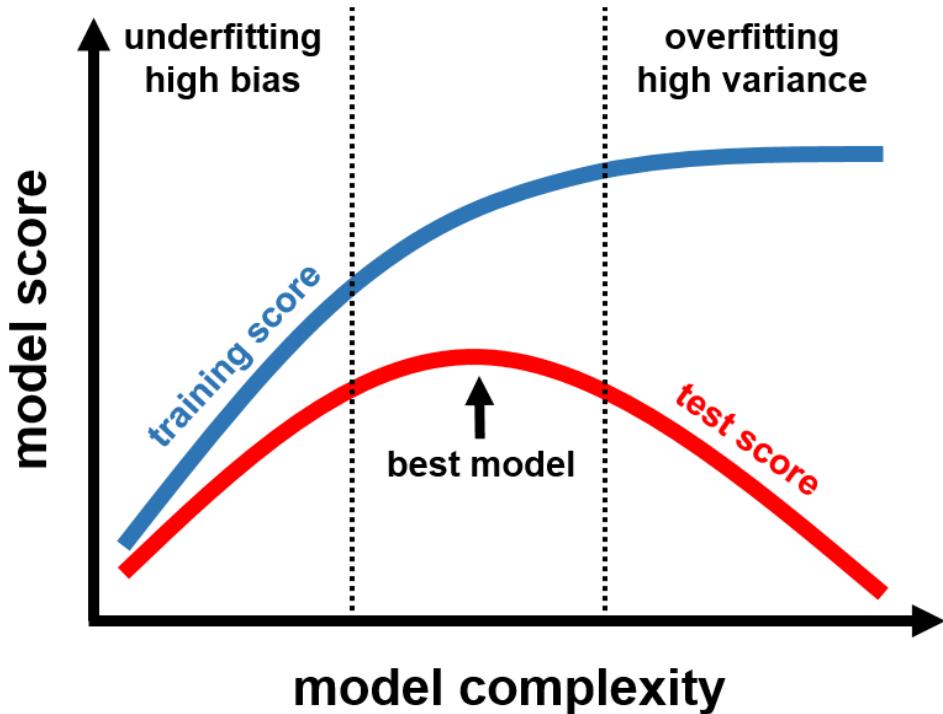
the model does not have enough complexity to account for all the features in the data. Hence, the model has a high bias.

The other decision tree (thin line, depth 5) makes a different mistake. This model has enough flexibility to nearly perfectly account for the fine structures in the data. However, at some points, the model seems to follow the particular pattern of the noise; we added to the sin function rather than the sin function itself. You can see that on the right-hand side of the graph, where the blue curve (thin line) would jitter a lot. Such a model is said to overfit the data. In other words, the model is so complex that it ends up accounting for random errors in the data. Hence, the model has a high variance.

Long story short—here's the secret sauce: fundamentally, selecting the right model comes down to finding a sweet spot in the trade-off between bias and variance.

*The amount of flexibility a model has (also known as the **model complexity**) is mostly dictated by its hyperparameters. That is why it is so important to tune them!*

Let's return to the kNN algorithm and the Iris dataset. If we repeated the procedure of fitting the model to the Iris data for all possible values of k and calculated both training and test scores, we would expect the result to look something like the following:



The preceding image shows model score as a function of model complexity. If there is one thing I would want you to remember from this chapter, it would be this diagram. Let's unpack it.

The diagram describes the model score (either training or test scores) as a function of model complexity. As mentioned in the preceding diagram, the model complexity of a neural network roughly grows with the number of neurons in the network. In the case of kNN, the opposite logic applies—the larger the value for k , the smoother the decision boundary, and hence, the lower the complexity. In other words, kNN with $k=1$ would be all of the way to the right in the preceding diagram, where the training score is perfect. No wonder we got 100% accuracy on the training set!

From the preceding diagram, we can gather that there are three regimes in the model complexity landscape:

- Very low model complexity (a high-bias model) underfits the training data. In this regime, the model achieves only low scores on both the training and test set, no matter how long we trained it for.
- A model with very high complexity (or a high-variance) overfits the training data, which indicates that the model can predict on the training data very well but fails on the unseen data. In this regime, the model has started to learn intricacies or peculiarities that only appear in the training data. Since these peculiarities do not apply to unseen data, the training score gets lower and lower.
- For some intermediate value, the test score is maximal. It is this intermediate regime, where the test score is maximal, that we are trying to find. This is the sweet spot in the trade-off between bias and variance!

This means that we can find the best algorithm for the task at hand by mapping out the model complexity landscape. Specifically, we can use the following indicators to know which regime we are currently in:

- If both training and test scores are below our expectations, we are probably in the leftmost regime in the preceding diagram, where the model is underfitting the data. In this case, a good idea might be to increase the model complexity and try again.

- If the training score is much higher than the test score, we are probably in the rightmost regime in the preceding diagram, where the model is overfitting the data. In this case, a good idea might be to decrease the model complexity and try again.

Although this procedure works in general, there are more sophisticated strategies for model evaluation that proved to be more thorough than a simple train-test split, which we will talk about in the following sections.

Understanding cross-validation

Cross-validation is a method of evaluating the generalization performance of a model that is generally more stable and thorough than splitting the dataset into training and test sets.

The most commonly used version of cross-validation is **k-fold cross-validation**, where k is a number specified by the user (usually five or ten). Here, the dataset is partitioned into k parts of more or less equal size, called **folds**. For a dataset that contains N data points, each fold should hence have approximately N / k samples. Then, a series of models are trained on the data, using $k - 1$ folds for training and one remaining fold for testing. The procedure is repeated for k iterations, each time choosing a different fold for ...

Manually implementing cross-validation in OpenCV

The easiest way to perform cross-validation in OpenCV is to do the data splits by hand.

For example, to implement two-fold cross-validation, we would perform the following procedure:

1. Load the dataset:

```
In [1]: from sklearn.datasets import load_iris  
...     import numpy as np  
...     iris = load_iris()  
...     X = iris.data.astype(np.float32)  
...     y = iris.target
```

2. Split the data into two equally sized parts:

```
In [2]: from sklearn.model_selection import model_selection  
...     X_fold1, X_fold2, y_fold1, y_fold2 = train_test_split(  
...             X, y, random_state=37, train_size=0.5  
...         )
```

3. Instantiate the classifier:

```
In [3]: import cv2  
...      knn = cv2.ml.KNearest_create()  
...      knn.setDefaultK(1)
```

4. Train the classifier on the first fold, then predict the labels of the second fold:

```
In [4]: knn.train(X_fold1, cv2.ml.ROW_SAMPLE, y_fold1)  
..., y_hat_fold2 = knn.predict(X_fold2)
```

5. Train the classifier on the second fold, then predict the labels of the first fold:

```
In [5]: knn.train(X_fold2, cv2.ml.ROW_SAMPLE, y_fold2)  
..., y_hat_fold1 = knn.predict(X_fold1)
```

6. Compute accuracy scores for both folds:

```
In [6]: from sklearn.metrics import accuracy_score  
...      accuracy_score(y_fold1, y_hat_fold1)  
Out[6]: 0.9200000000000004  
In [7]: accuracy_score(y_fold2, y_hat_fold2)  
Out[7]: 0.88
```

This procedure will yield two accuracy scores, one for the first fold (92% accuracy) and one for the second fold (88% accuracy). On average, our classifier hence achieved 90% accuracy on unseen data.

Using scikit-learn for k-fold cross-validation

In scikit-learn, cross-validation can be performed in three steps:

1. Load the dataset. Since we already did this earlier, we don't have to do it again.
2. Instantiate the classifier:

```
In [8]: from sklearn.neighbors import KNeighborsClassifier  
...      model = KNeighborsClassifier(n_neighbors=1)
```

3. Perform cross-validation with the `cross_val_score` function.
This function takes as input a model, the full dataset (`x`), the target labels (`y`), and an integer value for the number of folds (`cv`). It is not necessary to split the data by hand—the function will do that automatically depending on the number of folds. After the cross-validation is completed, the function returns the test scores:

```
In [9]: from sklearn.model_selection ...
```

Implementing leave-one-out cross-validation

Another popular way to implement cross-validation is to choose the number of folds equal to the number of data points in the dataset. In other words, if there are N data points, we set $k=N$. This means that we will end up having to do N iterations of cross-validation, but in every iteration, the training set will consist of only a single data point. The advantage of this procedure is that we get to use all-but-one data points for training. Hence, this procedure is also known as **leave-one-out** cross-validation.

In scikit-learn, this functionality is provided by the `LeaveOneOut` method from the `model_selection` module:

```
In [11]: from sklearn.model_selection import LeaveOneOut
```

This object can be passed directly to the `cross_val_score` function in the following way:

```
In [12]: scores = cross_val_score(model, X, y, cv=LeaveOneOut())
```

Because every test set now contains a single data point, we would expect the scorer to return 150 values—one for each data point in the dataset. Each of these points we get could be either right or wrong. Hence, we expect `scores` to be a list of ones (`1`) and zeros (`0`), which corresponds to correct and incorrect classifications, respectively:

```
In [13]: scores
Out[13]: array([ 1.,  1.,  1.,  1.,  1.,  1.,  1.,  1.,  1.,  1.,  1.,  1.,
```

```
1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,  
1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,  
1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,  
1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,  
1., 1., 1., 1., 0., 1., 0., 1., 1., 1., 1., 1., 1.,  
1., 1., 1., 1., 0., 1., 1., 1., 1., 1., 1., 1., 1.,  
1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,  
1., 1., 0., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,  
1., 1., 0., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,  
1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,  
1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.]
```

If we want to know the average performance of the classifier, we would still compute the mean and standard deviation of the scores:

```
In [14]: scores.mean(), scores.std()  
Out[14]: (0.9599999999999996, 0.19595917942265423)
```

We can see this scoring scheme returns very similar results to five-fold cross-validation.

You can learn more about other useful cross-validation procedures at http://scikit-learn.org/stable/modules/cross_validation.html.

Estimating robustness using bootstrapping

An alternative procedure to k-fold cross-validation is **bootstrapping**.

Instead of splitting the data into folds, bootstrapping builds a training set by drawing samples randomly from the dataset. Typically, a bootstrap is formed by drawing samples with replacement. Imagine putting all of the data points into a bag and then drawing randomly from the bag. After drawing a sample, we would put it back in the bag. This allows for some samples to show up multiple times in the training set, which is something cross-validation does not allow.

The classifier is then tested on all samples that are not part of the bootstrap (the so-called **out-of-bag** examples), and the procedure is repeated a large number of times ...

Manually implementing bootstrapping in OpenCV

Bootstrapping can be implemented with the following procedure:

1. Load the dataset. Since we already did this earlier, we don't have to do it again.
2. Instantiate the classifier:

```
In [15]: knn = cv2.ml.KNearest_create()
...
    knn.setDefaultK(1)
```

3. From our dataset with N samples, randomly choose N samples with replacement to form a bootstrap. This can be done most easily with the `choice` function from NumPy's `random` module. We tell the function to draw `len(X)` samples in the `[0, len(X)-1]` range with replacement (`replace=True`). The function then returns a list of indices, from which we form our bootstrap:

```
In [16]: idx_boot = np.random.choice(len(X), size=len(X),
...
    replace=True)
...
    X_boot = X[idx_boot, :]
...
    y_boot = y[idx_boot]
```

4. Put all samples that do not show in the bootstrap in the out-of-bag set:

```
In [17]: idx_oob = np.array([x not in idx_boot  
...      for x in np.arange(len(X))], dtype=np.bool)  
...      X_oob = X[idx_oob, :]  
...      y_oob = y[idx_oob]
```

5. Train the classifier on the bootstrap samples:

```
In [18]: knn.train(X_train, cv2.ml.ROW_SAMPLE, y_boot)  
Out[18]: True
```

6. Test the classifier on the out-of-bag samples:

```
In [19]: _, y_hat = knn.predict(X_oob)  
...      accuracy_score(y_oob, y_hat)  
Out[19]: 0.9285714285714286
```

7. Repeat steps 3-6 for a specific number of iterations.
8. Iteration of the bootstrap. Repeat these steps up to 10,000 times to get 10,000 accuracy scores, then average the scores to get an idea of the classifier's mean performance.

For our convenience, we can build a function from *steps 3-6* so that it is easy to run the procedure for some `n_iter` number of times. We also pass a model (our kNN classifier, `model`), the feature matrix (`x`), and the vector with all class labels (`y`):

```
In [20]: def yield_bootstrap(model, X, y, n_iter=10000):
...     for _ in range(n_iter):
```

The steps within the `for` loop are essentially *steps 3-6* from the code mentioned earlier. This involved training the classifier on the bootstrap and testing it on the out-of-bag examples:

```
...             # train the classifier on bootstrap
...             idx_boot = np.random.choice(len(X), size=len(X),
...                                         replace=True)
...             X_boot = X[idx_boot, :]
...             y_boot = y[idx_boot]
...             knn.train(X_boot, cv2.ml.ROW_SAMPLE, y_boot)
...
...             # test classifier on out-of-bag examples
...             idx_oob = np.array([x not in idx_boot
...                                for x in np.arange(len(X))],
...                                dtype=np.bool)
...             X_oob = X[idx_oob, :]
...             y_oob = y[idx_oob]
...             _, y_hat = knn.predict(X_oob)
```

Then, we need to return the accuracy score. You might expect a `return` statement here. However, a more elegant way is to use the `yield` statement, which turns the function automatically into a generator. This means we don't have to initialize an empty list (`acc = []`) and then append the new accuracy score at each iteration (`acc.append(accuracy_score(...))`). The bookkeeping is done automatically:

```
...             yield accuracy_score(y_oob, y_hat)
```

To make sure we all get the same result, let's fix the seed of the random number generator:

```
In [21]: np.random.seed(42)
```

Now, let's run the procedure for `n_iter=10` times by converting the function output into a list:

```
In [22]: list(yield_bootstrap(knn, X, y, n_iter=10))
Out[22]: [0.9833333333333328,
          0.93650793650793651,
          0.92452830188679247,
          0.92307692307692313,
          0.94545454545454544,
          0.94736842105263153,
          0.98148148148148151,
          0.96078431372549022,
          0.93220338983050843,
          0.96610169491525422]
```

As you can see, for this small sample we get accuracy scores anywhere between 92% and 98%. To get a more reliable estimate of the model's performance, we repeat the procedure 1,000 times and calculate both mean and standard deviation of the resulting scores:

```
In [23]: acc = list(yield_bootstrap(knn, X, y, n_iter=1000))
...      np.mean(acc), np.std(acc)
Out[23]: (0.95524155136419198, 0.022040380995646654)
```

You are always welcome to increase the number of repetitions. But once `n_iter` is large enough, the procedure should be robust to the randomness of the sampling procedure. In this case, we do not expect to see any more changes to the distribution of score values as we keep increasing `n_iter` to, for example, 10,000 iterations:

```
In [24]: acc = list(yield_bootstrap(knn, X, y, n_iter=10000))
...      np.mean(acc), np.std(acc)
Out[24]: (0.95501528733009422, 0.021778543317079499)
```

Typically, the scores obtained with bootstrapping would be used in a *statistical test* to assess the *significance* of our result. Let's have a look at how that is done.

Assessing the significance of our results

Assume for a moment that we implemented the cross-validation procedure for two versions of our kNN classifier. The resulting test scores are—92.34% for Model A and 92.73% for Model B. How do we know which model is better?

Following our logic introduced here, we might argue for Model B because it has a better test score. But what if the two models are not significantly different? These could have two underlying causes, which are both a consequence of the randomness of our testing procedure:

- For all we know, Model B just got lucky. Perhaps we chose a really low k for our cross-validation procedure. Perhaps Model B ended up with a beneficial train-test split so that the model had no problem classifying ...

Implementing Student's t-test

One of the most famous statistical tests is **Student's t-test**. You might have heard of it before: it allows us to determine whether two sets of data are significantly different from one another. This was a really important test for William Sealy Gosset, the inventor of the test, who worked at the Guinness brewery and wanted to know whether two batches of stout differed in quality.

Note that "Student" here is capitalized. Although Gosset wasn't allowed to publish his test due to company policy, he did so anyway under his pen name, Student.

In practice, the t-test allows us to determine whether two data samples come from underlying distributions with the same mean or *expected value*.

For our purposes, this means that we can use the t-test to determine whether the test scores of two independent classifiers have the same mean value. We start by hypothesizing that the two sets of test scores are identical. We call this the *null hypothesis* because this is the hypothesis we want to nullify, that is, we are looking for evidence to *reject* the hypothesis because we want to ensure that one classifier is significantly better than the other.

We accept or reject a null hypothesis based on a parameter known as the **p-value** that the t-test returns. The p-value takes on values between 0 and 1. A p-value of 0.05 would mean that the null hypothesis is right only 5 out of 100 times. A small p-value hence indicates strong evidence that the hypothesis can be safely rejected. It is customary to use $p=0.05$ as a cut-off value below which we reject the null hypothesis.

If this is all too confusing, think of it this way: when we run a t-test to compare classifier test scores, we are looking to obtain a small p-value because that means that the two classifiers give significantly different results.

We can implement Student's t-test with SciPy's `ttest_ind` function from the `stats` module:

```
In [25]: from scipy.stats import ttest_ind
```

Let's start with a simple example. Assume we ran five-fold cross-validation on two classifiers and obtained the following scores:

```
In [26]: scores_a = [1, 1, 1, 1, 1]
...       scores_b = [0, 0, 0, 0, 0]
```

This means that Model A achieved 100% accuracy in all five folds, whereas Model B got 0% accuracy. In this case, it is clear that the two results are significantly different. If we run the t-test on this data, we should hence find a really small p-value:

```
In [27]: ttest_ind(scores_a, scores_b)
Out[27]: Ttest_indResult(statistic=inf, pvalue=0.0)
```

And we do! We actually get the smallest possible p-value, $p=0.0$.

On the other hand, what if the two classifiers got exactly the same numbers, except during different folds? In this case, we would expect the two classifiers to be equivalent, which is indicated by a really large p-value:

```
In [28]: scores_a = [0.9, 0.9, 0.9, 0.8, 0.8]
...       scores_b = [0.8, 0.8, 0.9, 0.9, 0.9]
...       ttest_ind(scores_a, scores_b)
Out[28]: Ttest_indResult(statistic=0.0, pvalue=1.0)
```

Analogous to the aforementioned, we get the largest possible p-value, $p=1.0$.

To see what happens in a more realistic example, let's return to our kNN classifier from an earlier example. Using the test scores obtained from the ten-fold cross-validation procedure, we can compare two different kNN classifiers with the following procedure:

1. Obtain a set of test scores for Model A. We choose Model A to be the kNN classifier from earlier ($k=1$):

```
In [29]: k1 = KNeighborsClassifier(n_neighbors=1)
...
...     scores_k1 = cross_val_score(k1, X, y, cv=10)
...     np.mean(scores_k1), np.std(scores_k1)
Out[29]: (0.9599999999999996, 0.0533333333333323)
```

2. Obtain a set of test scores for Model B. Let's choose Model B to be a kNN classifier with $k=3$:

```
In [30]: k3 = KNeighborsClassifier(n_neighbors=3)
...
...     scores_k3 = cross_val_score(k3, X, y, cv=10)
...     np.mean(scores_k3), np.std(scores_k3)
Out[30]: (0.96666666666666656, 0.044721359549995787)
```

3. Apply the t-test to both sets of scores:

```
In [31]: ttest_ind(scores_k1, scores_k3)
Out[31]: Ttest_indResult(statistic=-0.2873478855663425,
pvalue=0.77712784875052965)
```

As you can see, this is a good example of two classifiers giving different cross-validation scores (96.0% and 96.7%) that turn out to be not significantly different! Because we get a large p-value ($p=0.777$), we expect the two classifiers to be equivalent 77 out of 100 times.

Implementing McNemar's test

A more advanced statistical technique is **McNemar's test**. This test can be used on paired data to determine whether there are any differences between the two samples. As in the case of the t-test, we can use McNemar's test to determine whether two models give significantly different classification results.

McNemar's test operates on pairs of data points. This means that we need to know, for both classifiers, how they classified each data point. Based on the number of data points that the first classifier got right but the second got wrong and vice versa, we can determine whether the two classifiers are equivalent.

Let's assume the preceding Model A and Model B were applied to the same five data points. Whereas Model ...

Tuning hyperparameters with grid search

The most commonly used tool for hyperparameter tuning is *grid search*, which is basically a fancy term for saying we will try all possible parameter combinations with a `for` loop.

Let's have a look at how that is done in practice.

Implementing a simple grid search

Returning to our kNN classifier, we find that we have only one hyperparameter to tune: k . Typically, you would have a much larger number of open parameters to mess with, but the kNN algorithm is simple enough for us to manually implement a grid search.

Before we get started, we need to split the dataset as we have done before into training and test sets:

1. Here we choose a 75-25 split:

```
In [1]: from sklearn.datasets import load_iris...      import
numpy as np...      iris = load_iris()...      X =
iris.data.astype(np.float32)...      y = iris.targetIn [2]:
X_train, X_test, y_train, y_test = train_test_split(...,
X, y, random_state=42...)
```

2. Then, the goal is to loop over all possible values of k . As we do this, we want to keep ...

Understanding the value of a validation set

Following our best practice of splitting the data into training and test sets, we might be tempted to tell people that we have found a model that performs with 97.4% accuracy on the dataset. However, our result might not necessarily generalize to new data. The argument is the same as earlier on in this book when we warranted the train-test split that we need an independent dataset for evaluation.

However, when we implemented a grid search in the last section, we used the test set to evaluate the outcome of the grid search and update the hyperparameter, k . This means we can no longer use the test set to evaluate the final data! Any model choices made based on the test set accuracy would leak information from the test set into the model.

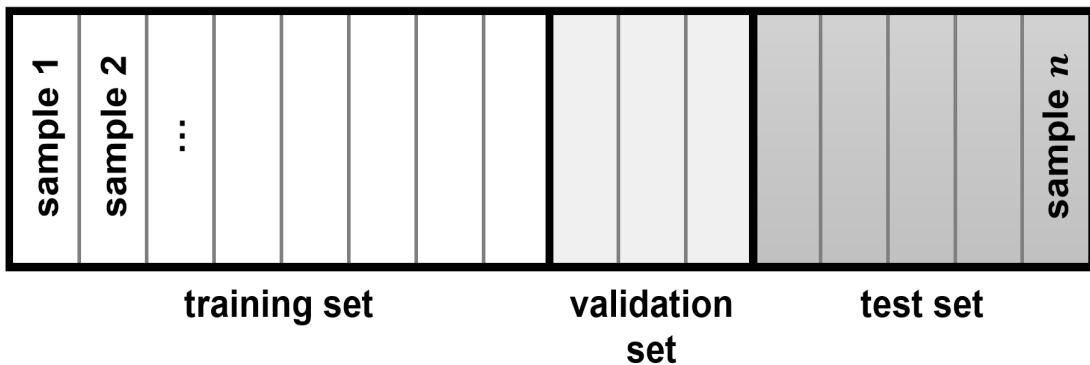
One way to resolve this data is to split the data again and introduce what is known as a **validation set**. The validation set is different from the training and test set and is used exclusively for selecting the best parameters of the model. It is a good practice to do all exploratory analysis and model selection on this validation set and keep a separate test set, which is only used for the final evaluation.

In other words, we should end up splitting the data into three different sets:

- A **training set**, which is used to build the model
- A **validation set**, which is used to select the parameters of the model

- A **test set**, which is used to evaluate the performance of the final model

Such a three-way split is illustrated in the following diagram:



The preceding diagram shows an example of how to split a dataset into training, validation, and test sets. In practice, the three-way split is achieved in two steps:

1. Split the data into two chunks: one that contains training and validation sets and another that contains the test set:

```
In [6]: X_trainval, X_test, y_trainval, y_test =
...         train_test_split(X, y, random_state=37)
In [7]: X_trainval.shape
Out[7]: (112, 4)
```

2. Split `x_trainval` again into proper training and validation sets:

```
In [8]: X_train, X_valid, y_train, y_valid = train_test_split(
...         X_trainval, y_trainval, random_state=37
...     )
```

```
In [9]: X_train.shape  
Out[9]: (84, 4)
```

Then, we repeat the manual grid search from the preceding code, but this time, we will use the validation set to find the best k (see code highlights):

```
In [10]: best_acc = 0.0  
...      best_k = 0  
...      for k in range(1, 20):  
...          knn = cv2.ml.KNearest_create()  
...          knn.setDefaultK(k)  
...          knn.train(X_train, cv2.ml.ROW_SAMPLE, y_train)  
...          _, y_valid_hat = knn.predict(X_valid)  
...          acc = accuracy_score(y_valid, y_valid_hat)  
...          if acc >= best_acc:  
...              best_acc = acc  
...              best_k = k  
...      best_acc, best_k  
Out[10]: (1.0, 7)
```

We now find that a 100% validation score (`best_acc`) can be achieved with $k=7$ (`best_k`)! However, recall that this score might be overly optimistic. To find out how well the model really performs, we need to test it on held-out data from the test set.

To arrive at our final model, we can use the value for k we found during grid search and re-train the model on both the training and validation data. This way, we used as much data as possible to build the model while still honoring the train-test split principle.

This means we should retrain the model on `X_trainval`, which contains both the training and validation sets and score it on the test set:

```
In [25]: knn = cv2.ml.KNearest_create()  
...      knn.setDefaultK(best_k)  
...      knn.train(X_trainval, cv2.ml.ROW_SAMPLE, y_trainval)  
...      _, y_test_hat = knn.predict(X_test)
```

```
| ...      accuracy_score(y_test, y_test_hat), best_k  
| out[25]: (0.94736842105263153, 7)
```

With this procedure, we find a formidable score of 94.7% accuracy on the test set. Because we honored the train-test split principle, we can now be sure that this is the performance we can expect from the classifier when applied to novel data. It is not as high as the 100% accuracy reported during validation, but it is still a very good score!

Combining grid search with cross-validation

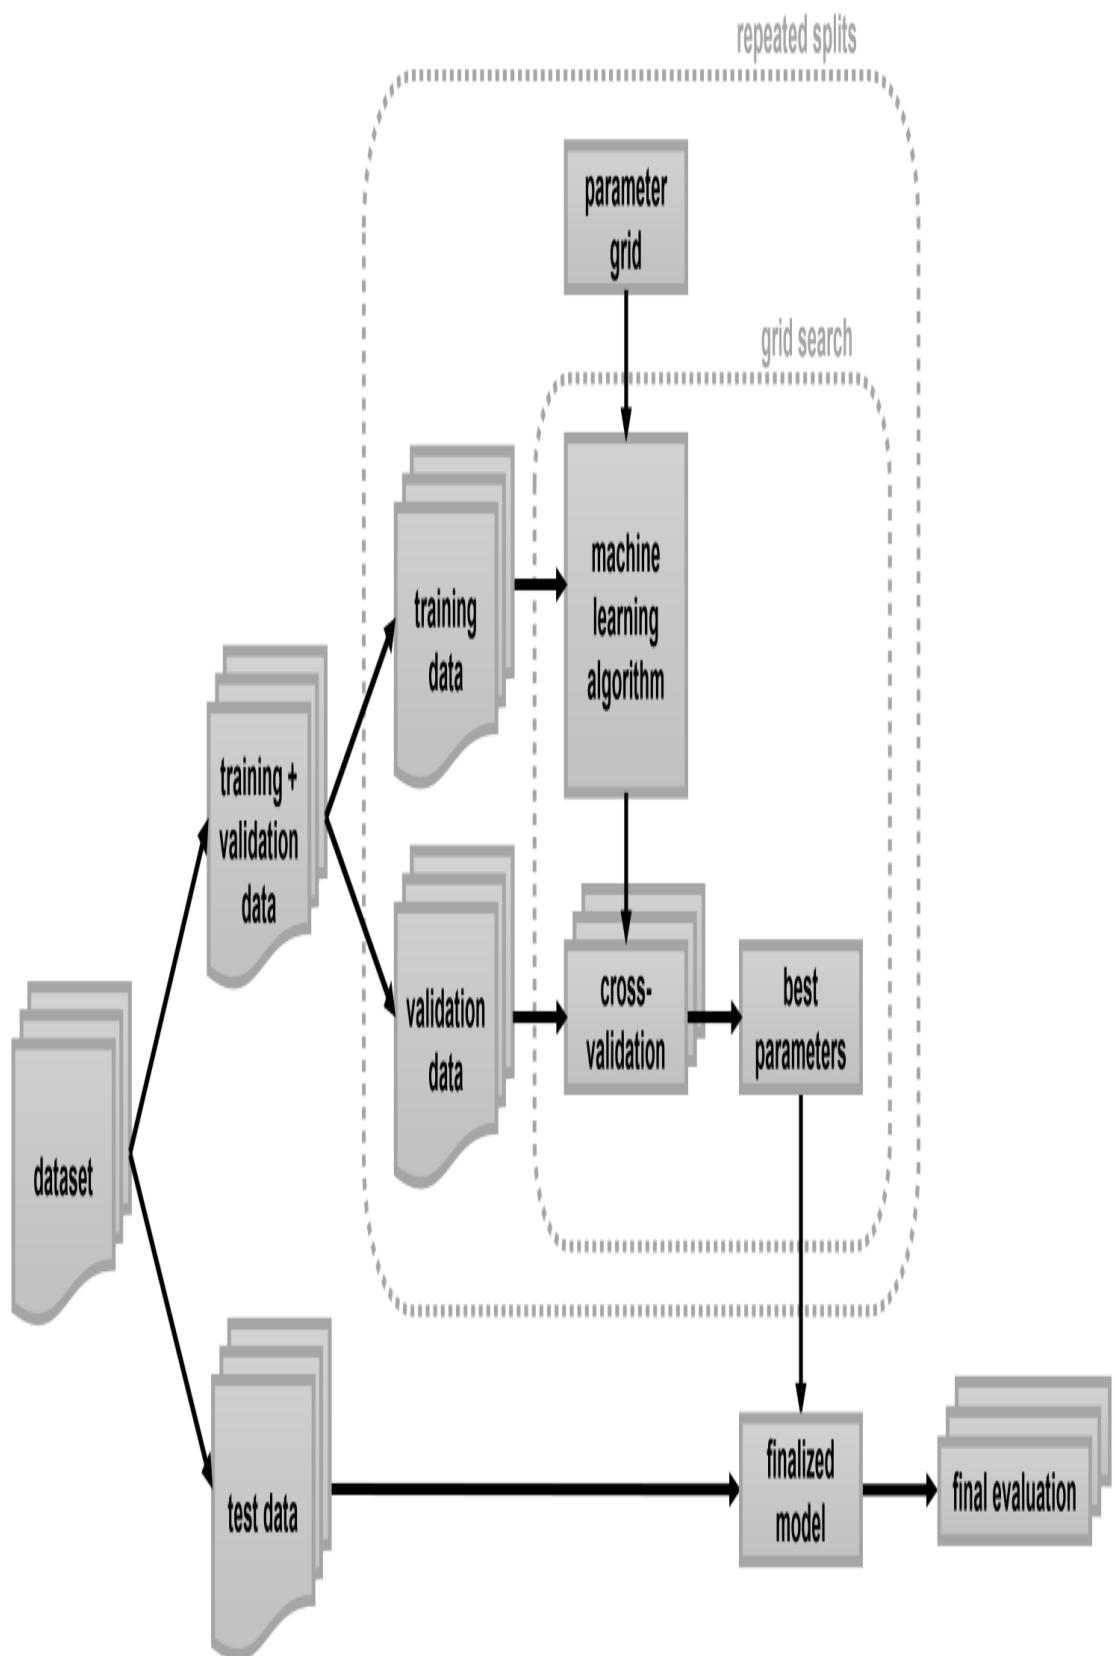
One potential danger of the grid search we just implemented is that the outcome might be relatively sensitive to how exactly we split the data. After all, we might have accidentally chosen a split that put most of the easy-to-classify data points in the test set, resulting in an overly optimistic score. Although we would be happy at first, as soon as we tried the model on some new held-out data, we would find that the actual performance of the classifier is much lower than expected.

Instead, we can combine grid search with cross-validation. This way, the data is split multiple times into training and validation sets, and cross-validation is performed at every step of the grid search to evaluate ...

Combining grid search with nested cross-validation

Although grid search with cross-validation makes for a much more robust model selection procedure, you might have noticed that we performed the split into training and validation sets still only once. As a result, our results might still depend too much on the exact training-validation split of the data.

Instead of splitting the data into training and validation sets once, we can go a step further and use multiple splits for cross-validation. This will result in what is known as **nested cross-validation**, and the process is illustrated in the following diagram:



In nested cross-validation, there is an outer loop over the grid search box that repeatedly splits the data into training and validation sets. For each of these splits, a grid search is run, which will report back a set of best parameter values. Then, for each outer split, we get a test score using the best settings.

Running a grid search over many parameters and on large datasets can be computationally intensive. A particular parameter setting on a particular cross-validation split can be done completely independently from the other parameter settings and models. Hence, parallelization over multiple CPU cores or a cluster is very important for grid search and cross-validation.

Now that we know how to find the best parameters of a model, let's take a closer look at the different evaluation metrics that we can use to score a model.

Scoring models using different evaluation metrics

So far, we have evaluated classification performance using accuracy (the fraction of correctly classified samples) and regression performance using R^2 . However, these are only two of the many possible ways to summarize how well a supervised model performs on a given dataset. In practice, these evaluation metrics might not be appropriate for our application, and it is important to choose the right metric when selecting between models and adjusting parameters.

When selecting a metric, we should always have the end goal of the machine learning application in mind. In practice, we are usually interested not just in making accurate predictions but also in using these predictions as part of a larger ...

Choosing the right classification metric

We talked about several essential scoring functions in [Chapter 3, First Steps in Supervised Learning](#). Among the most fundamental metrics for classification were the following:

- **Accuracy:** This counts the number of data points in the test set that have been predicted correctly and returns that number as a fraction of the test set size (`sklearn.metrics.accuracy_score`). This is the most basic scoring function for classifiers, and we have made extensive use of it throughout this book.
- **Precision:** This describes the ability of a classifier not to label a positive sample as a negative (`sklearn.metrics.precision_score`).
- **Recall (or sensitivity):** This describes the ability of a classifier to retrieve all of the positive samples (`sklearn.metrics.recall_score`).

Although precision and recall are important measures, looking at only one of them will not give us a good idea of the big picture. One way to summarize the two measures is known as the **f-score** or **f-measure** (`sklearn.metrics.f1_score`), which computes the harmonic mean of precision and recall as $2(precision \times recall) / (precision + recall)$.

Sometimes we need to do more than maximize accuracy. For example, if we are using machine learning in a commercial application, then the decision-making should be driven by the business goals. One of these goals might be to guarantee at least 90% recall. The challenge then becomes to develop a model that still has reasonable accuracy while satisfying all secondary requirements. Setting goals like this is often called **setting the operating point**.

However, it is often not clear what the operating point should be when developing a new system. To understand the problem better, it is important to investigate all possible trade-offs of precision and recall them at once. This is possible using a tool called the **precision-recall curve** (`sklearn.metrics.precision_recall_curve`).

Another commonly used tool to analyze the behavior of classifiers is the Receiver Operating Characteristic (ROC) curve.

The ROC curve considers all possible thresholds for a given classifier similar to the precision-recall curve, but it shows the false positive rate against the true positive rate instead of reporting precision and recall.

Choosing the right regression metric

Evaluation for regression can be done in similar detail as we did for classification. In [Chapter 3](#), *First Steps in Supervised Learning*, we also talked about some fundamental metrics for regression:

- **Mean squared error:** The most commonly used error metric for regression problems is to measure the squared error between the predicted and true target value for every data point in the training set, averaged across all data points (`sklearn.metrics.mean_squared_error`).
- **Explained variance:** A more sophisticated metric is to measure to what degree a model can explain the variation or dispersion of the test data (`sklearn.metrics.explained_variance_score`). Often, the amount of explained variance is measured using ...

Chaining algorithms together to form a pipeline

Most machine learning problems we have discussed so far consist of at least a preprocessing step and a classification step. The more complicated the problem, the longer this *processing chain* might get. One convenient way to glue multiple processing steps together and even use them in grid search is by using the `Pipeline` class from scikit-learn.

Implementing pipelines in scikit-learn

The `Pipeline` class itself has a `fit`, a `predict`, and a `score` method, which all behave just like any other estimator in scikit-learn. The most common use case of the `Pipeline` class is to chain different preprocessing steps together with a supervised model like a classifier.

Let's return to the breast cancer dataset from [Chapter 5, *Using Decision Trees to Make a Medical Diagnosis*](#). Using scikit-learn, we import the dataset and split it into training and test sets:

```
In [1]: from sklearn.datasets import load_breast_cancer...      import
numpy as np...      cancer = load_breast_cancer()...      X =
cancer.data.astype(np.float32)...      y = cancer.targetIn [2]: X_train,
X_test, y_train, y_test = train_test_split(... X, y, random_state=37
...
...
```

Using pipelines in grid searches

Using a pipeline in a grid search works the same way as using any other estimator.

We define a parameter grid to search over and construct `GridSearchCV` from the pipeline and the parameter grid. When specifying the parameter grid, there is, however, a slight change. We need to specify for each parameter which step of the pipeline it belongs to. Both parameters that we want to adjust, `c` and `gamma`, are parameters of `svc`. In the preceding section, we gave this step the name "`svm`". The syntax to define a parameter grid for a pipeline is to specify for each parameter the step name, followed by `_` (a double underscore), followed by the parameter name.

Hence, we would construct the parameter grid as follows:

```
In [8]: param_grid = {'svm__C': [0.001, 0.01, 0.1, 1, 10, 100],  
...                  'svm__gamma': [0.001, 0.01, 0.1, 1, 10, 100]}
```

With this parameter grid, we can use `GridSearchCV` as usual:

```
In [9]: grid = GridSearchCV(pipe, param_grid=param_grid, cv=10)  
...      grid.fit(X_train, y_train);
```

The best score in the grid is stored in `best_score_`:

```
In [10]: grid.best_score_  
Out[10]: 0.97652582159624413
```

Similarly, the best parameters are stored in `best_params_`:

```
In [11]: grid.best_params_
Out[11]: {'svm__C': 1, 'svm__gamma': 1}
```

But recall that the cross-validation score might be overly optimistic. To know the true performance of the classifier, we need to score it on the test set:

```
In [12]: grid.score(X_test, y_test)
Out[12]: 0.965034965034965
```

In contrast to the grid search we did before, now, for each split in the cross-validation, `MinMaxScaler` is refit with only the training splits, and no information is leaked from the test split into the parameter search.

This makes it easy to build a pipeline to chain together a whole variety of steps! You can mix and match estimators in the pipeline at will, you just need to make sure that every step in the pipeline provides a `transform` method (except for the last step). This allows an estimator in the pipeline to produce a new representation of the data, which, in turn, can be used as input to the next step.

The `Pipeline` class is not restricted to preprocessing and classification but can, in fact, join any number of estimators together. For example, we could build a pipeline containing feature extraction, feature selection, scaling, and classification, for a total of four steps. Similarly, the last step could be regression or clustering instead of classification.

Summary

In this chapter, we tried to complement our existing machine learning skills by discussing best practices in model selection and hyperparameter tuning. You learned how to tweak the hyperparameters of a model using grid search and cross-validation in both OpenCV and scikit-learn. We also talked about a wide variety of evaluation metrics and how to chain algorithms into a pipeline. Now, you are almost ready to start working on some real-world problems on your own.

In the next chapter, you will be introduced to an exciting and a new topic, that is, OpenVINO toolkit, which was one of the key releases in OpenCV 4.0.

Using OpenVINO with OpenCV

In the first chapter, we discussed various new additions in the OpenCV 4.0 release. One of the key releases to note is the OpenVINO toolkit. It's also interesting to note that the OpenVINO toolkit was selected as the 2019 Developer Tool of the Year by Embedded Vision Alliance.

In this chapter, we will focus only on how to use the OpenVINO toolkit with OpenCV. We will begin by installing the OpenVINO toolkit and then proceed to an interactive face detection demo with it. We will also learn to use OpenVINO Model Zoo with OpenCV and the **OpenVINO Inference Engine (IE)** with OpenCV. At the end of this chapter, we will also learn how to carry out image classification using OpenCV with OpenVINO IE.

In this chapter, we will cover the following topics:

- OpenVINO toolkit installation
- Interactive face detection demo
- Using OpenVINO Model Zoo with OpenCV
- Using OpenVINO IE with OpenCV
- Image classification using OpenCV with OpenVINO IE

Technical requirements

You can refer the code for this chapter at the following link: <https://github.com/PacktPublishing/Machine-Learning-for-OpenCV-Second-Edition/tree/master/Chapter12>.

Here is a summary of the software and hardware requirements:

- You will need OpenCV version 4.1.x (4.1.0 or 4.1.1 will both work just fine).
- You will need Python version 3.6 (any Python version 3.x will be fine).
- You will need Anaconda Python 3 to install Python and the required modules.
- You can use any OS—macOS, Windows, and Linux-based OSes—with this book. We recommend you have at least 4 GB RAM in your system.
- You don't need to have a GPU to run the code provided with this book.

Introduction to OpenVINO

OpenVINO (short for **Open Visual Inferencing and Neural Network Optimization**). It is designed to optimize various neural networks to speed up the inference stage. Inference, as we have discussed in previous chapters, is the process in which a trained neural network is used to generate results with unseen input data. For example, if a network is trained to classify a dog or cat, then if we feed the image of Tuffy (our neighbor's dog), it should be able to infer that the image is of a dog.

Considering that images and videos have become so common in today's world, there are a lot of deep neural networks trained to perform various operations, such as, multilabel classification, and motion tracking. Most of the inference performed in the world occurs on CPUs since GPUs are very expensive and usually do not suit the budget of individual AI engineers. In these cases, the speedup provided by OpenVINO toolkit is very crucial.

The speedup provided by OpenVINO toolkit consists of two steps. The first step focuses on the hardware specifications; it optimizes the network in a hardware-agnostic way using a Model Optimizer, which ships along with OpenVINO toolkit. The next step involves hardware-specific acceleration using OpenVINO IE.

The OpenVINO toolkit was developed by Intel, which is known for its optimized tools as well as hardware, and is focused on deep learning and artificial intelligence. It's not surprising to know that VPUs, GPUs, and FPGAs are also manufactured by Intel.

OpenVINO also provides optimized calls for the OpenCV and OpenVX libraries—the two most well-known computer vision libraries.

OpenVINO toolkit installation

In this section, we will use Intel's official instructions to install OpenVINO toolkit:

1. To start, first visit the OpenVINO toolkit download page (<https://software.intel.com/en-us/openvino-toolkit/choose-download>) and, depending on your system specifications, select and download the installer. You will have to first register your copy of the toolkit.
2. Use the installation instructions (<https://docs.openvinotoolkit.org/latest/index.html>) to install OpenVINO toolkit on your system.

OpenVINO toolkit will also install its own Intel-optimized version of OpenCV. If you already have OpenCV installed on your system, the installer will show that another version of OpenCV is already installed. It's better to install the ...

OpenVINO components

OpenVINO toolkit consists of the following main components:

- The **Deep Learning Deployment Toolkit (DLDT)** consists of Model Optimizer, the IE, pre-trained models and some tools to help you measure accuracy for your models.
- There's an optimized version of OpenCV compiled for Intel libraries (which are also optimized).
- There are OpenCL libraries.
- You get Intel's media SDK to speed up video processing.
- There's an optimized version of OpenVX.

Interactive face detection demo

The OpenVINO toolkit installation also provides various demo and sample applications. Just to test the installation, let's see whether we can run the interactive face detection demo.

First, we will move to the `samples` directory in the `deployment_tools/inference_engine` folder. You will find the various demo applications here, such as image classification and inference pipeline.

The interactive face detection demo takes in video as input and performs face detection coupled with age, gender, head-pose, emotion, and facial landmark detection. Depending on the kind of detection you want to perform, you can use a model from the list of the following pre-trained models:

- You can perform face detection only with

Using OpenVINO Inference Engine with OpenCV

In the previous section, we discussed how to run the interactive face detection demo. That's all good, but the question that still remains is how to harness the power of OpenVINO with your already existing OpenCV codes. Note that, here, we are emphasizing the utilization of the strength of OpenVINO with minimal changes in the code. This is very important because OpenVINO was not present in the earlier versions of OpenCV, including the more commonly used version 3.4.3. As a good developer, it's your job to make sure that your program supports the maximum number of systems and libraries.

Luckily for us, all it takes is just one line of code to start using OpenVINO Inference Engine for your OpenCV model's inference code, as shown in the following snippet:

```
cv::dnn::setPreferableBackend(DNN_BACKEND_INFERENCE_ENGINE); // C++  
setPreferableBackend(cv2.dnn.DNN_BACKEND_INFERENCE_ENGINE) # Python
```

And that's it! In a complete working example, this is how you will be using it:

```
net = cv2.dnn.readNetFromCaffe(prototxt, model)  
net.setPreferableBackend(cv2.dnn.DNN_BACKEND_INFERENCE_ENGINE)
```

Here, you can use any other method of reading your neural network. In this case, we are reading a Caffe model from the `.prototxt` and `.caffemodel` files.

Similarly, in the case of C++, this is how we can use it:

```
Net net = readNetFromCaffe(prototxt, model);
net.setPreferableBackend(DNN_BACKEND_INFERENCE_ENGINE);
```

Using OpenVINO Model Zoo with OpenCV

In the previous sections, we briefly discussed OpenVINO Model Zoo and how we can use OpenVINO IE with OpenCV. In this section, we will learn more about Model Zoo and what it offers.

OpenVINO Model Zoo is a collection of optimized pre-trained models that can be directly imported into OpenVINO for inference. The importance of this feature lies in the fact that one of the major reasons behind OpenVINO's speedup is the optimized model file that it takes for inference. The underlying inference principle is still the same as most deep learning inference toolkits and languages, such as OpenCV. OpenCV's `dnn` module uses this speedup principle of OpenVINO by using it as the default backend for all inference tasks.

...

Image classification using OpenCV with OpenVINO Inference Engine

The final topic that we will discuss in this chapter is how to carry out image classification using OpenCV with OpenVINO Inference Engine.

Before we jump into the details, let's have a brief look at an image classification problem. Image classification, also known as **image recognition**, is a part of the deep learning set of tasks and is perhaps one of the most common ones. In this task, a set of images are provided as input to the model and the model outputs the class or the label of the input images.

A common example of this is the dog and cat classification problem where a model is trained on a large number of images of cats and dogs and then, during the testing phase, the model predicts whether the input image is an image of a cat or a dog.

While this might look like a very naive problem, image classification carries a lot of importance in industrial applications. For example, if your camera boasts of having AI powers, what it means is that it can recognize the objects present in the image and change the image settings accordingly—whether it is an image of natural scenery or an Instagram-worthy snap of some food. The following image shows the output of an AI mobile phone camera:



Consider the preceding image I took of the roof of my house. Note that the camera, when switched to AI mode, was able to detect that I was taking a picture of plants and automatically changed the settings to match that. All of this is possible only because of image classification. Now, consider that you, as a computer vision engineer, are trying to train a model that can recognize whether the image is of a plant, waterfall, or human.

All of the efforts that you put in training the model will be wasted if your model cannot infer the class or the label of the image within a few milliseconds. No user wants to wait for even a couple of seconds for the camera to detect the object and change the settings.

This brings us back to the importance of OpenVINO's Inference Engine. OpenVINO has its own version of an image classification toolkit that can be used as follows.

Image classification using OpenVINO

Let's see how we can use the image classification demo present in OpenVINO's installation directory:

1. First, move to the `deployment_tools/demo` directory in your OpenVINO installation directory.
2. Next, let's run image classification on a demo image that's already present in the directory:

```
./demo_squeezeenet_download_convert_run.sh
```

Here's the result that I obtained:

```
Top 10 results:
```

```
Image /opt/intel/computer_vision_sdk_fpga_2018.2.298/deployment_tools/demo/./demo/car.png
```

```
817 0.8363345 label sports car, sport car
511 0.0946488 label convertible
479 0.0419131 label car wheel
751 0.0091071 label racer, race car, racing car
436 0.0068161 label beach wagon, station wagon, wagon, estate car, beach waggon, station waggon, waggon
656 0.0037564 label minivan
586 0.0025741 label half track
717 0.0016069 label pickup, pickup truck
864 0.0012027 label tow truck, tow car, wrecker
581 0.0005882 label grille, radiator grille
```

```
[ INFO ] Execution successful
```

```
#####
#
```

```
Demo completed successfully.
```

Let's also run another demo that uses the same image, the inference pipeline demo, which shows the speed of OpenVINO's Inference Engine quite well:

```
./demo_security_barrier_camera.sh
```

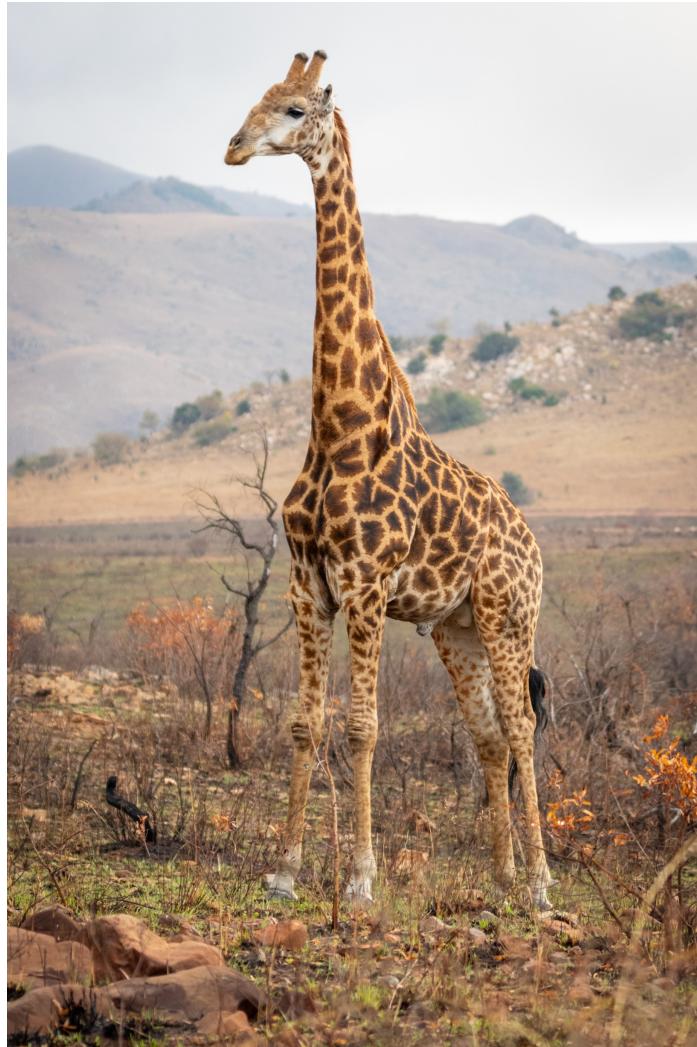
Here's the output image:

Since we are using the same image in ...

Image classification using OpenCV with OpenVINO

Let's first create an image classification inference code using OpenCV. Since we are only concerned with inference, we will use a pre-trained model:

1. First, let's download the Caffe model files, `deploy.prototxt` and `bvlc_reference_caffenet.caffemodel`, which can be obtained from Berkley Visions' repository (https://github.com/BVLC/caffe/tree/master/models/bvlc_reference_caffenet). Make sure that you download both files in your current working directory. We will also need a text file with the class labels mentioned. You can get it from https://github.com/torch/tutorials/blob/master/7_imagenet_classification/synset_words.txt.
2. Let's also use a sample image of a giraffe for image classification:



Next, let's start writing some code for image classification using OpenCV with OpenVINO.

3. Let's start off by importing some modules:

```
import numpy as np  
import cv2
```

4. Next, let's specify the model files:

```
image = cv2.imread("animal-barbaric-brown-1319515.jpg")
labels_file = "synset_words.txt"
prototxt = "deploy.prototxt"
caffemodel = "bvlc_reference_caffenet.caffemodel"
```

5. Now, let's read the labels from the labels text file:

```
rows = open(labels_file).read().strip().split("\n")
classes = [r[r.find(" ") + 1:].split(",")[0] for r in rows]
```

6. Let's specify the backend that we will use for inference:

```
net = cv2.dnn.readNetFromCaffe(prototxt, caffemodel)
net.setPreferableBackend(cv2.dnn.DNN_BACKEND_INFERENCE_ENGIN
net.setPreferableTarget(cv2.dnn.DNN_TARGET_CPU)
```

7. Let's carry out some basic image processing on the input image:

```
blob = cv2.dnn.blobFromImage(image, 1, (224, 224), (104, 117, 123))
```

8. Finally, let's pass this image to the model and get the output:

```
net.setInput(blob)
predictions = net.forward()
```

9. Let's obtain the top 10 predictions for the image of a giraffe we passed to the model:

```
indices = np.argsort(predictions[0])[::-1][:5]
```

10. Finally, let's display the top 10 predictions:

```
for index in indices:  
    print("label: {}, prob.: {:.5}".format(classes[index],  
                                           predictions[0][index]))
```

And, surprisingly, here is the result we obtain:

```
label: cheetah, prob.: 0.98357  
label: leopard, prob.: 0.016108  
label: snow leopard, prob.: 7.2455e-05  
label: jaguar, prob.: 4.5286e-05  
label: prairie chicken, prob.: 3.8205e-05
```

Notice that our model thought that the `giraffe` image we passed as input was actually a `cheetah` image. Why do you think that's the case? That's because `giraffe` was not present in the list of classes we had. So, the model came up with the closest match, which was because of the similar colored spots present on cheetahs and giraffes. So, the next time you carry out image classification, make sure the class is actually present in the label list.

We can also carry out a comparison between various backends to see the speedup obtained using OpenVINO's Inference Engine as the backend. Here's how it can be done. We need to change just one line in the preceding code:

```
net.setPreferableBackend(cv2.dnn.DNN_BACKEND_INFERENCE_ENGINE)
```

We can choose between the following backends:

- `cv2.dnn.DNN_BACKEND_DEFAULT`: This is if you have OpenVINO installed and will use it as the default backend.
- `cv2.dnn.DNN_BACKEND_HALIDE`: This requires OpenCV to be built with Halide. You can find detailed documentation for this at https://docs.opencv.org/4.1.0/de/d37/tutorial_dnn_halide.html.
- `cv2.dnn.DNN_BACKEND_OPENCV`: This is the best choice for carrying out a comparison between both the backends.

So, all you need to do is run the same code but replace the preceding line of code with the following:

```
net.setPreferableBackend(cv2.dnn.DNN_BACKEND_OPENCV)
```

And that's it! You can now carry out a comparison to see the speedup you obtain by using OpenVINO's Inference Engine as the backend.

You won't be able to see much difference in speed. To get a noticeable difference, use a `for` loop to carry out the inference 100 times, add up the total time taken during each step, and then divide it by 100 to obtain an average.

Summary

In this chapter, we briefly looked at the OpenVINO toolkit—what it is, what it's used for, and how we can install it. We also looked at how to run the demos and samples provided with the toolkit to understand and witness the power of OpenVINO. Finally, we saw how to harness this power in our pre-existing OpenCV codes by just adding one line specifying the backend to be used for model inference.

You might have also noticed that we didn't cover much hands-on content in this chapter. That's because OpenVINO is more suited for deep learning applications, which are not in the scope of this book. If you are a deep learning enthusiast, you should definitely go through the documentation provided by Intel on the OpenVINO toolkit and get started. ...

Conclusion

Congratulations! You have just made a big step toward becoming a machine learning practitioner. Not only are you familiar with a wide variety of fundamental machine learning algorithms, but you also know how to apply them to both supervised and unsupervised learning problems. Moreover, you were introduced to a new and exciting topic, OpenVINO Toolkit. In the previous chapter, we learned how to install OpenVINO and run an interactive face detection and image classification demo, among others. I am sure you have enjoyed learning about those topics.

Before we part ways, I want to give you some final words of advice, point you toward some additional resources, and give you some suggestions on how you can further improve your machine learning and data science skills. In this chapter, we will learn how to approach a machine learning problem and build our own estimator. We will learn how to write our own OpenCV-based classifier in C++ and a scikit-learn-based classifier in Python.

In this chapter, we will cover the following topics:

- Approaching a machine learning problem
- Writing your own OpenCV-based classifier in C++
- Writing your own scikit-learn-based classifier in Python
- Where to go from here

Technical requirements

You can refer to the code for this chapter from the following link: <https://github.com/PacktPublishing/Machine-Learning-for-OpenCV-Second-Edition/tree/master/Chapter13>.

Here is a summary of the software and hardware requirements:

- You will need OpenCV version 4.1.x (4.1.0 or 4.1.1 will both work just fine).
- You will need Python version 3.6 (any Python version 3.x will be fine).
- You will need Anaconda Python 3 for installing Python and the required modules.
- You can use any OS—macOS, Windows, and Linux-based OSes along with this book. We recommend you have at least 4 GB RAM in your system.
- You don't need to have a GPU to run the code provided along with this book.

Approaching a machine learning problem

When you see a new machine learning problem in the wild, you might be tempted to jump ahead and throw your favorite algorithm at the problem—perhaps the one you understood best or had the most fun implementing. But knowing beforehand which algorithm will perform best on your specific problem is not often possible.

Instead, you need to take a step back and look at the bigger picture. Before you get in too deep, you will want to define the actual problem you are trying to solve. For example, do you already have a specific goal in mind, or are you just looking to do some exploratory analysis and find something interesting in the data? Often, you will start with a general goal, such as detecting spam email messages, making movie recommendations, or automatically tagging your friends in pictures uploaded to a social media platform. However, as we have seen throughout this book, there are often several ways to solve a problem. For example, we recognized handwritten digits using logistic regression, k-means clustering, and deep learning. Defining the problem will help you to ask the right questions and make the right choices along the way.

As a rule of thumb, you can use the following five-step procedure to approach machine learning problems in the wild:

1. **Categorize the problem:** This is a two-step process:

1. Categorize by input: Simply speaking, if you have labeled data, it's a supervised learning problem. If you have unlabeled data and want to find structure, it's an unsupervised learning problem. If you want to optimize an objective function by interacting with an environment, it's a reinforcement learning problem.

2. Categorize by output: If the output of your model is a number, it's a regression problem. If the output of your model is a class (or category), it's a classification problem. If the output of your model is a set of input groups, it's a clustering problem.

2. Find the available algorithms: Now that you have categorized the problem, you can identify the algorithms that are applicable and practical to implement using the tools at our disposal. Microsoft has created a handy algorithm cheat sheet that shows which algorithms can be used for which category of problems. Although the cheat sheet is tailored toward **Microsoft Azure**, you might find it generally helpful.

The machine learning algorithm cheat sheet PDF (by Microsoft Azure) can be downloaded from <http://aka.ms/MLcheatsheet>.

3. Implement all of the applicable algorithms (prototyping): For any given problem, there are usually a handful of candidate algorithms that could do the job. So, how do you know which one to pick? Often, the answer to this problem is

not straightforward, so you have to resort to trial and error.

Prototyping is best done in two steps:

1. You should aim for a quick and dirty implementation of several algorithms with minimal feature engineering. At this stage, you should mainly be interested in seeing which algorithm behaves better at a coarse scale. This step is a bit like hiring: you're looking for any reason to shorten your list of candidate algorithms. Once you have reduced the list to a few candidate algorithms, the real prototyping begins.

2. Ideally, you would want to set up a machine learning pipeline that compares the performance of each algorithm on the dataset using a set of carefully selected evaluation criteria (see [chapter 11](#), *Selecting the Right Model with Hyperparameter Tuning*). At this stage, you should only be dealing with a handful of algorithms, so you can turn your attention to where the real magic lies: feature engineering.

4. **Feature engineering:** Perhaps even more important than choosing the right algorithm is choosing the right features to represent the data. You can read all about feature engineering in [chapter 4](#), *Representing Data and Engineering Features*.

5. **Optimize hyperparameters:** Finally, you also want to optimize an algorithm's hyperparameters. Examples might include the number of principal components of PCA, the

parameter, k , in the k-nearest neighbor algorithm, or the number of layers and learning rate in a neural network. You can look at [chapter 11, *Selecting the Right Model with Hyperparameter Tuning*](#), for inspiration.

Building your own estimator

In this book, we visited a whole variety of machine learning tools and algorithms that OpenCV provides straight out of the box. And, if for some reason, OpenCV does not provide exactly what we are looking for, we can always fall back on scikit-learn.

However, when tackling more advanced problems, you might find yourself wanting to perform some very specific data processing that neither OpenCV nor scikit-learn provide, or you might want to make slight adjustments to an existing algorithm. In this case, you may want to create your own estimator.

Writing your own OpenCV-based classifier in C++

Since OpenCV is one of those Python libraries that does not contain a single line of Python code under the hood (I'm kidding, but it's close), you will have to implement your custom estimator in C++. This can be done in four steps:

1. Implement a C++ source file that contains the main source code. You need to include two header files, one that contains all the core functionality of OpenCV (`opencv.hpp`) and another that contains the machine learning module (`ml.hpp`):

```
#include <opencv2/opencv.hpp>
#include <opencv2/ml/ml.hpp>
#include <stdio.h>
```

Then, an estimator class can be created by inheriting from the `StatModel` class:

```
class MyClass : public cv::ml::StatModel
{
public:
```

Next, you define constructor and destructor of the class:

```
MyClass()
{
    cout("MyClass constructor\n");
```

```
    }
~MyClass() {}
```

Then, you also have to define some methods. These are what you would fill in to make the classifier actually do some work:

```
int getVarCount() const
{
    // returns the number of variables in training samples
    return 0;
}

bool empty() const
{
    return true;
}

bool isTrained() const
{
    // returns true if the model is trained
    return false;
}

bool isClassifier() const
{
    // returns true if the model is a classifier
    return true;
}
```

The main work is done in the `train` method, which comes in two flavors (accepting either `cv::ml::TrainData` OR `cv::InputArray` as input):

```
bool train(const cv::Ptr<cv::ml::TrainData>& trainData,
           int flags=0) const
{
    // trains the model
    return false;
}

bool train(cv::InputArray samples, int layout,
```

```
        cv::InputArray responses)
    {
        // trains the model
        return false;
}
```

You also need to provide a `predict` method and a `scoring` function:

```
float predict(cv::InputArray samples,
              cv::OutputArray results=cv::noArray(),
              int flags=0) const
{
    // predicts responses for the provided samples
    return 0.0f;
}

float calcError(const cv::Ptr<cv::ml::TrainData>& data,
                bool test, cv::OutputArray resp)
{
    // calculates the error on the training or test dataset
    return 0.0f;
}
};
```

The last thing to do is to include a `main` function that instantiates the class:

```
int main()
{
    MyClass myclass;
    return 0;
}
```

2. Write a CMake file called `CMakeLists.txt`:

```
cmake_minimum_required(VERSION 2.8)
project(MyClass)
find_package(OpenCV REQUIRED)
```

```
add_executable(MyClass MyClass.cpp)
target_link_libraries(MyClass ${OpenCV_LIBS})
```

3. Compile the file on the command line by typing the following commands:

```
$ cmake
$ make
```

4. Run the executable `MyClass` method, which was generated by the last command and which should lead to the following output:

```
$ ./MyClass
MyClass constructor
```

Writing your own scikit-learn-based classifier in Python

Alternatively, you can write your own classifier using the scikit-learn library.

You can do this by importing `BaseEstimator` and `ClassifierMixin`. The latter will provide a corresponding `score` method, which works for all classifiers:

1. Optionally, first, you can overwrite the `score` method to provide your own metric `score` method:

```
In [1]: import numpy as np...      from sklearn.base import  
BaseEstimator, ClassifierMixin
```

2. Then, you can define a class that inherits from both `BaseEstimator` and `ClassifierMixin`:

```
In [2]: class MyClassifier(BaseEstimator, ClassifierMixin):...  
"""An example classifier"""
```

3. You need to provide a constructor, `fit` and `predict` methods.
The constructor defines all the parameters ...

Where to go from here

The goal of this book was to introduce you to the world of machine learning and prepare you to become a machine learning practitioner. Now that you know everything about the fundamental algorithms, you might want to investigate some topics in more depth.

Although it is not necessary to understand all of the details of all of the algorithms we implemented in this book, knowing some of the theory behind them might just make you a better data scientist.

If you are looking for more advanced material, then you might want to consider some of the following classics:

- *Stephen Marsland, Machine Learning: An Algorithmic Perspective, Second Edition, Chapman and Hall/Crc, ISBN 978-146658328-3, 2014*
- *Christopher M. Bishop, Pattern Recognition and Machine Learning. Springer, ISBN 978-038731073-2, 2007*
- *Trevor Hastie, Robert Tibshirani, and Jerome Friedman, The Elements of Statistical Learning: Data Mining, Inference, and Prediction. Second Edition, Springer, ISBN 978-038784857-0, 2016*

When it comes to software libraries, we already learned about two essential ones—OpenCV and scikit-learn. Often, using Python is great for trying out and evaluating models, but larger web services and applications are more commonly written in Java or C++.

For example, the C++ package is **Vowpal Wabbit** (VW), which comes with its own command-line interface. For running machine learning algorithms on a cluster, people often use `mllib`, a **Scala** library built on top of **Spark**. If you are not married to Python, you

might also consider using R, another common language of data scientists. R is a language designed specifically for statistical analysis and is famous for its visualization capabilities and the availability of many (often highly specialized) statistical modeling packages.

No matter which software you choose going forward, I guess the most important advice is to keep practicing your skills. But you already knew that. There are a number of excellent datasets out there that are just waiting for you to analyze them:

- Throughout this book, we made great use of the example datasets that are built into scikit-learn. In addition, scikit-learn provides a way to load datasets from external services, such as mldata.org. Refer to <http://scikit-learn.org/stable/datasets/index.html> for more information.
- Kaggle is a company that hosts a wide range of datasets as well as competitions on their website, <http://www.kaggle.com>. Competitions are often hosted by a variety of companies, nonprofit organizations, and universities, and the winner can take home some serious monetary prizes. A disadvantage of competitions is that they already provide a particular metric to optimize and usually a fixed, preprocessed dataset.
- The OpenML platform (<http://www.openml.org>) hosts over 20,000 datasets with over 50,000 associated machine learning tasks.
- Another popular choice is the UC Irvine Machine Learning Repository (<http://archive.ics.uci.edu/ml/index.php>), hosting over 370 popular and well-maintained datasets through a searchable interface.

Finally, if you are looking for more example code in Python, a number of excellent books nowadays come with their own GitHub repository:

- *Jake Vanderplas, Python Data Science Handbook: Essential Tools for Working with Data. O'Reilly, ISBN 978-149191205-8, 2016, <https://github.com/jakevdp/PythonDataScienceHandbook>*
- *Andreas Muller and Sarah Guido, Introduction to Machine Learning with Python: A Guide for Data Scientists. O'Reilly, ISBN 978-144936941-5, 2016, https://github.com/amueller/introduction_to_ml_with_python*
- *Sebastian Raschka, Python Machine Learning. Packt, ISBN 978-178355513-0, 2015, <https://github.com/rasbt/python-machine-learning-book>*

Summary

In this chapter, we learned how to approach a machine learning problem and built our own estimator. We learned how to write our own OpenCV-based classifier in C++ and scikit-learn-based classifier in Python.

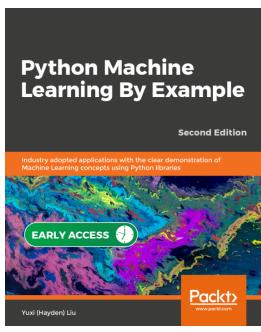
In this book, we covered a lot of theory and practice. We discussed a wide variety of fundamental machine learning algorithms, both supervised or unsupervised, and illustrated best practices as well as ways to avoid common pitfalls, and we touched upon a variety of commands and packages for data analysis, machine learning, and visualization.

If you made it this far, you have already made a big step toward machine learning mastery. From here on out, I am confident you will do just fine on your own.

All that's left to say is farewell! ...

Other Books You May Enjoy

If you enjoyed this book, you may be interested in these other books by Packt:

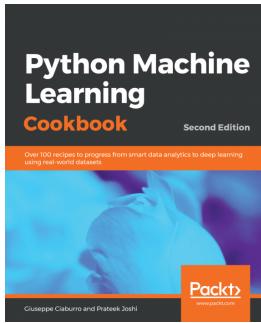


Python Machine Learning By Example - Second Edition
Yuxi (Hayden) Liu

ISBN: 978-1-78961-672-9

- Understand the important concepts in machine learning and data science
- Use Python to explore the world of data mining and analytics
- Scale up model training using varied data complexities with Apache Spark
- Delve deep into text and NLP using Python libraries such NLTK and gensim
- Select and build an ML model and evaluate and optimize its performance

- Implement ML algorithms from scratch in Python, TensorFlow, and scikit-learn



Python Machine Learning Cookbook - Second Edition

Prateek Joshi, Giuseppe Ciaburro

ISBN: 978-1-78980-845-2

- Use predictive modeling and apply it to real-world problems
- Explore data visualization techniques to interact with your data
- Learn how to build a recommendation engine
- Understand how to interact with text data and build models to analyze it
- Work with speech data and recognize spoken words using Hidden Markov Models
- Get well versed with reinforcement learning, automated ML, and transfer learning
- Work with image data and build systems for image recognition and biometric face recognition

- Use deep neural networks to build an optical character recognition system

Leave a review - let other readers know what you think

Please share your thoughts on this book with others by leaving a review on the site that you bought it from. If you purchased the book from Amazon, please leave us an honest review on this book's Amazon page. This is vital so that other potential readers can see and use your unbiased opinion to make purchasing decisions, we can understand what our customers think about our products, and our authors can see your feedback on the title that they have worked with Packt to create. It will only take a few minutes of your time, but is valuable to other potential customers, our authors, and Packt. Thank you!